

March 14, 2025

Suzanne H. Plimpton
Reports Clearance Officer
National Science Foundation

Via ostp-ai-rfi@nitrd.gov

Re: ABA Comment Letter responding to Request for Information on the Development of an Artificial Intelligence (AI) Action Plan

Dear Ms. Plimpton,

The American Bankers Association¹ (ABA) appreciates the opportunity to provide input from the financial services industry into the Development of an Artificial Intelligence (AI) Action Plan (“Plan”).² This Plan, as directed by President Trump’s Executive Order issued on January 23, 2025, will define the priority policy actions needed to sustain and enhance America’s AI dominance, and to ensure that unnecessarily burdensome requirements do not hamper private sector AI innovation.³ We hope these comments are helpful in formulating the Plan and look forward to working with the policymakers tasked under it on more specific deliverables.

The world has been captivated by the reality and promise of generative AI since its public unveiling in late 2022. It has powered investment, fueled dreams, and inspired new ways of thinking. It has the potential to improve the daily lives of individual Americans even while impacting geopolitics on the grandest scale. Currently, the leading AI developers call the United States of America home. But if US leadership in the field of AI is to be preserved, it will be necessary for the federal government to foster an environment of trust and consistency. Without meaningful standards to channel efforts in a constructive manner, companies scrambling for dominance would be so many “crabs in a bucket” seeking an edge over competitors.

Such an arrangement simply will not work given the interconnectedness of the AI ecosystem, in which Big Tech companies develop large language models (LLMs), downstream companies deploy these LLMs for particular use cases, and end users (agents or customers of the deployer) utilize these applications. The reliance of companies on third party entities in this arrangement is clear, driving home the need for baseline standards and due diligence. Complicating the picture

¹ *The American Bankers Association is the voice of the nation’s \$24.1 trillion banking industry, which is composed of small, regional and large banks that together employ approximately 2.1 million people, safeguard \$19.2 trillion in deposits and extend \$12.7 trillion in loans.*

² <https://www.federalregister.gov/documents/2025/02/06/2025-02305/request-for-information-on-the-development-of-an-artificial-intelligence-ai-action-plan>

³ <https://www.federalregister.gov/documents/2025/01/31/2025-02172/removing-barriers-to-american-leadership-in-artificial-intelligence>

even further is the fact that many LLM developers train their models by screen scraping data, potentially in contravention of website terms of service. This behavior can throttle internet performance and could, in time, disincentivize content creation.

But where to find these standards? Fortunately for the AI ecosystem, they already exist. Banks are a model for how other industries can explore AI-enabled use cases in a fruitful and sustainable manner. The compliance requirements, model risk management expectations, and supervision by specialized regulators has resulted in an environment of trust and responsible innovation, a prerequisite for prosperity. Accordingly, the financial services sector can be a leader to encourage the proliferation of these features. The framework under which banks operate can, with some clarifications and adjustments, effectively manage risk without endangering American dominance. We urge policymakers to extend similar standards to nonbank participants in the AI ecosystem, and we look forward to collaborating with the Administration, agencies, Congress, and other stakeholders to ensure that outcome.

To that end, this document contains a number of high-level policy recommendations. These are:

1. Congress must pass comprehensive laws establishing an AI risk management framework with strong preemptions of state requirements and which do not create duplicative or inconsistent obligations for banks.
2. Agencies should identify clear regulatory outcomes and objectives, while enabling regulated entities the ability to deploy effective risk management techniques based on common sense and best practices. Agencies must have perspective on the entire AI ecosystem rather than merely isolating the regulated entities, using existing authorities to address risks wherever they arise. Additionally, regulators must be transparent about the ways they utilize AI themselves.
3. The Federal Reserve (Fed), the Office of the Comptroller of the Currency (OCC), and the Federal Deposit Insurance Corporation (FDIC) should update model risk management guidance, subject to a notice and comment period. Moreover, field examiners must be trained not to focus on granular matters such as LLM's code but rather on the inputs, outputs, and outcomes of models.
4. Policymakers should encourage adoption of voluntary standards and frameworks where possible in order to encourage cross-sector collaboration. For example, this could include model cards to allow for explainability without divulging confidential commercial information, or certifications to help demonstrate sufficient maturity of policies and procedures.

These recommendations will be explored in additional detail below.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Introduction

Financial institutions have a long history of deploying and controlling risk related to novel technologies, and this includes both the traditional and generative iterations of the AI spectrum. Moreover, banks' use of AI is subject to a strong foundation of compliance to manage attendant risks. This foundation includes: three lines of internal risk-management defenses; application of technology-neutral laws, regulations, and guidance; and validation of the effectiveness of the framework through regular examinations by the bank regulatory agencies.

The following comments are based on discussions with ABA's AI Working Group, which is an interdisciplinary body of data scientists, technologists, compliance, legal, risk, security, audit, and human resources experts representing banks of all sizes. This group routinely comes together to formulate policy positions as well as to discuss operational best practices. Our comments are also informed by the work of the Financial Services Sector Coordinating Council's R&D Committee, which the ABA co-chairs.⁴

Our comments are organized into four sections: (I) Current State of AI Regulation; (II) How Banks Use AI; (III) Risk Management of AI; and (IV) Recommendations for Policymakers.

I. Current State of AI Regulation

AI is an umbrella term for various enabling technologies and is not a standalone program, product, or service. Rather, it is embedded into various bank operations, including security and products and services provided to customers. While it is useful to attach definitions to AI for purposes of framing the discussion, ABA strongly believes that any requirements should be rooted in use cases and the resultant risks presented, not in the technology that helps deliver them.

There is a complex overlay of applicable laws, regulations, and supervisory guidance that is relevant to AI usage by banks. The most important of these are model risk management

⁴ The FSSCC R&D committee convened a series of meetings with financial sector experts to explore the role of AI in cybersecurity and fraud, which resulted in a six page report that the US Treasury included in its March 2024 paper on AI and cybersecurity. See, *Appendix*, <https://home.treasury.gov/system/files/136/Managing-Artificial-Intelligence-Specific-Cybersecurity-Risks-In-The-Financial-Services-Sector.pdf>

expectations issued by the Fed, OCC, and FDIC.⁵ The guidance applies to models, which are defined as:

a quantitative method, system, or approach that applies statistical, economic, financial, or mathematical theories, techniques, and assumptions to process input data into quantitative estimates. A *model* consists of three components: an information input component, which delivers assumptions and data to the model; a processing component, which transforms inputs into estimates; and a reporting component, which translates the estimates into useful business information... The definition of *model* also covers quantitative approaches whose inputs are partially or wholly qualitative or based on expert judgment, provided that the output is quantitative in nature.”⁶

Further, even “more qualitative approaches used by banking organizations— i.e., those not defined as models according to this guidance—should also be subject to a rigorous control process.”⁷ Thus, treatment is scaled based on the use case and risk appetite. In 2021, the Federal Reserve, the OCC, and the FDIC issued interagency guidance addressing model risk management to support Bank Secrecy Act/Anti-Money Laundering and Office of Foreign Assets Control Compliance (BSA/AML and OFAC).⁸

Banks of all sizes use third-party sources to develop AI models, and the challenges associated with managing the risk posed by these models are universal across the banking industry. Accordingly, the interagency guidance issued by the Fed, the OCC, and the FDIC on third-party risk management is a crucial tool in conducting due diligence.⁹

During the prior Administration, several agencies released materials regarding AI.¹⁰ Although these materials did not purport to create new policy, entities subject to the agencies’ jurisdiction carefully monitor such activity in order to divine regulatory attitudes. A more open process that

⁵ SR 11-7, OCC Bulletin 2011-12, and FIL-22-2017, respectively. The OCC also released a booklet for its examiners to use as an aid when supervising banks’ model risk management programs; see <https://www.occ.treas.gov/publications-and-resources/publications/comptrollers-handbook/files/model-risk-management/index-model-risk-management.html>.

⁶ SR 11-7, page 3.

⁷ Id.

⁸ SR 21-8, OCC Bulletin 2021-19, and FIL-27-2021, respectively.

⁹ Guidance on Third-Party Relationships: Risk Management, <https://www.federalregister.gov/documents/2023/06/09/2023-12340/interagency-guidance-on-third-party-relationships-risk-management>.

¹⁰ See, e.g., Joint Statement on Enforcement Efforts Against Discrimination and Bias in Automated Systems, https://files.consumerfinance.gov/f/documents/cfpb_joint-statement-enforcement-against-discrimination-bias-automated-systems_2023-04.pdf; CFPB, Advisory on chatbots in consumer finance, <https://www.consumerfinance.gov/data-research/research-reports/chatbots-in-consumer-finance/chatbots-in-consumer-finance/>; CFPB, Circular 2023-03, “Adverse action notification requirements and the proper use of the CFPB’s sample forms provided in regulation B,” https://files.consumerfinance.gov/f/documents/cfpb_adverse_action_notice_circular_2023-09.pdf.

gathers and considers industry feedback would lead to superior outcomes—a process exemplified by this RFI.

II. How Banks Use AI

Banks have a long history of using traditional AI within the risk management framework outlined above. “Traditional AI” can be thought of as a system designed to respond to a particular set of inputs. It “learns” from the data and makes decisions or recommendations based upon the data, but does not create anything new. Traditional AI is generally predictable and has been vetted through years of usage and supervisory feedback.

Generative AI can be thought of as a form of AI that can create something new from prompts it receives. Generative AI models are “trained” on a set of data and learn the underlying patterns to generate new data that mirrors the training set. This can include images, music, and computer code, among other things.

Generative applications are still in the nascent phase; there is much more work to be done to fully understand and trust the technology, as well as the regulators’ attitudes towards it. Banks are paying particular attention to agentic AI, which acts as an assistant to employees who in turn function as a “human in the loop” control. Banks are proceeding very cautiously with generative AI, particularly with regard to customer-facing applications.

Examples of AI Use Cases

Banks are using traditional and generative forms of AI in the following ways (use cases utilize the traditional form of AI unless otherwise noted):

- Cybersecurity- AI is used to detect and respond to potential cyberattacks more quickly and efficiently than human intelligence could accomplish alone. AI-based network security software can monitor incoming and outgoing network traffic to identify suspicious patterns to aid security analysts in their initial classification by reducing the number of false positives. Banks may also utilize generative AI to pinpoint malicious code as well as aide internal developers in identifying vulnerabilities in their own code.
- Fraud Prevention - AI models using predictive analytics help banks proactively find anomalies in transactions and identify outliers that do not conform with customers’ past patterns or payment activity. AI models not only improve the performance of fraud detection capabilities, but also help catch fraudulent activity before it impacts customers.
- Lending- Banks use AI across lending processes to help identify accounts that can be approved for credit, as well as loan amounts and pricing. AI assists banks with evaluation

of creditworthiness and improves efficiency in decision-making. It can also provide metrics around key life indicators such as attrition rates for mortgages.

- Customer Service- AI assists banks with learning how customers interact with their products and services. In addition, AI can perform sentiment analysis to gain insight into satisfaction levels. This data can be used to better understand customer interactions and how to improve them. Generative AI's ability to summarize conversations with contact center personnel may also be used to help representatives grasp the backstory when customers have to re-engage on the same issue with a different individual and can streamline the process of finding answers to customer inquiries.
- Chatbots- Chatbots powered by traditional AI are commonly used and can respond with static responses to certain keywords. Customers gravitate towards these channels due to ease of use and preference for self-service. In fact, in certain circumstances (such as to obtain a routing number) many customers may prefer to interact with a chatbot rather than a human. Thus, routine questions could be handled by the chatbot, with referral to human beings for more complex issues. Moreover, they can also be leveraged for simulations to train employees on how to deal with complaints or other eventualities. Generative AI may further expand chatbot capabilities.
- Marketing- Traditional and generative AI may help automate the creation of various types of marketing content, including ad copy, social media posts, image/video generation, and product descriptions. AI can personalize content to optimize its relevance. Personalization must be balanced with privacy principles that protect personal information, respect consent, and comply with legal requirements.
- Back Office- Code generation, regulatory reporting, internal document/knowledge management, and agentic tools for compliance and audit functions are promising use cases for generative AI.

III. Risk Management of AI

A material feature of generative AI to bear in mind is its prompt-based nature, which lowers the barrier to access and expands the pool of potential users. This attribute makes it even more important to stand up effective AI governance.

On a broad scale, banks are undergoing a paradigm shift characterized by the breakdown of silos in favor of enterprise-wide functions. We expect this movement will continue. As part of this transition, banks are building interdisciplinary teams to identify, assess, and mitigate risks stemming from particular use cases. These cross-functional groups allow for superior risk awareness and mitigation versus compartmentalized approaches.

ABA members have been using AI technologies for years, while being subject to supervision and regulation. Their existing and mature governance frameworks are already mitigating AI-specific risks. They do not require a separate process for evaluating AI applications; rather, they ensure the existing risk management framework has the necessary personnel/controls to flag the inherent risk associated with AI-enabled use cases.

Further, ABA members are confident that banks will be able to adapt their existing risk management and governance frameworks to mitigate risks associated with increased use of emerging AI technologies. This is testament to the three lines of defense structure used by banks, coupled with agency supervision.

Three Lines of Defense

The “three lines of defense” refers to the division of roles and responsibilities within a bank in order to identify, assess, and mitigate risks. Among other sources, this expectation is provided in the model risk management guidance, SR 11-7:¹¹

- The first line are business units, which are generally responsible for the risk associated with their business strategies. They are ultimately accountable for the risk and performance within the framework set by bank policies and procedures, and are responsible for ensuring processes are properly developed, used, and evaluated.
- The second line is the control function. The responsibilities include risk measurement, limits, and monitoring. Other responsibilities include managing the independent validation and review process to ensure that effective challenge takes place. Control staff should have the authority to restrict business operations and order corrective action. Control work can exercise judgment in order to prioritize risk treatment.
- The third line is the bank's internal audit function. The third line's role is not to duplicate risk management activities but to evaluate whether risk management is comprehensive, rigorous, and effective. They should be independent and document findings. The third line should possess expertise but should not be directly involved in the first or second line of work. The third line should also verify that acceptable policies are in place, owners and control groups comply with those policies, validation work is conducted properly, and appropriate degrees of effective challenge is being carried out.

Examples of Risks

Some examples of the risks stemming from AI usage that financial institutions mitigate through the governance framework are:

¹¹ SR 11-7, see pp. 18-19.

- Cybersecurity and Fraud- banks must identify vulnerabilities in their own systems caused by AI (such as potential data leakage or backdoors). They must also defend against increasingly sophisticated AI usage by bad actors.
- Privacy and Data Governance- the training data must be of sufficient quality to ensure good model performance, and moreover usage of the data must be permissible under applicable law and terms of service. Banks seeking to leverage customer data must consider whether there is sufficient disclosure and transparency around AI processes.
- Bias and Fair Lending- Financial institutions are committed to ensuring that credit decisions are made in a fair and non-discriminatory manner. Banks utilize a rigorous validation process for models and conduct a statistical analysis to assess performance in meeting their business needs. An exhaustive compliance review is a crucial step prior to deployment.
- Third-Party Risk Management- Given that vendors are an essential element of the AI ecosystem, banks have to conduct proper due diligence to understand how third parties deliver AI-enabled services, as well as how AI affects the third parties' own operations. Banks must then negotiate contractual provisions to mitigate risk and delineate responsibilities. Banks must be able to understand and explain third-party solutions in order to inform their customers and satisfy regulatory expectations.
- Illicit Finance - Many banks use AI and machine learning technologies as part of their risk-based approach to BSA and sanctions compliance. In order to adopt an effective risk-based approach, banks must have an accurate understanding of the actual risks associated with their business practices. Banks also need refined and accurate models to avoid expending unnecessary resources investigating false positives; while ensuring they do not miss important red flags. In addition, generative AI can be used by bad actors to create realistic IDs, such as fraudulent passports, when attempting to open accounts online.

The risk management framework used by banks is able to address these and other issues; nonbanks and other ecosystem participants would benefit from it as well.

IV. Recommendations for Policymakers

ABA makes the following recommendations for inclusion in the Plan:

A. Legislation

Congress must assert its leadership over this issue. Momentum is building for legislative action; if Congress does not act, the states will (and some have already). This patchwork of requirements would be burdensome to national companies and have adverse consequences to American dominance in this space.

Instead, Congress must pass comprehensive laws establishing an AI risk management framework with strong preemptions. These laws should create baseline standards that are not duplicative or inconsistent with the model risk management obligations of banks.

Moreover, such a law should take a different path than laws such as the EU AI Act¹² or the vetoed California bill SB 1047,¹³ which are overly prescriptive and accordingly may discourage innovative practices such as fine-tuning existing LLMs. Instead, a flexible risk-based framework adapted from that used by banks would be far more effective at both achieving dominance and mitigating potential harm.

B. Regulation

Regulators should identify clear regulatory outcomes and objectives, while enabling regulated entities the ability to deploy effective risk management techniques based on common sense and best practices.

Requirements of banks to examine and monitor third-party AI algorithms, training data, or performance are not possible without third party cooperation. Any rulemaking related to AI should leverage existing authorities to focus on third-party nonbank AI models, tools, and platforms to impose the same obligations directly on such third parties and to require such providers to furnish sufficient credible, reliable information if used by financial institutions or in the financial services market.

At the same time, regulation of AI outside of federally regulated financial institutions needs to be strengthened. Far from being an impediment, this oversight would bolster innovation by creating a culture of trust and consistency that allows companies to experiment with confidence. Dominance is achieved on a level playing field; not when government tips the scales in a particular stakeholder's direction.

Finally, governmental agencies must be transparent about their own use of AI, especially the ways in which it is brought to bear on regulated entities.

C. Supervision

ABA encourages updates to model risk management guidance to be more reflective of bank operations and to make applicability to AI usage more obvious. Existing model risk guidance

¹² <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX%3A32024R1689>.

¹³ <https://legiscan.com/CA/bill/SB1047/2023>.

from the prudential agencies¹⁴ establishes standards for the use of third-party models used by banks, but banks may not have the ability to fully oversee all AI that third parties use to provide services. Updated guidance should clearly delineate the responsibilities when banks use third-party AI. However, overly prescriptive safeguards or prohibitions on the use of AI could become a barrier to valuable use cases if banks do not have appropriate flexibility to test and adapt new AI tools.

Critically, however, such updates must be developed through a notice and comment period to ensure that the proposed clarifications reflect the current state of technology, industry practice, consumer interests, etc.

In addition, many ABA members report that field examiners often stymie innovation efforts by needlessly focusing on granular matters (such as explaining the finer points of computer code) rather than areas of true risk (such as a model's inputs, outputs, and outcomes). Field examiners should be responsible for implementing an agency's policy, not making their own. This is especially true because field examiners exercise outsize power on how banks are able to function and adapt to their markets. Improper supervisory practices hurt the banking industry (often disproportionately affecting community banks) and drive potential customers into the arms of underregulated fintechs.

D. Other

The Plan should encourage the development of voluntary strategies for managing AI-related risk. One example would be a standardized disclosure template (model cards) shared for validation exercises which would not require the sharing of confidential commercial information. Another example are industry certifications to provide evidence of compliance for a baseline of fairness, transparency, and explainability; such programs would aid banks of all sizes and give AI developers incentives to build controls. Cooperation between industry and government via public/private partnerships is also needed to meet the challenges posed by advanced technologies.

The Plan should seek to develop an approach to explainability that transcends traditional methodologies by leveraging a suite of coordinated risk management practices, including but limited to data governance, weighted decision-making criteria, assurance and testing, and continuous risk monitoring. This can be achieved by mapping to the NIST AI Risk Management Framework¹⁵ (as well as the NIST Privacy Framework¹⁶) and/or creating sector-specific profiles. This holistic approach should include all participants in the AI ecosystem, including technology companies and non-financial industry actors, particularly because the economics of LLM development prevents internal development thereof and drives adoption of third-party offerings. Such work could be the foundation of interoperability across sectors and jurisdictions and would allow the entire ecosystem to innovate confidently.

¹⁴ SR 11-7, OCC Bulletin 2011-12, and FIL-22-2017, *supra*, note 5.

¹⁵ <https://www.nist.gov/itl/ai-risk-management-framework>.

¹⁶ <https://www.nist.gov/privacy-framework>.

Conclusion

As demonstrated above, the financial services sector can serve as an exemplar for responsible innovation of AI. While emerging technologies are exciting and can open new possibilities for the human condition, it is essential that they occur in an environment that fosters trust and accountability. Castles cannot be built on quicksand; dominance can only result from order.

We have made several high-level recommendations that the Plan should pursue. There are more particular actions to be taken, and ABA and our members would follow up directly with the entities tasked under the Plan on those deliverables.

ABA and our members are grateful for the opportunity to provide input into the Plan and stand ready to support policymakers as they work to implement it.

If you have any questions about this comment, please contact Ryan T. Miller.

Respectfully submitted,

Ryan T. Miller
Vice President & Senior Counsel, Innovation Policy