

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 03, 2025
Status: [REDACTED]
Tracking No. m7t-har8-s4dz
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1056
Comment on FR Doc # 2025-02305

Submitter Information

Name: Paula Allen

[REDACTED]
[REDACTED]

General Comment

In regards to Artificial Intelligence, let me be blunt. Artificial means fake, not real.

In order for America to reach it former Greatness and Strength, Americans need truth. Journalists need to get back to unbiased, ethical reporting. Our Scientists, Pharmaceutical, Medical professionals need to Follow the real Science, men are men (they have Penis's) and women are women (we have monthly Periods, bear children, and in our later years menopause.), medications with side effects worse than the illness they're trying to treat...It's all Artificial and I, for one, am tired of the stupidity of the fake.

President Trump is real and I love that, I get his sarcastic, blunt remarks. He says it like it is, and I love it. "Don't tell us how to feel!" Amen Mr. President!

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 04, 2025
Status: [REDACTED]
Tracking No. m7v-04wp-0lff
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1065

Comment on FR Doc # 2025-02305

Submitter Information

Name: STEPHEN CASPER

[REDACTED]

General Comment

AI is an emerging technology, and there is a great deal of uncertainty about how it will affect the world in the coming years. A lack of regulation in AI could miss opportunities to promote competitiveness and inform the public about what is consuming. Meanwhile, regulation that is too onerous could harm America's competitiveness.

As a middle ground, collaborators and I emphasize the value of process-based regulations, which do NOT limit what AI companies can do, but only serve to promote reporting and visibility. We believe that facilitating more public knowledge about AI developers and their systems is essential to advance the science and to ensure that our democratic Society is capable of making informed choices in the future.

In our recent paper (<https://arxiv.org/abs/2502.09618>), collaborators and I outline 15 evidence-seeking AI regulations. None of which place requirements on what developers can and cannot do. All of which are designed to help inform the public:

1. AI governance institute: A federal AI governance institute to research risks, evaluate systems, and curate best risk management practices that developers are voluntarily encouraged to adhere to.
2. Model registration: Maintaining a federal registry of frontier AI systems.
3. Model specification and basic info: Requiring developers to document intended use cases, behaviors, and basic information about frontier systems.
4. Internal risk assessments: Requiring developers to conduct and report on internal risk assessments of frontier systems.
5. Independent third-party risk assessments: Requiring developers to have an independent third-party conduct and produce a report (including access, methods, and findings) on risk assessments of frontier systems. Developers can also be required to document if and what "safe harbor" policies they have to facilitate independent evaluation and red-teaming.
6. Plans to minimize risks to society: Requiring developers to produce a report on risks posed by their frontier systems and risk mitigation practices that they are taking to reduce them.
7. Post-deployment monitoring reports: Requiring developers to establish procedures for monitoring and periodically reporting on the uses and impacts of their frontier systems.
8. Security measures: Given the challenges of securing model weights and the hazards of leaks, frontier developers can be required to document high-level noncompromising information about their security measures.
9. Compute usage: Given that computing power is key to frontier AI development, frontier developers can be required to document their compute resources including details such as the usage, providers, and the location of compute clusters.
10. Shutdown procedures: Requiring developers to document if and which protocols exist to shut down frontier systems that are under their control.
11. Documentation Availability: All of the above documentation can be made available to the public (redacted) and AI governing authorities (unredacted).
12. Documentation comparison in court: To incentivize a race to the top where frontier developers pursue established best safety practices, courts can be given the power to compare documentation for defendants with that of peer developers.
13. Labeling AI-generated content: To aid in digital forensics, content produced from AI systems can be labeled with metadata, watermarks, and notices.
14. Whistleblower protections: Regulations can explicitly prevent retaliation and offer incentives for whistleblowers who report violations

of those regulations.

15. Incident reporting: Frontier developers can be required to document and report on substantial incidents in a timely manner.

Best,

Stephen Casper, MIT

Attachments

2502.09618v3

PITFALLS OF EVIDENCE-BASED AI POLICY

Stephen Casper
MIT CSAIL

David Krueger
Mila

Dylan Hadfield-Menell
MIT CSAIL

At this very moment, I say we sit tight and assess.
– President Janie Orlean, Don't Look Up

ABSTRACT

Nations across the world are working to govern AI. However, from a technical perspective, there is uncertainty and disagreement on the best way to do this. Meanwhile, recent debates over AI regulation have led to calls for “evidence-based AI policy” which emphasize holding regulatory action to a high evidentiary standard. Evidence is of irreplaceable value to policymaking. However, holding regulatory action to too high an evidentiary standard can lead to systematic neglect of certain risks. In historical policy debates (e.g., over tobacco ca. 1965 and fossil fuels ca. 1985) “evidence-based policy” rhetoric is also a well-precedented strategy to downplay the urgency of action, delay regulation, and protect industry interests. Here, we argue that if the goal is evidence-based AI policy, the first regulatory objective must be to actively facilitate the process of identifying, studying, and deliberating about AI risks. We discuss a set of 15 regulatory goals to facilitate this and show that Brazil, Canada, China, the EU, South Korea, the UK, and the USA all have substantial opportunities to adopt further evidence-seeking policies.

CONTENTS

1	How do We Regulate Emerging Tech?	2
1.1	“Nope, I’m against evidence-based policy.”	2
1.2	A Broad, Emerging Coalition	2
1.3	A Vague Agenda?	3
2	The Evidence is Biased	4
2.1	Selective Disclosure	4
2.2	Easy- vs. Hard-to-Measure Impacts	4
2.3	Precedented vs. Unprecedented Impacts	5
2.4	Ingroups vs. Outgroups	5
2.5	The Culture & Values of the AI Research Community	5
2.6	Industry’s Entanglement with Research	6
3	Lacking Evidence as a Reason to Act	8
3.1	Substantive vs. Process Regulation	9
3.2	In Defense of Compute & Cost Thresholds in AI Regulation	9

4 We Can Pass Evidence-Seeking Policies Now	9
4.1 15 Evidence-Seeking AI Policy Objectives	10
4.2 Ample Room for Progress	10
4.3 The Duty to Due Diligence from Discoverable Documentation of Dangerous Deeds	12
4.4 Considering Counterarguments	12
5 Building a Healthier Ecosystem	13

1 HOW DO WE REGULATE EMERGING TECH?

Recently, debates over AI governance have been ongoing across the world. A common underlying theme is the challenge of regulating emerging technologies amidst uncertainty about the future. Even among people who strongly agree that it is important to regulate AI, there is sometimes disagreement about when and how. This uncertainty has led some researchers to call for “evidence-based AI policy.”

1.1 “NOPE, I’M AGAINST EVIDENCE-BASED POLICY.”

See how awful that sounds? This highlights a troublesome aspect of how things are sometimes framed. Of course, evidence is indispensable. But there is a pitfall of holding policy action to too high an evidentiary standard:

Postponing regulation that enables more transparency and accountability on grounds that it’s “not evidence-based” is counterproductive.

As we will argue, focusing too much on getting evidence before we act can paradoxically make it harder to gather the information we need.

1.2 A BROAD, EMERGING COALITION

Recently, there have been a number of prominent calls for evidence-based AI policy. For example, several California congressmembers and Governor Gavin Newsom recently argued against an AI regulatory bill in California by highlighting that it was motivated by mitigating future risks that have not been empirically observed:

There is little scientific evidence of harm of ‘mass casualties or harmful weapons created’ from advanced models.

– Zoe Lofgren et al. in an [open letter](#) to Gavin Newsom

[Our] approach. . . must be based on empirical evidence and science. . . [we need] AI risk management practices that are rooted in science and fact.

– Gavin Newsom in [his veto](#) of bill [SB 1047](#)

Others in academia have echoed similar philosophies of governing AI amidst uncertainty. For example, in their book *AI Snake Oil* (Narayanan & Kapoor, 2024), Princeton researchers Arvind Narayanan and Sayash Kapoor claim:

The whole idea of estimating the probability of AGI risk is not meaningful. . . We have no past data to calibrate our predictions.

– Narayanan & Kapoor (2024), *AI Snake Oil*

They follow this with [an argument](#) against the precautionary principle (Taleb et al., 2014), claiming that policymakers should take a noncommittal approach in the face of uncertainty and not act on speculative estimates of future AI risks.

Meanwhile, Jacob Helberg, a senior adviser at the Stanford University Center on Geopolitics and Technology, has argued that there just isn't enough evidence of AI discrimination to warrant policy action.

This is a solution in search of a problem that really doesn't exist. . . There really hasn't been massive evidence of issues in AI discrimination.

– Jacob Helberg on priorities for the current presidential administration

And Martin Casado, a partner at Andreesen Horowitz, recently argued in a post that we should hold off on taking action until we know the marginal risk:

We should only depart from the existing regulatory regime, and carve new ground, once we understand the marginal risks of AI relative to existing computer systems. Thus far, however, the discussion of marginal risks with AI is still very much based on research questions and hypotheticals.

– Casado (2024), Base AI Policy on Evidence, Not Existential Angst

And finally, the seventeen authors of a recent article titled, *A Path for Science- and Evidence-Based AI Policy*, argue that:

AI policy should be informed by scientific understanding. . . if policymakers pursue highly committal policy, the. . . risks should meet a high evidentiary standard.

– Bommasani et al. (2024a), A Path for Science- and Evidence-based AI Policy

Overall, the evidence-based AI policy coalition is diverse. It includes a variety of policymakers and researchers who do not always agree with each other. We caution against developing a one-dimensional view of this coalition or jumping to conclusions from quotes out of context. However, **this camp is generally characterized by a desire to avoid pursuing highly committal policy absent compelling evidence.**

1.3 A VAGUE AGENDA?

Calls for evidence-based policy are not always accompanied by substantive recommendations. However, Bommasani et al. (2024a) end their article with a set of four milestones for researchers and policymakers to pursue:¹

Milestone 1: A taxonomy of risk vectors to ensure important risks are well-represented

Milestone 2: Research on the marginal risk of AI for each risk vector

Milestone 3: A taxonomy of policy interventions to ensure attractive solutions are not missed

Milestone 4: A blueprint that recommends candidate policy responses to different societal conditions

These milestones are extremely easy to agree with. Unfortunately, they are also unworkably vague. It is unclear what it would mean for them to be accomplished. In fact, for these milestones, it is not hard to argue that existing reports reasonably meet them. For example, the AI Risk Repository (Slattery et al., 2024) predates Bommasani et al. (2024a) and offers a meta-review, taxonomy, and living database of AI risks discussed in the literature. If this does not offer a workable taxonomy of risks (Milestone 1), it is unclear what would.²

¹Bommasani et al. (2024a) also call for the establishment of a registry, evaluation, red-teaming, incident reporting, and monitoring but do not specify any particular role for regulators to play in these. They also make a nonspecific call for policymakers to broadly invest in risk analysis research and to investigate transparency requirements.

²For milestone 2, most relevant research is domain-specific; consider Metta et al. (2024), Sandbrink (2023), Musser (2023), and Cazzaniga et al. (2024) as examples. Note, however, that forecasting future marginal risks will always be speculative to some degree. See also Bengio et al. (2025). Meanwhile, milestones 3 and 4 essentially describe the first and second stages of the AI regulation process, so existing regulatory efforts already are working on these (e.g., Arda, 2024).

These milestones are an encouraging call to actively improve our understanding. However, absent more precision, we worry that similar arguments could be misused as a form of tokenism to muddy the waters and stymie policy action. In the rest of this paper, we will argue that holding regulatory action to too high an evidentiary standard can paradoxically make it harder to gather the information that we need for good governance.

2 THE EVIDENCE IS BIASED

In its pure form, science is a neutral process. But it is never done in a vacuum. Beneath the cloak of objectivity, there are subjective human beings working on problems that were not randomly selected (Kuhn, 1962). There is a laundry list of biases subtly shaping the evidence produced by AI researchers. A policymaking approach that fixates on existing evidence to guide decision-making will systematically neglect certain problems.

2.1 SELECTIVE DISCLOSURE

In February 2023, Microsoft announced Bing Chat, an AI-powered web browsing assistant. It was a versatile, semi-autonomous copilot to help users browse the web. It was usually helpful, but sometimes, it went off the rails. Users found that it occasionally took on shockingly angsty, deceptive, and outright aggressive personas. It would go so far as to threaten users chatting with it. Rest assured, everyone was fine. Bing Chat was just a babbling web app that could not directly harm anyone. But it offers a cautionary tale. Right now, developers are racing to create increasingly agentic and advanced AI systems (Chan et al., 2023; Casper et al., 2025). If more powerful future systems go off the rails in similar ways, it could spell trouble.

Following the Bing Chat incidents, Microsoft’s public relations strategy focused on patching the issues and moving on. To the dismay of many AI researchers, Microsoft never published a public report on the incident. If Microsoft had nothing but humanity’s best interests at heart, it could substantially help researchers by reporting on the technical and institutional choices that led to Bing Chat’s behaviors. However, it’s just not in their public relations interests to do so.

Historically, AI research and development has been a very open process. For example, code, models, and methodology behind most state-of-the-art AI systems were broadly available pre-2020. More recently, however, developers like Microsoft have been exercising more limited and selective transparency (Bommasani et al., 2024b). **Due to a lack of accountability in the tech industry, some lessons remain simply out of reach. There is a mounting crisis of transparency in AI when it is needed the most.**

2.2 EASY- VS. HARD-TO-MEASURE IMPACTS

The scientific process may be intrinsically neutral, but not all phenomena are equally easy to study. Most of the downstream societal impacts of AI are difficult to accurately predict in a laboratory setting. The resulting knowledge gap biases purely evidence-based approaches to neglect some issues simply because they are difficult to study.

Thoroughly assessing downstream societal impacts requires nuanced analysis, interdisciplinarity, and inclusion. . . there are always differences between the settings in which researchers study AI systems and the ever-changing real-world settings in which they will be deployed.

– Bengio et al. (2024), International Scientific Report on the Safety of Advanced AI

Differences in the measurability of different phenomena can cause insidious problems to be neglected. For instance, compare explicit and implicit social biases in modern language models. Explicit biases from LLMs are usually easy to spot. For example, it is relatively easy to train a language model against expressing harmful statements about a demographic group. But even when we do this to language models, they still consistently express more subtle biases in the language and concept associations that they use to characterize different people (Wan et al., 2023; Wan & Chang, 2024; Bai et al., 2024; Hofmann et al., 2024).

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 04, 2025
Status: [REDACTED]
Tracking No. m7v-04wp-0lff
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1065

Comment on FR Doc # 2025-02305

Submitter Information

Name: STEPHEN CASPER



General Comment

AI is an emerging technology, and there is a great deal of uncertainty about how it will affect the world in the coming years. A lack of regulation in AI could miss opportunities to promote competitiveness and inform the public about what is consuming. Meanwhile, regulation that is too onerous could harm America's competitiveness.

As a middle ground, collaborators and I emphasize the value of process-based regulations, which do NOT limit what AI companies can do, but only serve to promote reporting and visibility. We believe that facilitating more public knowledge about AI developers and their systems is essential to advance the science and to ensure that our democratic Society is capable of making informed choices in the future.

In our recent paper (<https://arxiv.org/abs/2502.09618>), collaborators and I outline 15 evidence-seeking AI regulations. None of which place requirements on what developers can and cannot do. All of which are designed to help inform the public:

1. AI governance institute: A federal AI governance institute to research risks, evaluate systems, and curate best risk management practices that developers are voluntarily encouraged to adhere to.
2. Model registration: Maintaining a federal registry of frontier AI systems.
3. Model specification and basic info: Requiring developers to document intended use cases, behaviors, and basic information about frontier systems.
4. Internal risk assessments: Requiring developers to conduct and report on internal risk assessments of frontier systems.
5. Independent third-party risk assessments: Requiring developers to have an independent third-party conduct and produce a report (including access, methods, and findings) on risk assessments of frontier systems. Developers can also be required to document if and what "safe harbor" policies they have to facilitate independent evaluation and red-teaming.
6. Plans to minimize risks to society: Requiring developers to produce a report on risks posed by their frontier systems and risk mitigation practices that they are taking to reduce them.
7. Post-deployment monitoring reports: Requiring developers to establish procedures for monitoring and periodically reporting on the uses and impacts of their frontier systems.
8. Security measures: Given the challenges of securing model weights and the hazards of leaks, frontier developers can be required to document high-level noncompromising information about their security measures.
9. Compute usage: Given that computing power is key to frontier AI development, frontier developers can be required to document their compute resources including details such as the usage, providers, and the location of compute clusters.
10. Shutdown procedures: Requiring developers to document if and which protocols exist to shut down frontier systems that are under their control.
11. Documentation Availability: All of the above documentation can be made available to the public (redacted) and AI governing authorities (unredacted).
12. Documentation comparison in court: To incentivize a race to the top where frontier developers pursue established best safety practices, courts can be given the power to compare documentation for defendants with that of peer developers.
13. Labeling AI-generated content: To aid in digital forensics, content produced from AI systems can be labeled with metadata, watermarks, and notices.
14. Whistleblower protections: Regulations can explicitly prevent retaliation and offer incentives for whistleblowers who report violations