

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 08, 2025
Status: [REDACTED]
Tracking No. m80-wo1f-1uuu
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1116
Comment on FR Doc # 2025-02305

Submitter Information

Name: Naive Skeptic
Email: [REDACTED]

General Comment

The fundamental problem of computer science has always been that we cannot make a machine do what we want completely reliably. Historically this has not mattered very much, but as computer systems become more powerful and influential this matters more and more. Consider the big CrowdStrike outage last year. AI systems are on a path towards ever-increasing power and influence. Systems like companies and governments are already far more powerful than individuals, so it seems unlikely that humans represent the upper limit of AI capability. As AI systems become more powerful the amount of harm they will be able to do will increase. Unless the government steps in the trend of increasing AI capability will not stop on its own. Eventually the power these systems will possess will mean that we will be capable of substantial harm. And because we cannot rely make them do what we want, they will try to do things we don't want. There are two fundamental challenges we face: AI Capabilities keep increasing and we do not know how to make AIs do what we want reliably. Government should address these areas, and hold AI companies to account in these areas.