

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 11, 2025
Status: [REDACTED]
Tracking No. m84-onse-vbdz
Comments Due: March 15, 2025
Submission Type: Web

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1138
Comment on FR Doc # 2025-02305

Submitter Information

Email: [REDACTED]
Government Agency Type: Foreign
Government Agency: Noema Research

General Comment

Please find our comment attached.

Attachments

Response to White House's RFI on AI Action

Strengthen State Capacity to Wield AI as a Geopolitical Resource

By [Noema Research](#)

In this response, we articulate policy objectives that we argue to be of strategic relevance to the US, and indeed to any global AI power that is determined to lead in the development of this remarkably powerful technology. This resource is structured as follows:

- **New Resource Emerging.** In this section, we describe AI as one of many geopolitically relevant resources, such as steel or energy.
- **Strengthen State Capacity.** In this section, we introduce two straightforward policy objectives concerning this new resource.
- **Leverage, Credibility, Security.** In this section, we reflect on the geopolitical impacts of the suggested initiatives.
- **Barriers Versus Setbacks.** Finally, we advance an intuitive framing for weighing the upsides and downsides of specific AI policies.

We conclude with pointers to further resources related to the arguments advanced in the rest of the document.

New Resource Emerging

An effective way of conceiving of AI development is as the creation of a qualitatively novel resource based on certain raw materials and manufacturing processes. Just as steel production requires iron ores and specialized furnaces, and energy generation demands fossil fuels and combustion equipment, AI development necessitates computational power and learning algorithms. Feed in larger quantities of raw materials and more efficient manufacturing processes, and you predictably obtain more "virtual labor" that can be used to extend what humans can achieve.

Perhaps more than any other resource that has preceded it, AI can be applied in an incredibly wide range of contexts. The intellectual faculties of these virtual entities may be harnessed to augment the productivity of the human workforce, may be directed

towards strengthening the cybersecurity of critical infrastructure, and may be recruited towards furthering the development of effective healthcare interventions.

Naturally, the remarkable generality of this intellectual labor forged out of computing power and learning algorithms may also be applied in ways detrimental to national interests. These virtual capabilities may disrupt the labor market in ways that cause societal unrest, may enable foreign adversaries to identify vulnerabilities in domestic infrastructure, and may empower rogue actors to interfere with public health in increasingly sophisticated ways.

Like steel, energy, and countless other resources being manufactured on an industrial scale, AI is a highly geopolitically relevant resource due to its dual-use nature. The unusual origin of this new asset — largely, the private sector — has led governments from around the world to scramble towards achieving adequate state capacity with varying degrees of success. However, we argue that bringing governments “in the loop” is a critical step for effectively wielding this emerging resource on the world stage, navigating domestic developments wisely, and enabling meaningful leadership in the development of this historical technology.

Strengthen State Capacity

To enact this critical step, we suggest two policy objectives:

- **Measure capability consumption domestically.** In order to credibly and effectively wield this technology to further national interests, governments need to first and foremost become aware of the quantity of these newly manufactured resources. Governments may especially benefit from gaining insight into types of virtual labor that are of high geopolitical relevance, such as automated hacking or automated research, among others. Governments should have knowledge of, for instance, how much autonomous hacking is being minted domestically, and how much of this capability gets exported to foreign adversaries. They may also benefit from, for instance, being aware of the consumption of economically valuable capabilities, in order to gain a deeper understanding of the impacts on the labor market.
- **Pursue associated domestic and international opportunities.** With this enhanced awareness of the synthesis of these resources domestically, governments may benefit from building capacity for actively investigating the strategic opportunities enabled by this understanding. We hint at particular directions in the following section, yet here stress the policy objective of building the capacity to investigate such interventions in the first place. This may be easiest to achieve as part of existing organizations with expertise in related

domains, such as relevant public bodies (e.g., [US AISI](#)), and think tanks with a history of navigating the game theory of emerging technologies to advance national interests (e.g., [RAND](#)).

Leverage, Credibility, Security

As governments gain awareness of the quantity of these resources and actively pursue strategies informed by this understanding, we expect the following benefits to materialize:

- **Leverage.** This awareness may help highlight the scope and scale of this virtual workforce as a resource to be used in bargaining across the world stage. As nations boast extensive natural resources, advanced military equipment, or thriving consumer economies, so may the intellectual horsepower of these virtual entities constitute an additional negotiation token to be used in securing strategic deals. When composed with established policy practices, such as for instance through tariffs on imported quantities of virtual labor, these resources may provide a net increase in the range of available foreign policy actions.
- **Credibility.** Demonstrating awareness of domestic AI development through measures of capability consumption may help boost credibility both internationally, as well as among citizens impacted by the synthesis of this resource. Deploying mechanisms for proving the correctness and validity of these measurements can further provide the foundation of credibility needed to pursue landmark bilateral or multilateral initiatives on AI, analogous to how many nuclear security agreements have been predicated on literal [warhead counts](#) or [tons](#) of dual-use materials.
- **Security.** Similar to managing the proliferation of chemical precursors or radioactive materials for national security reasons, monitoring the quantities of dual-use AI capabilities being consumed by domestic organizations may help governments effectively address emerging threats. Beyond the domestic setting, surgical international initiatives may help extend existing non-proliferation efforts in the domain of dual-use AI capabilities, such as autonomous hacking.

Barriers Versus Setbacks

This extensive optionality may be achievable at little cost to companies developing AI. Consider that computational resources used in AI development are estimated to increase by [~5x per year](#), and the learning algorithms are estimated to become [~3x more efficient](#) in the same period, for a total increase of ~15x per year in the amount of virtual labor that can effectively be "minted." Assume that metering capability

consumption time incurs a 5% overhead relative to the baseline computational power of serving frontier models. This would be equivalent to a one-time setback of less than a week for the development of AI capabilities. That said, AI companies based in the US are estimated to be much farther ahead, potentially making it palatable to invest in the powerful geopolitical instruments detailed above, in order to be able to collect their future dividends.

To conclude, strengthening state capacity in AI by equipping governments with this enhanced awareness would allow them to wield these emerging resources more effectively in their dealings. By pursuing these policy objectives, governments can maintain leadership and ensure these powerful technologies serve national interests. We lay out a more concrete vision for how these remote measurement capabilities may be developed and applied in a dedicated resource on ["virtual diplomacy."](#)

[Noema Research](#) is an R&D lab working on techniques for remotely measuring AI capabilities with a view towards averting Great Power conflict. Based in the EU, we work with frontier labs and institutions from around the world to understand and manage the capability surface of AI systems. To enshrine this mission, leadership has formally pledged to donate 100% of proceeds on liquidity to charity.

March 15, 2025

Trustible Comments on the Request for Information (RFI) on the Development of an Artificial Intelligence (AI) Action Plan

To the Office of Science and Technology Policy:

On behalf of Trustible, a leading technology company based in Virginia that helps build trust through AI governance software, we appreciate the opportunity to submit comments in response to the Office of Science and Technology Policy's RFI on developing an AI Action Plan.¹

Trustible provides a Software as a Service platform that leverages AI to help large and medium size organizations implement internal processes to manage and oversee their use of AI. Trustible supports the Administration's goal to sustain and enhance America's global AI competitiveness and innovation. In order for the U.S. to maintain and build upon its AI leadership, we should encourage an AI ecosystem that leverages our world-class technology infrastructure and build trust amongst innovative AI tools. Innovation and trust flourish when there are common sense, industry-driven standards available for technology companies to adopt at scale.

I. The Second Trump Administration Can Continue Its Work from the First Trump Administration by Promoting Common Sense, Pragmatic Standards.

We encourage the Trump Administration to convene pragmatic stakeholders from across industry, academia, and other faucets of civil society to create technical standards that build trust in AI technologies. Establishing a practical set of AI standards helps grow the AI ecosystem and economy because companies that adopt those standards can demonstrate a basic level of reliability for their AI tools – enhancing the marketability and procurement of these systems. As part of our work, we gain valuable insights from the private sector about its development and adoption of AI tools. We consistently hear from companies that they want certain assurances about AI technology before they adopt it.

President Trump understood the importance that standards have in building trust and growing economic opportunity when he signed the Executive Order on Maintaining American Leadership in AI in February 2019. We encourage the Administration to build upon the success it achieved with regards to issuing technical standards. The current standards landscape is more saturated with guidance for foundational model creators than companies that integrate or deploy these models for their own products and services. However, foundation model creators are vastly outnumbered by organizations that are not developing their own models. While these organizations may have strong subject matter expertise, they may lack the requisite talent to demonstrate trustworthiness in their systems that is found in frontier model labs.

¹ Our comments are approved for public dissemination and contain no business-proprietary or confidential information. We understand that the contents of these comments may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Therefore, standards can be extremely valuable for non-frontier model companies because it provides them with a baseline of scalable best practices.

II. The Trump Administration Can Learn from the Cybersecurity Ecosystem to Develop Scalable AI Standards.

The Administration should reference the success with cybersecurity standards as a roadmap for continuing its work on developing AI standards. Standards, such as the Service Organization Control 2 (SOC-2), Payment Card Industry Data Security Standard, and HITRUST, set attainable goals for companies to achieve while also helping them set a strong foundation for cybersecurity practices. These standards are particularly helpful for small and medium enterprises (SMEs) because they are market driven efforts that lower entry barriers for companies who may otherwise not have been able to demonstrate a baseline level of cybersecurity practices for their customers. In fact, SOC-2's scalability helped us build trust with potential and existing customers. An auditable standard tailored towards AI can make recommendations for companies, particularly SMEs, that deploy AI systems but not develop them. We should avoid the pitfalls of critics who assert that standards simply serve as "check the box exercise," when in practice these standards assist smaller enterprises like ourselves with understanding the types of cybersecurity controls to implement. Instead, the Administration should view AI standards as a means to help entrepreneurs and new SMEs incorporate best practices from rational industry AI experts.

III. The Trump Administration Should Lead on AI Standards to Promote American Values in the AI Ecosystem.

In the absence of continued U.S. leadership on AI standards, there is a heightened risk for other countries and international bodies to dictate heavy-handed protocols that are the antithesis of U.S. freedom, competitiveness, and innovation. There are many emerging standards that either impose unattainable requirements for SMEs or prevent operationalization due to overbroad and convoluted language. As a fast-growing startup company, we understand the unique challenges that entrepreneurs and SMEs face when trying to demonstrate trust in their products to prospective customers. The Administration can help avoid these barriers by encouraging the development of guidance or standards that are scalable for early AI startups or tools to increase the market adoption of these technologies.

Trustible appreciates and supports the Trump Administration's efforts to meaningfully engage stakeholders on how best to position the U.S. as a global leader in AI. Being the leader in AI standards will help achieve that goal, while also unlocking America's technological innovation and economic prosperity. Trustible looks forward to building a meaningful partnership with the Administration as it continues to pursue a robust AI policy agenda.

Respectfully,

Gerald Kierce
Co-Founder and CEO
Trustible

Andrew Gamino-Cheong
Co-Founder and CTO
Trustible