

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 13, 2025
Status: [REDACTED]
Tracking No. m88-6wbs-tkuk
Comments Due: March 15, 2025
Submission Type: Web

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1289

Comment on FR Doc # 2025-02305

Submitter Information

Name: Holly Elmore

[REDACTED]

General Comment

To the NITRD NCO and OSTP,

I am relieved to see the United States government treating frontier AI as the absolute force that it will be. The potential benefits are enormous, but our top priority at this critical juncture must be ensuring the loyalty and reliability of these frontier AI systems. The United States must take the reins while we have the lead to ensure that no country develops out-of-control superhuman AI. We must make an AI Deal to prevent superhuman AI development by all nations while we still have the lead.

I did my PhD work in bioinformatics. I still remember when GPUs for scientific applications were novelties in the compute cluster at Harvard only seven years ago. How far we have come. Five years after I graduated, frontier AI could have completed my PhD work in a fraction of the time. I marvel at the potential of this amazing technology to improve our lives. I became a scientist to improve our lives through knowledge and innovation. But now I fear that, without wise leadership, frontier AI will become a threat to humanity instead of the boon that it can be.

The first thing we need to steer frontier AI and keep it aligned with American interests is more time. Over the past few years, AI has gone from barely being able to write coherent sentences to rivaling physicists and mathematicians at elite-level problems. AI experts believe that truly human-level AI could be around the corner, and I agree. Without creating more time to steer and innovate in the direction we want, we will lose the ability to control the technology at all.

Superhuman AI could exploit cybersecurity vulnerabilities, cripple our infrastructure, work with our adversaries to induce armed conflict, or make synthetic superpathogens. We may be able to eliminate some of these risks by hardening our security— for instance, by mandating better biosecurity practices— but not all. By definition, any AI system with greater-than-human intelligence will be able to think up new ways to threaten our security that we ourselves had failed to consider. Even if we are lucky enough to avoid widespread loss of life, superhuman AI that escapes our control could nonetheless threaten our national interests, destabilize the economy, and plunge our country— and our world— into chaos. The world needs strong leadership to steer frontier AI away from catastrophe and into a better future.

The United States can provide that leadership, for ourselves and the world. America should have a great AI Plan that ensures that we build beneficial AI and lead the rest of the world in keeping it beneficial. Other nations are developing their own frontier AI, but the United States is ahead. This puts us in a position to lead the rest of the world in avoiding potentially fatal mistakes with this powerful technology. An AI Deal that limits how much any nation can experiment with AI over a certain level of power (achieved perhaps through limits on training compute) would allow the US to ensure the beneficial development of future AI and maintain its lead.

Thank you,

Holly Elmore, PhD