

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status: [REDACTED]
Tracking No. m88-z5wk-cutb
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1365

Comment on FR Doc # 2025-02305

Submitter Information

Email: [REDACTED]

Organization: Convergence Analysis

General Comment

Convergence Analysis is a US 501(c)(3) non-profit research organization focused on frontier AI governance.

We respectfully submit this response to the Request for Information on the Development of an AI Action Plan, focusing on information relevant for securing American leadership across 4 high-priority areas:

- (1) US Policy Leadership in the AI Economy;
- (2) National Security and AI Innovation in the USA;
- (3) AI Diplomacy and American Control;
- (4) Strategic Information on Powerful Dual-Use Models.

Please see the attached file for our full comment.

Attachments

Response by Convergence Analysis to the Request for Information on the Development of an Artificial Intelligence (AI) Action Plan

Response by Convergence Analysis to the Request for Information on the Development of an Artificial Intelligence (AI) Action Plan

March 15, 2025

AI Action Plan
Attn: Faisal D'Souza
NCO, 2415 Eisenhower Avenue
Alexandria, VA 22314

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Introduction

Convergence Analysis is a US 501(c)(3) non-profit research organization focused on frontier AI governance. We respectfully submit this response to the Request for Information on the Development of an AI Action Plan, focusing on information relevant for securing American leadership across 4 high-priority areas: **(1) US Policy Leadership in the AI Economy; (2) National Security and AI Innovation in the USA; (3) AI Diplomacy and American Control; (4) Strategic Information on Powerful Dual-Use Models.**

(1) US Policy Leadership in the AI Economy

The economic impacts of near-future AI systems are highly uncertain, and understudied, and yet there is a strong consensus among experts that they will be transformational. The economic policy frameworks of this administration's term will lay the foundations for Americans and the world for generations into the future.

In this section we highlight findings from our recent expert conference Threshold 2030¹ on AI's near-term economic impacts, convening specialists from cutting-edge organizations including Google, DeepMind, Stanford, MIT, Metaculus and OpenAI. The conference developed practical AI impact forecasting methods and mapped areas of expert consensus and disagreement. Below we list 7 key findings:

- 1. Within 5 years AI developments may lead to [a sharp increase in unemployment](#).**

¹ [Conference Summary: Threshold 2030 - Modeling AI Economic Futures - Convergence Analysis](#)

The impact of AI on unemployment by 2030 will be highly sector-specific, with transport and tech-heavy fields facing rapid automation, while skilled physical work, such as surgery and dentistry, and high-prestige roles seeing more AI augmentation than replacement. Expert predictions for 2030 ranged from modest unemployment with wage depression in creative industries, to 40% unemployment rates and significant government intervention.

2. Within 5 years AI developments are likely to lead to [increasing wealth inequality](#).

AI advancement will likely concentrate wealth among capital owners, tech companies, and skilled professionals, while automation-driven unemployment increases financial pressure on others, leading to rising strain on government welfare.

3. AI will disrupt [wage and labor share](#).

AI-driven automation is expected to decrease labor's share of national income as capital owners capture a greater portion of economic output, although some high-skill and service-based professions may command rising wages. Slow AI progress in the next 5 years may only slightly alter wage structures, while in scenarios where AI progress continues to advance quickly, AI systems could lead to a near-complete displacement of human labor in cognitive fields.

4. The [rate of diffusion](#) for AI Systems will likely be uneven across industries and regions.

Experts predicted infrastructure and institutional inertia will limit widespread deployment of AI, even in scenarios with advanced AI. While some economies may integrate AI rapidly, the global distribution of benefits is expected to be highly unequal, potentially leading to geopolitical tensions. Several experts predicted limited global AI diffusion by 2030, indicating that they expect significant infrastructural and implementation barriers.

5. [Public Responses to AI-Driven Economies](#) will be significant.

Experts predicted possible increased demands for wealth redistribution, shifts towards local labor, and growing political pressure for AI regulation. Alternative economic models such as universal basic income or cooperative labor structures may gain traction if conventional employment declines. In scenarios with high unemployment, there may be widespread crises around finding purpose as traditional work-based identities collapse.

6. There is an immediate need for new economic metrics.

There are significant gaps in our ability to measure and track AI's economic impacts. Conventional measures fail to capture the full scope of AI-driven economic transformation. The development of such metrics would be crucial for tracking AI's impact on both economic productivity and broader social outcomes.

7. There is a glaring lack of research on concrete AI economic policies.

A notable gap emerged in the discussions around policy interventions to address AI's economic impacts. While attendees discussed significant existing research describing potential economic impacts from AI systems, there was limited research referenced or discussion regarding concrete policy ideas to address these impacts.

Specific neglected [topics of research](#) include:

- Tax & revenue policy (corporate profit taxation, robot taxes, tax base distortions),
- Social safety nets (UBI, job transition funds, AI reskilling, negative income taxes),
- Predistributive policies & market incentives (public benefit corporations, insurance policies, procurement policies, due diligence), and
- Public / private coordination (AI as a public utility, government management / integration).

More details on these expert discussions, models of the near-future, and conclusions can be found in the [Threshold 2030 Conference Report](#).

(2) National Security and AI Innovation in the USA

While economic opportunity drives AI innovation, experts forecast that advanced AI will enable capabilities with significant implications for national security, such as cyberwarfare capabilities and autonomous weapons systems.

Private US labs are currently the leading organizations pushing the frontier of AI development, and substantial evidence² points towards the current and continued dominance of private US AI labs. As AI system capabilities are demonstrated to threaten US national security, we expect the US government to increase its involvement in AI technologies in order to maintain American superiority globally.

Some experts have predicted, or advocated for a ‘Manhattan Project for AI’, to consolidate and nationalize all existing frontier AI research in the interest of national security. However, such a project and other similar descriptions³ of nationalization represent only a narrow subset of possible scenarios of US government involvement, and are not likely to align with the objectives of this Action Plan of solidifying both US AI leadership and national security for three reasons:

- 1) Nationalization may be seen as a threat to innovation driven by private sector competition and diversity of approaches. The American innovation model, built on free-market competition, is arguably why the US leads the AI race today.
- 2) Nationalization of frontier AI labs would face unprecedented practical, legal, and political challenges. Companies like Microsoft, Google, and Meta have market

² [AI Index | Stanford HAI](#)

³ [AI and Geopolitics: How might AI affect the rise and fall of nations? | RAND](#)

capitalizations exceeding \$1 trillion each. Similarly, Nvidia—controlling 80% of the AI chip market with a \$3 trillion valuation—would require government ownership in many nationalization scenarios, making this financially and legally implausible.

- 3) Less intrusive alternatives exist: combinations of targeted policy levers could address national security concerns while allowing private AI development to continue.

We identify 13 sets of these policy levers that can meet both national security and innovation needs, as an alternative to nationalization projects. Each set of levers offers a series of options that afford the government increasingly more influence, on a spectrum ranging from standard regulations to more comprehensive government control.

The comprehensive overview of these levers can be found in **Section 2 of our report [Soft Nationalization](#)**.

We do not advocate or recommend for the application of any of these policy levers. This section is informative in nature – it is intended solely to describe the space of plausible policy levers that may occur.

(3) AI Diplomacy and American Control

Securing America's AI advantage requires not only domestic policy frameworks but also strategic international engagement. The competitive dynamics between the US and China generate a range of threats to American control of powerful AI systems, their creation and their production. Without strong diplomatic action, rapid progress may create new, undue challenges. In this section we discuss two of these challenges, and make recommendations for each.

1. Diplomacy to prevent pre-emptive attacks on the US

A national US project to develop Artificial Superintelligence⁴ (ASI) would incentivize pre-emptive strikes against the US and its AI infrastructure, as US adversaries will recognize that superintelligent AI threatens their ability to retaliate in nuclear or conventional warfare⁵.

Adversaries would also be incentivized to exfiltrate model weights to equal or surpass America's AI capabilities. The most extreme security measures are not guaranteed to prevent this, and come with a new set of associated challenges; for example, extreme centralization that such security measures likely entail would make a state's AI development more vulnerable to a single military attack by an adversary⁶.

⁴ We define Artificial Superintelligence as an AI system that vastly exceeds human performance across nearly all cognitive tasks.

⁵ [Artificial Intelligence and Nuclear Stability | CNAS](#)

⁶ [The Manhattan Trap | Convergence Analysis](#) 12

A cheap and effective way to reduce these risks is to open diplomatic channels to competing states to defuse competitive dynamics in the short term. We believe it is in the interest of the US government to:

- Develop metrics for distinguishing between mainstream AI development that generates no strategic risk, and developments that may generate degrees of strategic risk, such as a national ASI project.
- Initiate high-level dialogue with competitors, particularly China, focused specifically on superintelligent AI development.
- Establish secure communication channels to prevent misperception of regular AI development as national ASI projects.

2. Verification to maintain American Control

Hypercompetitive acceleration of AI development and deployment also increases vulnerability to bad actors, and threatens American control over deployed systems, both military and civilian. As AI becomes more advanced, it will by definition operate with less human oversight, doing more with more autonomy. This increases the surface area for control failures from technical errors, terrorist groups or other bad actors. Given the growing autonomy of these systems, the impact of control failures for widely deployed military systems or of civilian infrastructure (such as self-driving cars) would be devastating.

Maintaining American control of the most powerful AI systems requires a controlled scaling approach. To support this strategy, mechanisms for verifying the AI progress of competing states are essential. Visibility into adversaries' progress towards superintelligence would inform the speed of US development, allowing for US-led controlled scaling of AI systems while ensuring that adversaries do not race ahead. Unlike general AI development, an ASI project would likely require distinctive infrastructure, methods, and resources that creates detection opportunities⁷. These characteristics make verification mechanisms both strategic and feasible. We recommend that the USG:

- Research methods to distinguish between general AI development and ASI-specific development.
- Develop technical capabilities to monitor indicators of ASI development activities

(4) Strategic Information on Powerful Dual-Use Models

Effective diplomacy and national security both depend on a foundation of accurate information about AI capabilities. To maintain American leadership while addressing the dual-use nature of frontier AI, the United States Government must establish mechanisms for strategic visibility into the most powerful models being developed.

⁷ [The Manhattan Trap | Convergence Analysis](#) 22

Powerful dual-purpose AI models have emerged as a critical technological frontier with serious implications for national security and global strategic capabilities, including, but not limited to:

- **Biosecurity, Terrorism, and Weapons Proliferation** - AI could be used to develop biological or chemical weapons or even radiological or nuclear weapons⁸.
- **Cybersecurity Threats** - AI could lower the barrier to entry for sophisticated cyberattacks from state and non-state adversaries⁹.
- **Control failures** - Advanced AI systems or autonomous weapons systems may become extremely dangerous if they behave in unexpected ways, through technical failure or misuse by bad actors.

The US government currently has limited visibility into the development and capabilities of frontier dual-use models. This lack of information creates severe limitations on any efforts to detect, anticipate, and respond to potential opportunities and threats¹⁰. To meet these challenges, the Bureau of Industry and Security (BIS) has proposed a rule¹¹ establishing reporting requirements for the development of advanced AI models and computing clusters under the authority of the Defense Production Act. To add to the work being done by the US Government, Convergence Analysis has prepared the following proposal on lightweight registration of the most powerful dual-use models.

1. The Proposal for an Advanced Dual-Use AI Registry

A registry would serve as a central repository of information on national dual-use AI capabilities and vulnerabilities. It would:

- A. **Inform the creation of national strategies**, with accurate information about existing models, their capabilities and vulnerabilities.
- B. **Support implementation of national strategies**, ensuring that safety-critical interventions are carefully targeted so that all and only the relevant models are affected.
- C. **Build capacity in government** by providing structured information on AI for policymakers and encouraging the development of the technical knowledge required to manage AI systems.

To capitalize on these advantages, we propose that the USG implement a registry for the most powerful dual-use AI models under the appropriate authority.

2. Avoiding Excessive Burdens on AI Labs

⁸ [Skating to Where the Puck Is Going | CSET](#)

⁹ [Russia will not intimidate us with cyberthreats | Reuters](#)

¹⁰ [AI Model Registries: | Convergence Analysis](#)

¹¹ [Federal Register :: Establishment of Reporting Requirements for the Development of Advanced Artificial Intelligence Models and Computing Clusters](#)

A model registry should serve as a foundation for more robust policy development without imposing excessive burdens that could impede technological advancement or create unnecessary resistance from industry stakeholders. To ensure the proposed registry fosters AI innovation, we highlight several key design principles, including: (1) Minimal Reporting requirements that capture only essential information linked to governance goals, and (2) No Mandatory Standards or Licensing, as AI safety evaluations and standards remain immature. The registry should function as a lightweight information-gathering mechanism rather than an enforcement tool.

3. Defining the Most Powerful Models

Our proposal includes a holistic definition of model capabilities that engages with the latest findings in scaling laws and capitalizes on all readily available information about a given model's potential. This definition ensures that only the most powerful models are included in a registry, meeting national security objectives without placing burdens on innovation. This includes a combination of simple information on model training compute, as is already captured by BIS' work, but also model architecture, size and training data, in addition to findings from capability evaluations.

4. Optimized Reporting Mechanism

We propose a reporting schedule that builds on BIS' suggested quarterly update requirements. Rather than enforcing quarterly reports from labs, our proposed mechanism requires new reports from labs only when a registered model is meaningfully improved beyond a defined capability threshold. Complementing this capability-based reporting, a bi-annual update can keep information on model security and evaluation metrics up to date, since model security or other features may be changed and have national security implications without improving the model capability.

5. Information to be Reported

We recommend that **12 categories** of information are reported upon registration. These categories of information have been chosen to optimize USG knowledge and capacity building while creating negligible additional overhead on AI labs and AI innovation. Our full proposal details what specific information should be required within each category.

1. Basic Information	5. Training Data	9. Functions of the Model
2. Open-Source Status	6. Model Architecture	10. Post-Deployment Monitoring
3. Model Size & Parameters	7. Hardware Information	11. Model Physical Security
4. Training Compute	8. Model Evaluations	12. Model Cybersecurity

6. Enforcement

A registry can be effectively enforced through market mechanisms that encourage compliance. By making model registration a prerequisite for commercial operations in the US, developers of the most powerful dual-use models would have strong economic incentives to register. Similar examples in other industries include:

- Know-Your-Customer¹² (KYC) standards in banking protect financial institutions against fraud, money laundering, and terrorist financing;
- Worker registration systems such as the E-Verify system¹³ that requires employers to verify the eligibility of their employees to work in the United States. Employers face penalties for knowingly hiring or continuing to employ unauthorized workers.

See the full report here for the comprehensive proposal: [AI Model Registries: A Foundational Tool for AI Governance](#)

In Conclusion

The rapid advancement of artificial intelligence presents a strategic imperative for American leadership. The United States has an opportunity to act decisively to secure its technological edge while managing competitive dynamics with strategic rivals.

In this document, we have submitted information on four high-priority areas:

1. **US Policy Leadership in the AI Economy:** AI will transform the global economy with sector-specific disruptions. The U.S. must implement forward-looking policies to address potential unemployment, inequality, and uneven technology diffusion to capture benefits, maintain competitive advantage and become a future leader in AI Economic policy.
2. **National Security and Innovation:** Frontier AI technologies will redefine defense capabilities. Rather than nationalizing innovation, which threatens market competition and faces financial and legal hurdles, targeted policy levers can effectively harness private sector innovation while securing critical national interests.
3. **AI Diplomacy and American Control:** Establishing high-level communication channels with competitors will prevent misperception and reduce the risk of pre-emptive actions against U.S. AI infrastructure, preserving America's technological lead.
4. **Strategic Information on Powerful Dual-Use Models:** The proposed Advanced Dual-Use AI Model Registry will provide American leadership with foundational information on the most powerful AI models while supporting America's innovation ecosystem.

¹² [Know Your Customer \(KYC\) | Swift](#)

¹³ [E-Verify | USG](#)

For more information on any of the above, please contact [REDACTED]
Convergence Analysis stands ready to provide our continued support and additional insights and clarification on any aspects of this submission as the government develops the AI Action Plan.

Respectfully submitted,

Convergence Analysis