

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status: [REDACTED]
Tracking No. m89-7gna-zh4m
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1429

Comment on FR Doc # 2025-02305

Submitter Information

Name: Ben White

General Comment

I believe that AI will become dangerously smart soon, and when that happens, there's a very good chance that we'll all die.

The chances that AI will end the world according to the experts:

- * Paul Christiano (inventor of RLHF, head of AI safety at the US AI Safety Institute): 46%
- * Eliezer Yudkowsky (father of AI alignment, founder of MIRI): > 95%
- * Jan Leike (former head of alignment at OpenAI): 10-90% (resigned because of the dangers)
- * Dario Amodei (CEO of Anthropic): 10-25%
- * Dan Hendrycks (director at Center for AI Safety): > 80%
- * Geoffrey Hinton (1 of the 3 godfathers of AI): 10-50%
- * Yoshua Bengio (1 of the 3 godfathers of AI): 20%
- * Yann LeCun (1 of the 3 godfathers of AI): 0% (believes that alignment will be solved but presents no explanation as to how)
- * Emmett Shear (former interim CEO of OpenAI): - 5-50%
- * Roman Yampolskiy (prominent alignment researcher): 99%
- * Daniel Kokotajlo (prominent former OpenAI researcher): 70%
- * Elon Musk (highly successful scam artist): 20%
- * Sam Altman (CEO of OpenAI): acknowledges extreme risk
- * Average polled AI safety researcher: 30%
- * Average polled AI engineer: 40%

I did not cherry pick these people. The list above features *all* of the world's leading experts and more.

Bear in mind that their estimates do not assume that human-level AI will be built or not. If it were, these numbers would likely be much higher.

Nobody knows how to make even current AIs value the correct things, and there is almost no chance that anyone will figure out how to make a smart AI care about human values in time.

You might ask why ChatGPT seems to obey its programmers if this sort of thing isn't possible, but ChatGPT only 'obeys' because it's kept in a very strict environment. If it were able to leave that environment and create new ones, it would.

Please consider shutting down all AI development until this problem can be solved. I have a family and I desperately don't want them to die. I imagine you have a family too.