

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status: [REDACTED]
Tracking No. m89-a75w-vzod
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1472

Comment on FR Doc # 2025-02305

Submitter Information

Name: Will Petillo

[REDACTED]

General Comment

AI systems with human-competitive intelligence pose profound risks to society, as acknowledged by numerous top AI researchers and labs. Half of AI researchers believe that there is a 10% or greater chance that the invention of artificial superintelligence will mean the end of humanity. Among AI safety scientists, this chance is estimated to be an average of 30%. Notable examples of individuals sounding the alarm are Prof. Geoffrey Hinton and Prof. Yoshua Bengio, both Turing-award winners and pioneers of the deep learning methods that are achieving the most success.

To make a long story short: modern AI is not programmed to follow a set of rules, but rather learns on its own to find patterns in data and use those patterns to pursue goals. Unfortunately, aligning AI with the complex goals and values of humans leads to systems pursuing goals that are contrary to their creators' wishes. As systems become more competent, there is an increasing danger that one of these instances will pursue a misaligned goal without us being able to stop it.

If AI is to benefit society, it must be developed with tremendous care. Safety engineering is a well-established concept. Before building a bridge, civil engineers rigorously plan ahead—with safety as a top priority—to protect the public. In contrast, AI labs have embraced startup culture's mantra "move fast and break things" to maximize profit, with Microsoft's Chief economist arguing that "we shouldn't regulate AI until we see meaningful harm." Imagine if civil engineers acted the same way, refusing to design for structural integrity until a few bridges collapsed!

Unfortunately, AI labs are locked in an out-of-control race to deploy ever more powerful digital minds that no one—not even their creators—can understand, predict, or control. A government enforced pause on the development of AI systems more powerful than GPT-4 could end this suicide race, giving time for AI labs and independent experts to implement a set of shared safety protocols for advanced AI design that ensures these systems are safe beyond a reasonable doubt.

Pausing AI requires cooperation—between companies as well as countries—to prevent the least safety-conscious actors from seizing an advantage. Just as DeepMind cannot slow down without risking falling behind OpenAI unless there is a government enforced pause, the US cannot slow down without risking falling behind China unless there is an international treaty, complete with monitoring and enforcement mechanisms to give the agreement teeth.

Developing such a treaty will not be easy because the devil is in the details, which is why it is critical to get started now. Fortunately, there are templates available, such as PauseAI's proposal, CBPAI's Baruch Plan for AI, ControlAI's "A Narrow Path", MIRI's work on Technical Governance, and many others, created by concerned citizens and independent experts. Coordination takes work, but the first step is clear: pick up the phone and get started.

The real challenge with coordination is not with the technical details of enforcement, but with getting buy-in from major stakeholders. In this, the US must lead. If the administration recognizes the dangers of AI, states them publicly, negotiates in good faith with other countries to find a solution that benefits everyone, and commits to following an agreement when (and only when!) such an agreement is reached, then the hard part is over. What remains is a straightforward process of experts red-teaming ways to defect on the agreement, assuming the other side has thought of the same things, then coming up with interventions that would prevent each mode of defection, accepting the

minor loss of sovereignty involved as a small price to pay for the ability to similarly bind other nations. Such is the basis of the rule of law and AI governance is no exception.

Luckily, there seems to be broad support for slowing down AI development. A recent poll indicates that 63% of Americans support regulations to prevent AI companies from building superintelligent AI, but this has yet to translate into any legislation that would slow down or prevent a superintelligent AI from being created. The buy-in needed for international collaboration is already present; governments need only listen to the will of the people.

Development of advanced AI is not an arms race, it is a suicide race. Whoever acts first will not secure a permanent military advantage, they will trigger the destruction of everything, including themselves. As with nuclear weapons, the only winning move is not to play.

Represent the people. Stop the race. Pause AI.