PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status: [redacted]
Tracking No. m89-b2zh-0vn8
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1488
Comment on FR Doc # 2025-02305

Submitter Information

Email: [redacted]
Organization: Coalition for Health AI (CHAI)

General Comment

Please see attached file for the Coalition for Health AI (CHAI)'s response.

Attachments

CHAI AI Action Plan_Final

About the Coalition for Health AI

The Coalition for Health AI (CHAI) is an industry-led, non-profit public-private partnership that believes in the potential of Artificial Intelligence (AI) to advance the practice of medicine.

We bring together the broad spectrum of interdisciplinary stakeholders in the US health ecosystem to drive the development and appropriate use of Responsible AI in health. We create consensus among innovators, clinicians, health systems, payors, and expert stakeholders across the sector.

Our 4,000 members come from nearly 3,000 organizations, along with hundreds of individual experts and stakeholders actively participating in workgroups, bringing diverse perspectives to our community.

- **Health Systems:** Our community includes over 200 health system members, spanning regional providers like MedStar Health, Mercy, and Providence, as well as leading academic medical centers such as Duke Health, Mayo Clinic, and Stanford Medicine.
- **Professional Advocacy and Specialty Groups:** Core to our work is building on foundational knowledge and standards that ensure quality in health care. We have a strategic partnership with the National Health Council who help ensure patient and community advocacy groups lead our workgroups, and we have the President and CEO of the National Association of Community Health Centers on our board. Organizations like the American Society of Clinical Oncology, College of American Pathologists, National Council for Mental Wellbeing, HL7, the Patrick J. McGovern Foundation, and The SCAN Foundation are among more than 100 philanthropic, professional advocacy, medical specialty, and standards organizations that have joined the CHAI community.
- **Healthcare Industry:** We're thankful to have membership from leaders in the health industry from CVS and UnitedHealth Group to OCHIN and Solventum.
- **Startups:** 75% of members are from industry, of which 24% are startups with fewer than 50 employees. These startups are creating solutions used across healthcare. They include Abridge, Aidoc, Ambience Healthcare, Affineon, Bayesian Health, Bend Health, Healthvana, Innovaccer, and Suki.
- **Tech Infrastructure Providers:** We are grateful for the broad participation of the AI frontier model and infrastructure providers, including OpenAI, Oracle, AWS, Google, Anthropic.

Board of Directors

- **Chair: John Halamka, M.D., M.S**., President, Mayo Clinic Platform
- **Secretary and Treasurer: Michael Pencina, Ph.D.,** Chief Data Scientist, Duke Health; Professor of Biostatistics & Bioinformatics and Director of Duke AI Health, Duke University
- **Suchi Saria, Ph.D.,** Endowed Chair & AI Professor, Johns Hopkins; Founder & President, Bayesian Health; Advisor, National Academy of Medicine AI Code of Conduct
- **Nigam Shah, MBBS, Ph.D.,** Chief Data Scientist, Stanford Health Care; Professor of Medicine and Biomedical Data Science, Stanford University School of Medicine
- **Eric Horvitz, M.D., Ph.D.,** Chief Scientific Officer, Microsoft; member, President's Council of Advisors on Science and Technology (PCAST)
- **Morgan Cheatham,** Vice President, Bessemer Venture Partners; Medical Trainee, Brown University
- **Jennifer Goldsack, OLY,** CEO, Digital Medicine Society (DiMe)
- **Kyu Rhee MD, MPP**, President and CEO of National Association of Community Health Centres (NACHC)

Our AI Action Plan

Methodology

CHAI is the nation's leader at bringing together diverse perspectives to forge actionable, consensus-driven solutions in health AI. Our strength lies in translating high-level principles into practical tools that work and are adopted across the ecosystem.

In December 2022, we published our **Blueprint for Trustworthy AI**, developed through extensive stakeholder collaboration to define foundational principles for responsible AI in healthcare. We then refined this work with the **Responsible AI Guide and Checklist**, published in June 2024, bringing together hundreds of experts to create a structured playbook for AI development and deployment - offering concrete guidance on ethics, quality assurance, and evaluation.

Having defined principles of fairness, transparency, usefulness, security, privacy, and safety in these playbooks and guidelines, CHAI then created, with similar consensus-building, its **Applied Model Card.** This is a practical tool that balances the needs of our vendor community (for example, protecting their intellectual property) and health systems (for example, their liability for technology used in their health system) to produce an approach that has been widely adopted and is easy to use. For example, 36 leading health systems and AI solution providers, have publicly pledged the rapid adoption of the Applied Model Card. Often compared to a **nutrition label for algorithms**, CHAI has established the Applied Model Card as a trusted standard for AI transparency, procurement, and governance and is freely available on GitHub, complete with a fillable template, instructions, an example, and supporting resources. It's a testament to CHAI's ability to turn broad agreement into something tangible, scalable, and impactful.

Similarly, we have developed this AI Action Plan alongside **150** of our founding partners and members. It reflects that there is, in fact, broad agreement across stakeholders on the core principles of responsible AI in healthcare. Our technology industry and health system members **resoundingly align** on key priorities, including the need for **transparency, preferred routes for AI assurance, trust-building mechanisms, disclosure of failures, and proactive bias mitigation**. CHAI continues to lead in transforming these shared priorities into practical, scalable solutions that we put forward today.

More detail on respondents and our survey findings are in Appendix 2.

Our Vision: Leading the Future of AI in Healthcare to Foster Human Flourishing

We are in a transformative era of technology, poised to redefine healthcare and unlock new opportunities for human flourishing.

AI, especially frontier models in Generative AI, offers an unprecedented opportunity to advance health, empower individuals, and drive well-being across physical, mental, and social dimensions.

We envision a future where America is the global leader in deploying AI in healthcare, ensuring that AI is embedded in every health system (regardless of resources), and benefits every American. Over the next four years, we want to see notable and tangible progress made in addressing healthcare's most persistent challenges - access, cost, and quality - driven by AI. Our goal is for everyone to be using AI confidently, because the AI-enabled solutions in healthcare are not only widely available but also high-quality, safe, easy-to-use and effective. We envision AI enabling America, its businesses and its entire population to truly flourish.

To realize this vision, we set out a plan to:

- **Accelerate AI Innovation –** Investing in infrastructure that ensures the development of safe, effective, and reliable AI solutions capable of addressing healthcare's greatest challenges.
- **Promote Responsible AI** – Embedding external and local evaluation, monitoring, and transparency reporting into AI governance to drive responsible, and high-quality AI deployment.
- **Democratize Access to Data –** Upholding free-market principles to foster secure, robust, and widely accessible data while incentivizing ethical, patient-centric data sharing and protection.
- **Strengthen U.S. Leadership in AI** – Advancing public-private partnerships that unite government, industry, and research institutions to unlock cutting-edge AI advancements while upholding quality, safety, and reliability.
- **Cultivate World-Class Health AI Talent** – Empowering the people behind healthcare's most critical decisions to harness AI's potential safely and effectively. Just as past medical revolutions required new skills, the AI era demands a workforce fluent in its capabilities, limitations, and ethical considerations.

CHAI's AI Action Plan: Enabling Responsible & Scalable AI in Healthcare

1. Accelerate AI Innovation by building National Health AI Innovation and Solution Infrastructure

The health AI market is expanding rapidly, yet adoption remains slow; of 1,016 FDA-approved AI products, only around 2% have been adopted.[1] A key challenge is ensuring AI performs reliably in real-world clinical settings, where models can subtly degrade over time or fail to generalize beyond lab conditions. Health systems have historically struggled to distinguish effective AI from unproven or low-quality solutions, making standardized evaluation critical for broader adoption.

A National AI Innovation and Solution infrastructure is essential to address these challenges. While some health systems may have the capacity to assess AI performance internally, many will need independent, standardized evaluation resources. Our survey of 150 CHAI members confirmed that **standardized AI performance benchmarking** is the top priority, by a long way, across all stakeholder groups, including startups. Health customers want standards to compare apples with apples, and industry want standards to perform against.

AI applications vary widely in risk. High-risk applications, such as models used in diagnosis or direct patient care, should be subject to stronger oversight, while lower-risk AI, such as administrative tools that indirectly impact patient safety and care, should require fewer reporting requirements. Unlike the EU's blanket classification of most health AI as high-risk that will likely stifle innovation, a nuanced, risk-based framework would allow models to be de-risked through mitigations like human oversight, increasing adoption without unnecessary regulatory burden. This is widely agreed upon in our community, with 94% of our survey respondents supporting a tiered AI risk classification framework.

Just as risk mitigation in health AI should use a tiered approach, so should monitoring. Some domains of healthcare are inherently high-risk, and identifying when AI has failed - and when harm could have been prevented is challenging, but imperative. A model used nationwide might benefit thousands while harming others in ways too subtle to detect at individual sites. Mandatory post-market monitoring was identified as the top accountability priority by 72% of our survey's respondents, further justifying real-world performance tracking. A national framework for AI monitoring and post-market surveillance would provide the scale and visibility needed to detect systemic failures and improve safety.

Ultimately, standardized evaluation and a tiered assurance framework would enhance trust in AI, allowing health systems to deploy proven solutions with confidence, helping innovators demonstrate success while preventing harm from ineffective models. To achieve this, CHAI urges the Administration to establish a federated AI quality assurance network, leveraging public-private partnerships to provide data access, tools, benchmarks, and real-world monitoring to evaluate AI performance, safety, and impact.

- **Develop AI Assurance Resources for Trust & Transparency**

[1] Dr Keith Dreyer, CDSO Mass General Brigham, testimony to FDA Digital Health Advisory Committee

- Establish a federated AI quality assurance network leveraging public-private partnerships to provide testing frameworks, benchmark datasets, and performance metrics.
- Enable AI validation against standardized evaluation frameworks before deployment in clinical settings.
- Enable AI model performance disclosures detailing training data, limitations, and risks where it is relevant to safety.
- Establish a national framework for real-world monitoring & AI pre- and post-market surveillance to ensure reliability and prevent unintended harm.

- **Develop AI Safety & Risk Stratification**
    - Implement tiered risk classification of AI solutions based on their impact on patient care, and level of human oversight.
    - Establish robust risk mitigation mechanisms to address emerging safety issues.

---

2. Promoting Responsible AI by Strengthening AI Governance & Liability Protections

Healthcare has historically operated within a clear, well-established legal framework; one that was not designed for AI. Trying to fit all AI into the current state of regulation forces providers and health systems to take on incalculable legal risks, especially pertaining to clinical AI solutions, leading to stagnated adoption.

AI has the potential to vastly improve diagnostic speed and accuracy, outperforming human experts in areas such as radiology while dramatically reducing the administrative time and cost involved in care delivery:

- **Performance Gains:** 40% improvement in high-skilled task performance [2]
- **Time Savings:** 30% reduction in time to complete documentation [3]
- **Productivity Improvements:** 25% increase in patient encounters from using AI tools in clinic [4]
- **Lower Costs**: 20x cheaper cost per support interaction [5]

As AI advances, the risk paradox grows more complex: holding clinicians accountable for a failure in AI-driven pattern recognition that goes far beyond human ability and explainability seems at odds with fostering broad-scale innovation. Yet, currently it is assumed that providers and health systems shoulder the legal burden of AI liability, despite neither developing nor testing how these models operate.

[2]Dell'Acqua F, McFowland E, Mollick E, et al. Navigating the Jagged Technological Frontier: Field Experimental Evidence of the Effects of AI on Knowledge Worker Productivity and Quality. Social Science Research Network. 2023;(24-013). doi:https://doi.org/10.2139/ssrn.4573321
[3]Oscar Health. AI Use Case: Messaging Encounter Documentation - Oscar Tech - Medium. Medium. Published January 23, 2024. Accessed March 14, 2025. https://medium.com/oscar-tech/ai-use-case-messaging-encounter-documentation-5f47380e000f
[4]Bali E. Primary Care Costs Explained. Carbonhealth.com. Published June 5, 2023. Accessed March 14, 2025.https://carbonhealth.com/blog-post/introducing-hands-free-charting-ai-enabled-note-taking-for-more-personal-visits
[5]AI: The Coming Revolution. Coatue. Published November 16, 2023. https://www.coatue.com/blog/perspective/ai-the-coming-revolution-2023

Our survey findings show that respondents strongly prefer real-world oversight mechanisms over pre-deployment validation alone. Our community also resoundingly (92%) support mandatory fairness and bias disclosure, which suggests assurance frameworks should explicitly include demographic fairness evaluations. Ensuring that AI decisions remain explainable and auditable is critical for both patient trust and legal clarity.

Without a regulatory framework that contemplates AI and its rapid transformation, health providers will remain hesitant to adopt the AI solutions that stand to make the greatest impact and enable even the smallest rural hospital to have access to specialized insights. This Administration has the opportunity to establish clear accountability structures that balances liability protections for both developers and implementers of AI solutions.

- **Define AI Liability & Accountability Mechanisms**
    - Clarify accountability structures for AI developers, healthcare institutions and providers, and users.
    - Develop clear protocols for AI failure reporting, liability claims, and enforcement mechanisms.
    - Ensure AI solutions are fair by enabling access to testing to better understand variability in performance in AI tools, and mitigate accordingly.
    - Ensure model performance is documented and standardized, including considerations on: limitations, and decision-making processes to ensure accountability and informed decision-making.

---

3. Democratize Access to Data by Ensuring Secure, Interoperable, and Fair AI Data Ecosystems

Data powers the AI revolution, and building effective health AI requires access to vast amounts of patient data, safely and securely, and with adequate privacy controls in place.

Advances in privacy-preserving data methods, such as Federated Learning and Confidential Compute, offer potential solutions, allowing AI to be trained on health system data without risking patient information leakage and unnecessary disclosure..

Federated learning - a decentralized AI training approach - has significant stakeholder support, with 71% of survey respondents endorsing it as a viable method for privacy-sensitive healthcare AI. However, 74% expressed concerns about the lack of standardization across healthcare systems, underscoring the need for clear, industry-wide frameworks for interoperability. Similarly, more than two-thirds of respondents emphasized the need for stronger ethical AI data-sharing guidelines, reinforcing the importance of ensuring the current regulatory approach contemplates the importance of enabling better use of data for appropriate purposes..

The nuclear industry was stalled for decades due to a few high-profile disasters - health AI faces the same risk if cybersecurity and privacy aren't prioritized. A handful of breaches could erode public trust, derailing AI adoption in healthcare for years. Robust privacy protections are critical to ensuring AI's future.

- **Establish federated & privacy-preserving AI data infrastructure**
  - Support decentralized data-sharing models that protect patient privacy while enabling AI innovation.
  - Build upon TEFCA, FHIR, USCDI+, and other efforts to ensure data standardization continues to prioritize easier facilitation of data sharing and insight generation.
  - Establish or promote privacy preserving guidelines for primary and secondary health data use.
- **Enhance AI Cybersecurity & Digital Trust**
  - Enable health systems to use privacy-preserving technologies to better leverage internal data.
  - Develop cross-border cloud and cyber harmonization strategies to align AI policies with international best practices.

---

4. Strengthen US Leadership by Aligning AI Investment with Healthcare Transformation Goals

AI is a means, not an end - a tool to make healthcare higher-quality, more affordable, and accessible. To maximize its impact, investment must be aligned with healthcare's most pressing challenges. Governmental guidance and leadership is critical to ensuring that developers and implementers of AI solutions have the confidence and guidance necessary to ensure continued investment and deployment of AI at a rapid pace. .

The HITECH Act revolutionized healthcare with the deployment of EHRs, but also by linking investment to the most valuable transformations via Meaningful Use. This directed innovation toward real system needs. Similarly, healthcare delivery providers find themselves hesitant to experiment because differentiation of AI products is unclear and return on investment has not been demonstrated in many cases.

This Administration can accelerate AI adoption by defining key areas of opportunity and backing them with investment - ensuring AI delivers real, transformative value where it's needed most.

- **Create clear federal AI priorities**

- Fund infrastructure advancements for healthcare delivery providers to ensure that they are ready for rapid adoption of AI solutions that improve access, quality, and cost-efficiency in healthcare.

- Direct AI investment toward high-priority clinical areas, such as chronic disease management and rural healthcare access.

- **Integrate AI into value-based care models**
  - Align AI adoption with CMS reimbursement structures that incentivize safe, effective, and fair AI use.
  - Establish economic incentives for responsible AI deployment.
  - Support AI-driven care coordination and predictive analytics to enhance value-based care and alternative payment models.

---

5. Cultivate World-Class Health AI Talent by Building AI Literacy & Workforce Readiness

AI can transform healthcare, but people remain the most valuable asset. Maximizing AI's potential means helping healthcare workers adapt - from administrators streamlining appointments to physicians using AI for early disease detection.

To work effectively with AI, clinicians must understand its strengths and weaknesses. AI models can fail in specific scenarios, such as data shifts or population changes, making AI literacy essential for safe and confident adoption. We have big gaps in training clinicians to use AI tools effectively. We know that clinicians can over-rely on the technology, so knowing when to use it and to what extent is critical. For example, in a study of 50 physicians' diagnostic reasoning, physicians using GPT4 scored 76.3, physicians using conventional resources scored 73.7, but GPT4 alone scored 92.1.[6] AI literacy is foundational to responsible adoption, and there is near-universal agreement (90% of survey respondents) that investing in workforce training is essential for safe and effective AI integration.

As trust grows, AI in healthcare will become as unremarkable as Electronic Health Records - a routine part of care. However, patients must trust that AI serves their interests, not hidden agendas. While they don't need to understand AI's technical details, they should know when AI is used in their care - especially in critical decisions like prior authorization.

- **Launch AI Training & Upskilling Programs**
  - Develop AI literacy programs for healthcare providers, with a primary focus on the nursing workforce, to ensure proper AI integration into clinical practice.

---

[6]Goh E, Gallo R, Hom J, et al. Influence of a Large Language Model on Diagnostic Reasoning: A Randomized Clinical Vignette Study. medRxiv (Cold Spring Harbor Laboratory).doi:https://doi.org/10.1101/2024.03.12.24303785

- Expand training opportunities for AI developers, regulators, and policymakers in healthcare AI evaluation.

- **Empower Patients & Providers with AI Education**
    - Foster trust in AI through public education campaigns.
    - Create AI explainability resources for clinicians and patients to enhance transparency.

Appendix 1: Definitions

- "Artificial intelligence" means a machine-based system that can, for a given set of human-defined objectives, make predictions, recommendations, or decisions influencing real or virtual environments;
- "Quality Assurance resources" means a set of tools, frameworks, methodologies, and guidelines that support the training, validation, testing and evaluation of Health AI performance, ensuring it meets specified standards for accuracy, clinical utility, reliability, safety, and ethical considerations.
- "Quality assurance network" means a coordinated system of organizations, institutions, and regulatory bodies that collaborate to establish, maintain, and improve the evaluation, monitoring, and validation of Health AI systems, ensuring their ongoing compliance with safety, efficacy, and ethical standards.
- "Responsible AI" means the design, development, deployment, and governance of artificial intelligence systems in a manner that prioritizes fairness, transparency, accountability, reliability, and the minimization of harm, particularly in high-stakes domains such as healthcare.
- "Human flourishing" means the holistic well-being of individuals and communities, encompassing physical health, mental well-being, social connection, and the ability to pursue meaningful lives, which can be enhanced or undermined by the responsible deployment of Health AI.
- "Post-market surveillance" means the continuous monitoring, assessment, and regulation of Health AI systems after deployment, ensuring that they maintain safety, accuracy, and effectiveness over time, while identifying and mitigating risks that may emerge in real-world clinical settings.
- "Federated learning" is a decentralized approach to training machine learning models. It doesn't require an exchange of data from client devices to global servers. Instead, the raw data on edge devices is used to train the model locally, increasing data privacy.
- "Confidential compute" is a security and privacy-enhancing computational technique focused on protecting data in use. Confidential computing can be used in conjunction with storage and network encryption, which protect data at rest and data in transit respectively.

Appendix 2: CHAI Survey Responses

We received 150 respondents from senior stakeholders across our ecosystem:

- 40% Health Systems
- 36% Tech Industry (Enterprise: 13%; Small to medium enterprise 10%; Startup 13%)
- 19% Academic and/or Research Institutions
- 13% Professional Services
- 7% Associations or Foundations
- 3% Government (State or Federal)
- 3% Health Insurer

Examples of respondents in each category are below:

- Health Systems
  - Providence, Duke, Stanford, Northwell, Mayo Clinic, UC Davis, The Permanente Medical Group, Ochsner, Community Health Network, UMass Memorial, RUSH, Sutter Health, Emory, University of Kentucky Healthcare, John Hopkins, Mount Sinai, Mass General Brigham
- Tech Industry
  - Medtronic; Samsung, Bayer, Hyro, Atropos Health, Lyric AI, Lifelink Systems, Affineon Health, Veradigm LLC, NTT Data, Globant, Evidentli, Cognome, Healthvana
- Academic and/or Research Institutions
  - Johns Hopkins University; Creighton; SSM Health; Mayo Clinic; Drexel University; Emory Pediatric Institute; University of California, Elsevier
- Professional Services
  - NCQA, Health Information Alliance, Slalom, C4i, Global Solutions Group, Wolters Kluwer, Ascent Strategy Group, FINN Partners, Metis Consulting Services,
- Associations or Foundations
  - Civitas Networks for Health, Massachusetts Health Data Consortium, American Speech–Language–Hearing Association, The SCAN Foundation, DiMe, National Assoc. for Behavioral Healthcare, American Nurses Association, National Alliance against Disparities in Patient Health, American Psychiatric Association
- Government (State or Federal)
  - FDA, VA
- Health Insurer:
  - Optum; CVS; Blue Cross Blue Shield (MN), SCAN Health Plan

Executive Summary:

Across stakeholders, there is broad agreement on foundational principles for responsible AI in healthcare. Both healthcare systems and the technology industry align on key priorities, including the need for transparency, preferred routes for AI assurance, trust-building mechanisms, disclosure of failures, and proactive bias mitigation.

However, a notable point of divergence arises around **AI risk classification**. While most stakeholders support a tiered framework to differentiate AI risks, the technology sector does not agree on a tiered classification of risk framework.

When it comes to **AI data governance**, there was low agreement across the groups. Health systems, the technology industry and academic research institutes (72% of respondents) agreed on mandatory post-market monitoring as the priority for accountability.

Interestingly, **government respondents (though only 5 respondents) were less supportive of transparency measures** compared to other organizational groups.

On the issue of **AI literacy**, there is widespread agreement on the necessity of training for healthcare providers.

Summary of results:

Quality Assurance Resources

- **Standardized AI performance benchmarks** are the top priority for CHAI members when it comes to AI assurance resource development, significantly leading other categories. This was consistent across all subgroups.
- The survey results indicate a strong consensus among CHAI members that **transparency in AI systems is critical for responsible adoption in healthcare.** ~90% of respondents rated key transparency measures as **somewhat important or very important**, with particularly high support for public disclosures of AI model limitations and risks (111 respondents rating it very important) and real-world bias testing & demographic fairness evaluations (95 rating it very important).
- 94% agree with a **tiered AI risk classification framework,** indicating strong enthusiasm for a risk-based oversight model rather than a one-size-fits-all regulatory approach.
- When asked for further quality assurance resource development, themes emerged around:
    - **Data Quality & Governance** – tools to track data provenance to ensure auditability, and frameworks for standardized data quality assessment.
    - **AI Testing & Evaluation Resources** – open-access benchmarking datasets, independent third-party verification, and sector-specific AI assurance labs to ensure continuous validation and generalizability.
      Fairness & Bias Mitigation – clear fairness guidelines, demographic testing for wide-ranging clinical populations, and mechanisms to prevent AI-driven healthcare inequities.
    - **Regulatory & Policy Guidance** – AI policy toolkits, liability frameworks, and AI incident reporting systems akin to adverse event tracking in pharma.

- **Continuous Monitoring & Post-Deployment Oversight** – real-world monitoring of AI performance, structured reporting of failures, and AI readiness criteria for safe clinical adoption.
- **Security & Privacy Protections** – AI-specific cybersecurity protocols, multi-factor authentication, controlled data access, and privacy safeguards to protect both individuals and businesses.
- **Education & Stakeholder Collaboration** – AI literacy programs, clinician engagement in AI development, and clear guidance on AI implementation and limitations.
- **Ethical AI Use & Human Oversight** – clear guidelines for when AI should and should not be used, ensuring human-in-the-loop oversight in critical decisions, and setting usability and governance ethical standards.
- **Transparency & Explainability** – clear, contextual AI decision explanations, standardized dataset descriptions, public access to evaluation benchmarks, and AI labeling requirements.

Fairness and Bias Mitigation

- There is consensus (92%) that supports mandatory disclosure, with a split between requiring it for all AI models (55%) vs. only high-risk AI (45%). Voluntary disclosure is largely rejected (2%).
- 92% of respondents agree that **prioritizing fairness and bias mitigation** in AI across diverse populations is important. This suggests AI assurance frameworks should include strong bias evaluation methods and fairness criteria, as well as inclusion of fairness in transparency efforts (e.g. model card).

Federated Learning: where AI models are trained on decentralized data rather than centralized patient data.

- Federated learning **is widely supported (71%), making it a viable direction for AI model training in privacy-sensitive domains like healthcare.** The 26% neutral group presents an opportunity to educate stakeholders on how federated learning works, its trade-offs, and potential industry adoption.
- However, 74% of respondents cited **lack of standardization across different healthcare systems** as their top concern with federated learning. Stakeholders need clear guidelines and industry-wide frameworks to harmonize data formats, protocols, and governance practices

AI Governance

- The responses suggest a preference for **ongoing, real-world oversight and clear accountability structures** rather than just pre-deployment validation. This reinforces the need for a multi-layered governance approach that combines regulatory reporting, liability frameworks, and direct user feedback to ensure AI safety and effectiveness.
- Over two thirds of respondents said **standardized guidelines for ethical AI data-sharing** was the most critical priority. Over half (53%) opted for stronger patient data access & consent controls, highlighting the importance of individual agency and trust in AI governance.

Workforce

- 90% of respondents support either significant or moderate AI training investment, making **workforce development** a non-negotiable element of responsible AI integration.