PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status: [redacted]
Tracking No. m89-btf9-2uz6
Comments Due: March 15, 2025
Submission Type: Web

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1495
Comment on FR Doc # 2025-02305

Submitter Information

Email: [redacted]
Organization: PauseAI NYC

General Comment

AI systems with human-competitive intelligence pose profound risks to society, as acknowledged by numerous top AI researchers and labs. Half of AI researchers believe that there is a 10 percent or greater chance that the invention of artificial superintelligence will mean the end of humanity, not including the risk of other catastrophic harms. Among AI safety scientists, this chance is estimated to be an average of 30 percent. Notable examples of individuals sounding the alarm are Prof. Geoffrey Hinton and Prof. Yoshua Bengio, both Turing-award winners and pioneers of the deep learning methods that are achieving the most success.
To make a long story short: modern AI is not programmed to follow rules, but rather learns on its own to find patterns in data and use those patterns to pursue goals. Unfortunately, attempting to align AI with the complex goals and values of humans leads to systems frequently pursuing goals that are unknowable, unsafe, unethical or all three. As AI systems become more competent, there is a significant danger that one of these instances will pursue a harmful goal. Once this happens, we humans may very quickly find out that we are unable to stop it.
If AI is to benefit society, it must be developed with tremendous care. Safety engineering is a well-established concept. Before building a bridge, civil engineers rigorously plan ahead - with safety as a top priority - to protect the public. In contrast, AI labs have embraced startup culture's move fast and break things mantra to maximize profit, with Microsoft's Chief economist arguing that "we shouldn't regulate AI until we see meaningful harm." Imagine if civil engineers acted the same way, refusing to even think about structural integrity until a few bridges collapsed and killed thousands!
Unfortunately, AI labs are locked in an out-of-control race to deploy ever more powerful digital minds that no one - not even their developers - can understand, predict, or control. A government enforced pause on the development of AI systems more powerful than GPT-4 could end this suicide race, giving time for AI labs and independent experts to implement a set of shared safety protocols for advanced AI design that ensures these systems won't cause the catastrophic harm their own creators expect.
Pausing AI requires cooperation - between companies as well as countries - to prevent the least safety-conscious actors from seizing an advantage. Just as DeepMind cannot slow down without risking falling behind OpenAI unless there is a government enforced pause, the US cannot slow down without risking falling behind China unless there is an international treaty, complete with monitoring and enforcement mechanisms. Catastrophic harm from previous technologies like nuclear weapons has already been limited by such treaties (even with weak enforcement), so such a treaty for AI is not unusual. Even politically unaligned nations like China have demonstrated interest in global regulations.
Developing such a treaty will not be easy because the devil is in the details, which is why it is critical to get started now. Fortunately, there are templates available, such as PauseAI’s proposal, CBPAI’s Baruch Plan for AI, ControlAI’s A Narrow Path and many others, created by concerned citizens and independent experts. Coordination takes work, but the first step is clear: pick up the phone and get started.
The real challenge with coordination is not with the technical details of enforcement, but with getting buy-in from major stakeholders. In this, the US must lead, even if it is to our political detriment. If this administration recognizes the dangers of AI, states them publicly, negotiates in good faith with other countries to find a solution that benefits everyone, and commits to following an agreement when such an agreement is reached, then the hard part is over. What remains is a straightforward process of experts red-teaming ways to defect on the agreement, assuming the other side has thought of the same things, then coming up with interventions that would prevent each mode of defection, accepting the minor loss of sovereignty involved as a small price to pay for the safety of its people.
Luckily, there is broad support for slowing down AI development. A recent poll indicates that 63% of Americans support regulations to prevent AI companies from building superintelligent AI, but this has yet to translate into any legislation that would slow down or prevent a superintelligent AI from being created. The buy-in needed for international collaboration is already present; governments need only listen to the will of the people.
Development of advanced AI is not an arms race, it is a sprint to trigger the end of the world. Whoever reaches the finish line first will be

responsible for the destruction of all humanity. As with nuclear weapons, the best move is not to play, while convincing others to do the same.
Protect your people. Stop the race. Pause AI