As of: March 21, 2025
Received: March 14, 2025
Status: [redacted]
Tracking No. m89-e434-b2uh
Comments Due: March 15, 2025
Submission Type: API

PUBLIC SUBMISSION

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1540
Comment on FR Doc # 2025-02305

Submitter Information

Email: [redacted]
Organization: AI Futures Project

General Comment

See attached file(s)

Attachments

AI Futures RFI Response (2)

AI Futures RFI Response

March 14, 2025

To: Faisal D'Souza, NCO
Office of Science and Technology Policy
2415 Eisenhower Avenue
Alexandra, VA, 22314

From: Thomas Larsen, on behalf of the AI Futures Project

Re: Request for Information (RFI) on the Development of an Artificial Intelligence (AI) Action Plan (“Plan”)

The AI Futures Project is a non-profit dedicated to understanding and forecasting the development of AGI.

We have a good track record in predicting the future of AI. Daniel Kokotajlo, our executive director, was a governance researcher at OpenAI, where he worked on scenario planning. In 2021, he wrote a scenario forecast called "What 2026 Looks Like", whose predictions for the 2021 - 2025 period have been quite accurate. He was recently named one of Time Magazine's "100 Most Influential People In AI 2024". Eli Lifland, a researcher at the AI Futures Project, is one of the world's top competitive forecasters. He ranks first on the RAND Forecasting Initiative all-time leaderboard.

We think that AGI companies are substantially *underhyping* the level of AI capabilities that is achievable over the coming years. Artificial General Intelligence (AGI)—defined as AI systems surpassing human performance across all cognitive tasks—will likely be developed within the next decade.[1] While talk of AGI has become mainstream, very few people are acknowledging that soon after the development of AGI, the automation of AI research will likely precipitate a rapid transition to superintelligence: systems outperforming humans in every domain. These superintelligences will have a massive effect on the world — much larger than the industrial revolution.

The development of superintelligence poses a substantial chance of existential risk — outcomes where humanity is permanently disempowered by highly capable AI systems. The

[1]Our team has different views: I (Thomas) have a median AGI timeline of 2030. Meanwhile, Daniel's median is 2028, and Eli’s is 2032. All of us have substantial uncertainty: it could easily be later or earlier.

two central threats are a) the risk of AI takeover, and b) AI-driven concentration of power into a tiny group of elites. Both of these risks are irreversible, and so we need to take proactive measures to avoid them.

We make the following brief recommendations:

1. Build capacity in government to understand AGI. There should be a body in government trying to understand the timeline to consequential AI capabilities (including extreme capabilities, such as fully general superhuman researchers). This should involve regularly interacting with AGI companies like OpenAI, Anthropic, and Google DeepMind, to understand their current levels of capabilities.
2. Ensure that the technical alignment problem is solved before superintelligence is built. The US government should attempt to recognize when superintelligence is imminent. At that point, it is likely that AI alignment will not be solved. In this case, the US government should attempt to recognize the situation and pause or substantially slow down frontier AI development. This will be difficult, likely requiring a deal with China. The other option — directly building misaligned superintelligences — is substantially more dangerous, and must be avoided.
3. Aim to align superintelligence to the interests of all of humanity. On the current trajectory, even if we solve the alignment problem, the most likely outcome is that superintelligence is built to pursue the interests of a tiny group of elites. The US government should take urgent measures to distribute power over AGI projects. A good first step here is to require transparency over the model spec (the intended goals of the AI system).

Thank you for your time. We are happy to discuss or expand on any of the above proposals as desired.

March 14, 2025

To: Faisal D'Souza, NCO
Office of Science and Technology Policy
2415 Eisenhower Avenue
Alexandra, VA, 22314

From: Thomas Larsen, on behalf of the AI Futures Project

Re: Request for Information (RFI) on the Development of an Artificial Intelligence (AI) Action Plan ("Plan")

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

The AI Futures Project is a non-profit dedicated to understanding and forecasting the development of AGI.

We have a good track record in predicting the future of AI. Daniel Kokotajlo, our executive director, was a governance researcher at OpenAI, where he worked on scenario planning. In 2021, he wrote a scenario forecast called "What 2026 Looks Like", whose predictions for the 2021 - 2025 period have been quite accurate. He was recently named one of Time Magazine's "100 Most Influential People In AI 2024". Eli Lifland, a researcher at the AI Futures Project, is one of the world's top competitive forecasters. He ranks first on the RAND Forecasting Initiative all-time leaderboard.

We think that AGI companies are substantially *underhyping* the level of AI capabilities that is achievable over the coming years. Artificial General Intelligence (AGI)—defined as AI systems surpassing human performance across all cognitive tasks—will likely be developed within the next decade.[1] While talk of AGI has become mainstream, very few people are acknowledging that soon after the development of AGI, the automation of AI research will likely precipitate a rapid transition to superintelligence: systems outperforming humans in every domain. These superintelligences will have a massive effect on the world — much larger than the industrial revolution.

The development of superintelligence poses a substantial chance of existential risk — outcomes where humanity is permanently disempowered by highly capable AI systems. The

[1]Our team has different views: I (Thomas) have a median AGI timeline of 2030. Meanwhile, Daniel's median is 2028, and Eli's is 2032. All of us have substantial uncertainty: it could easily be later or earlier.

two central threats are a) the risk of AI takeover, and b) AI-driven concentration of power into a tiny group of elites. Both of these risks are irreversible, and so we need to take proactive measures to avoid them.

We make the following brief recommendations:

1. Build capacity in government to understand AGI. There should be a body in government trying to understand the timeline to consequential AI capabilities (including extreme capabilities, such as fully general superhuman researchers). This should involve regularly interacting with AGI companies like OpenAI, Anthropic, and Google DeepMind, to understand their current levels of capabilities.
2. Ensure that the technical alignment problem is solved before superintelligence is built. The US government should attempt to recognize when superintelligence is imminent. At that point, it is likely that AI alignment will not be solved. In this case, the US government should attempt to recognize the situation and pause or substantially slow down frontier AI development. This will be difficult, likely requiring a deal with China. The other option — directly building misaligned superintelligences — is substantially more dangerous, and must be avoided.
3. Aim to align superintelligence to the interests of all of humanity. On the current trajectory, even if we solve the alignment problem, the most likely outcome is that superintelligence is built to pursue the interests of a tiny group of elites. The US government should take urgent measures to distribute power over AGI projects. A good first step here is to require transparency over the model spec (the intended goals of the AI system).

Thank you for your time. We are happy to discuss or expand on any of the above proposals as desired.