

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 14, 2025
Status [REDACTED]
Tracking No. m89-esrd-fgw7
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-1557
Comment on FR Doc # 2025-02305

Submitter Information

Email: [REDACTED]
Organization: Salesforce

General Comment

See attached file(s)

Attachments

20250314_OSTP RFI Response on AI Action Plan



March 14, 2025

Faisal D'Souza, NCO,
2415 Eisenhower Avenue,
Alexandria, VA 22314, USA

Re: Request for Information on the Development of an Artificial Intelligence (AI) Action Plan.

Comments of Salesforce

Salesforce appreciates the opportunity to provide written comments on the development of an Artificial Intelligence (AI) Action Plan in response to the Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO) and National Science Foundation's request for information on behalf of the Office of Science and Technology Policy (OSTP).

About Salesforce

Founded in 1999, Salesforce is a global leader in cloud enterprise software for customer relationship management, providing software-as-a-service and platform-as-a-service offerings to businesses, governments, and other organizations around the world. Our customers represent companies of all sizes and across all sectors. Our business model is cloud-based and low-code, allowing for faster technology deployment and greater agility. We help our customers connect with their customers-or employees or citizens-in new ways using cloud, data, and AI technologies.

Salesforce has been active in the development and research of AI for more than a decade. Our latest AI innovation is Agentforce, a suite of out-of-the-box AI agents and a set of tools to build, customize, and deploy them across any industry. Agentic AI is a technological revolution-the third wave of artificial intelligence after predictive and generative AI. AI agents go beyond traditional automation, being capable of searching for relevant data, analyzing it to formulate a plan, and then putting the plan into action. This technology has incredible potential to enhance speed, responsiveness, and personalization for [private enterprises and the public sector](#).

Enterprise AI

At Salesforce, our commitment to trusted AI stands out through our focus on the enterprise AI space's unique characteristics. The business model of enterprise AI is trust. Both our commercial and government customers demand the highest levels of trust, reliability, and accuracy. As a trusted enterprise AI provider, we tailor our approach to focus on the use-case and context-specific applications of AI tools optimized for our customer's business



needs. Unlike consumer AI, which is often generalized and susceptible to misuse, Salesforce develops AI models and applications that are highly specialized and relevant to our customers' needs. These models are designed to perform everyday work tasks such as generating emails, summarizing sales or service calls, and suggesting actions for case resolution during customer support interactions, among many other agentic AI use cases. Our AI systems are built and trained on curated data with stringent privacy, security, and accuracy requirements. Whether or not mandatory regulations are in place, our enterprise customers will expect high standards of trust, responsibility, and safety from Salesforce's enterprise AI, and Salesforce aspires to exceed those expectations as a U.S.-headquartered global leader.

AI Opportunity

For AI to deliver on its promise to boost economic growth, raise living standards, and improve the delivery of public services, it needs to be widely adopted and used by organizations of all sizes operating across different economic sectors. Policies that support the uptake of enterprise AI have the potential to dramatically accelerate and pull forward these gains, as companies deploy AI at scale. AI agents are poised to become an increasingly important part of the way that enterprise AI is adopted across the economy. To unlock this opportunity, the US government should both lead in the uptake of these technologies as well as ensure a policy environment that encourages the development and use of responsible AI. Taking action on these issues is both a matter of national security as well as economic leadership for the United States.

Salesforce employs AI to help governments:

- Automate routine tasks-By offloading tedious manual tasks to AI or AI-powered agents, staff can spend more time focusing on complex tasks and strategic initiatives.
- Improve government efficiency-Streamlining processes and workflows allows agencies to complete tasks faster, often with fewer resources and at lower operating costs.
- Make better decisions-Real-time data and insights that AI can provide help agencies make more informed decisions faster.

No matter if it is agentic AI or some other AI innovation, the fundamental principles of sound AI public policy remain unchanged: risk-based approaches, with clear delineation of the roles of different actors in the ecosystem, supported by robust privacy, transparency, and safety rules. With this in mind, it is time to think beyond just regulating how AI is built and used. Policymakers should focus on creating the right conditions to support the wide adoption of trusted agentic AI across industries and geographies. Salesforce has previously shared thoughts in our white papers, [The Next Frontier in Enterprise AI: Shaping Public](#)



[Policies For Trusted AI Agents](#) and [Shaping the Future: A Policy Framework for Trusted Enterprise AI](#).

Trust as a Foundation For AI Adoption

Companies that use our products entrust their sensitive data, customer relationships, and, ultimately, their credibility in the marketplace, to us. For these customers, many of whom operate around the world and are subject to a wide array of rules, ensuring that our enterprise AI products are safe, reliable, and trustworthy is table stakes—our business depends on getting these issues right. For Salesforce to retain its market positioning, we must be always looking for means to meet and exceed the expectations of our customers regarding the trust, responsibility, and safety of our enterprise AI. One way we operationalize these commitments is by designing experiences that help make AI more reliable and that clearly explain the AI's action while delivering clear and correct information.

We are committed to making AI's benefits accessible to everyone while ensuring these technologies are safe, reliable, and adaptable to various environments. We provide our employees, customers, and partners with the necessary tools to develop and use AI safely and accurately. To achieve this, we prioritize Trust as our core value and we have established our [Trusted AI Principles](#) and our [Guiding principles for responsible agentic AI](#) that apply to both agentic and generative AI. Further, Salesforce proactively engages with governments, industry, academia, and civil society to advance trusted, risk-based, and globally applicable AI norms. By focusing on trusted AI, Salesforce aims to encourage faster adoption of this technology, leading to significant economic and productivity gains.

Policy Recommendation: Consumer and business trust in AI remains [elusive](#). To close this gap, public-private collaboration will be key. Salesforce supports, through an organization like the National Institute of Standards and Technology (NIST), the means for businesses to demonstrate their efforts to develop trusted AI. Tools like guidelines, principles, and frameworks can create a path to closing the trust gap and increasing AI adoption.

Policy Recommendation: As the federal government considers rules for the use and adoption of AI, a collaborative approach is critical. Salesforce supports the creation of a presidentially appointed committee to bring together the public sector, academics, civil society, and the private sector to develop innovative policies and approaches to AI. This multi-stakeholder approach will help us advance adoption and maintain national competitiveness.

Policy Recommendation: Testing is a crucial part of the [responsible AI lifecycle](#). Salesforce conducts rigorous AI testing in simulated environments. However, there is a need for a global, science-backed approach to evaluations, as AI testing is still a developing field with few established global best practices. Therefore, Salesforce



supports empowering an organization to create a forum where researchers can collaborate on scientific approaches and solutions to AI testing.

Procurement & Government Adoption

The US Government has a key role in fostering AI adoption. It doesn't just set rules for the private sector but can lead the way in deployment. AI in the public sector promises to increase government efficiency, improve decision-making, and streamline citizen service interactions. Federal leadership also gives important signals to state and local authorities. To move ahead rapidly and adopt AI with confidence, agencies need clear signals, incentives, and appropriate guidance from policy.

Policy Recommendation: Federal, state, and local governments should prioritize the adoption of AI applications with high-level strategies, administrative guidance, and resources that facilitate faster learning, experimentation, and scalable adoption of high-impact public sector use cases. In particular, Salesforce recommends the following:

- The Office of Management and Budget (OMB) should issue refreshed guidance leveraging the NIST AI Risk Management Framework for AI, ensuring a common framework for agencies to manage their AI systems and conduct AI risk assessments - rather than agency-by-agency approaches.*
- Streamline reporting and impact assessment requirements.*
- Maintain agency Chief AI Officers (CAIOs) and channels for cooperation and coordination across agencies.*
- Whenever possible, state-of-the-art commercially available solutions should be prioritized to reduce development and maintenance costs.*
- Maintain a central AI use case library of public sector applications.*
- Create a central point for agencies to consult with experts (government and private sector) to identify the best opportunities and paths forward to leverage AI for more efficient work process transformation.*

Maximizing the public benefit from AI requires a government that can move as fast as the private sector, while still respecting high standards for security, privacy, and responsible management of government funds. Public sector procurement rules need to be fit for purpose so that lack of clarity and outdated ways of managing government technology doesn't slow the ability to experiment with and deploy the latest tools. This includes accelerating and taking an appropriately risk-based approach to privacy and security evaluations so they don't become a bottleneck to the way AI applications are developed and used.

FedRAMP has been critical to facilitating the secure adoption of cloud tools by agencies that are not equipped to independently evaluate the security of technology products. However,



while "do once, use many times" is the right approach, FedRAMP is not today moving with the speed required to get a range of innovative AI tools in the hands of federal agencies.

***Policy Recommendation:** Modernize and update procurement and security evaluation processes across government so that individual agencies are better empowered to flexibly experiment with and deploy AI, including:*

- *Fully resource and streamline the FedRAMP evaluation process to reduce paperwork and speed evaluation timelines.*
- *Narrow the scope of systems for which FedRAMP approval is required, e.g. to exclude low-risk/ public information systems.*
- *Recognize commercial-grade certifications when appropriate to the level of risk.*
- *Extend FedRAMP's Agile Delivery Pilot program to streamline approvals for iterative projects that are frequently updated with new features.*

Data Privacy & Data Flows

Continued AI innovation and adoption of AI depends on trust, and trust starts from the foundation of clear privacy rules. Today, data privacy and security concerns are top barriers to AI adoption. Improving the protection of consumers' privacy rights in the United States will accelerate adoption by fostering trust between consumers and companies.

Personal data protection supports AI adoption by setting clear consumer and business expectations for protection of privacy; and preventing and establishing accountability for misuse. It also supports innovation by simplifying compliance for businesses of all sizes and putting clear rules around when and how personal data can be leveraged to train and use AI. When incentivized by clear rules, AI with the right guardrails is a [privacy enabler](#) that can minimize human error, excessive collection, or inappropriate access and use.

***Policy Recommendation:** To ensure trust and overcome an increasingly complex state-by-state patchwork, a federal law protecting the privacy rights of consumers should be adopted, based on the principles of transparency, participation or choice of the consumer, and accountability. Where highly sensitive personal or commercial data are used, Federal rules should continue to support and incentivize the use of privacy-enhancing technologies that protect privacy while allowing data to be productively used.*

Additionally, US leadership in AI means ensuring that US tech can be used around the world. This depends on a stable environment in which to conduct digital commerce and transfer data across borders, free of unreasonable restrictions. American AI providers seeking to engage with international markets can face a number of requirements that restrict the ability to transmit and store the data that underpins AI. These requirements can be costly, and impair the ability of American AI companies to access markets and compete.



Policy Recommendation: The US should take leadership in promoting and negotiating bilateral and multilateral agreements that provide durable protections for American AI and digital businesses to access markets and transfer data across borders seamlessly.

The US has worked with its European partners to ensure uninterrupted and compliant transfers of personal information from Europe to the US. The EU-U.S. Data Protection Framework (DPF) has been successful in facilitating more than USD 1 trillion in trade and investment annually, including a flourishing trade in AI services. Its continuance will ensure that US companies can fairly access European markets and ensure that European businesses leverage cutting-edge American AI technology.

Policy Recommendations: The US should work with European authorities to ensure the continued operation of the EU-U.S. Data Privacy Framework for seamless provision of services from US businesses to European markets.

Workforce

It is [estimated that 75%](#) of AI's ultimate value lies in the "front office," where businesses directly interact with their customers. Yet, [41 % of employee](#) time is currently lost to repetitive, low-value tasks that have little to do with this essential front-office work. AI has the potential to ease these burdens, allowing workers to focus on more meaningful and strategic tasks, making work not only more productive but also more fulfilling. The global workforce must be equipped to use this technology safely and confidently.

As AI reshapes the way Americans work, the focus should be on training, creativity, and critical thinking skills that are uniquely human. Equitable access to AI is essential so that all individuals, regardless of their background or location, can benefit from these advancements and contribute to a more prosperous future.

Governments should prioritize collaboration with civil society and industry to equip workers with the right skills to thrive in the new opportunities created by AI. It is imperative to expand and invest in public-private partnerships for workforce reskilling: Governments, industry, NGOs, and civil society each bring necessary perspectives and capabilities. Policymakers should focus on public-private partnerships specifically geared towards reskilling the workforce to adapt to the unique challenges and opportunities posed by AI.

Policy Recommendation: The Federal government should promote public-private partnerships focusing on upskilling workers in industries at risk of AI disruption, helping them transition to higher-value roles that involve managing and collaborating with AI systems. Understanding the industries and jobs at risk of AI



disruption is another opportunity for the government to partner with relevant stakeholders to produce quantitative information about the effects of AI disruption.

Policy Recommendation: Preparing the next generation for an AI-first future is essential. Education and vocational curriculums should include ongoing training in AI concepts, skills, ethical considerations, and its potential future integration as the discipline evolves. Educators should also be enabled and reskilled to integrate AI into the curriculum and meet the needs of future classrooms.

Policy Recommendation: Collaboration among AI companies, industry, universities, and labor unions to connect AI-ready talent with suitable job opportunities, facilitate the transition into the workforce, and develop AI skills frameworks that are transferable within and between industries.

Competition

Salesforce believes that markets and innovation thrive when customers have free and fair choices to find the solutions that best meet their needs. However, when a few dominant enterprises lock up critical inputs-directly or through partnerships-the choice is restricted and innovation is harmed. Regulators and policymakers must work to ensure that the AI space remains open and fair and does not tip in favor of incumbent cloud infrastructure providers who control critical AI inputs, including infrastructure, leading LLM models, data, and AI chips. Control of these components by these companies, coupled with vertical integration down the technology stack, has the potential to lead to durable market power across AI products and adjacent technologies, such as cloud infrastructure, office productivity software, and communication/collaboration software. Rather than AI leading to greater competition and innovation, as seen in past technological advancements, the AI revolution risks further entrenching existing dominant incumbent firms.

A core guiding value of Salesforce's generative AI strategy is to maximize choice for customers by being open, extensible, and model-agnostic in its integrations with the emerging AI ecosystem. Salesforce aims to enable customers to use the LLMs of their choice with Salesforce's services-whether that be a model provided by a third party, Salesforce, or the customer itself.

Policy Recommendation: Globally, competition authorities have been working to understand the complex implications of generative AI on market dynamics and, in some cases, assessing the health of competition in this space presenting an opportunity for global coordination and collaboration. Salesforce believes that the US has robust tools to meet the demands of a rapidly evolving landscape and would urge the use of this authority to conduct market studies, maintain an ongoing



dialogue with industry about market health, and take enforcement action where needed.

US AI Leadership

As AI technology rapidly evolves, the United States must implement an approach that reinforces its position as a global leader in innovation. Maintaining US leadership in AI means not only considering policies that drive economic growth or enhance national security but also those that create a foundation of confidence in the technology. This is best achieved by adopting a comprehensive and forward-thinking public policy approach that results from collaboration, seeks to build on global best practices, and is agile. Given the global nature of many actors in the AI ecosystem, the United States should look to continually engage globally to shape the future of technology.

The US has many tools at its disposal to ensure AI leadership, among them NIST. NIST has proven to be one of the most effective organizations for establishing voluntary, globally recognized standards that are comprehensive, collaborative, and dynamic. Their frameworks, which are internationally influential, have helped companies large and small evaluate AI practices, which contributes to faster adoption of this technology. Salesforce has engaged with and hopes to continue engaging with NIST's AI RMF and Playbook, which include many of the measures we have internalized as we develop our own AI technology.

Salesforce believes in a sensible approach to AI frameworks, rules, or guidelines that are interoperable with global best practices. This sensible approach must also extend to the establishment of standards and means for AI testing. This work is already being done in AI safety institutes around the world and we would encourage the US to engage in this global collaboration. They represent additional opportunities for the US to contribute to the creation of a globally agreed-upon scientific approach to the evaluation of AI models. In *President Trump's Executive Order 13859 on Maintaining American Leadership in Artificial Intelligence (2019)* one of the goals was prioritizing AI R&D and collaboration with academia and the private sector.

The US has the opportunity to lead in finding a path toward interoperability between international standards, schemes, and protocols. The emergence of fragmented policies and approaches to the responsible development and deployment of AI systems can raise the cost of operations for US AI companies and undermine consumer protections. To ensure US interests, global forums provide a means for the US to shape the global policy discourse to ensure an approach that allows for US market access and competitiveness.

***Policy Recommendation:** NIST is highly relevant to the AI governance conversation in the US. The agency should continue to be effectively funded and empowered to enact comprehensive AI frameworks. By establishing and championing rigorous AI standards, benchmarks, and best practices, via NIST, the U.S. can shape global*



regulatory frameworks so that American innovations become the benchmark for security, accuracy, safety, interoperability, and ethical governance.

Policy Recommendation: Any approach to AI rules promoted by the Administration should be both durable and flexible, focused on risk-based guidelines to ensure the trusted development and deployment of AI systems. Rules for AI should focus on the outcomes posed by AI systems that result in concrete harms instead of just the AI systems and models themselves. These rules should aim to assign responsibility and accountability based on companies' position on the value chain, which is rapidly growing both in size and complexity.

Policy Recommendation: Given that the US is home to some of the largest AI companies in the world, it is important for the country to be actively involved in developing better methods for testing and measuring the safety and security of AI systems. If an American perspective is absent from the work of safety institutes abroad, other viewpoints will fill the void, potentially lacking the same experience and insight, or not aligning with the same interests and priorities. Salesforce believes this effort aligns with the mission of the Select Committee on Artificial Intelligence established under the previous Trump Administration. This committee is responsible for coordinating AI research and policy across federal agencies to maintain the United States' leadership in AI. Salesforce urges the Administration to continue this vital work, but on a global scale and in collaboration with the various members of the AI Institutes network.

Policy Recommendation: As the US considers a range of possible responses to the opportunities and challenges presented by AI, it should prioritize finding global consensus in venues like the G7, the Organization for Economic Co-operation and Development (OECD), or the Association of Southeast Asian Nations (ASEAN).

Conclusion

It is evident that a cohesive and forward-thinking public policy approach is crucial for the United States to establish and maintain its leadership in the rapidly evolving field of artificial intelligence. To secure America's AI edge, the federal government should lead by example in standards development, governance, and adoption. Abroad, we must work with allies that share our values to deter the adoption of detrimental policies that restrict commercial access and impede data flows. Doing so without haste will boost economic growth and help safeguard national security.

Salesforce values the opportunity to share perspectives and contribute to the U.S. government's policy development process. We have long believed it is the responsibility of industry leaders to lend expertise for the public good. We are proud of our collaborations and investments in the necessary tools, training, and support for our customers' AI



transformations. Finally, Salesforce is grateful for the White House's work to foster a policy and business environment for a bright future with trustworthy, innovative, and competitive AI.