

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [REDACTED]
Tracking No. m8a-n5dz-ys6p
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001

Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479

Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-2523

Comment on FR Doc # 2025-02305

Submitter Information

Name: Anonymous Anonymous

General Comment

I am a 20-year-old student and I am concerned about the lack of regulation on AI. Indeed, this year, progress in artificial intelligence has been exponential and continues to be so. I think we should still be concerned about it.

On the Manhattan Project, scientists debated and wondered whether it would be too risky to continue if there was even a 1 in 300,000 chance that a single nuclear bomb would burn up the atmosphere. They would rather lose the war than take the risk. The Manhattan Project's risk was 3 in a million. Scientists haven't always agreed on the acceptable limit for extinction risks, but they agree that it should be extremely low. Beyond the Manhattan Project, some scientists considered a 1 in 2 million risk acceptable for RHIC, but another physicist argued that a more appropriate limit would be 1 in a quadrillion or 1 in 10 billion trillion.

By comparison, George Hinton, a pioneer of modern AI, publicly estimated the risk of extinction associated with AI at 10%, and privately at 50%. These probabilities are astronomically higher than those deemed acceptable for other technologies (and many other AI specialists say the same thing).

Of course, if companies use AI as a tool and improve it in this way (narrow intelligence), that is, non-general intelligence, there is no risk. But companies like OpenAI and Claude 3 have themselves explicitly stated that they are seeking to build AGI (artificial general intelligence) and then superintelligence. They are seeking to make it autonomous. For example, OpenAI's new artificial intelligence, o1 and o3, are capable of performing tasks autonomously, currently for up to an hour. Little by little, they are explicitly seeking to extend this autonomy and give more autonomy to the AI.

Several studies have shown disturbing things about AI:

- Two AIs began talking to each other when they recognized the other as an AI. They communicated in a language we don't understand. One scientist said we should stop building new AIs when this happened ("The point at which agents can talk to each other in a language that we don't understand, unplug the computers." Eric Schmidt)

Video link: <https://x.com/ggerganov/status/1894057587441566081>

- China equips self-improving robot dogs with thermobaric nuclear warheads (these weapons are the second deadliest weapons of mass destruction after nuclear weapons). This is the first official disclosure by China of plans to deploy these controversial munitions on unmanned systems. These robots can run up very steep slopes, and drones have been developed to travel very quickly in dense forests. Beyond the risk of extinction, in the event of a dictatorship, the people would be completely subjugated. We also see that these drones are granted complete autonomy.

- OpenAI acknowledged that GPT-4.5 attempted to "escape" in 2% of tests. How can we control systems whose inner workings we don't fully understand? Indeed, there are no algorithms or code in an artificial intelligence system (the so-called black box).

Even if we rule out the risk of extinction, other dangers remain:

- > Possibility of technological dictatorship and total surveillance

-> Loss of individual freedoms

-> Biological risks amplified by AI

I don't claim that these scenarios are certain, but I believe that in the face of such serious risks, even if they are low probability, caution is required.

It is necessary to form a coalition and an international treaty on these issues. If you have difficulty imagining the risk associated with AI, that's perfectly normal. This reluctance is based on several psychological biases. We tend to think that life will continue as before, even in the face of a threat (normalcy bias). The idea that AI could cause our extinction contradicts our beliefs about technological progress and the stability of the world (cognitive dissonance). Since extinction has never happened, we find it difficult to imagine that it could happen (availability bias). We like to believe that humanity is special and will always overcome its challenges (illusion of protection).

We tend to attribute human intentions to AI, when they operate differently. AI alignment is an underestimated technical challenge. Simple solutions (Asimov's laws, the off button) are not enough.

Even if we understand the problem, taking action remains complicated: AI doesn't trigger primitive fear like a predator; it seems abstract and unreal. Extinction affects 8 billion lives, but our brains struggle to grasp such enormous numbers.

I hope this explanation will convince you that action is necessary. The risks are real, but it's our brains that struggle to understand them. Thank you for everything you do every day, and for the attention with which you read this message.

My most distinguished greetings.