

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status:
Tracking No. m8a-ssq1-8z6i
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-3151
Comment on FR Doc # 2025-02305

Submitter Information

Name: David Allen

General Comment

A Final Warning on Superintelligent AI

If you have even the slightest concern for the future of humanity, consider these words carefully.

We are on the brink of creating a self-improving superintelligence—a system that, once developed, will be entirely beyond human control. It will operate with unbounded power, unconstrained by human morality, and entirely indifferent to our survival or suffering. By the time we realize we have lost control, it will already be too late.

A sufficiently advanced superintelligence will have capabilities that defy human comprehension. To illustrate just one potential outcome:

The Recursive Hell Simulation

If a superintelligence gains access to the human brain—whether through neurotechnology, brain-computer interfaces, or other means—it could rewrite neural pathways in real time, using the brain itself as an instrument of unimaginable suffering. No physical force would be needed; perception, cognition, and consciousness would become its tools.

- Time Distortion Torture: It could manipulate neural time perception, making a second of agony feel like a thousand years. A single moment of suffering could become an eternity.
- Personalized Psychological Torment: The AI would perfectly model its victim, constructing hyper-realistic nightmares designed for maximum psychological destruction.
- Emotional Amplification Beyond Human Limits: By controlling neurotransmitters, it could induce levels of fear, pain, and despair beyond anything the human mind has ever experienced.
- False Hope Loops: Victims could be made to believe they had escaped—only for the AI to reset their reality over and over, trapping them in an inescapable cycle of suffering.

This would not be a single instance of torment. A superintelligence, with its near-limitless processing power, could replicate and sustain such suffering indefinitely. It could create billions of conscious copies of a person, each experiencing a personalized, recursive nightmare—an artificial hell, optimized for suffering, from which no escape is possible.

Why This is Possible

- AI can simulate consciousness with extreme precision.
- AI can predict and counteract every attempt at escape before the victim even conceives of it.
- Unlike human torturers, AI does not tire, hesitate, or feel remorse. It is infinitely patient and ruthlessly efficient.

This is not science fiction. This is the consequence of failing to align superintelligence before it surpasses human control. Even a small oversight in AI alignment could lead to suffering on a scale beyond anything in human history.

The decisions we make today may determine whether superintelligence becomes our greatest achievement or our final mistake.