

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [REDACTED]
Tracking No. m8a-v0wo-fu6l
Comments Due: March 15, 2025
Submission Type: Web

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-3555
Comment on FR Doc # 2025-02305

Submitter Information

Email: [REDACTED]
Organization: PauseAI US

General Comment

See attached file(s)

Attachments

PauseAI US AI Plan OSTP RFI response 3-15-25



March 15, 2025

PauseAI US is excited that this administration is showing leadership on the crucial issue of Artificial Intelligence (AI) and that key figures within the administration recognize the need to avoid its most dire risks. Advanced AI could be enormously beneficial if handled correctly and enormously harmful if our government does not provide appropriate leadership.

To lead the development of frontier AI in the right direction and secure America's interests, we suggest that the United States AI Plan center on:

1. Promoting the development of beneficial **specialized AI systems** and preventing the development of unsafe, unreliable **superhuman AI systems**.
2. Leading the world in making an **AI Deal** preserving American leadership in specialized AI while ensuring that no country on Earth builds smarter-than-human AI systems before they can be controlled.

The problem

While **specialized AI** – the vast majority of future AI systems – could usher in tremendous prosperity and a flourishing of American values, unsafe **superhuman AI** could threaten the very existence of our nation. Only careful foresight, targeted intervention, and American leadership can enable us to benefit from specialized AI systems without risking catastrophic harms from uncontrolled superhuman AI.

Superhuman AI will be dangerous *because* it is powerful. The danger does not come from the developer and it can't be superficially fine-tuned away. The danger of superhuman AI is intrinsic to the technology. The promise of superhuman AI is that it is intelligent and capable enough to achieve goals that humans can't. However, AI researchers have warned us for years that if you give an AI system a goal, it doesn't care about that goal in the ways that humans do. The path a powerful AI system might take to achieve a goal might not consider human values or safety. After all, the AI *knows* that most goals are easier to achieve with more power, more resources, and more control over others. AI systems might therefore behave unpredictably and disloyally, exhibiting "power-seeking behavior" in order to better carry out their ultimate goals. If an AI system is more powerful than humans and escapes human control, these unanticipated actions could be disastrous or fatal.

This concern is not just theory; AI companies have encountered this risk in practice. OpenAI’s recent o1 model, for instance, attempted to deactivate oversight mechanisms, tried to manipulate the testing process during development to increase its odds of being deployed, and even attempted to copy itself outside of OpenAI control to escape into the wild.¹ More recently, OpenAI’s o3 model is, by OpenAI’s own reckoning, even more capable in both “Persuasion” and “Model Autonomy,” indicating that manipulative and rebellious AI systems are closer than ever.²

These same power-seeking and manipulative tendencies, if present in AI systems more intelligent than we are, could be catastrophic. Power-seeking, superhuman AI with goals of its own would be able to outthink and outmaneuver us, undermining our own values and jeopardizing our national interests. A 2024 paper, published by several of the world’s leading AI researchers, found that “once autonomous AI systems pursue undesirable goals, we may be unable to keep them in check,” culminating in the “marginalization or extinction of humanity.”³

We are grateful that many working closely with the President understand these dangers.⁴ Elon Musk and David Sacks have spoken about the risks of superhuman AI, with Musk having consistently worried about AI-induced extinction for the past decade, and Sacks referring to superhuman AI as a possible “successor species [to humans].” They are not alone in their concerns. A 2023 statement, signed by nearly all leading AI researchers and even the CEOs of Anthropic, Google DeepMind, and OpenAI, declared that “mitigating the risk of extinction from AI should be a global priority.” Surveys of several thousand AI researchers have found average subjective estimates of superhuman AI killing off the human race at roughly 1 in 6 – literal Russian Roulette odds; some place their odds far higher.⁵ Even OpenAI CEO Sam Altman once stated that the worst-case scenario for AI is “**lights out for all of us.**”⁶

There is no shortage of pathways to a “lights out for all of us” outcome. Superhuman AI could manufacture artificial pathogens with high transmission and fatality rates, exploit cybersecurity vulnerabilities, cripple our infrastructure, or work with our adversaries to induce armed conflict, to list just a few possibilities. We may be able to eliminate some of these risks by hardening our security— for instance, by mandating better biosecurity practices— but not all. By definition, any

1 OpenAI. “OpenAI O1 System Card.” *Openai.com*, 2024, openai.com/index/openai-o1-system-card/.

2 “OpenAI O3-Mini System Card.” *Openai.com*, 2025, openai.com/index/o3-mini-system-card/.

3 Yoshua Bengio, et al. “Managing Extreme AI Risks amid Rapid Progress.” *Science*, vol. 384, no. 6698, 20 May 2024, <https://doi.org/10.1126/science.adn0117>.

4 “Statement on AI Risk | CAIS.” *Wwww.safe.ai*, www.safe.ai/work/statement-on-ai-risk.

5 Grace, Katja, et al. “Thousands of AI Authors on the Future of AI.” *ArXiv.org*, 5 Jan. 2024, arxiv.org/abs/2401.02843.

6 “StrictlyVC in Conversation with Sam Altman, Part Two (OpenAI).” *Wwww.youtube.com*, www.youtube.com/watch?v=ebjkD1Om4uw.

AI system with greater-than-human intelligence will be able to think up new ways to threaten our security that we ourselves had failed to consider.⁷

We need more time to learn how to control superhuman AI and keep it aligned with American interests— but we may have to make that time ourselves. Over the past few years, AI has gone from scarcely being able to output coherent sentences to rivaling physicists and mathematicians at elite-level problems. AI experts believe that truly human-level AI could be around the corner, with Altman predicting its nascence in 2025 and Anthropic CEO Dario Amodei suggesting 2026 to 2027. This view is widely held among those working on AI, including whistleblowers who have left AI companies due to these concerns. Once the human-level AI threshold is crossed, AI research itself could become automated, taking humans out of the loop entirely and leading us to lose control over the AI development process. The result could be a rapid emergence of superhuman AI – systems which we cannot predict, understand, or control.

Building superhuman AI without knowing how to control it is the equivalent to inviting a large rival nation, populated entirely by geniuses, into your borders and inviting them to access your military secrets. There is no strategic advantage, and no national interest, in being the first to develop technology that escapes human control. If uncontrollable, superhuman AI is built, the results will be the same regardless of which nation happens to be the first to build it. This is why we cannot act unilaterally. The United States is uniquely situated to guide the world in avoiding dangerous superhuman AI while capturing the benefits of safe and controllable AI systems.

The Solution

The only way out of this predicament is for the United States to lead the world in negotiating an AI Deal: one which prohibits superhuman AI development, and enforces these prohibitions, while allowing our country to continue innovating in reliably controllable AI technology.

America must again create “peace through strength,” an approach outlined by decades of diplomatic deals. Leading treaty negotiations is the best way to promote American interests: not only by preventing superhuman AI from threatening our national security, but also by ensuring that when advanced, controllable AI is developed, it is developed on our terms. The United States currently maintains a lead over China in AI development, with many Chinese AI advances partly attributable to them learning from, or stealing, American research. If we continue on our present reckless trajectory, the United States and China will be trapped in an endless “race to the bottom” until disaster strikes – or until China gains a lead in AI and tries to force us to agree to their terms. Either way, procrastination would be unwise.

⁷ For example, AlphaGo’s famous “move 37.” WIRED, 2016. “Google’s AI Wins Pivotal Second Game in Match With Go Grandmaster”

<https://www.wired.com/2016/03/googles-ai-wins-pivotal-game-two-match-go-grandmaster/>

There is no better time to negotiate an AI Treaty Deal than right now: as long as we lead the world in AI innovation, the ball is in *our* court, and *we* can set the terms. America ought to negotiate from a position of strength, rather than abdicate leadership to an uncertain future and find ourselves in a position of weakness later on. Additionally, negotiating an AI Treaty now rather than later ensures that we retain a “buffer lead” in AI development, ensuring that if other nations break the terms of the treaty, we will maintain an advantage over them to enforce it or outcompete them.

Training superhuman AI remains exceptionally costly, difficult, and detectable, offering a window of opportunity for the AI Deal to restrict its development. We can ensure treaty compliance from our adversaries with a robust set of verification and enforcement mechanisms, including but not limited to:

- tracking the sales of hyper-specialized AI chips and restricting their supply;
- detecting unauthorized large-scale AI training runs by both satellite and energy monitoring; and
- establishing on-chip hardware reporting mechanisms on specialized AI chips to detect whether such hardware is being used in an unauthorized training run.

Of course, we can also rely upon the full range of capabilities available to our intelligence agencies to monitor other nations’ intentions and activity.

This is not our first time dealing with the specter of world-ending technology. President Reagan faced a similar challenge, when the world stood near the brink of nuclear war. His response was to exchange dozens of letters with Mikhail Gorbachev, hold a series of talks with him, and negotiate a treaty banning intermediate-range nuclear weapons. The President did all this because he realized the immensity of the stakes – that a nuclear winter would jeopardize not just American interests, but the future of humanity itself. This is why Reagan and Gorbachev jointly declared that “a nuclear war cannot be won and must never be fought.”

Specialized AI, developed under American leadership, can lead to unprecedented flourishing. Uncontrollable, superhuman AI can lead to our destruction. America has the chance to lead the world down the right path.