PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [redacted]
Tracking No. m8a-ywdg-16go
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-5664
Comment on FR Doc # 2025-02305

Submitter Information

Email: [redacted]
Organization: Arm, Inc.

General Comment

See attached file(s)

Attachments

OSTP_AI_RFI_February2025

Arm's Response to the Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO), National Science Foundation Request for Information on the Development of an AI Action Plan

Submitted to: Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO), National Science Foundation
Date: 15 March 2025
From: Vince Jesaitis, Head of Global Government Affairs, Arm Inc.
Subject: Response to Request for Information on the Development of an Artificial Intelligence (AI) Action Plan

Dear Faisal D'Souza, NCO,

Arm welcomes the opportunity to respond to the U.S. Department of Commerce's Request for Information (RFI) on the development of an Artificial Intelligence (AI) Action Plan. As a global leader in semiconductor and computing architecture, intellectual property (IP) and solutions, Arm is a critical enabler of the modern digital economy providing the world's most widely used computing architecture. Arm designs high-performance, power-efficient, scalable, and secure computing architecture and central processing units (CPUs) that serve as the "brain" in virtually every form of computing, from smartphones and laptops to cloud data centers, supercomputers, automobiles, and embedded systems, enabling AI across all forms of computing. Unlike traditional chip manufacturers, Arm does not fabricate semiconductors; instead, it licenses its architecture and designs to a vast ecosystem of technology companies, enabling a competitive, diverse, and globally resilient supply chain. This unique, open business model allows semiconductor companies—including U.S. industry leaders—to build customized, highly optimized AI, mobile, and infrastructure chips, driving innovation across multiple sectors.

The Arm architecture is at the heart of 99% of the world's smartphones, as well as a rapidly growing share of data centers, AI workloads, automotive, and edge computing applications. Each year, Arm partners ship more than 30 billion Arm-based chips (this equates to more than 900 Arm-based chips being produced each second), bringing the total to over 310 billion chips shipped globally. These chips power critical infrastructure, from cloud computing and 5G networks to industrial automation, automotive, and aerospace applications. As AI demands more efficient and scalable computing architectures, high-performance, power-efficient Arm-based processors are becoming essential to AI acceleration, from training large language models in hyperscale data centers to running real-time AI inference at the edge.

As one example, Arm is involved as a lead technology partner in the launch of the Stargate AI project, a collaboration between Arm, OpenAI, Nvidia, and Oracle to develop advanced AI infrastructure. The project requires high-performance and energy-efficient computing to support large-scale AI model training and inference. As part of this effort, Nvidia's Grace Blackwell Superchip, a central component of Stargate, incorporates Arm-based CPUs to optimize efficiency and performance. The Grace CPU, built on Arm's architecture, is designed to manage the increasing computational demands of AI while maintaining lower power consumption compared to traditional architectures. By enabling more efficient processing and scalability, Arm's technology plays a critical role in supporting the next generation of AI systems and ensuring that AI development remains both performant and cost-effective.

With its globally adopted, open computing architecture, Arm plays a key role in semiconductor supply chain resilience, enabling a broad ecosystem of U.S. and international chipmakers, device manufacturers, and cloud providers to innovate without relying on a single supply source. This diverse and decentralized ecosystem strengthens technological security and economic competitiveness, ensuring that AI, high-performance computing, and next-generation technologies remain accessible, efficient, and secure. As policymakers seek to enhance U.S. leadership in AI and semiconductor technology, fostering an open, competitive, and innovation-driven semiconductor landscape—one that leverages Arm's foundational architecture—will be critical to long-term economic growth and national security.

Given the strategic importance of AI to U.S. economic growth, national security, and technological leadership, it is critical that the AI Action Plan recognizes the role of semiconductor innovation in AI development and prioritizes investment in AI-optimized compute architectures. We appreciate the Department's efforts to shape a comprehensive AI strategy and offer the following insights and recommendations:

The Foundational Role of Semiconductor Design in AI

AI advancements depend on highly efficient, scalable, and specialized computing architectures. As AI models become larger and more complex, high-performance and energy-efficient computing is essential to sustaining AI innovation and deployment at scale.

Arm is at the forefront of next-generation AI computing, providing power-efficient, high-performance processor architectures that power AI across cloud, edge, and mobile platforms. This is why Arm is a key technology contributor to the Stargate project, a cutting-edge AI supercomputing initiative aimed at developing the next generation of AI infrastructure. Stargate represents a major leap forward in AI compute capabilities, requiring unprecedented levels of efficiency, scalability, and performance—areas where Arm's architecture excels.

The central processing unit, or CPU, is often referred to as the brain of a computing device because it manages and coordinates almost every operation. It processes instructions, controls data flow, and ensures that different software and hardware components work together. While CPUs are versatile and capable of handling many different tasks, they are not always the most efficient at processing the complex calculations required for artificial intelligence. AI workloads involve vast amounts of data and mathematical operations that need to be processed quickly and in parallel, which is beyond what a general-purpose CPU is designed to do.

To handle the demands of AI, modern computing systems incorporate specialized processors that work alongside the CPU. Each of these components has a distinct role in improving performance and efficiency.

Graphics processing units, or GPUs, were originally designed for rendering images and animations. However, they are now widely used in AI computing because they excel at handling multiple calculations at the same time. AI models, particularly those that rely on deep learning, require extensive computations to recognize patterns and make predictions. GPUs process these tasks efficiently by working on many pieces of data simultaneously.

Neural processing units, or NPUs, are specifically designed for AI-related tasks. While GPUs are powerful for general parallel processing, NPUs go further by optimizing calculations that are unique to AI models, such as recognizing images, text, and speech. NPUs are commonly found in devices that need to perform AI computations efficiently while consuming less power, such as smartphones, cameras, and embedded systems.

Tensor-based processors are another category of specialized AI chips. These processors are built to accelerate the complex mathematical functions needed for machine learning, making them particularly useful for large-scale AI applications, such as training deep learning models or processing real-time AI tasks.

Other types of specialized chips, such as field-programmable gate arrays (FPGAs) and application-specific integrated circuits (ASICs), offer additional flexibility or efficiency depending on the specific AI workload. FPGAs can be reconfigured for different tasks, making them adaptable for evolving AI applications, while ASICs are custom-designed for specific functions, maximizing efficiency but limiting flexibility.

Memory and data movement also play a crucial role in AI computing. High-bandwidth memory, or HBM, is used to ensure that AI processors can quickly access large amounts of data without delays. AI systems require rapid processing of vast datasets, and having efficient memory solutions allows processors to keep up with the demands of advanced AI models.

All of these components work together to power modern AI systems. The CPU acts as the central coordinator, directing data flow and managing system resources. GPUs, NPUs, and tensor-based processors accelerate AI computations, while specialized memory like HBM ensures that data is readily available for processing.

As AI continues to evolve, the need for specialized processors will continue to grow. Each new generation of AI models requires greater computing power, and the underlying hardware must advance to keep pace. The combination of CPUs, AI accelerators, and optimized memory solutions ensures that AI can drive innovation across a wide range of industries, from powering intelligent assistants and autonomous systems to enabling breakthroughs in scientific research and automation.

To ensure continued leadership in AI compute, U.S. policy should incentivize investment in AI-optimized semiconductor design for a range of computing needs, support energy-efficient AI compute architectures to balance performance with sustainability, and recognize the critical role of semiconductor design in AI leadership by integrating it into national AI strategies.

AI's Dependence on Specialized Compute Architectures

As discussed above, AI workloads are increasingly demanding specialized compute architectures optimized for efficiency, power consumption, and data movement. Large-scale AI training, such as the models supported by the Stargate project, requires highly optimized hardware architectures that maximize performance while minimizing energy costs.

Arm's power-efficient processor architectures are uniquely suited to support AI workloads at hyperscale. As AI systems become more compute-intensive, the industry must prioritize heterogeneous computing architectures that integrate CPUs, GPUs, NPUs, and AI accelerators. The success of initiatives like Stargate underscores the need for continued public-private collaboration in AI infrastructure development, where industry leads the technology development and government creates an environment conducive to innovation and investment to enable rapid buildout and upgrades of these projects.

To enable AI's continued evolution, U.S. policy should protect and promote open and interoperable AI computing architectures to encourage flexibility and competition, support next-generation semiconductor research to advance AI-specific memory, networking, and accelerator technologies, and encourage collaboration between AI research labs, semiconductor firms, and infrastructure providers to co-optimize AI hardware and software development.

AI Hardware Development Must Accelerate to Keep Pace with Software

The rapid advancement of AI software is outpacing hardware development, creating critical bottlenecks in AI compute performance, energy efficiency, and scalability. Developing cutting-edge semiconductor technology requires long lead times, often spanning four to five years from initial design to commercialization. The AI models powering today's breakthroughs are running on chips that were conceived half a decade ago, meaning the hardware innovations needed for future AI applications must be prioritized now. Without sustained investment in next-generation AI hardware, the U.S. risks falling behind in AI capabilities, limiting innovation and increasing reliance on foreign semiconductor supply chains.

The next generation of AI models will require:

- Custom AI silicon optimized for efficiency and performance, enabling faster model training and inference with reduced power consumption.
- Advanced memory and interconnect technologies to efficiently manage data flow in large AI models, reducing latency and bandwidth constraints.
- Low-power AI processing to enable broader AI deployment across edge and mobile platforms, ensuring AI can run efficiently in real-world applications, from autonomous vehicles to smart infrastructure.
- Flexible and reconfigurable computing architectures, such as chiplet-based designs and AI accelerators, to provide modular, scalable performance for a wide range of AI applications.

The U.S. must take proactive steps to secure its leadership in AI computing by prioritizing AI-specific semiconductor research and development (R&D). Policymakers should:

- Ensure CHIPS Act incentives account for domestic AI semiconductor design and manufacturing.
- Support AI chip startups and fabless semiconductor firms through incentives, tax credits, and access to national AI research infrastructure.
- Develop a national AI hardware roadmap, aligning industry, government, and academic research to accelerate advancements in AI-specific computing architectures.
- Streamline regulatory and permitting processes to ensure the rapid development of new AI hardware fabrication facilities and testing centers in the U.S.

Without a sustained focus on AI hardware, the U.S. risks a bottleneck in AI progress, hindering the economic and national security benefits that AI innovation can deliver. A comprehensive, long-term AI hardware strategy will ensure the country remains at the forefront of AI computing, semiconductor design, and next-generation digital infrastructure.

Ensuring Secure and Resilient AI Hardware

AI security begins at the hardware level, as the integrity, reliability, and trustworthiness of AI systems depend on the security of the underlying compute platforms. Without secure hardware foundations, AI models and applications are vulnerable to a range of threats, including data breaches, adversarial manipulation, and supply chain disruptions. As AI is increasingly deployed in critical sectors such as defense, healthcare, finance, and infrastructure, ensuring the security of AI compute platforms is essential for maintaining national security, economic stability, and public trust.

Hardware-based security measures provide a robust foundation for mitigating AI-related risks, as software protections alone are insufficient to defend against sophisticated attacks. Secure-by-design semiconductor architectures can help prevent tampering, unauthorized access, and adversarial exploits by embedding security mechanisms directly into AI processing units. Features such as secure enclaves, cryptographic accelerators, trusted execution environments, and real-time anomaly detection at the hardware level can significantly enhance AI system resilience. By addressing security at the hardware level, there is an opportunity to eliminate entire classes of vulnerabilities before they can be exploited at the software layer, creating a more robust security posture for AI-driven applications.

A secure AI ecosystem also depends on a transparent and resilient semiconductor supply chain. The increasing globalization of semiconductor manufacturing introduces risks related to counterfeit components, firmware backdoors, and potential disruptions in access to essential AI chips. Strengthening domestic and allied semiconductor production capabilities, enhancing traceability and verification of AI hardware components, and reducing reliance on high-risk supply chain nodes are all necessary steps to ensure long-term AI security.

To enhance AI security and trustworthiness, we recommend the following policy actions:

- Increase federal focus on cybersecurity at the hardware level by prioritizing research and investment in AI-specific, secure-by-design semiconductor architectures that embed safety and security features directly into AI compute platforms.
- Support industry-driven hardware security standards for AI computing, ensuring that AI hardware meet rigorous cybersecurity benchmarks for integrity, confidentiality, and resilience against tampering and adversarial attacks.
- Encourage domestic production of secure AI hardware by incentivizing semiconductor companies to integrate advanced security features into AI chips and processors though targeted CHIPS Act projects and funding, particularly under the Department of Defense Microelectronic Commons programs.
- Expand supply chain risk mitigation programs by enhancing semiconductor traceability, establishing stronger verification mechanisms for AI hardware components, and working with allies to reduce dependence on high-risk supply chain nodes.
- Require government AI systems to adopt hardware-level security protections to ensure that federal AI deployments, particularly in defense and critical infrastructure, are built on a foundation of secure compute architectures.

By prioritizing cybersecurity at the hardware level, the U.S. government can proactively address systemic AI vulnerabilities, safeguard national security, and strengthen trust in AI-powered systems across public and private sectors

---

Global Supply Chain Resilience

AI hardware development is deeply interconnected, requiring contributions from multiple countries at various stages, including design, fabrication, packaging, and testing. The semiconductor industry, which underpins AI computing, is one of the most complex supply chains in the world. Unlike traditional manufacturing industries where production can be localized, semiconductor production relies on highly specialized expertise, raw materials, and advanced equipment that are distributed across a global network.

Given this complexity, ensuring the resilience of AI hardware development requires more than just securing access to critical materials or increasing domestic manufacturing capacity. It also depends on enabling cross-border collaboration between companies that contribute to different parts of the supply chain. Many semiconductor firms and AI hardware developers operate globally, with research teams, design centers, and production facilities spread across multiple countries. Intra-company collaboration across borders is essential for optimizing production, accelerating innovation, and maintaining competitiveness.

Developing and producing an advanced AI chip requires an immense amount of engineering effort, with millions of total engineering hours spanning research, design, fabrication, testing, and packaging. From initial architecture and circuit design to advanced fabrication processes that require precision at the atomic level, thousands of highly specialized engineers contribute across multiple disciplines. This includes semiconductor physicists working on transistor scaling, electrical engineers optimizing power efficiency, software engineers developing firmware and AI model optimizations, and materials scientists refining the chemical compositions needed for high-performance chips. The sheer scale of expertise involved, combined with the need for continuous iteration and quality control, makes AI chip development one of the most labor-intensive and knowledge-intensive engineering feats in the world. Leading U.S. multinational semiconductor companies build and maintain this expertise by hiring critical talent in the U.S. and allied countries. Limiting these collaborations through overly restrictive policies could slow down innovation and increase costs.

While safeguarding national security is critical, policies that indiscriminately restrict technology sharing or impose rigid export controls on AI hardware risk fragmenting the supply chain in ways that could ultimately weaken the U.S. position in AI and semiconductor development. Export control measures must be carefully designed to prevent sensitive technologies from falling into the hands of adversaries without unnecessarily restricting American companies from participating in the broader global semiconductor market. Overly broad restrictions could isolate U.S. firms from essential partnerships, making it more difficult to source the best and brightest minds and the most advanced IP, designs, materials, components, and manufacturing processes needed for AI chip development.

A balanced policy approach focusing on strengthening strategic alliances with trusted international partners, encouraging collaboration among companies within and across borders, particularly between firms in allied nations with shared security and economic interests, will help ensure that AI hardware development remains robust, competitive, and innovative. At the same time, incentivizing investment in domestic capabilities—such as advanced semiconductor design, manufacturing, packaging, and testing—can reduce vulnerabilities without undermining the benefits of international cooperation.

Ultimately, the future of AI hardware development will depend on the ability of policymakers to recognize the interconnected nature of the semiconductor industry and refine policies that protect national security while fostering an environment where innovation can continue to thrive through global collaboration. Instead of viewing cross-border collaboration as a risk, policymakers should work to enhance partnerships that strengthen the semiconductor supply chain, ensuring that the U.S. and its allies remain at the forefront of AI hardware developments.

The Need for Government Investment in AI Infrastructure for Sensitive Use Cases

Many government applications require dedicated AI infrastructure to support sensitive and classified workloads that cannot rely solely on commercial cloud providers. National security and intelligence agencies increasingly depend on AI models for real-time threat analysis, situational awareness, and decision-making in complex operational environments. Similarly, AI plays a growing role in cybersecurity, helping to detect and mitigate cyber threats before they can compromise critical systems. Public sector services, including emergency response, fraud detection, and law enforcement, also benefit from AI-driven automation and predictive analytics, enabling faster and more effective government operations. However, the commercial AI marketplace is advancing at an unprecedented pace, rapidly outpacing government AI systems in capability and deployment speed. Without decisive action, this gap will continue to grow, making it more difficult for government agencies to adopt cutting-edge AI solutions and ensure they remain competitive in critical national security and public sector applications.

Commercial cloud providers offer valuable scalability, cost efficiency, and rapid deployment capabilities, making them well-suited for many government AI applications, particularly those that involve unclassified data, collaboration across agencies, or public service delivery. However, for highly sensitive workloads, reliance on commercial infrastructure presents risks related to data sovereignty, supply chain security, and potential access limitations in times of geopolitical tension or crisis. Additionally, certain national security and intelligence operations require real-time processing with minimal latency, which can be difficult to achieve with cloud-based solutions that rely on geographically distributed data centers. In these cases, government-owned and operated AI infrastructure is necessary to ensure secure, uninterrupted access to mission-critical computing resources.

To address these challenges, the U.S. government should:

- Invest in secure, government-owned AI data centers to handle classified and national security workloads.
- Expand access to AI-optimized supercomputing for national defense, intelligence, and cybersecurity operations.
- Ensure critical AI workloads can run on-premises where necessary, particularly for applications requiring low latency and continuous availability.

- Establish a hybrid AI infrastructure strategy that leverages commercial cloud solutions for non-sensitive applications while maintaining dedicated, government-controlled infrastructure for the most critical use cases.

By implementing these measures, the U.S. can ensure that its AI infrastructure remains secure, resilient, and capable of supporting evolving national security and public sector requirements.

The Need for U.S. Government Investment in AI Infrastructure and AI at the Edge

In addition to large-scale AI compute infrastructure, AI at the edge—where AI models run directly on local devices rather than centralized cloud servers— can play a critical role in enhancing efficiency, reducing latency, and improving security across government applications. Edge AI is particularly valuable for defense, critical infrastructure, and public safety, enabling real-time decision-making in remote or high-security environments. Investing in sovereign AI edge capabilities will allow the U.S. to fully leverage AI in secure communications, battlefield intelligence, and industrial automation.

A key example of this approach is the AI-RAN (Artificial Intelligence for Radio Access Networks) Alliance, a public-private initiative aimed at integrating AI-driven wireless networking, 5G, and Open RAN (Radio Access Networks) to enhance network efficiency, security, and performance. AI-powered RAN technologies improve spectrum management, optimize network traffic, and enable more resilient communications for government and military applications. The U.S. government should support AI-RAN and other edge AI initiatives to maintain technological leadership in next-generation networks and ensure secure, AI-optimized communications infrastructure.

To maximize the benefits of AI infrastructure and AI at the edge, policymakers should:

- Expand federal funding for AI edge computing research to accelerate innovations in secure, low-power AI processing.
- Integrate AI-powered edge solutions into federal agencies for cybersecurity, emergency response, and defense operations.
- Provide incentives for AI-driven Open RAN deployment to enhance secure, software-defined wireless networks.
- Ensure CHIPS Act funding supports edge AI development alongside large-scale AI infrastructure computing.
- Promote public-private partnerships to advance AI hardware and software co-optimization for high-efficiency, real-time AI applications.

By prioritizing AI infrastructure investments at both the cloud and edge levels, the U.S. can strengthen national security, industrial competitiveness, and public sector efficiency, ensuring that AI innovation remains a strategic asset for the country's long-term economic and geopolitical leadership.

Reducing Environmental and Permitting Barriers to Accelerate AI Deployment

As AI adoption expands across industries, data centers and semiconductor facilities are critical infrastructure for AI model training, inference, and deployment. However, lengthy environmental reviews and permitting processes create significant delays and cost burdens for data center construction, semiconductor fabrication plants, and advanced packaging facilities. To maintain U.S. leadership in AI and ensure timely deployment of computing infrastructure, policymakers should streamline environmental and permitting requirements without compromising essential protections.

Specifically, the U.S. government should:

- Expedite permitting for AI-critical infrastructure by creating a fast-track approval process for data centers and semiconductor facilities, recognizing them as national strategic assets.
- Reform the National Environmental Policy Act (NEPA) process to reduce unnecessary delays, including setting clear time limits for environmental reviews and eliminating redundant state and federal assessments.
- Standardize and simplify state-level permitting for AI data centers and chip manufacturing to ensure consistency across jurisdictions.
- Incentivize sustainable AI infrastructure by providing tax credits or incentives for facilities that meet high energy-efficiency and carbon-reduction benchmarks, rather than imposing restrictive regulations.

Without streamlined permitting and regulatory clarity, AI infrastructure projects risk delays of years, undermining U.S. technological competitiveness and economic growth.

The Need for Competitive Tax Policy to Support Semiconductor R&D

Semiconductor development is among the most R&D-intensive industries in the world, requiring sustained investment in chip design, fabrication technologies, AI acceleration, and energy-efficient computing. Developing next-generation AI chips involves multi-billion-dollar investments over multiple years, with increasing costs driven by the complexity of semiconductor design, advanced packaging techniques, and the integration of AI-specific accelerators. For example, the average cost to design a 2nm chip, which is the most advanced process commercially available today, is nearly three times higher than it was to design a 7nm chip, which was the leading process commercially available in 2018.

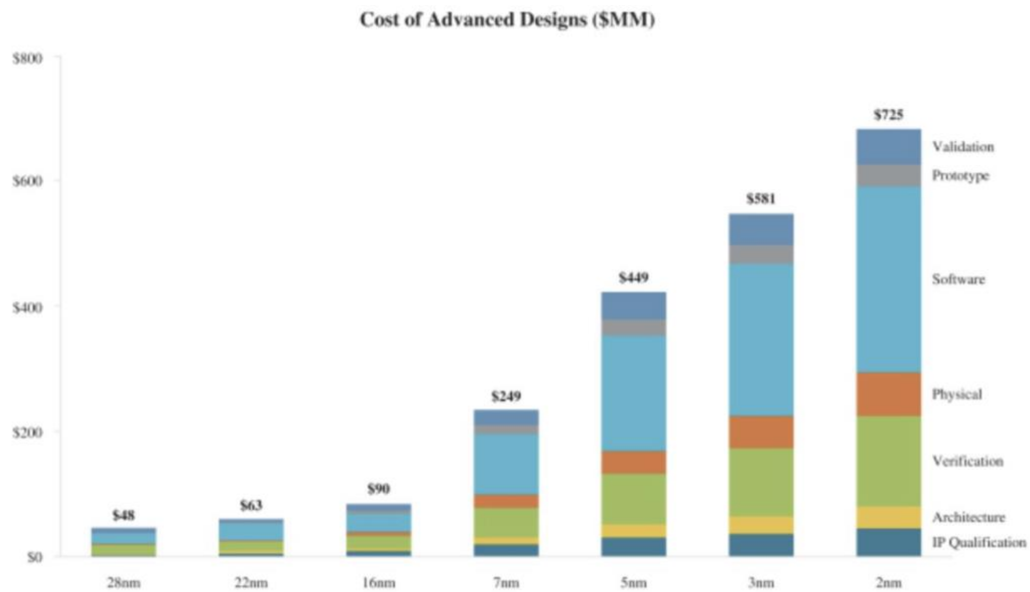

Source: IBS July 2022.

Given the diverse computing needs, and the financial demands on the private sector to continue this innovation cycle, a competitive tax policy is critical to ensuring that the U.S. remains a global leader in AI computing and semiconductor innovation. To sustain leadership in AI hardware (including semiconductor IP development and design), the U.S. must provide strong, predictable, and globally competitive tax incentives for semiconductor and AI-related R&D. Policymakers should:

- Restore immediate expensing of R&D costs (as opposed to multi-year amortization) to maintain investment levels and accelerate innovation.
- Expand tax credits for AI and semiconductor research to offset the rising costs of advanced chip design, prototyping, and software-hardware co-optimization.
- Ensure that tax policy encourages U.S.-based semiconductor design, preventing innovation from shifting to jurisdictions with more favorable R&D incentives.
- Align U.S. tax incentives with global competitors to retain AI chip design leadership, as many nations—including China, South Korea, and EU member states—offer more aggressive semiconductor tax incentives.

Without a globally competitive tax structure, U.S. semiconductor firms risk being outpaced in AI chip development, undermining both economic competitiveness and national security objectives. A robust and forward-looking tax policy will allow companies like Arm and its unparalleled ecosystem of semiconductor partners to continue investing in AI-driven chip innovation, ensuring that the U.S. leads in next-generation computing, AI acceleration, and secure, efficient semiconductor architectures.

Building a Strong, Competitive AI and Semiconductor Workforce

A highly skilled and adaptable workforce is vital to maintaining America's leadership in AI and semiconductor innovation. As demand for specialized talent grows, incentivizing education, training, and industry partnerships will be critical to economic growth, national security, and long-term competitiveness. Without a strong domestic talent pipeline, the U.S. risks falling behind in semiconductor design, AI development, and advanced manufacturing, ceding technological leadership to global competitors.

The Arm-led Semiconductor Education Alliance, founded in collaboration with key industry leaders, is working to address these challenges by bridging academia and industry through hands-on technical training, curriculum development, and workforce upskilling initiatives. This collaborative approach ensures that American workers remain at the forefront of innovation by equipping them with the cutting-edge skills needed to advance AI and semiconductor technologies.

To build a world-class workforce that sustains U.S. leadership in AI and semiconductor innovation, we recommend:

- Incentivizing STEM education and workforce training programs aligned with industry needs, particularly in semiconductor design, AI model optimization, and hardware-software integration.
- Strengthening partnerships between businesses, universities, and technical colleges to ensure that education programs are directly linked to real-world applications and industry demands.
- Expanding federal programs for AI and semiconductor apprenticeships to create hands-on learning opportunities and reduce barriers to entry into these fields.
- Encouraging private-sector leadership in workforce development through public-private partnerships, tax incentives for AI and semiconductor training programs, and investment in reskilling initiatives.
- Reforming immigration policies to attract and retain top AI and semiconductor talent, ensuring that the U.S. remains the preferred destination for the world's leading engineers, researchers, and technologists.

By leveraging industry expertise, targeted incentives, and market-driven solutions, the U.S. can cultivate a highly skilled workforce ready to sustain AI and semiconductor innovation for decades to come. Strategic action now will ensure that the U.S. remains the global hub for advanced computing, chip design, and AI-powered technologies.

Conclusion

AI innovation depends on advancements in semiconductor design and computing infrastructure. Arm is committed to collaborating with government and industry partners to drive the development of efficient, secure, and scalable AI computing architectures that enable AI everywhere. We appreciate the opportunity to provide input on the AI Action Plan and look forward to continued engagement.

Sincerely,
Vince Jesaitis
Head of Global Government Affairs
Arm Inc.