

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [REDACTED]
Tracking No. m8a-z68b-xow5
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-5899
Comment on FR Doc # 2025-02305

Submitter Information

Name: Cody Rushing

Email: [REDACTED]

General Comment

See attached file(s)

Attachments

aiaction

**Response to the National Science Foundation's and
Office of Science & Technology Policy's Request for Information
on the Development of an Artificial Intelligence (AI) Action Plan**

90 Fed. Reg. 9088 (Feb. 6, 2025)

Docket No. NSF_FRDOC_0001

March 15, 2025

Cody Rushing | [REDACTED]

As an AI security researcher, I'm submitting this comment to inform the AI Action Plan's development¹. We are witnessing unprecedented acceleration in AI capabilities, with major labs projecting transformative AI systems within the current administration's term. Anthropic, for instance, anticipates that by 2027, Claude will "find breakthrough solutions to challenging problems that would have taken teams years to achieve."

To ensure AI development promotes "human flourishing, economic competitiveness, and national security" as stated in Executive Order 14179, it's paramount that we address important AI Security issues. While many groups will comment on the potential of these AI models, the innovation must be supported by policies that address two crucial security concerns:

- We lack robust methods to ensure powerful AI systems reliably act in accordance with security protocols. This is reflected both in expert opinion, and the growing body of empirical examples of AI's attempting to subvert security protocols or self-exfiltrate². The technical challenge of preventing advanced AI systems from attempting to circumvent security measures remains unsolved.
- As AI labs produce more capable models, adversaries including nation-state actors will be increasingly incentivized to steal these models. The consequences of such theft would be severe and multifaceted, introducing a variety of national security risks

We need to be sure that labs are equipped to handle this challenge, and that governments establish incentives and resources to support them in this endeavor. To ensure AI development promotes "human flourishing, economic competitiveness, and national security" as stated in Executive Order 14179, there must be improved transparency between labs and government.

To maintain America's AI leadership while safeguarding national security, the AI Action Plan should have targeted transparency requirements that facilitate important information transfer

¹ The views expressed in this submission represent my professional assessment and do not necessarily reflect the official positions of my employer.

² See <https://www.anthropic.com/research/alignment-faking> as a particularly strong example of this.

between AI labs and governments. These measures address portions of critical security threats identified above:

i) Disclosure of In-Development Capabilities

Frontier AI labs should be required to notify federal authorities when they achieve significant capability breakthroughs, such as when they pose severe CBRN risks, or have the ability to automate important processes of AI development. Many such important capabilities are currently tracked by AI labs through their respective scaling plans (such as OpenAI's Preparedness Framework), but it's important that the performance of these models with respect to these evaluations is disseminated to the government, such that the government can be prepared to make important decisions. A good system would:

- Enable the AI Safety Institute (and/or other federal government institutions) to maintain awareness of capability advances without compromising the labs proprietary methodologies and secrets.
- Allow the government to proactively assess security implications of new capabilities rather than reacting after deployment.
- Create an early warning system for capabilities that could pose national security risks if stolen or misused.
- Help prioritize federal resources toward securing the most advanced systems from nation-state theft attempts.

ii) Robust Whistleblower Protections

Strong protections for AI researchers and engineers are essential for identifying security vulnerabilities before they can be exploited. It would be very beneficial to introduce whistleblower protections such that AI labs can be held accountable to the American people. These protections should:

- Establish secure channels for reporting severe security concerns, including model behavior that attempts to circumvent security controls.
- Reinforce America's commitment to safe AI development while deterring practices that could undermine security.
- Help identify potential insider threats before they result in data breaches or model theft.
- Create mechanisms for reporting cases where models demonstrate deceptive behaviors or failures that could manifest as security vulnerabilities.

In the global race for AI leadership, unnecessary regulatory friction would disadvantage American companies. However, certain minimal guardrails are not burdensome but rather essential infrastructure for sustainable progress.

These transparency measures are some of many that would require minimal resources from AI labs while providing crucial information that enables the government to identify where intervention is needed and—equally important—where it is not. These recommendations can help ensure that AI promotes the interests of the American people.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.