

PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [REDACTED]
Tracking No. m8b-2j9a-w6ke
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-8358
Comment on FR Doc # 2025-02305

Submitter Information

Name: Steven
Kaas **Email:**

[REDACTED]

General Comment

My main hope for the AI Action Plan is that it gives serious attention to the risks of losing control over future, highly advanced AI systems — risks that could lead to catastrophic or even existential consequences. Many submissions likely outline concrete policy proposals to address this. I don't claim special expertise in that area, but I want to highlight one approach that may not yet have been raised.

AI experts differ sharply on the severity of existential risk from misaligned AI. Researchers like Geoffrey Hinton, Yoshua Bengio, and Stuart Russell have warned that current AI advances are placing us on a highly dangerous trajectory. Others, like Yann LeCun and Andrew Ng, argue the risk is overstated.

Sound policy depends on rigorously assessing how serious an extinction-level threat AI poses, understanding the nature of that threat, and identifying how it can be addressed.

To that end, I suggest supporting structured, public dialogues between leading researchers on AI and its risks — perhaps modeled after the Dutch government's "Climate Dialogue" project (2012–2015) (<https://mwenb.nl/climate-dialogue/>), which brought together climate scientists with opposing views. Such a forum could foster more transparent, systematic, and constructive expert conversations, grounding AI policy in well-vetted ideas and making them accessible for public scrutiny.