PUBLIC SUBMISSION

As of: March 21, 2025
Received: March 15, 2025
Status: [redacted]
Tracking No. m8b-3shb-9tad
Comments Due: March 15, 2025
Submission Type: API

Docket: NSF_FRDOC_0001
Recently Posted NSF Rules and Notices.

Comment On: NSF_FRDOC_0001-3479
Request for Information: Development of an Artificial Intelligence Action Plan

Document: NSF_FRDOC_0001-DRAFT-9583
Comment on FR Doc # 2025-02305

Submitter Information

Email: [redacted]
Organization: Equity and Access in Algorithms, Mechanisms, and Optimization

---

General Comment

See attached file(s)

Attachments

OSTP_AI_comment

Final Submission:

Recommendation: Strengthening AI Literacy and Ethical Education

Education is a cornerstone for mitigating AI risks. By raising public awareness and fostering AI literacy, education helps users and developers recognize ethical pitfalls and safety issues early in the AI development process. Integrating AI ethics, AI risks and challenges, and risk management into higher education curricula and AI workforce training encourages responsible AI development and deployment. Both public awareness campaigns and formal education can nurture a culture of accountability, making AI technologies and AI developers align with human, societal, economic, and governmental values and well-being.

To achieve this, we encourage the following actions:

- Integrate AI ethics and risk awareness into curricula
    - Schools and universities should embed AI ethics, safety, and governance in STEM, social sciences, and humanities programs.
    - Educational initiatives should cover core AI concepts, biases, and responsible AI use, equipping students to navigate the ethical challenges of AI.
- Expand public AI literacy programs
    - Government should incentivize tech companies and media organizations to develop accessible AI education initiatives, such as online courses, workshops, and interactive platforms.
    - Governments should incentivize libraries, NGOs, and local governments should facilitate AI literacy programs to reach underrepresented communities and ensure inclusive access to AI knowledge.
- Support AI research and academic initiatives
    - Public institutions and academia should receive funding to advance research on AI safety, robustness, and ethical considerations.
    - Governments should fund universities to integrate AI governance principles and technical safety measures into educational programs to prepare the next generation of developers and policymakers.

Governments should promote open research on AI risk mitigation and establish ethical guidelines for AI development.

Recommendation: Categorizing AI Risks

A risk-based policy framework is essential for regulating AI. The EU's AI Act, for example, defines four levels of AI risk: minimal risk (free use, e.g. AI spam filters), limited risk (subject to transparency obligations, e.g. AI-powered chatbots), high risk (strict requirements before deployment, e.g. AI in hiring or critical infrastructure control), and unacceptable risk (banned applications, e.g. AI for social scoring). We believe that U.S. policies should take a similar tiered approach identifying prohibited or "high-impact" AI use cases that pose risks to national security, democratic values, or safety (such as Framework to Advance AI Governance and Risk Management in National Security).

We believe that U.S. policies should adopt a categorization of high-risk AI applications as those that could seriously affect people's rights, safety, or critical systems. These might include AI for biometric surveillance and identification, AI managing critical infrastructure (energy grids, traffic control), or algorithms making high-stakes decisions in healthcare, law enforcement, or employment similarly as in EU AI Act. Such systems should demand rigorous oversight, including robust risk assessments, transparency, human oversight, and regulatory approval before deployment. However, to maintain economic competitiveness and relief of unnecessarily burdensome requirements that hamper private sector AI innovation applications that do not affect people's rights and safety could be categorized as a low-risk AI application. These applications still require some notion of ethical guidelines (such as steps to mitigate bias), but generally can be less governed. We believe that educational institutions and industry groups through governmental-aided and instructed processes can develop best practices for low-risk AI, ensuring these innovations promote economic competitiveness and convenience without infringing on rights.

However, for high-risk AI, strong governmental control and regulation are warranted. Governments should set stringent standards, perform audits, and, where necessary, restrict or ban uses of AI that imperil human rights or national security. This aligns with global policy trends. Therefore, we ask if the United States' national security frameworks that emphasize controlling AI uses that could undermine democratic values or safety are going to be employed (Framework to Advance AI Governance and Risk Management in National Security), as well as if the EU AI Act that mandates strict compliance for high-risk systems are going to be adapted. We believe that such oversight ensures AI development serves humanity progress and flourishing, while protecting human dignity. Just not to be focused on U.S. and EU documents, China's AI ethical governance position paper warns that unregulated AI misuse may "undermine human dignity and equality". Majority of AI regulation documents, including Russia, articulate that AI advancement must safeguard human values. China advocates a people-centered, "AI for good" approach aimed at enhancing the common well-being of humanity, and Russian AI principles likewise state that protecting the interests and rights of human beings is the main priority of AI development (for example, Safe and Ethical AI (SEA) Platform Network · Linking Artificial Intelligence Principles (LAIP)). These perspectives reinforce that high-risk AI must be kept under governmental control to uphold human dignity, societal values, and security.

In contrast, low-risk AI systems can be guided more by public and educational institutions than by heavy regulation. Schools, universities, and civil society organizations could lead in creating ethical guidelines and literacy programs for everyday AI tools. By doing so, they ensure widespread understanding of AI and promote innovation that is competitive yet responsible. This lighter governance for low-risk applications – through codes of conduct, voluntary standards, and education – allows economic and technological advancement while still respecting privacy and fairness. It enables entrepreneurs and educators to experiment and deploy benign AI applications (like learning assistants or personalized recommendations) under general best practices, without the bottleneck of extensive regulation, as long as these applications remain in the low-risk category.

Recommendation: Preserve public data infrastructure (including NIST AI Institute)

Historically, much of the development of more advanced AI and ML has progressed via the so-called "common-task framework", the practice of a subfield declaring performance on a common challenge task as the barometer of AI progress, such as detecting handwritten digits from images (MNIST), categorizing natural-language labels from images (ImageNet), radiologist annotations of chest X-rays (CheXpert). The compilation of such domain-specific data is a necessary step to evaluate AI systems and thereby improve them, especially since often many innovations comprise of changes in design that are a priori difficult to compare which potential change must be best. Benchmarking on public data is crucial to advance AI. Public data and benchmarks also crucially accelerate the development of open-source AI models, which in turn accelerate AI development in general and lower barriers to adoption via enabling innovations in required compute resources, increasing their potential economic impacts.

Even private-sector data innovations within companies also rely on publicly available data. Federal public statistics provide vetted, reliable, and uncertainty-quantified information about the nation's population, aka feasible markets. Private firms conducting market research, for example, also rely on publicly available data on demographics to size potential market opportunities and compare trends based on their current customer base to project future operations under potential expansions. Keeping this public data accessible, rather than leaning on commercial private-sector data brokers, also enables data-driven decision-making for small businesses and other economic enterprises, spurring additional growth.

Additionally, the gamification of current standard benchmarks has led to the massive overfitting to benchmarking measurements and thereby deceptive and untrustworthy performance results that may not even generalize to the real world. However, by providing public benchmarking data, the US positions itself in standardizing more trustworthy evaluations; specifically, **we advocate for public benchmarking data to be dynamically updated and with proper documentation (e.g. datasheets)**. With the ability to dynamically update benchmarks, we can better ensure that models are evaluated on datasets that continuously reflect deployed contexts. With the feature of being properly documented, benchmarks can be more easily open to public scrutiny, which keeps in check how these datasets embody real world situations and values.

Recommendation: Strengthen Scientific Evaluation Frameworks for AI

As AI systems become more integrated into society, it is critical that their development and evaluation are grounded in **rigorous scientific standards**. We urge the AI Action Plan to incorporate policies that enhance the **measurement, transparency, and diversity of AI research**, ensuring that progress is both meaningful and responsible. Specifically, I recommend:

1. **Reducing Data Leakage:** Establishing clear guidelines to prevent data contamination in AI training and evaluation. Many reported AI advancements are undermined by data leakage, leading to overestimated capabilities. Standardized methodologies for dataset construction and rigorous validation protocols are essential.
2. **Converging on Evaluation Frameworks:** Developing standardized, transparent benchmarks for assessing AI capabilities, particularly for advanced AI models with general-purpose reasoning (e.g., AGI-like claims). Without consensus on evaluation metrics, it is difficult to distinguish between genuine advancements and misleading results.
3. **Encouraging the Publication of Negative Results:** AI research currently lacks incentives for publishing failures, leading to redundant work and an incomplete scientific record. Encouraging the disclosure of negative results will promote accountability, prevent wasted resources, and accelerate genuine innovation.
4. **Supporting Alternative Research Directions:** The dominant paradigm of scaling AI models has produced significant advances, but it should not be the sole focus of AI research. Investment in **alternative approaches**—such as neurosymbolic methods, causal reasoning, and energy-efficient AI—will help ensure a more **robust and sustainable** AI ecosystem.

By incorporating **scientific best practices** into the AI Action Plan, we can build a foundation for **trustworthy, reproducible, and equitable** AI progress.