

Securing American Leadership:

Recommendations and Implications for a National AI Action Plan

Parv Mahajan¹, Cyra Alesha, Jaehun Baek, Merritt Hornbuckle, Adelaida Jaramillo, Bratee Podder, Kishor Shanmugavel, Seth Shoneman, Emalee Collins, Andrew Wei, Dwayne Wilkes²

Executive Summary

This authorship group was organized by and represents the AI Safety Initiative³, a community of technical and policy researchers at Georgia Institute of Technology and Georgia Tech Research Institute interested in mitigating risks to American and human interests from advanced artificial intelligence.

Frontier AI models represent a critical national security asset requiring immediate action to ensure America's economic competitiveness and geopolitical dominance. The recommended dual approach involves classifying frontier models as vital security assets while fostering commercial applications that drive innovation and economic growth. Success requires substantial investment in infrastructure, intelligence networks, and public-private partnerships, alongside comprehensive education and workforce development initiatives. In our response, we split AI models into two main classes - (1) frontier AI models and (2) their commercial applications - and provide four classes of recommendations:

(1) National Security, Defense, and Research Dominance: Secure frontier AI development through establishing classified partnerships between AI labs and DoD/IC government agencies, implement multi-layered defense strategies, secure human capital through

¹ Corresponding author. [REDACTED] [REDACTED]

² Parv organized, led, and managed the authoring group, and wrote the introduction, among other sections. All other authors are ordered alphabetically.

AI Safety Initiative at Georgia Tech: Parv, Cyra, Jaehun, Bratee, Seth, Andrew, Dwayne

School of Public Policy, Georgia Institute of Technology: Adelaida, Emalee, Merritt

School of Cybersecurity and Privacy, Georgia Institute of Technology: Kishor

³ AI Safety Initiative. (2025). *Research*. <https://www.aisi.dev/research>

competitive retention programs, prioritize explainable AI over AGI, and strengthen IP protection and procurement controls.

(2) Technical Infrastructure and Model Development: Develop resilient model evaluation tools, expedite domestic hardware manufacturing through the CHIPS Act, secure critical mineral supply chains beyond China, and invest in energy-efficient computation technologies.

(3) Innovation and Education: Support industry research on AI adoption best practices, develop comprehensive workforce training programs, accelerate government AI adoption, and build AI literacy across K-12 education.

(4) Balancing Autonomy and Safety: Enforce strict data handling policies, establish NIST as the primary regulatory body for commercial AI, incorporate Probabilistic Risk Assessment methodologies, and support market-driven solutions like incident reporting systems.

Introduction

AI dominance should be a vital national security priority. Through this lens, we provide concrete US government policy recommendations and examine the implications of these policies. Our recommendations encompass defense applications, infrastructure development, innovation strategy, education policy, and research funding for sustaining American leadership and strengthening the US economy in the years to come.

Our discussions are broadly split into two focus areas: (1) *frontier AI*, or highly capable foundation models with never-before-seen capabilities, such as GPT 4o; and (2) *commercial applications* of these frontier models which interface with consumers, businesses, or government partners directly, such as ChatGPT. The US government should rapidly promote an **open, economically competitive environment for commercial applications**, while treating the most sensitive components of **frontier AI as critical national security assets**.

The price of ineffective action is high: Extremely transformative AI is likely to come during this presidential term, and, unless managed responsibly, will result in acute negative consequences for American interests globally across almost every sector. We acknowledge the risks involved and provide recommendations to mitigate them in four (4) sections:

(1) National Security, Defense, and Research Dominance: *Why are frontier models so vital to national security, and what can be done to protect American dominance?*

(2) Technical Infrastructure and Model Development: *What resources and infrastructure does frontier AI dominance require, and how can the United States rapidly protect these resources?*

(3) Innovation and Education: *What educational and economic systems are necessary to ensure commercial AI application dominance, and how can these be innovatively promoted?*

(4) Balancing Autonomy and Safety: *What are the risks of integrating commercial AI applications in businesses and governments, and how can these risks be managed responsibly?*

(1) National Security, Defense, and Research Dominance

Why frontier AI is a national security asset: Advanced AI systems can perform highly complex tasks, such as analyzing vast datasets, automating decision-making, and even simulating scenarios for defense and intelligence purposes, with efficiency and quality rivaling human experts, and will continue to gain in intelligence, likely *exceeding* human experts sometime this decade⁴. Nations that develop and control these human-exceeding frontier models will gain an insurmountable edge in everything from cybersecurity and military strategy to economic forecasting. For these reasons, **frontier AI models should be treated with the same level of classification as advanced nuclear research and intelligence-gathering systems.**

Implications of underinvestment: Inaction and underinvestment in AI will have devastating economic, military, and geopolitical consequences as America risks losing its global AI dominance. If geopolitical counterparts such as China and the EU begin to outpace US

⁴ Grace, K., Stewart, H., Sandkühler, J. F., Thomas, S., Weinstein-Raun, B., & Brauner, J. (2024, January 5). *Thousands of AI Authors on the Future of AI*. ArXiv.org. <https://arxiv.org/abs/2401.02843>

investment and innovation in AI-driven national defense, the US military will be technologically outmatched and unable to withstand catastrophic security threats by adversaries. With other nations controlling the most advanced AI systems and critical electronic components, the US would find itself increasingly dependent on adversaries such as China, heightening its vulnerability to cyber warfare and strategic coercion⁵.

Emerging threats to frontier AI models: American frontier AI models face constant threats from nation-state adversaries, particularly China, who aim to steal model weights, corrupt training pipelines, or embed malicious code at the supply-chain level. These sophisticated threats extend beyond traditional cyberattacks, employing covert surveillance of research facilities, exploitation of open-source libraries, and injection attacks designed to manipulate system outputs. **As frontier AI capabilities increase, these models become even more valuable targets**, attracting increasingly advanced and persistent threats. If not proactively addressed, foreign adversaries will erode stakeholder trust in our systems and technological dominance.

Mitigating threats: To safeguard America's frontier AI capabilities against sophisticated threats, the federal government must implement a coordinated, multi-layered defense strategy spanning the entire AI lifecycle. Since cutting-edge models predominantly reside with private corporations, we recommend establishing formalized partnerships between leading AI laboratories and key government agencies, **with frontier model research and data classified at the SECRET/TS level**. This classification necessitates streamlined security clearance processes for qualified researchers while maintaining rigorous standards.

Within this classified framework, we propose a division of agency responsibilities: Cybersecurity and Infrastructure Security Agency (CISA) should lead the implementation of Zero Trust architecture across national AI research facilities, ensuring continuous verification of all access requests with advanced cyber-physical protection measures implemented at data centers housing frontier AI systems. Concurrently, Defense Advanced Research Projects Agency

⁵ Special Competitive Studies Project. (2022). *Mid-Decade Challenges to National Competitiveness*. <https://www.scsp.ai/wp-content/uploads/2022/09/SCSP-Mid-Decade-Challenges-to-National-Competitiveness.pdf>

(DARPA) and Intelligence Advanced Research Projects Activity (IARPA) should allocate funding for comprehensive adversarial testing and vulnerability assessments, utilizing established frameworks such as MITRE ATLAS⁶ and OWASP⁷ to systematically identify and remediate security vulnerabilities.

For enhanced threat detection capabilities, CISA should develop and deploy advanced anomaly-monitoring systems throughout federal and contractor-operated facilities, creating an early warning network for potential intrusions. By integrating these security measures under federal coordination and aligning them with existing intelligence capabilities, the United States can establish a dynamic, adaptive defense posture that evolves alongside emerging AI risks, thereby preserving America's technological advantage in this critical domain.

Securing human capital: US dominance in artificial intelligence depends on advanced technologies and the skilled professionals behind that innovation. Yet, a major concern is the potential loss of talent when researchers and engineers are drawn to foreign competitors or rival companies⁸. To counter this risk, **the federal government should work with public universities and University-Affiliated Research Centers to develop strong AI workforce retention programs.** These programs should include competitive grants, fellowships, and clearly defined career paths designed to attract and keep top experts. In addition, the State Department should simplify visa and immigration processes to help recruit leading AI professionals from allied nations. At the same time, Intelligence Community (IC), such as the National Security Agency (NSA) and Federal Bureau of Investigation (FBI), must enhance personnel evaluations and insider-threat detection to prevent any adversarial impact. By consistently investing in skill

⁶ MITRE. (2024). Atlas Matrix. <https://atlas.mitre.org/matrices/ATLAS>

⁷ Sotiropoulos, J., Rosario, R. F. D., & Rosario, S. C. (2024, December 16). Announcing the OWASP LLM and Gen AI Security Project Initiative for securing agentic applications. OWASP Top 10 for LLM & Generative AI Security. <https://genai.owasp.org/2024/12/15/announcing-the-owasp-llm-and-gen-ai-security-project-initiative-for-securing-agentic-applications/>

⁸ Marks, D., & Frenkel, O. (2025, February 11). *America's Technology Lead Depends on High-Skilled Workers*. American Leadership Initiative. <https://www.american-leadership.org/issues/americas-techology-lead-depends-on-highskilled-workers-7yytz>

development, professional growth, and international partnerships, the federal government can build an attractive pipeline of AI talent. This strategy not only reduces the risk of intellectual capital draining to non-ally entities but also supports a secure and innovative ecosystem that benefits US agencies, private enterprises, and the overall national interest.

Explainability, not AGI: For the most strategic pathway forward in AI development, the US should steer clear of initiating a single, massive project focused on AGI development and pursue a more holistic framework. An aggressive, monolithic push towards developing superintelligent AI could destabilize global relations and provoke similar initiatives by our adversaries, particularly nations like China. Instead, fostering innovation through public-private partnerships can promote responsible AI advancement without escalating geopolitical risks. Additionally, the inner workings of all frontier models are currently opaque, reducing their reliability in defense and intelligence applications. Enhancing Explainable AI (xAI) is crucial for building trust and facilitating informed decision-making, as it focuses on developing models that provide clear, understandable explanations for their outputs. **We encourage research investment from DoD/IC stakeholders, specifically into mechanistic and non-mechanistic interpretability of frontier AI models.**

Intellectual property and procurement controls: To protect intellectual property and optimize procurement in AI, the government should adopt strategic measures that target both security and efficiency. This includes tightening counterintelligence and law enforcement in coordination with LEOs to combat state-sponsored theft of AI technologies, prosecuting espionage attempts, imposing trade sanctions on recipients of stolen IP, and enhancing cybersecurity for defense contractors and AI firms⁹. **Additionally, classified research programs will help safeguard critical innovations by ensuring sensitive projects are conducted in secure facilities with strict access controls.** On the procurement front, streamlining acquisition processes by employing agile, tech-friendly approaches, such as using fast-track acquisition

⁹Godefroy, J. (2024, March 7). *How Does Artificial Intelligence Affect Intellectual Property Protection?* Rouse.com. <https://rouse.com/insights/news/2024/how-does-artificial-intelligence-affect-intellectual-property-protection>

models and Other Transaction Authority (OTA) agreements, can reduce development timelines and facilitate rapid integration of AI solutions. Enhanced unclassified public–private channels for less sensitive national security applications, including the establishment of AI test beds, innovation hubs, and exchange programs between industry and federal agencies, will further ensure that cutting-edge AI technology is effectively deployed to meet agency needs.

Opportunities for stronger export controls: Global governance of AI demands a balanced approach that promotes allied innovation while strictly controlling the export of sensitive technologies; US investments in chips manufacturing represent a unique opportunity to effectively enforce export regulation through technical controls. **We encourage work from the Bureau of Industry and Security to study and test geo-tagging American exports to Tier 1/2 Countries**, and recommend that National Institute of Standards and Technology (NIST) codify geo-tagging standards, which will make it efficient to identify any attempts of refurbishment or smuggling. These protections may be complemented with end-use agreements to bestow legal onus for smuggling on responsible parties.

(2) Technical Infrastructure and Model Development

The need for evaluations: As emerging race dynamics between AI labs have shown, frontier AI advancements are extremely rapid, even outpacing the capacity of traditional means to share research effectively. We expect this to continue most acutely around evaluations of top AI models. For this reason, the Administration must **immediately develop resilient model evaluation tools and benchmarks in key fields of national security**, including CBRN threats, intelligence-gathering, cybersecurity, and strategic planning. These efforts must be shared and developed across DoD/IC agencies, allowing for accurate assessment of US model capabilities and data-backed model alignment with US interests.

Hardware as a key resource: The recently enacted CHIPS and Science Act of 2022 initially served to incentivize domestic chip manufacturing by providing \$53 billion to US manufacturing facilities; however, due to bureaucratic delays, the federal government was slow

to allocate grants. Many companies did not receive the funding that they were promised; therefore, in order to promote timely critical hardware manufacturing and development, legislators should aim to expedite bureaucratic grant-delivering processes. By facilitating efficient public-private collaboration, the US will accelerate its path toward increased domestic hardware protection and become less reliant on other countries, thereby strengthening national security. **Establishing data centers in allied countries is essential to address domestic constraints on compute resources.** We recommend the US establish partnerships to maintain data centers in allied states while ensuring the security and integrity of these facilities.

In order to ensure long-term sustainability and efficiency, the US must secure its supply chain of critical minerals for AI development. Currently, the US relies heavily upon China for minerals that are crucial for semiconductor production, such as germanium, gallium, palladium, and polysilicon¹⁰. This allows China to exert control over the American AI sector through restricting exports, as was done in December 2024¹¹. Any disruptions to the AI supply chain could compromise national security. Therefore, the US should bolster strategic international relationships with countries such as Australia, Canada, South Africa, Namibia, and various South American nations—all of whom have access to rare earth elements and other equally important minerals. Additionally, **policymakers should develop bilateral trade agreements for other countries to export their mineral resources to the US.** In doing so, the US can enhance production efficiency by minimizing supply chain vulnerabilities, ultimately fortifying national security.

Efficiency increases: Emerging technologies and operational practices are offering promising pathways to mitigate the energy intensity of AI infrastructure. AI-driven optimization systems are themselves becoming critical tools for improving data center efficiency, as they can

¹⁰ Menell, B. (2025, March 3). China's critical mineral monopoly threatens U.S. national security. The Washington Times.<https://www.washingtontimes.com/news/2025/mar/3/chinas-critical-mineral-monopoly-threatens-us-national-security/>

¹¹ Ministry of Commerce notice 2024 no. 46: Notice concerning strengthening controls on exports of relevant dual-use items to the United States. (n.d.). Center for Security and Emerging Technology. Retrieved March 15, 2025, from <https://cset.georgetown.edu/publication/china-rare-earth-export-ban/>

monitor, analyze, and optimize energy usage in real time, precisely identifying inefficiencies and reducing wasted power^{12 13}. At the hardware level, fundamental innovations in chip architecture are yielding significant efficiency gains. In order to incentivize the production, distribution, and usage of these emerging technologies, **the NSF should allocate additional funds toward research institutions as well as prominent software/hardware companies for the study of energy-efficient computation.**

(3) Innovation and Education

It is not enough for the United States to host and protect frontier models - the Administration should actively support an open environment for their commercial applications.

Support industry research on best practices and adoption methods: As AI is able to automate more jobs, productivity advantages will accrue disproportionately to AI-proficient individuals and organizations, underscoring the importance of effective AI adoption strategies for American companies to stay competitive¹⁴. Key research to be federally funded through the Small Business Administration should include: (1) comprehensive impact assessments across sectors, (2) evidence-based AI utilization frameworks, (3) research on effective teaching methods for building both practical skills and conceptual understanding of AI (including its limitations, risks, and benefits), (4) developmental frameworks examining AI's effects on youth cognitive and social development, and (5) frameworks for Human-AI partnerships that amplify human strengths and creativity rather than attempting to substitute for human capabilities.

Workforce AI training and development: A workforce prepared for AI integration demands both technical proficiency and enhanced human-centric skills. In government, AI training needs to be a standard requirement for both new and existing employees to ensure a

¹² Singh, A., & Martinez, C. (2024). AI-driven optimization for sustainable data centers. *Journal of Green Computing*, 5(2), 112-128. <https://doi.org/10.1145/3542895.3542901>

¹³ Meta AI Research. (2023). Machine learning for data center efficiency (Research bulletin 2023-07). <https://ai.meta.com/research/publications/ml-for-data-centers>

¹⁴ Marquis, Yewande & Oladoyinbo, Tunbosun & Olabanji, Samuel & Olaniyi, Oluwaseun & Ajayi, Samson. (2024). Proliferation of AI Tools: A Multifaceted Evaluation of User Perceptions and Emerging Trend. *Asian Journal of Advanced Research and Reports*. 18. 30-55. 10.9734/AJARR/2024/v18i1596.

workforce proficient in AI tools. In the public sector, **the Department of Labor, in partnership with research institutions, should conduct thorough studies on occupations at risk of AI disruption.** This research should guide the development of targeted training programs that build on existing worker expertise while also enhancing digital skills through online, industry-aligned micro-credentials and college courses. Additionally, the Department of Commerce should establish a national skills database to track evolving job requirements across sectors, enabling educational institutions and individuals to adapt rapidly to market demands.

Private sector applications: AI adoption in highly controlled industries like healthcare, finance, and energy should be guided by clear standards set by specific regulatory bodies, such as the FDA for healthcare, the SEC for finance, and the FERC for energy. The use of regulatory sandboxes, as seen in the UK's Financial Conduct Authority¹⁵ model, would allow these sectors to safely test AI systems under supervision, ensuring their effectiveness without compromising safety or integrity. Key applications—such as AI for diagnostic analysis in healthcare, financial reporting in finance, and predictive maintenance in energy grids—should be selectively scrutinized by these agencies to maintain industry standards while promoting innovation.

To further encourage AI innovation, the Small Business Administration (SBA) should provide targeted subsidies and clear guidelines for cloud computing and AI usage to small and medium-sized enterprises. These incentives could include grants, tax breaks, and educational resources on integrating AI into business operations, helping smaller firms more quickly adapt to technological advancements. This would help bridge the gap between large and small businesses, fostering innovation across a wider range of industries while ensuring that ethical AI practices are upheld.

Bureaucratic adoption of AI: To maximize AI's impact in government, agencies must take strategic steps to enhance efficiency, decision-making, and citizen engagement. Government agencies should expand AI-powered public interfaces—including chatbots and virtual

¹⁵ Financial Conduct Authority. (2024). *AI Update*. <https://www.fca.org.uk/publication/corporate/ai-update.pdf>

assistants—to improve citizen engagement and simplify access to public services. AI-driven decision-making tools - building on GSaI¹⁶ - should be embedded in government operations to enable real-time data analysis, refine policy outcomes, and enhance democratic governance. **We recommend the OPM prioritize AI-driven automation of routine administrative tasks to reduce inefficiencies and free up public servants for higher-value work throughout the federal government.** Automating data entry, scheduling, and document processing will improve operational efficiency and service delivery to the American people.

Build AI capabilities among youth: AI literacy should be integrated across K-12 curricula, merging technical skills, responsible use, and critical evaluation of AI systems. Building on the foundation set by the 18 states that mandate media literacy classes, efforts should expand to include AI competencies.¹⁷ **We recommend establishing an AI Education Coordination Office within OSTP.** This board should consist of educators, AI researchers, ethicists, and industry representatives responsible for developing comprehensive AI literacy standards. Given AI's rapid evolution, the advisory board should conduct quarterly curriculum relevance reviews, feeding into a structured curriculum update every 18 months—significantly faster than traditional cycles. Additionally, OSTP should oversee a specialized rapid response team responsible for promptly creating supplemental educational materials addressing major AI breakthroughs. These materials may be made available through a digital repository, continuously accessible to schools and students nationwide, and tailored to different audience ages. Assessment methods should transition to portfolio-based evaluations, enabling students to demonstrate AI competencies practically, highlighting technical, nuanced understanding, and reasoning skills on AI.

¹⁶ Kelly, M., & Schiffer, Z. (2025, March 7). *DOGE Has Deployed Its GSaI Custom Chatbot for 1,500 Federal Workers*. WIRED. <https://www.wired.com/story/gsa-chatbot-1500-federal-workers/>

¹⁷ Justin Klawans. (2024, April 2). The push for media literacy in education amid the rise of AI. The Week. <https://theweek.com/tech/media-literacy-AI-schools>

(4) Balancing Autonomy and Safety

Data Handling and Protection: Government agencies must enforce strict policies on data usage and cross-border data sharing that clearly define permissible data flows, limit inference risks, and comply with privacy regulations, while emergency data-erasure protocols can be used to swiftly remove compromised information in the event of a breach^{18 19}. By aligning these measures with international standards and rigorous screening processes, the federal government can safeguard data integrity and privacy while enabling secure AI operations throughout their lifecycle.

Risk assessments & regulatory Frameworks: Establishing effective risk assessments and regulatory frameworks is essential to ensuring that AI technologies are developed and deployed responsibly without stifling innovation. Existing regulatory bodies, such as the Federal Trade Commission (FTC), have already taken steps to address AI-related risks, but regulations should be streamlined further to avoid overburdening developers, especially smaller companies and startups, which may struggle with compliance costs.²⁰ **We recommend that NIST take a leading role in creating and overseeing commercial AI regulations, acting as the primary body for monitoring AI development and working with the FTC to ensure that companies are held accountable for their commercial AI applications.** Additionally, the OECD's Principles on Artificial Intelligence²¹, which were established under the previous Trump administration, provide valuable guidance for crafting regulations that promote innovation while protecting public safety and ensuring fairness. Probabilistic Risk Assessment (PRA)²² is one widely accepted methodology that leverages data modeling to evaluate the likelihood of risks

¹⁸ Pugh, B., & Ward, S. (2024). *International Association of Privacy Professionals*. iapp.org.

<https://iapp.org/news/a/what-does-ai-need-a-comprehensive-federal-data-privacy-and-security-law>

¹⁹ Zafir-Fortuna, G. (2025). *What to Expect in Global Privacy in 2025*. Future of Privacy Forum.

<https://fpf.org/blog/what-to-expect-in-global-privacy-in-2025/>

²⁰ Staff in the Office of Technology and the Division of Advertising Practices. (2024, December 31). *AI and the Risk of Consumer Harm*. Federal Trade Commission.

<https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2025/01/ai-risk-consumer-harm>

²¹ OECD. (2019). *AI principles*. OECD. <https://www.oecd.org/en/topics/sub-issues/ai-principles.html>

²² NRC. (2020, July 7). *Probabilistic Risk Assessment (PRA)*. NRC Web.

<https://www.nrc.gov/about-nrc/regulatory/risk-informed/pr.html>

occurring and their potential impact. In the context of AI, PRA could address issues like algorithmic bias, security vulnerabilities, and ethical concerns early in the development process. To further enhance the responsible development of AI, we recommend prioritizing the integration of PRA into FTC AI policy frameworks.

In addition to structured risk assessments, market-driven solutions like AI incident reporting systems are effective tools for identifying and addressing emerging threats. Voluntary initiatives like the AI Incident Database (AIID)²³ serve as independent fact-checking resources that document real-world AI failures, fostering transparency and promoting industry-wide learning. To strengthen the reliability of these systems, we recommend efforts to improve their data quality and standardization, recognizing that while there are challenges²⁴ encouraging the development of similar initiatives can still play a valuable role in promoting accountability without the need for heavy-handed regulations. By supporting the FTC's leadership and building on these frameworks, the US can foster a balanced regulatory environment that encourages growth while safeguarding societal interests.

Conclusion

Frontier AI models represent a critical national security asset requiring immediate action—they will determine America's economic competitiveness and geopolitical stability in the coming years. We advocate a dual approach: classifying and securing frontier models as vital national security assets while simultaneously fostering commercial applications that drive innovation and economic growth. This strategy demands substantial investment in physical infrastructure, robust intelligence-sharing networks, and public-private partnerships with frontier AI labs.

The Administration must mobilize the full DoD/IC apparatus to mitigate security risks while integrating AI capabilities across defense applications, bureaucratic systems, and industry

²³ AI Incident Database. (2025). *Artificial Intelligence Incident Database*. Incidentdatabase.ai. <https://incidentdatabase.ai/>

²⁴ Stanley, J. C., & Dorton, S. L. (2023). Exploring Trust With the AI Incident Database. *Proceedings of the Human Factors and Ergonomics Society Annual Meeting*, 67(1), 489–494. <https://doi.org/10.1177/21695067231198084>

sectors. The US cannot neglect its citizens in its journey to AI dominance - we must establish educational programs for youth, hone the technical workforce, and protect existing human capital. The technological revolution before us offers unprecedented opportunity, but only if America acts decisively now—because in the race for AI dominance, there are no silver medals.

Acknowledgements

We especially thank, in no particular order, Olajide Olugbade and Sue Bae from the School of Public Policy at the Georgia Institute of Technology for high-level mentorship and feedback on an earlier draft of this submission.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

The arguments expressed and policies recommended in this document do not necessarily reflect the views of the Georgia Institute of Technology, a public R1 research university, or the Georgia Tech Research Institute, the nonprofit, applied research organization of the Georgia Institute of Technology.