

March , 2025

AI Action Plan
Attn: Faisal D'Souza
NCO
2415 Eisenhower Avenue
Alexandria, VA 22314

Subject: AI Action Plan Request for Information 2025-02305 (90 FR 9098)

Dear Sir

We are writing in response to the RFI published on behalf of the Office of Science and Technology Policy (OSTP) on the Development of an Artificial Intelligence (AI) Action Plan. This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Alignment AI is a company focused on the governance, risk management, measurement, and transparency for the use of AI in healthcare. Alignment AI works with both healthcare AI developers, and healthcare AI procurers/implementors/users on matters including certification, accreditation, governance development, risk assessment, and measurement of AI.

We offer comments on the following topics raised by the RFI.

General Comments

Healthcare offers a more fraught environment for the use of AI than many industry sectors. It offers not only the challenges of trust, privacy of personal information, validity, appropriateness, and fairness of AI present in most, but it has additional profound concerns that accrue to patient safety. In many healthcare use cases, the role of AI is to support clinical decision makers often in critical or emergency care situations. Healthcare has used clinical decision support algorithms for decades based on more predefined reference data-driven models that have been developed based on empirically researched protocols and clinical trials. Unlike these models, the use of AI that can learn and adapt asks for new approaches to be taken to ensure they stay true to the intended purpose of use, and adaptations based on new learning are treated with appropriate rigor to assure the AI continues to perform in manners that are valid, safe, and accurate. Our concern is especially to argue for the adoption of robust governance models predicated on a risk-based approach to ensure that AI can be used safely with sufficient transparency to enable implementors and users to understand how they have been designed, tested, and deployed.

Use

AI represents an evolution of the use of decision support algorithms in healthcare that have long been used to represent the application of empirical best practices for supporting clinical decision making. Healthcare information technology (HIT) developers have been working to respond to the need for reliable, transparent, effective, and safe AI applications to support use cases for clinical, financial, and operational needs. AI procurers/implementors need to be able to understand how the AI they place into use has been designed, developed, trained, and tested. They need to understand as the AI models evolve how they do so with fidelity to intended use, and that modifications are subject to the same rigor as original development. Users need to be able to have transparency into how the AI algorithms function. There is not a current par level of such requirements across all AI use cases within healthcare. Instead, there is partial coverage offered by regulation between the [guidance](#) published by the Food and Drug Administration (FDA) for software as a medical device, and the Office of the National Coordinator for Health Information Technology (ONC) in its HTI-1 final rulemaking which adopted a [set of requirements](#) for HIT and AI developers seeking certification with the ONC. However, HIT of other kinds such as for use in revenue cycle processes and operational processes or by payers and providers who are not required to use Certified EHR Technology (CEHRT) may be using AI not bound by such requirements.

The healthcare vertical needs a common baseline approach to provide transparency into how the AI was designed, developed, trained, and tested, and what kind of risk management approach was taken to support assurance of performance. Implementors and users need to be able to have the same expectation for transparency no matter the type of use case at hand, and regardless of whether current regulation demands it of the developer. Whether this is through a federally driven regulatory approach or through development of industry best practices that form a code of conduct for governance based on standards such as [ISO 42001:2023](#), we believe that such requirements should be supported by this Administration for healthcare to set a baseline for AI developers and adopters alike to expect as a minimum bar to support insight into the transparency of the AI that is procured and used.

Explainability and Assurance of AI Development and Deployment

Significant to the matter of providing transparency into healthcare AI is how it supports and stays accurate to its purpose of use. The current ONC regulations require that certified HIT developers identify any source data attributes that serve as inputs into a given algorithm's processing, and source reference information into the logic of the algorithm. The requirements also require information to be shared by the developer as to how the AI was trained, what kind of considerations went into the design of the training data set, how the AI was tested, and how representativeness was preserved through the AI's development to assure it held fidelity to its purpose. Healthcare is unique in that the concerns for patient safety must always be met, which introduces a real human impact that may not be present in a lot of other industries. The reliability, validity, and efficacy of AI in terms of the impact on patient safety is top of mind. This means that the design and performance of the AI must be understood by procurers and implementors to make good decisions about introducing any given algorithm into use in a safe, reliable manner. The risks to safe and effective AI use are too great not to have strong governance and measurement systems in place.

To exemplify the points above we offer a couple illustrative examples. ALIGNMT AI, through programmatic red-teaming, identified a critical flaw in an ambient AI solution that mistakenly transcribed patient consultations as if a procedure had been performed, rather than merely discussed. Similarly, ALIGNMT AI discovered a design pattern in an ambient AI solution that prevented the summarization of objective clinical findings, leading to medical charts missing clinical encounter data required for risk-adjustment.

Mitigating risks in healthcare requires a strong, risk-based governance framework for both AI developers and implementers. This recommendation is an element of both the FDA and the ONC regulatory approaches for AI for the medical devices and software products for which they have regulatory purview. For the developer, the interest is in assuring that the AI remains of fidelity to its purpose of use throughout the development lifecycle, and that what has been done to engage in risk management to assure AI performance is understood by procurers and users. For the implementor, this means that they have established a firm strategic purpose of why they seek to adopt AI, how they are going to measure its performance to assure not only that it supports its purpose of use, but also that any potential risks to patient safety are evaluated on a continuous basis. To establish this, information such as what ONC requires be provided by certified HIT developers must be made available to implementors to design and develop their own AI governance and risk management approaches.

Healthcare providers like integrated health systems already have in place policies, procedures, and practices for governance and risk management in areas such as clinical research, cybersecurity, medical staff credentialing, and similar demands. The use of AI represents another key area of requirement that necessitates transparent inputs, and measurement systems. Implementors should be urged to connect their AI risk management approaches to the same approaches they use for similar needs across the enterprise.

In the real world, AI risk-mitigation platforms like ALIGNMT AI have started to support the FAVES framework—Fairness, Appropriateness, Validity, Effectiveness, and Safety—originally required by the ONC. This framework ensures that AI systems in healthcare are equitable, clinically relevant, scientifically sound, impactful, and secure, aligning with regulatory expectations and best practices for AI management during pilot and post deployment in clinical settings.

We urge the Administration to consider how it can set the expectation with AI developers and procurers/implementors to develop governance and risk management approaches that depend on transparent sharing of information about AI required to build a sound foundation of trust and awareness of AI performance. This should start with continuing to support what has been done under ONC's HTI-1 rule, urging the adoption of a common floor across areas of the industry that may not be subject to federal regulation where an industry driven best practice is in order, and to strongly support the need to develop measurement capabilities to support continuous monitoring of AI performance similar to the FAVES approach adopted by ONC.

Risks

In its framework of requirements for healthcare AI, ONC leveraged the concepts of FAVES (Fair, Appropriate, Valid, Effective, and Safe) to describe the risk areas that are of high concern to healthcare uses of AI. Each of these represents an area of material risk for what may go wrong with AI and are risks that undermine AI's use in healthcare.

By implementing a risk management framework, healthcare enterprises can proactively mitigate AI risks—ensuring AI models are validated, appropriately deployed, and continuously monitored to protect patient safety and promote equitable, effective care. We offer below a few examples illustrating the types of risks that can be detected by the FAVES framework:

- **Fairness** – A predictive readmissions model may systematically underestimate risk for Black and Hispanic patients due to historical disparities in healthcare access, leading to inequitable allocation of preventive resources.
- **Appropriateness** – An AI sepsis detection tool designed for ICU patients might be applied in general hospital wards without validation, triggering false alarms or missed diagnoses in lower-acuity settings.
- **Validity** – A breast cancer screening AI trained primarily on data from Caucasian women may fail to detect cancer as accurately in Asian or African American women, leading to misdiagnoses.
- **Effectiveness** – An AI-powered chatbot for mental health screening may perform well in controlled studies but fail to engage real-world patients effectively, leading to low adoption and missed opportunities for early intervention.
- **Safety** – A robot-assisted surgical system with AI-driven guidance could misinterpret anatomic structures due to poor image quality, increasing the risk of surgical errors.

AI use may be impacted by the risk that the AI was developed given a limited training data set that does not support extension to a patient population that is dissimilar to that represented by the training data set. Without good oversight and understanding to monitor performance in production use, AI may not provide fair outcomes for a patient population that differs in its demographic makeup from that which informed the training database. Bias in the outcome can result.

Similarly, if the AI is adapted for use to solve for clinical conditions that vary from the original purpose of use, or for a patient population that differs in its composition from that which it was trained upon, it may experience “drift” from its purpose of use that renders it inappropriate to be used the way the implementor intends. Without the ability to measure such drift and again lacking in understanding as to the AI's original design and intended purpose, the AI may not be appropriate to be used in the manner attempted.

The AI may also be at risk for reaching incorrect or inaccurate conclusions for the intended purpose. If the production use experience differs because of the types of issues that come from deviating from the intended patient population or from the clinical use cases for which it was designed, inaccurate outcomes could result from hallucination of the AI.

Any of these issues can render the AI to be ineffective for its intended use. If it is being used for a purpose deviating from its intended use, for a patient population not

represented by its training, or using inputs it has not been trained to encounter, the effectiveness of the AI may be undermined.

Finally, the accumulative effect of any of these risk areas endangers patient safety or introduces potential for detrimental effects on those individuals who are subjects of its processing.

Without measurement systems that are designed to monitor AI performance against risk areas represented by the FAVES when those risks are inadequately defined, the realization of those risks in adverse ways is likely and especially problematic given the overriding concern for patient safety.

We urge the Administration to take the lead in emphasizing the significance in designing governance, risk management approaches, and measurement of AI performance informed by design principles such as the FAVES to instill confidence and assurance into the use of AI. We urge a harmonization of approach between the FDA and ONC, and a collaboration with the healthcare industry to extend these principles and requirements into industry best practices such that no matter the intended purpose of the AI, whether clinical, operational, administrative or financial, there is a par level of requirement for supporting sound AI risk management.

Data Privacy and Security

Corollary to the risk areas represented by the FAVES are risks that accrue to the privacy and security of the data sets used to train and test AI as well as the integrity of the AI algorithms themselves. The Administration should support efforts to require the use of de-identified data sets for training and testing purposes where such data sets may be developed derived from identifiable personal information. For healthcare, the Administration should develop guidance outlining how HIPAA Security and Privacy rules apply to the systems and data sets that underlie AI development, deployment, and use. Beyond the risks that may accrue from handling identifiable personal electronic health information, there are risks to adulteration of the AI algorithm itself by hostile actors using malware or other cyberattack methods as well as denial of service attacks against environments that support AI provision as software as a service. In the new [proposed HIPAA Security Rule](#), HHS proposes that covered entities must test their cybersecurity methods such as encryption on a regular basis to assure contemporary capabilities are maintained to ensure continued effectiveness for the purpose intended. We urge the Administration to support similar approaches for the security of healthcare AI especially should this rule be made final and include proposed requirements such as removing the addressable nature of encryption. This may include providing guidance to reinforce with covered entities and business associates that the requirements of any final Security Rule from the current rulemaking extends to healthcare AI as “technology assets” as defined in the proposed rule. This would help raise the significance and clarity of how AI is included in the inventory of technology assets as systems that are material to HIPAA Security Rule compliance. Healthcare AI bears a high degree of relevance to the “confidentiality, integrity, and availability” of electronic personal health information.

Regulation

We have observed that there is a need for a sound consistent set of regulating principles for AI use in healthcare, to help enterprises who are now struggling to comply with a myriad of rules emerging at the state level. We are not suggesting that it must be achieved solely through federal regulation. However, we argue for preserving the current regulatory framework established by the ONC for AI use in decision support. We also argue for this Administration to take the lead in strongly encouraging industry best practices built on the approaches taken by the ONC to inform industry codes of conduct for developers, and by international standards such as the ISO standard referenced earlier on the matter of building governance frameworks for the oversight of AI. Further, we urge the Administration to consider how these can extend to the procurer/implementor of healthcare AI to help inform a framework for how they may in turn develop sound governance practices consistent with industry standards. The Administration should require these kinds of practices to be in evidence in federal contracting for AI across the board, and specifically to healthcare, adopt as requirements the need for contractors to prove they have in place robust systems of governance and measurement of AI performance.

ISO standards can play a key role in guiding AI procurement in healthcare and use without requiring additional federal regulation. ISO frameworks, such as ISO 42001 (AI Management Systems) and ISO 23894 (AI Trustworthiness), offer structured, internationally recognized methodologies for assessing and mitigating AI risks—making them valuable tools for healthcare enterprises and AI developers alike.

Furthermore, positioning ONC as the leader in AI oversight for the healthcare industry, rather than the FDA, aligns well with ONC's broader scope of AI in healthcare beyond just medical devices. The ONC's approach to algorithmic transparency, data quality, and real-world performance validation is more adaptable for AI-driven decision support tools, clinical workflows, and administrative automation. This could foster an AI innovation-friendly ecosystem with clear but flexible governance standards rather than strict regulatory mandates.

Outside of federal-sponsored certification programs, AI risk-adjustment companies like ALIGNMT AI work with the industry on voluntary certification of AI best practices. For example, ALIGNMT AI partners with Drummond Group (a certification agency) to help AI developers certify to an independent benchmark for AI governance best practices. While this helps AI developers demonstrate compliance with AI best practices and standards, it also helps showcase to the implementors of AI the reliability of the solution they are looking to purchase.

Governance

We have recommended that the Administration support robust governance practices by AI developers and implementors alike, and to consider using its authorities to require this from federal contractors and agencies. For healthcare, we have urged consideration of the program elements ONC has incorporated into its certification requirements for decision support interventions or to consider promoting a par level industry best practice that can be based on recognized standards such as ISO 42001:2023. Further, the HIPAA Security Rule requires security risk assessment of all information systems considered in scope of its compliance requirements. This also is a risk-driven exercise that complements the

approach that should be central to any governance effort. We urge that any industry code of conduct for AI development and use include measurement systems for accountability against the principles of the FAVES we discussed earlier. The main need is for a par-level governance approach across all industry use cases for AI regardless of the presence of federal regulation mandating its existence.

We appreciate the opportunity to offer these comments, and we are happy to offer our expertise to assist efforts in any way to help the Administration understand healthcare segment specific considerations further regarding the governance and measurement of AI deployment and performance.

Sincerely,

Andreea Bodnari

Chief Executive Officer

ALIGNMT AI