

Policy Recommendations for Zero-Knowledge Proofs (ZKP)

in AI:

A Response to the NITRD AI Action Plan RFI

Author: Am Asif Haider
Date: March 15th 2025
Cage Code: 90SD8
UEI: TB4NVKFNHKZ4
Multiple Contract since 2015

Submission Contact:

Asif Haider

SARZAD LLC

[REDACTED]

Submitted to:

[REDACTED]

Public Disclosure Statement:

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Introduction & Statement of Intent

Artificial Intelligence (AI) stands at a critical juncture where continued innovation must be balanced with robust privacy, security, and transparency measures. The Executive Order emphasizes the need to **remove barriers to AI innovation** while maintaining the **highest standards of data privacy and security**. This document addresses those objectives by proposing **Zero-Knowledge Proofs (ZKPs)** as a **lightweight, scalable** security and privacy framework that enhances AI without imposing the heavy compliance and regulatory burdens often associated with traditional oversight models.

The National Science and Technology Council (NSTC) Networking and Information Technology Research and Development (NITRD) Program's AI Action Plan underscores key areas:

- **Data Privacy:** Safeguarding sensitive information.
- **Security:** Protecting AI systems against threats.
- **Efficiency:** Ensuring large-scale AI operations remain cost-effective.
- **Regulatory Harmonization:** Encouraging solutions that **do not** unnecessarily hinder innovation or create duplicative oversight.

ZKPs allow one party to prove a statement about data **without revealing** the underlying data itself. Originally developed within cryptography, ZKPs are well-suited to **verifying AI model performance, compliance, and integrity** in a manner that preserves **confidentiality** and **minimizes compliance overhead**.

Statement of Intent:

This policy proposal aims to show how **federal agencies, industry, and academia** can leverage ZKPs in AI development, deployment, and governance. We propose **non-regulatory, incentive-based** approaches to encourage ZKP adoption—aligning

with the **Executive Order's goal** of removing barriers to innovation while safeguarding national leadership in AI.

AI Privacy & Security Without Regulatory Burdens

As AI systems scale, so do concerns over data privacy and system integrity. Whether it is GDPR in Europe or various data protection laws in the United States, organizations often face a **patchwork** of regulatory requirements. This complexity can **slow innovation**:

- **Excessive Audits & Compliance:** Traditional compliance frameworks often demand **detailed disclosures** of datasets and model internals.
- **Interoperability Gaps:** Varying regulations create barriers to deploying AI solutions across borders or industries.
- **Resource Constraints:** Smaller firms find the compliance burden disproportionately heavy, discouraging them from investing in cutting-edge AI.

Zero-Knowledge Proofs are cryptographic protocols that allow a prover to **demonstrate the truth of a statement** to a verifier without revealing *how* that statement is true or exposing **sensitive underlying data**. In the AI context, this capability offers several benefits:

1. **Privacy Assurance:** Organizations can **prove** they handle data according to compliance standards—e.g., HIPAA, FISMA—without revealing patient information or proprietary IP.
2. **Regulatory Streamlining:** Instead of requiring manual audits of raw datasets, regulators could accept ZKP-based certifications. This drastically lowers the compliance burden.

3. **Cross-Border Data Sharing:** ZKPs can **abstract away** jurisdictional complexities by focusing on verifiable compliance claims (e.g., “This model meets EU privacy standards”), which fosters **international AI collaboration**.

Recommendations for Policy

- **Adopt Voluntary, ZKP-Based Compliance Pilots:** Encourage federal agencies to implement pilot programs where **ZKP proofs** can substitute for detailed data-sharing in compliance audits.
- **Promote Public-Private Collaboration:** Industry, academia, and government labs can co-develop **ZKP frameworks** that align with existing regulatory standards like ISO 27001 or NIST Privacy Framework.
- **Tax Incentives for ZKP R&D:** Consider offering **research tax credits** specifically for developing or integrating ZKPs into AI workflows, thereby accelerating adoption.

By placing **ZKPs at the core** of privacy and security strategies, federal policy can **reduce red tape** and **unleash AI innovation**—all while upholding rigorous data protection principles.

ZKP for AI Security and Model Assurance

Modern AI systems face threats on multiple fronts:

- **Data Poisoning:** Malicious actors can insert corrupted or misleading data into training sets, compromising model outputs.
- **Adversarial Attacks:** Sophisticated techniques can manipulate input data (e.g., subtle pixel changes in images) to mislead AI systems.

- **Inference Extraction:** Attackers may query an AI model to reverse-engineer its underlying parameters, exposing proprietary information.

How ZKPs Strengthen Model Integrity

ZKPs address these vulnerabilities by enabling **provable assurances** of AI model integrity:

1. **Model Integrity Proofs:** Developers can create **cryptographic commitments** tied to a model's parameters at deployment. ZKPs then allow users to verify that a model has **not been tampered with** or replaced.
2. **Secure Model Updates:** When a model is updated or retrained, ZKP can confirm that **legitimate updates** were performed according to authorized processes—no hidden modifications allowed.
3. **Ownership and Licensing:** ZKPs can prove a party holds valid rights to a model's weights or architecture without revealing the sensitive details themselves, curbing **IP theft** risks.

Policy Recommendation: Incentivize ZKPs in AI Security Frameworks

A key step toward **voluntary, non-regulatory** adoption is the creation of **incentives** and technical guidance:

- **Guidance Documents:** NIST, in collaboration with industry, can publish **best-practice guidelines** for using ZKPs to fortify AI models.
- **Voluntary Certification:** An official “ZKP-Ready” seal or certification could be introduced, signaling a model meets certain security thresholds **without** forcing regulatory mandates.
- **Procurement Preferences:** Federal agencies could grant *priority* in procurement decisions to AI solutions that use

ZKPs for model assurance, incentivizing broader industry adoption.

Such measures would bolster AI's security posture while **preventing heavy-handed regulation** that could stifle the growth of America's AI ecosystem.

Technology Readiness Levels (TRL)

The proposed AI and ZKP technologies are currently at TRL 4-7, indicating that they are near-commercialization but require additional development and testing. By focusing on technologies that are close to market, this proposal ensures that the AI and ZKP solutions will be deployed rapidly in real-world production environments.

To bring these technologies to full deployment, the project will emphasize pilot testing, scaling, and optimization across multiple manufacturing sites. The goal is to move from initial development and testing at TRL 4 to broader deployment and integration, with a clear roadmap for commercialization and industry-wide adoption.

Energy Efficiency & AI Infrastructure

AI's exponential growth requires massive computation, from **data center operations** to **model training**. This demand raises concerns about energy consumption, carbon footprint, and **associated operational costs**. Many U.S. data centers already account for a **significant percentage** of national energy usage, creating an urgent need for **sustainable solutions**.

Although ZKPs were traditionally considered **computationally expensive**, modern breakthroughs—such as **zk-SNARKs** (Zero-Knowledge Succinct Non-Interactive Arguments of Knowledge)

and **zk-STARKs**—are **dramatically reducing** computational overhead. When properly integrated, ZKPs can *optimize* AI operations:

1. **Selective Disclosure**: Instead of running full-blown audits or duplicating data, AI systems can produce **ZKP proofs** that reduce the need for repeated computations.
2. **Batch Verification**: ZKP schemes can verify **multiple transactions or model inferences** in one succinct proof, saving energy across large-scale AI services.
3. **Edge Computing**: ZKPs can facilitate **lightweight verification** on edge devices, minimizing back-and-forth communication with data centers.

Policy Recommendation: R&D Investment in ZKP for Sustainable AI

- **Targeted Funding**: Federal R&D grants for **ZKP-based AI systems** that explicitly emphasize **energy-efficient protocols** (e.g., HPC labs collaborating with cryptographic researchers).
- **Infrastructure Upgrades**: Incentivize the **retrofitting** of existing data centers to incorporate hardware-level acceleration for ZKPs, ensuring alignment with the Department of Energy's efficiency standards.
- **Public-Private Partnerships**: Encourage **demo projects** that showcase real-world improvements in **energy usage** and **carbon footprint** when ZKPs are part of AI workflows.

Aligning energy policy with advanced cryptographic techniques helps sustain AI's growth while respecting environmental imperatives.

ZKP in AI Governance & International Collaboration

AI governance frameworks—both domestic and international—struggle with balancing transparency and **corporate confidentiality**. Nations differ in privacy laws, creating **barriers** to cross-border AI deployments. Meanwhile, global

concerns about **AI's impact** on jobs, security, and data sovereignty intensify demands for **common governance standards**.

ZKPs for Trust-Building Across Borders

By design, ZKPs enable stakeholders to **verify** compliance and model claims **without disclosing** sensitive or proprietary details. This capability is vital for building **mutual trust** among nations with varying regulatory and cultural norms:

- **Regulatory Sandboxes:** Countries can collaborate on **ZKP-based "sandbox" environments** where AI tools from multiple jurisdictions are tested under shared verification protocols.
- **Uniform Verification Layer:** A common ZKP-based layer could unify compliance checks for data protection (e.g., verifying no personally identifiable information is mishandled) even if local laws differ.
- **Interoperable Standards:** International bodies such as ISO could adopt **ZKP guidelines** for AI governance, streamlining cross-border data flows and regulatory acceptance.

Policy Recommendation: Integrate ZKP in AI Governance Policies

- **Bilateral & Multilateral Agreements:** Encourage the integration of **ZKP-based auditing** in trade agreements focused on technology and digital goods.
- **Global Forums:** Use organizations like the **OECD, WTO, and G7** to advocate for **cryptographic proof mechanisms** as a standard best practice in AI governance.
- **Funding Mechanisms for Joint Research:** Allocate specific funds within **international science cooperation** budgets to support cross-border ZKP research (e.g., U.S.-E.U. collaborative grants).

Such initiatives help position the United States as **both** a global leader in AI **and** a pioneer in cryptography-driven governance.

AI Workforce Development & Education

Despite AI's growth, a notable **skills gap** persists: not enough professionals are trained in **cryptography, privacy engineering, and secure AI**. This shortage hampers the adoption of advanced privacy solutions like ZKPs.

Integrating ZKP into AI Education

To keep the U.S. AI workforce globally competitive, **federal and state educational policies** must adapt:

1. **Curriculum Updates:** Revise AI-related courses (e.g., computer science, data science, cybersecurity) to **include modules on ZKPs**—covering theory and practical deployment.
2. **Hands-On Labs:** Encourage the development of **open-source** ZKP libraries and lab exercises, allowing students to experiment with **real-world cryptographic code**.
3. **Cross-Disciplinary Training:** Combine AI, cryptography, and **policy** studies, producing graduates who can navigate technical, legal, and governance aspects of AI adoption.

Policy Recommendation: Federal Funding for ZKP Education

- **Grant Programs:** Direct agencies like the NSF to **prioritize grants** for academic programs that expand ZKP courses or labs in AI curricula.
- **Fellowships & Internships:** Establish **joint fellowships** between government agencies (e.g., DARPA) and industry to train students specifically on **privacy-preserving AI**.
- **Professional Certification:** Work with industry associations (e.g., CompTIA, ISC²) to create a **ZKP specialization** certification that meets both commercial and government needs.

By boosting ZKP literacy among AI professionals, we ensure the **sustained growth** of a privacy- and security-focused AI ecosystem.

Conclusion & Call to Action

Summary of Key Recommendations

This document outlines how **Zero-Knowledge Proofs** can serve as a **transformative, non-regulatory tool** to enhance the U.S. AI sector. Our policy proposals include:

1. **Voluntary, ZKP-based privacy and security standards** to streamline compliance.
2. **Incentivized adoption of ZKPs** for AI model integrity and assurance.
3. **Federal R&D support** for ZKP-enhanced energy-efficient AI infrastructures.
4. **Global collaboration** to standardize ZKP usage, fostering international AI trust.
5. **Educational investments** in ZKP to close the AI skills gap.

Call to Action

The federal government, in partnership with industry and academia, should **immediately** begin pilot programs that integrate ZKPs into AI compliance, model security, and infrastructure. Doing so will:

- **Spur Innovation** by eliminating the need for heavy-handed regulatory mandates.
- **Strengthen National Security** by reducing AI vulnerabilities.
- **Position the U.S. as a Global Leader** in both AI and cryptographic innovation.

We encourage agencies responding to the NITRD AI Action Plan to adopt these recommendations, **ensuring the United States remains at the forefront of trusted, privacy-preserving AI.**

Appendices (For Future Request)

- **Appendix A:** Technical Diagram of ZKP-Enabled AI Model Verification
- **Appendix B:** Industry Case Studies (Healthcare, Finance, Defense)

(Note: If additional technical depth is required, these appendices can be developed as supplementary materials without exceeding the main document's page limit.)

Final Note on Submission

This document, tailored as a policy recommendation for the **NITRD AI Action Plan RFI**, remains well within the **15-page limit**. It addresses **key areas** of AI privacy, security, infrastructure efficiency, governance, and workforce development—**all through the lens of ZKPs**.

Thank you for considering these recommendations. We look forward to continued engagement with the federal government, industry leaders, and the broader AI community to realize the potential of **Zero-Knowledge Proofs** as a **cornerstone technology** for privacy-preserving and innovation-friendly AI.

Asif Haider

Asif Haider is a visionary entrepreneur and AI expert who specializes in making complex technological concepts accessible for businesses. As a founding member of the S.M.A.R.T. AI Community and co-founder of AxeGENAI, Asif bridges the gap between AI innovations and real-world applications. His background includes service in the U.S. Air Force and a Master's degree in Negotiation and Conflict Resolution.

With experience spanning logistics, manufacturing, and government contracting, Asif has developed AI solutions that enhance productivity and cybersecurity. As the host of the podcast **AYAYAYAI**, Asif breaks down AI concepts for non-technical audiences and offers actionable insights for industries such as healthcare, retail, and transportation.

Allen Hall

Allen Hall has over 30 years of experience in industrial operations and consulting across sectors such as aerospace, steel fabrication, and medical devices. He holds a Black Belt certification in Lean Six Sigma and has led numerous improvement projects while training thousands of employees in process optimization. An Air Force veteran, Allen holds an MBA from Oklahoma City University and specializes in driving operational excellence through data-driven decision-making.

Jamal Awan

Jamal Awan is an accomplished Information Technology Executive with a focus on IT Service Management, solutions architecture, and digital transformation. As Chief Technology Officer at

YurTech Federal, Jamal leads cloud adoption, automation, and data analytics initiatives for public and commercial clients.

With experience managing federal contracts worth over \$300 million, Jamal has driven IT transformation for clients including NASA, the Navy, and the Department of State. He holds a Master's in Business Administration, a Data Analytics Certification from Cornell, and certifications in PMP and ITIL Expert.
