

From: [REDACTED]
To: [ostp-ai-rfi](#)
Subject: [External] AI Action Plan
Date: Monday, March 17, 2025 9:21:12 AM

CAUTION: This email originated from outside your organization. Exercise caution when opening attachments or clicking links, especially from unknown senders.

Introduction

I'm a software engineer who has observed the transformative impacts of artificial intelligence firsthand. Language models have become increasingly powerful multipliers of developer productivity, at first as learning tools since they're better than internet search engines for many topics, and increasingly as autonomous agents that write code directly. At the same time, however, language models have hobbled the software industry as a whole by making recruitment harder. Employers used to be able to find great engineers by posting jobs online, but now, public job postings are reliably flooded with hundreds or thousands of AI-generated applications optimized to the point of being undifferentiable. Employers are increasingly returning to older, less efficient ways to find candidates. Teams are becoming less productive as they start to feel the effects. As AI programs for writing job applications continue to proliferate and become easier for nontechnical job seekers to use, other industries will become less productive as well. It's too late to stop this trend, but it can serve as an early warning for greater unintended harms that can still be prevented. Artificial intelligence is advancing quickly and its advances cannot be assumed to be beneficial to America.

Background

The AI Action Plan is a government initiative of extreme importance, even compared to the other government initiatives making the news. Its highest priority must be mitigating the downside risks of artificial intelligence.

Analysis

Training transformers directly on text from the internet made them generally knowledgeable. The latest frontier models are specifically trained to reason by writing thoughts before they act, making them much better at self-correction, and Sam Altman and Dario Amodei are promising much larger advances in a few months. I suspect the next wave will be models specifically trained to read and write memories as they think, making them much better at learning from their own experience. It's hard to predict what happens then, but it's safe to say that

increasingly destructive actions will become increasingly accessible as AI capabilities continue to proliferate. Much of society relies on certain types of crime being inherently difficult to commit, an assumption which will break down when anyone can command human-level autonomous agents to work tirelessly on any task. Large organizations, especially those relying on legacy systems, may collapse as autonomous agents efficiently uncover their security vulnerabilities and poison their data. Automated identity theft could cause large banks to lose everyone's accounts with no reliable way to correct the error. Designs provided by AI could make it easy to build counterfeit currency printers. An advanced cyberattack targeting a vehicle manufacturer could cause millions of cars to crash at once. Software written by AI could allow turning a cheap quadcopter with a haphazardly attached gun into a viable sniper drone. AI will even be capable of designing deadly viruses. This is far from an exhaustive list. These crises will happen without warning, deploying AI defensively won't prevent all of them, and many organizations won't upgrade their security soon enough regardless. Finally, at some unknown threshold, autonomous agents will surpass their creators' ability to develop more powerful intelligence. Calling recursive self improvement a national security risk is an understatement: it is America's final deadline.

Recommendations

The only way to prevent these disasters is to control the proliferation of frontier artificial intelligence. Frontier models must not be open sourced, and even closed source models must be limited, since they are all trivial to jailbreak and probably always will be. The US government must procure exclusive access to the most powerful models to ensure that its defense stays ahead of the strongest publicly available offense. AI research labs must have state-level security. AI chips must be tracked as WMD inputs. Since superhuman intelligence is inherently uncontrollable by humans even under ideal circumstances, a hard stopping point must be defined in advance. Once all of this is in place, America can focus on ensuring that other nations, particularly China, follow its lead. Last week's Superintelligence Strategy paper by Dan Hendrycks, Eric Schmidt, and Alexandr Wang is a good starting point for considering the risks that America must be prepared for and the measures that it can take: <https://www.nationalsecurity.ai>

Conclusion

Artificial intelligence is dual use and advancing rapidly. The dangers are unprecedented. Without government policies to control proliferation, AI will become analogous to the weapons HG Wells hypothesized in *The World Set Free*: "It was a matter of common knowledge that a man could carry about in a handbag

an amount of latent energy sufficient to wreck half a city." If America doesn't become the world leader in AI security, then America won't survive AI.

Austin Henrie

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

All e-mails to and from this account are for NITRD official use only and subject to certain disclosure requirements.

If you have received this e-mail in error, we ask that you notify the sender and delete it immediately.