

United States Office of Science and Technology Policy (OSTP) - Development of an Artificial Intelligence (AI) Action Plan Input

[Center for Existential Safety](#)

March 15, 2025

1. Introduction

The [Center for Existential Safety](#) believes that the continued development of artificial intelligence without extremely stringent safeguards will soon lead to humanity's extinction or permanent disempowerment.

We've known since at least [1863](#) with Samuel Butler's "Darwin Among the Machines" that machines might one day supplant humans. In [1951](#), amidst the digital computer revolution, Alan Turing cogently outlined artificial intelligence's existential risks. And for the last 20+ years AI safety pioneers have been attempting to help our species come to terms with this risk.

We are not alone in our concerns.

The Nobel laureate Geoffrey Hinton [says](#), "I actually think the existential risk [AI causing human extinction] is more than 50%."

Stuart Russell, the President of the International Association for Safe and Ethical Artificial Intelligence, [says](#), "The development of highly capable AI is likely to be the biggest event in human history. The world must act decisively to ensure it is not the last event in human history."

Sam Altman, the CEO of OpenAI, [said](#), "AI will probably most likely lead to the end of the world, but in the meantime, there'll be great companies."

To date, over 1,000 [risks](#) and 3,500 reports of [harms](#) (or 14,000 reports of [incidents and hazards](#)) from AI have been cataloged. And capacity is [growing](#) at extraordinary rates by [tech companies](#) that prioritize profits over safety.

In response, essentially no realistic plans have been developed, funded, and enacted on an international level to mitigate those risks. As a species, we spend far more on [candy](#) than we do on preventing our own extinction.

We believe we have collectively failed to prepare for this crucial moment in humanity's history. But with strong leadership from the United States, we can change that now.

2. Recommendation

In the spirit of this [open letter](#) “calling on world leaders to show long-view leadership on existential threats”, we urge the United States to now lead us with a laser focus on ensuring humanity's long-term existential safety.

Our primary recommendation for the Trump Administration is to join the [International AI Governance Alliance](#). If the United States joins and champions this alliance, we believe it will be leading humanity toward a much safer, brighter future for all. We believe this could be the biggest win-win in all of human history. The Trump Administration could champion this, with the President himself playing a defining role.

We consider this to be strongly in America's national interests. Indeed, it is one of the only ways to ensure its national security. If nearly all people don't agree to cooperate on AI, then everyone in the world is likely to die or be disempowered soon. This is because any one person anywhere might intentionally or unintentionally unleash a rogue superintelligence that quickly supplants humanity as the most powerful force on Earth. Or they may develop and deploy narrow AI systems to cause catastrophic damage or obtain unprecedented power, disrupting the geopolitical order.

3. Details

Full details are on the International AI Governance Alliance [website](#). Below we include some key parts.

The Need

Many experts [agree](#) that humanity faces a plausible risk of extinction due to the rapidly advancing capabilities of artificial intelligence.

The explicit near-term goal of many Big Tech companies is to create godlike AI systems far smarter and more capable than humanity itself.

These companies [acknowledge](#) the extinction risks posed by their products, yet continue developing them without the consent of the eight billion people they are affecting. This is entirely unacceptable.

For everyone's [existential safety](#), all large-scale frontier development of AI systems must [pause](#) immediately. We must regroup as a species and collectively decide what we want to do about AI.

The Response

We need a [supranational organization](#) capable of effectively mitigating extinction risks from AI and fairly distributing its economic benefits to all. We can call this the International Artificial Intelligence Governance Alliance (IAIGA). For simplicity, it is pronounced 'eye-guh'.

This organization should do three things:

1. Develop the Global AI Safety Treaty that regulates research and development of AI systems, ensuring they are provably safe and aligned to human values.
2. Manage the enforcement of this treaty.
3. Manage the fair distribution of the economic value that these AI systems generate.

To be effective, this organization must operate at a speed and scope commensurate with the technology it aims to govern. It must be proactive, not reactive.

It must also operate openly, in conjunction with all signatories and alongside extensive public and expert review.

It cannot be "politics as usual" with endless talk and limited action. Humanity's very existence is at stake, which means we all must transcend "normal" politicking.

The Commitment

We propose a Collective Commitment on AI for humanity. We commit to:

1. Adhere to internationally accepted safety standards for AI research and development.

2. Adhere to an internationally accepted distribution system for AI-derived economic value.

Individuals and organizations can make this commitment now. After completion, nations can adopt the Global AI Safety Treaty, which would put this commitment into law.

Committing now indicates your support for this safer, fairer future state of the world. Everyone can win, if we develop AI with extraordinary care. Otherwise, everyone will lose.

How will you determine the AI safety, enforcement, and distribution policies and standards of the Global AI Safety Treaty?

We believe it is too important for any one individual, company, organization or country to determine these policies and standards on their own, as they profoundly affect all life on Earth.

Therefore, we will invite all signatories to participate in an open door meeting to collaboratively make these decisions.

Prior to the meeting, we will run an open call for proposals from policymakers and the general public. Anyone can submit their ideas, even non-signatories.

In the meeting signatories will be invited to debate all known approaches to AI governance. See some examples [here](#). Proposals will be evaluated by signatories for their ability to adequately address humanity's near-term risk of extinction and the fair allocation of resources derived from AI. A prediction market for each proposal will also be created to capture collective insight into their viability.

Given the extreme urgency for immediate action, the meeting will not conclude until agreement on a single proposal is reached. We will then share that with world leaders.

We will reconvene at least annually to reaffirm or update our policies and standards.

4. Endorsements

In addition to our recommendation above, we also endorse the Machine Intelligence Research Institute's [Response to RFI on the Development of an Artificial Intelligence \(AI\) Action Plan](#).

—

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.