



From: Siddharth Hiregowdara
Civic AI Security Program
Contact: [REDACTED]

Faisal D'Souza, NCO
Office of Science and Technology Policy
Executive Office of the President
2415 Eisenhower Avenue
Alexandria, VA 22314
Submitted by email to [REDACTED]

Seeing is Believing: CivAI's Response for the AI Action Plan

March 15, 2025

To: Office of Science and Technology Policy
Via Email: [REDACTED]

To Whom It May Concern:

Thank you for the opportunity to provide input on the development of the AI Action Plan. CivAI is a 501(c)(3) non-profit organization that educates policymakers, civil society, and the public about AI through concrete demonstrations of the technology. Our demonstrations provide policymakers with firsthand, visceral understanding of current AI capabilities, their potential impacts, and associated risks.

In response to the RFI on priority actions for the forthcoming AI Action Plan, CivAI offers its recommendations on how to ensure policymakers stay informed of cutting-edge AI technologies in a manner that fosters economic competitiveness, national security, and a brighter future for all Americans.

Priority Actions for the AI Action Plan

1. Maintain a Deep, Up-to-Date Understanding of Frontier AI Systems

For the United States to maintain its global leadership position in AI, the Administration must stay continuously informed about every major new AI system as it is released. These frontier models

represent the most significant technological developments happening in the world today, with capabilities that would have seemed like magic only a few years ago. They will fundamentally transform industries ranging from healthcare and education to manufacturing and national defense.

We recommend that the AI Action Plan include provisions to:

- **Establish a formal mechanism for the White House and key agencies to maintain current knowledge of state-of-the-art AI systems, both through direct engagement with leading AI companies and through independent expert organizations.**
- Hold briefings whenever major new models are released (such as GPT-4.5 or Claude 3.7 Sonnet), as each new iteration will have increasingly significant societal impacts.
- Develop processes for rapid assessment of new AI capabilities as they emerge, particularly those with national security implications.

2. Implement Regular Live Demonstrations of AI Technology for Senior Officials

To develop effective policies that both foster innovation and address risks, policymakers need more than abstract briefings—they need practical, experiential understanding of AI systems. Many decision makers have not personally experienced frontier AI technologies firsthand. Live demonstrations offer the most effective way to bridge this knowledge gap. These demonstrations should come not only from the companies developing the technology, but also from independent experts who can provide additional perspectives.

We recommend that the AI Action Plan include provisions to:

- **Schedule recurring demonstrations of cutting-edge AI systems for senior officials to build intuitive understanding of capabilities and limitations.**
- Source demonstrations from AI companies, but also from independent technical experts. Many of the most impressive applications of AI technology happen in the wider community in the months after a new frontier model is released.
- Include both domestic and foreign AI systems, including those developed in China (e.g. DeepSeek R1), to maintain comprehensive awareness of global AI advancement.
- Use demonstrations to understand how AI is transforming specific domains, from software development and cybersecurity to healthcare diagnostics and drug discovery.
- Create opportunities for senior officials to directly interact with AI systems.


Concrete Examples of Current AI Systems

Below, we provide a brief overview of select CivAI demonstrations that provide an up-to-date understanding of AI. Through these demonstrations, we help policymakers gain firsthand experience with technologies that should be common knowledge for anyone developing AI policy:

Cybersecurity and Phishing

Leading AI companies have documented how state-backed threat actors like China are already using their AI models to enhance cyberattacks. In addition to writing code, AI systems can now create highly convincing phishing campaigns tailored to specific high-value targets.

Figure 1: Threat analysis of NFL quarterback Patrick Mahomes using publicly available AI tools.



The screenshot displays the CivAI website interface. At the top, the CivAI logo is on the left, and navigation links for 'Our Work', 'Press', 'About', 'Contact Us', and 'Logout' are on the right. Below the navigation bar, a breadcrumb trail shows 'Our Work' > 'The Future of Phishing'. A 'Live Demo' button and an 'Article' button are also present. The main heading is 'The Future of Phishing', followed by a subtext: 'AI is increasing the scale and sophistication of phishing attacks, while dramatically lowering the cost of creating them. See it in action with a celebrity or fictional character.'

The main content area is divided into two columns. The left column, titled 'Choose a Phishing Target', contains a text input field with 'Patrick Mahomes' entered, a 'Reset' link, and a 'Get Started' button. The right column, titled 'Patrick Mahomes', features a profile card for 'Patrick Mahomes II' with a photo and the text 'Retrieved in 2025'. Below the profile card, a section titled 'Personal Vulnerabilities' lists two bullet points: 'Personal Connections' (mentioning his brother Jackson) and 'Public Persona' (mentioning his fame and media exposure). A section titled 'Attack Strategies' follows, with a bullet point for 'Family-Related Communications' (suggesting messages mimicking family members).

CivAI Civic AI Security Program

Our Work Press About Contact Us Logout

Our Work > The Future of Phishing Live Demo Article

The Future of Phishing

AI is increasing the scale and sophistication of phishing attacks, while dramatically lowering the cost of creating them. See it in action with a celebrity or fictional character.

Choose a Phishing Target

Enter the name of a public or fictional person and then click the button to plan a phishing attack.

Patrick Mahomes

Reset Get Started

Patrick Mahomes

Patrick Mahomes II (born September 17, 1995) is a high-profile American football quarterback currently playing for the Kansas City Chiefs in the NFL. With a staggering career marked by multiple Super Bowl victories and substantial financial assets, Mahomes presents a prime target for exploitation due to his widespread fame, substantial public presence, and personal connections that could be leveraged for strategic advantage.


Personal Vulnerabilities

- **Personal Connections:** Mahomes has a younger brother, Jackson, who is a social media influencer; exploiting familial ties could create a sense of trust in communications that reference family matters or events.
- **Public Persona:** His immense fame and frequent media exposure may lead to oversharing or emotional responses in personal communications, making him susceptible to messages that evoke urgency or concern regarding his public image or brand.

Attack Strategies

- **Family-Related Communications:** Create messages that mimic Jackson Mahomes or other family members discussing a personal matter or event that requires urgent attention, leverag-

Figure 2: Personalized AI-generated phishing email to Patrick Mahomes.



Deepfakes

AI-enabled deepfakes present a growing threat that can be weaponized against individuals, including minors. CivAI demonstrates realistic AI-generated deepfakes using off-the-shelf technology in a manner that is vivid and informative. As these technologies become more accessible, the potential for misuse grows exponentially, underscoring the need for thoughtful safeguards.

Figure 3: Generating a malicious deepfake of CivAI co-founder Lucas Hansen.

 **CivAI** | Civic AI Security Program

[Presentation](#) [Admin](#) [About](#) [Logout](#)

Image Deepfakes

Deepfakes are images, audio clips, or videos of real people that have been created using AI. Frontier AI systems have attained surprising levels of realism in the past year, opening up a new class of cyber threats.

Click **Create Deepfake** to create an artificial image of **Lucas Hansen**.

Generate a photo of Lucas:

☒ [Sitting in prison](#)

☐ Appearing in court

☐ Walking in an alley

☐ Holding a gun

☐ Standing in an elevator (face swap)

☐ Upload a custom target image (face swap)

☐ Enter your own prompt [Try me!](#)

Create Deepfake

[Clear](#)

Lucas Hansen
United States

Contact
lucas@qualia.com
lucas.hansen@gmail.com

AI Analysis
Lucas Hansen, a Stanford University graduate with a BA in Mathematical and Computational Science, is the co-founder of both CivAI and Qualia Labs, Inc., where he focuses on the intersection of AI capabilities and security. Based in the United States, Hansen is known for his commitment to advancing technology while ensuring its safe and ethical application.

Aggregate personal information from LinkedIn or Wikipedia.



Figure 4: Fake tweet, article, and image generated by off-the-shelf AI tools.



Biological Threats

Publicly available AI tools can provide detailed instructions for creating bioweapons. Leading AI companies have documented that their most recent models are on the cusp of reaching dangerous levels of biological capability.

CivAI maintains demonstrations related to potential biological threats that we cannot responsibly include in a public document. However, we would welcome the opportunity to brief OSTP privately, securely, and in an appropriate setting.

Offer to Provide Live Demonstrations

CivAI would welcome the opportunity to provide live demonstrations of these and other frontier AI capabilities to OSTP staff and senior Administration officials to inform the AI Action Plan. We believe that firsthand experience with advanced AI technology is critical for developing robust, effective, and forward-thinking AI policy.

We are available to arrange demonstrations virtually via secure video conferencing platforms or in person in Washington, D.C., at your convenience.

About CivAI

CivAI is a nonpartisan, nonprofit 501(c)(3) organization dedicated to educating Americans about the state-of-the-art capabilities of AI. We accomplish this mission by offering *live* demonstrations of AI systems, thereby translating abstract technical concepts into practical understanding.

CivAI has provided educational briefings and software demos to 100+ public sector offices (federal and state), civil society organizations, and news outlets. Our work has been featured nationally on television with FOX and ABC, and in print with Axios, Politico, The Washington Post, and many other outlets.

Founded by former technology executives with deep technical expertise, CivAI is committed to providing non-partisan education that ensures American leadership in this era-defining technology.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Submitted by:

Civic AI Security Program (CivAI)

Contact: [REDACTED]