March 11, 2025

Faisal D'Souza, NITRD National Coordination Office
2415 Eisenhower Avenue
Alexandria, VA 22314

Submitted Electronically via ostp-ai-rfi@nitrd.gov

Re: AI Action Plan

Dear Mr. D'Souza,

Thank you for the opportunity to respond to your RFI regarding the Development of an Artificial Intelligence (AI) Action Plan. Epic is a global healthcare technology company that works with hospitals and clinics in every state across the nation, providing an integrated software suite to manage both clinical care and back-office operations. Our patient portal, MyChart, puts medical records in the pockets of patients, enhancing medical literacy, communication, and wellness.

AI has the potential to ease administrative burdens for providers and help patients get well faster. Epic has over 100 use cases of generative AI under active development, and more than half of our customers are already using at least one generative AI feature. However, there are several unnecessarily burdensome requirements in U.S. regulation that have either held back our development efforts or made our customers question their willingness to implement our features. In this response, we provide multiple categories of reform for your consideration and provide specific examples within each.

We would be happy to answer any questions you may have and would welcome further collaboration on the development of the AI Action Plan. This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Sincerely,

Ladd Wiley
Senior Vice President for Global Corporate Affairs, Public Policy, and Advocacy

Define AI Appropriately, and Don't Regulate Beyond It

Government regulation of AI will be most effective when it is grounded in an understanding of how the technology works and what specific risks it presents. Healthcare organizations have used rules-based, deterministic decision support to improve patient care for decades. Predictive and generative AI are different, using statistics and probability to draw inferences from data and recreate patterns in ways that vary across organizations and even from one use to the next. AI-specific rules should not inadvertently be applied to non-AI software.

It is important for the federal government to set the standard on core, definitional issues in AI regulation because, at present, many states are taking matters into their own hands and causing confusion. For example, the California Privacy Protection Agency (CPPA) recently issued a proposed rule regulating, among other things, "automated decision-making technology (ADMT)." While California has adopted an industry-standard definition of AI, the proposed rule defines ADMT broadly as "any technology that processes personal information and uses computation to… substantially facilitate human decision making."[1] This definition effectively encompasses all computer software. Even a bipartisan group of California legislators submitted a public comment imploring the CPPA to "scale back" the ADMT regulations and "avoid the general regulation of AI."[2]

This example demonstrates the importance of adopting an appropriately narrow and precise definition of AI. Other common, flawed definitions often invoke "human-like capabilities," which undermines regulatory certainty by introducing a moving target; presumably, at some point, simple pocket calculators were perceived as performing tasks typically associated with human intelligence.

The federal definition of AI currently established in 15 USC § 9401(3) is an outdated version of the definition drafted by the Organization for Economic Co-operation and Development (OECD) and was put in force by the William M. (Mac) Thornberry National Defense Authorization Act for Fiscal Year 2021.[3] The definition is decent, but OECD has since improved it further, and an increasing number of jurisdictions (including several U.S. states and other countries) have been adopting the new version. The federal government could better promote interstate commerce by updating its definition of AI to "a machine-based system that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments."

---

[1] CPPA, "Proposed Regulations on CCPA Updates, Cybersecurity Audits, Risk Assessments, Automated Decisionmaking Technology (ADMT), and Insurance Companies," https://cppa.ca.gov/regulations/ccpa_updates.html

[2] Assm. Jacqui Irwin, et. al, "Public Comment on ADMT Regulations," https://acrobat.adobe.com/id/urn:aaid:sc:VA6C2:688f11b5-83dc-48b3-b066-4a568b6e6b69

[3] Congress.gov, "Public Law 116-283—Jan. 1, 2021," https://www.congress.gov/116/plaws/publ283/PLAW-116publ283.pdf

It would be similarly helpful for the federal government to establish a risk stratification of AI that other parties could adopt and follow. For example, AI that supports human decision-making but leaves the human ultimately in control of any outcome is fundamentally lower risk than completely automated decision-making. Similarly, AI that is used to determine insurance premiums is higher risk than AI that makes shopping recommendations. However, such a risk stratification must not conflate sector with use case. Healthcare, for example, is of critical importance to the wellbeing of people and society, but not all applications of AI within healthcare present the same risk. AI that helps clinicians take notes during appointments so they can focus their attention on the patient in front of them is lower risk than AI that develops complex treatment plans.

Ensure Executive Agencies Do Not Interfere with the Practice of Medicine

In one notable case where Congress has asserted itself on AI in healthcare, the Food and Drug Administration (FDA) seems to be actively thwarting Congressional intent. As part of the 21st Century Cures Act of 2016, Congress exempted clinical decision support (CDS) software that meets certain criteria from regulation as a medical device. However, in September 2022, FDA issued final guidance that narrows the exemption considerably. For example, FDA asserts that if an advisory makes a single suggestion instead of providing a list of options, it fails the exemption criterion requiring CDS to be "for the purpose of supporting or providing recommendations," relying on the argument that the pluralized word "recommendations" is not inclusive of its singular form.

This questionable interpretation of statute threatens to sweep a tremendous number of CDS tools into FDA's oversight. Among just Epic's U.S. customers using a specific CDS feature known as OurPractice Advisories, there are more than 150,000 unique advisories in active use. Many advisories take the form of a "single suggested next step" because a team of clinical leaders at the healthcare organization has concluded a clinician should consider a specific action in that specific clinical situation. FDA has limited resources and is unlikely to be prepared for this kind of volume of new "medical devices."

In November 2024, members of Congress from both parties submitted a letter to FDA indicating that "FDA's guidance does not reflect its typical risk-based approach or consider the significant clinical oversight under which these tools are configured."[4] In a February 2025 article in JAMA Health Forum, Scott Gottlieb, former FDA commissioner, argued that "the current regulatory posture of the FDA for classifying tools with…

[4] Healthcare IT News, "Lawmakers ask CDRH to revisit its CDS guidance," https://www.healthcareitnews.com/news/lawmakers-ask-cdrh-revisit-its-cds-guidance

advanced analytical capabilities as medical devices could impede or even block their integration into [electronic medical record] systems."[5]

FDA provided some clarification via FAQs in December 2024, but did not resolve the fundamental concerns about the CDS exemption. It would be appropriate for FDA to completely rescind its CDS final guidance and rewrite it, adhering more closely to the Congressional intent of the 21st Century Cures Act and promoting safety in healthcare in a way that does not stymie the integration of AI into this important sector.

Align Policy with the Current State of AI Capabilities

The significant leap forward in AI capabilities manifested by the introduction of ChatGPT in late 2022 merits a reevaluation of AI regulations that had been in place before that time. Certain regulations may no longer be as appropriate or relevant as they once were.

For example, the 2016 rule implementing section 1557 of the Patient Protection and Affordable Care Act (ACA), which prohibits discrimination in health programs and activities, included a requirement regarding meaningful access for individuals with limited English proficiency that a covered entity "use a qualified [human] translator when translating written content in paper or electronic form." Such a requirement may well have been appropriate in 2016, when the capabilities of AI were much more limited. Today, however, large language models have been shown to perform translation, even of complex and highly technical content, excellently.

Since 2016, several subsequent rules have revised the implementation of section 1557 of the ACA, but the concept of using only humans for language translation has persisted through each iteration. Most recently, the 2024 rule introduced 45 CFR § 92.201(c)(3), which stipulates that "if a covered entity uses machine translation when the underlying text is critical to the rights, benefits, or meaningful access of an individual with limited English proficiency, when accuracy is essential, or when the source documents or materials contain complex, non-literal or technical language, the translation must be reviewed by a qualified human translator."[6]

While this revision represents clear progress in that it acknowledges the existence of machine translation for the first time, the requirement prevents, in effect, any widespread adoption of AI translation services. If a healthcare provider is going to have to compensate a human translator anyway, there will be little justification for also

[5] JAMA, "New FDA Policies Could Limit the Full Value of AI in Medicine," https://jamanetwork.com/journals/jama-health-forum/fullarticle/2830189

[6] Code of Federal Regulations, "45 CFR § 92.201(c)(3)," https://www.ecfr.gov/current/title-45/subtitle-A/subchapter-A/part-92/subpart-C/section-92.201

investing in AI services to accomplish the same task. Finding a way to keep humans in the loop without undermining the value proposition of AI translation would represent meaningful progress. For example, it may now be more appropriate to require that qualified human translators provide initial validation of machine translation tools but then not necessarily be involved in each individual subsequent translation activity.

It is also worth noting that the concept of "translation" should not be limited to converting text from one language to another. Another common barrier to access for patients from all language backgrounds is clinical instructions filled with medical jargon that a specific patient may not understand. Using AI to "translate" such text to an appropriate reading level for the patient could significantly improve adherence, which could in turn greatly contribute to overall population health.

Focus On and Provide Support for Local Validation of AI Features

Over the past few years, a key focus for AI policymakers has been on how to validate the performance of AI. The training process used for modern AI introduces the potential for bias and discrimination, and the probabilistic nature of AI's output makes it particularly tricky to pinpoint the future behavior of an AI feature with confidence.

In the healthcare space, groups like the Coalition for Health AI (CHAI) proposed a nationwide network of "assurance labs" to perform premarket validation of AI tools. While outsourcing this work to the private sector likely would be more efficient than a centralized FDA premarket review process, any mandate of such a structure would represent a significant expansion of the scope of healthcare tools subject to premarket approval. (Importantly, using AI for health does not inherently make it a medical device. Designating something as a medical device reflects its intended use; AI, on the other hand, is a technology that may be applied to accomplish such a use.)

Not only is premarket approval a framework that should be used only when truly necessary (because it can add *significant* delays to the time to market for innovative applications of technology), but also it would not be particularly well-suited to the specific challenges of validating AI. AI performance is sensitive to factors such as patient acuity mix and even the workflow into which it is integrated. For example, organizations with dedicated remote monitoring teams for conditions such as sepsis or deterioration may set the alert threshold on a predictive model much lower than those who can only rely on frontline nurses managing a large caseload of patients. While such adjustments make perfect sense considering workflow, they fundamentally alter the "performance" of the AI models on paper, shifting the proportion of false positives and false negatives.

As a result, the exact same AI product that functions well in an academic medical center in New Haven, CT may not be appropriate for a critical access hospital in rural Iowa. Furthermore, some AI models continue to learn and evolve over time, meaning that premarket validation will always be out of date. Thus, while baseline performance testing by AI developers is always important, ongoing post-deployment monitoring by deployers (e.g., healthcare delivery organizations) may be even more important.

The challenge for post-deployment validation, in healthcare at least, is that small or rural healthcare organizations often struggle to afford the necessary data science resources and expertise to perform effective monitoring. Yet because AI performance will vary from organization to organization, post-deployment validation is not something that can be outsourced in practice. Without appropriate support for “local validation” activities, AI could widen, rather than close, the digital divide in healthcare.

AI developers can help; in 2024, for example, Epic released an open-source tool called the AI Trust and Assurance Suite that automates the most resource-intensive aspects of post-deployment monitoring. The suite's real-time evaluations use the deployer's local data and workflows and can monitor both Epic and third-party AI features, using the deployer's choice of evaluation methodologies.

However, even with access to the Suite, many lower-resourced healthcare organizations say they don't have the time and staffing required to engage in local validation and, thus, to fully embrace AI. One opportunity for government investment in AI worth considering would be a grant program for small and rural hospitals and clinics to enable proper local validation and post-deployment monitoring of AI performance.