



Comments on the Development of an Artificial Intelligence (AI) Action Plan

March 15, 2025

The Hacking Policy Council (“HPC”) submits the following comments in response to the Office of Science and Technology Policy (OSTP) and Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO)’s Request for Information (“RFI”) on the Development of an Artificial Intelligence (AI) Action Plan.¹ We thank OSTP and NCO for the opportunity to provide input towards this important proposal.

HPC is a group of industry experts dedicated to creating a more favorable legal, policy, and business environment for vulnerability management and disclosure, good faith security research, penetration testing, bug bounty programs, and independent repair for security.² Many of our members are deeply involved in AI system deployment, testing, and red teaming.

HPC’s comments focus on the critical role of AI testing and red teaming. Even as we use AI to drive innovation in cybersecurity, we also recognize that its evolution introduces significant risks, particularly in cybersecurity and broader security domains. Given the growing reliance on AI technologies across industries, it is essential to prioritize AI security and resilience. A comprehensive AI strategy must include testing AI vulnerabilities early in development and throughout their lifecycle to ensure public safety and mitigate potential threats. It is imperative that these risks be understood and addressed to safeguard against unintended consequences.

As the White House works to define its priorities around AI, below are our recommendations:

¹ Office of Science and Technology Policy (OSTP) and Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO), Request for Information, Development on an Artificial Intelligence (AI) Action Plan, Feb. 6, 2025, <https://www.federalregister.gov/documents/2025/02/06/2025-02305/request-for-information-on-the-development-of-an-artificial-intelligence-ai-action-plan>

² Hacking Policy Council, <https://hackingpolycouncil.org>.

1. Encourage red-teaming to evaluate cybersecurity and security risks

Given that the Trump administration's priority is to strengthen the U.S. AI leadership and competitiveness, particularly in comparison to other nations such as the People's Republic of China, we believe it is crucial for the government to actively promote and support the adoption of AI red-teaming practices. By rigorously testing AI systems, red-teaming helps to identify vulnerabilities and weaknesses, ensuring the security and resilience of AI technologies.

In cybersecurity, red-teaming typically involves simulating adversarial attacks to expose system vulnerabilities and improve defensive strategies.³ However, as AI technologies become increasingly integrated into different sectors, it is important to recognize that AI red-teaming can go beyond traditional adversarial testing. It should address both adversarial and non-adversarial risks, evaluating a wider range of potential threats. This will help identify and mitigate risks that may not stem from direct attacks but still pose serious threats to public safety, national security, and societal well-being.

For example, aside from cybersecurity risks, red-teaming could assess unintended outputs and wider security concerns, such as chemical, biological, radiological, and nuclear (CBRN) threats. These areas represent critical infrastructure where AI could have profound and potentially dangerous consequences if misused or poorly implemented. Testing systems in these domains could help prevent hazardous chemical releases, biological terrorism or pandemics, and nuclear incidents.

Given the diverse nature of AI testing and the range of associated risks, we urge the development of a nuanced framework that pushes for the adoption of these practices, but also distinguishes between different types of AI testing. This framework should provide clarity on when and how each type of testing should be employed to effectively manage these risks.

2. Advocate and establish secure communication channels for reporting AI-related vulnerabilities.

Much like vulnerability disclosure policies (VDPs) for security vulnerabilities, AI developers should be equipped with structured processes for receiving, analyzing, and addressing AI-related vulnerabilities. We strongly encourage the government to advocate for and establish secure communication channels dedicated to the

³ Hacking Policy Council, Comments to Request for Information Related to NIST's Assignments Under the Executive Order Concerning Artificial Intelligence, Feb 2, 2024, https://assets-global.website-files.com/62715f02a51b614ce64867fd/65c53257140b8d81b6f27422_Hacking%20Policy%20Council%20-%20comments%20to%20NIST%20re%20AI%20red%20teaming%20-%2020240202.pdf

responsible disclosure of AI vulnerabilities. These channels will ensure that vulnerabilities are identified and remediated in a timely and efficient manner, reducing the risks to public safety and national security.

However, we'd like to note that it may be helpful to have different processes for AI cybersecurity disclosures and disclosures for other unintended weaknesses. Disclosures related to CBRN or other risks will require different mitigations than those for traditional security vulnerabilities and may involve different expert teams and resources.

3. Provide clear protections for external researchers conducting AI-related research.

To further promote responsible AI research, the government should establish clear legal protections for external researchers conducting AI-related testing. These protections will ensure that independent researchers are shielded from potential legal risks while performing both AI-related research. By providing these safeguards, the government can encourage more widespread engagement from the research community.

In particular, we urge the government to update the Computer Fraud and Abuse Act (CFAA) to extend protections for good faith security research to good faith AI research. While the Department of Justice announced its intent to support this, no substantial action has been taken. We encourage this administration to follow through on this important initiative, ensuring that researchers conducting responsible AI testing are clearly protected under the law.

4. Ensure results of red-teaming are sensitive and place appropriate protections in place.

Finally, HPC strongly supports reducing regulatory burdens and believes that the mandatory reporting of physical and cybersecurity protections, as well as AI red-teaming test results, should not be required. As mentioned in previous comments⁴ to the Department of Commerce's Bureau of Industry & Security (BIS), we believe that the detailed reporting of red-team testing outcomes could inadvertently create significant security risks, particularly for dual-use foundation models and the organizations developing them.⁵

⁴ Hacking Policy Council, Comments to BIS on the Establishment of Reporting Requirements for the Development of Advanced Artificial Intelligence Models and Computing Clusters, Oct. 10, 2024, https://cdn.prod.website-files.com/660ab0cd271a25abeb800453/670e9bb356f840554f037131_Hacking%20Policy%20Council%20-%20Comments%20to%20BIS%20re%20reporting%20requirements%20for%20advanced%20AI%20-%2020240924.pdf

⁵ Hacking Policy Council, Comments to BIS on the Establishment of Reporting Requirements for the Development of Advanced Artificial Intelligence Models and Computing Clusters, Oct. 10, 2024,

Overly detailed reports have the potential to expose unmitigated vulnerabilities, testing methods, or approaches that could be exploited to disable, manipulate, or compromise the integrity of AI systems. If such sensitive information were to be accessed by unauthorized parties, it could provide malicious actors with a clear blueprint for targeting these systems. Therefore, it is critical that results from red-teaming exercises be handled with the utmost discretion, ensuring that only necessary information is shared with relevant stakeholders, and that sensitive findings are protected to prevent them from being misused.⁶

*

*

*

Thank you for the opportunity to provide input to the AI Action Plan. If we can be of additional assistance, please contact Heather West, at [REDACTED] and Tanvi Chopra, at [REDACTED]

https://cdn.prod.website-files.com/660ab0cd271a25abeb800453/670e9bb356f840554f037131_Hacking%20Policy%20Council%20-%20Comments%20to%20BIS%20re%20reporting%20requirements%20for%20advanced%20AI%20-%2020240924.pdf

⁶ Hacking Policy Council, Comments to NIST on Managing Misuse Risk for Dual-Use Foundation Models, Sept. 9, 2024,

http://cdn.prod.website-files.com/660ab0cd271a25abeb800453/66e08a055a60774fbf66e0ce_HPC%20Comments%20to%20NIST%20on%20Managing%20Misuse%20Risk%20for%20Dual%20Use%20Foundation%20Models%20-%2020240924.pdf