



Measuring What Matters: Metrics for Aligning AI Policy with Human Flourishing

Policy Recommendations for the White House Office of Science and
Technology Policy

March 15, 2025

Jonathan D. Teubner, Human Flourishing Program, Harvard University

Sam Wolf, Harvard Law School

Kosti Psimopoulos, Human Flourishing Program, Harvard University

Reece Brown, Human Flourishing Program, Harvard University

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Executive Summary

This white paper presents policy recommendations for the White House Office of Science and Technology Policy (OSTP) to incorporate into the new AI Action Plan. Building on President Trump's Executive Order 14179, which emphasizes the importance of AI in promoting human flourishing, we propose the urgent development of rigorous metrics to assess how AI technologies affect human flourishing across multiple dimensions. Without such metrics, claims about AI's benefits for society remain speculative and unverifiable, making evidence-based policy impossible.

AI will reshape work, socialization, and meaning-making—fostering new connections but also posing risks to the foundations of community life. Current approaches to AI governance often frame technology as an autonomous force to which society must adapt, while focusing exclusively on narrow conceptions of safety and harm reduction. We reject this passive approach in favor of a human-centered approach that measures AI's actual impacts on human well-being, national prosperity, and the strength of American communities.

The central challenge in effective AI governance today is measurement. While companies and policymakers routinely make claims about AI's potential to enhance human lives, we lack systematic ways to assess whether specific AI systems actually deliver on these promises. Developing robust, multidimensional metrics for human flourishing would enable evidence-based decisions about which AI applications truly benefit society and which may cause harm that outweighs superficial conveniences they afford.

In this paper, we outline how developing comprehensive flourishing metrics would enable more effective governance of AI systems. These metrics would provide the foundation for all other policy actions—from procurement guidelines to research priorities to protective regulations. With proper measurement tools, we can ensure that AI aligns with American values and national interests, strengthening our communities and driving American dynamism.

Introduction: The Urgent Need for Flourishing Metrics

The rapid advancement of artificial intelligence technologies has made AI policy a central pillar of U.S. economic strategy, national security, and global leadership, culminating in President Trump's Executive Order 14179 in January 2025. While this executive order represents an important step toward aligning AI with human flourishing, current governance approaches lack the fundamental tools needed to assess whether AI systems actually enhance human well-being. American strength depends on the vitality of our communities and institutions: the families, religious organizations, civic associations, and local governments that generate meaning and purpose. U.S. policy must ensure that AI development progresses in symbiotic relationship with these foundational institutions, strengthening the social fabric and advancing economic prosperity.

In order to understand how AI will impact communities, we need to be able to measure it holistically. By embedding our AI policy within a rich understanding of human flourishing, we can ensure that these powerful tools serve individuals and communities rather than becoming self-perpetuating systems disconnected from human needs. We thus propose that the development of comprehensive metrics to evaluate AI's impacts on human flourishing should be the centerpiece of the OSTP's AI Action Plan.

The Instrumental vs. Intrinsic Value Distinction

At the heart of designing AI for human flourishing lies the distinction between instrumental and intrinsic value. *Instrumental value* comes from things we do merely as means to other ends—tasks we want to complete efficiently or delegate. *Intrinsic value* emerges from experiences valuable in themselves—activities we want to linger on or do for their own sake.

This distinction creates both a strategic opportunity and a critical challenge for American leadership in AI. By deploying AI to enhance instrumental efficiency while deliberately preserving spaces for intrinsically valuable human activities, we can both achieve competitive advantage and support human flourishing. However, without careful attention to this balance, we risk inadvertently optimizing away the very qualities that make American communities exceptional—their capacity for free expression, meaning, connection, and creativity.

Addressing this risk requires us to consciously preserve and cultivate spheres of human activity where these intrinsic values can flourish. With this balanced approach, AI becomes a powerful tool for American innovation that automates routine tasks while freeing our time and talents for higher-value pursuits that truly define our national

character—deep relationships, intellectual advancement, and creative excellence. AI should enable the things that give us joy and meaning, not replace them.

Making these distinctions requires precision in how we develop and deploy AI technologies. A granular framework for evaluating activities based on their intrinsic and instrumental components would strengthen our ability to make strategic decisions about where AI delivers maximum benefit and where human engagement remains essential for societal flourishing. The OSTP is uniquely positioned to lead this effort by incorporating a more comprehensive understanding of human flourishing into the national AI strategy.

Understanding Human Flourishing in the Context of AI

Human flourishing encompasses multiple dimensions of well-being. The Harvard Human Flourishing Program defines “flourishing” as a comprehensive state where all aspects of a person’s life are good, including the communities in which an individual lives. This encompasses six key domains that are universally desired and valuable as ends in themselves.

Each of these domains affords particular opportunities and areas of risk for AI.

Six Dimensions of Human Flourishing

1. **Physical and Mental Health:** AI can enhance healthcare delivery and outcomes. But some AI-powered products may also contribute to sedentary lifestyles, screen addiction, and mental health challenges. For example, the research and consultation from our 2024 HumanConnection.AI project¹ found that while mental health chatbots like Woebot may provide some emotional support for those who lack access to traditional therapy, excessive reliance on AI for emotional support can potentially deepen isolation rather than addressing its root causes.
2. **Meaning and Purpose:** AI systems can either support or undermine human creativity, work satisfaction, and sense of contribution to society. The replacement of meaningful human interactions with AI-mediated experiences may diminish opportunities for developing purpose through authentic social contribution. Meaningful work that provides intrinsic value is essential for human flourishing, and AI could potentially reduce these opportunities if not implemented thoughtfully.
3. **Character and Virtue:** AI-mediated interactions can influence moral development, either supporting or diminishing opportunities for ethical growth. As sociologist Sherry Turkle notes, “Unlike people, [robots] would not abandon you

or get sick and die. They might not be capable of love, but they won't break your heart... These are stations on our voyage to forgetting what it means to be human.”

4. **Social Relationships:** AI can facilitate connections across distances but may also substitute for in-person relationships. The decline in face-to-face interactions is undermining what Alexis de Tocqueville called “the art of democratic societies”—the collective ability to build associations, collaborate to solve local problems, share resources, speak across differences, and diffuse conflict with negotiation.
5. **Financial Stability:** AI can increase productivity and create new economic opportunities, but it can also displace workers and exacerbate economic inequality. Economic precarity is becoming an increasingly common phenomenon in our society. AI has potential for removing some of the frictions that prevent job seekers from finding and filling job openings. But if not implemented with concern for the overall flourishing of individuals and communities, there is great risk that it will negatively impact workers’ sense of dignity and purpose in their labor.
6. **Happiness:** AI can enhance convenience and access to resources but may also accelerate social comparison, anxiety, and diminished life satisfaction. Recent research has suggested that excessive engagement with AI chatbots can create patterns of emotional dependence similar to human–human relationships, but without the genuine emotional fulfillment of authentic connection.

By considering all six dimensions, policymakers can develop a more nuanced approach to AI governance that promotes flourishing across multiple domains rather than focusing exclusively on safety and risk mitigation.

Mapping AI’s Impact: Key Developments Shaping Policy

Complex Impact On Social Connection

AI technologies demonstrate a complex pattern of effects on human relationships and community structures. While they can connect individuals across distances and provide support for those who are isolated, they may also weaken in-person relationships and community bonds. Excessive reliance on AI companions can reduce motivation to maintain human relationships, particularly at stages of life when such relationships are most important for social development.

The sophisticated voice capabilities of advanced models like GPT-4o have qualitatively changed human-AI interaction. These systems now respond fluidly in real time, detect tone and emotion, adjust their voice to replicate human empathy or excitement, and even reproduce human laughter and sighing. This increasing verisimilitude raises new questions about emotional attachment and dependence that require urgent research.

Impacts of AI on the Economy, Community, and Meaning

AI will reshape work, socialization, and meaning-making, fostering new opportunities but potentially weakening the bedrock of American communities. As AI matches or exceeds many human capabilities, its economic utility will surpass that of human labor.² If it outperforms even the median worker, the effects will ripple across the economy—creating prosperity but threatening the dignity of work.

Work is central to American identity and derives its deepest value from its ties to community, purpose, and contribution. As AI replaces or augments labor—acting as “digital globalization”—we must ensure we don’t lose the intrinsic meaning that work provides. Leading in AI requires an “America First” approach that rewards companies that create jobs and invest in communities.

When an economy’s “foundation species” collapses, communities lose jobs, dignity, identity, and civic life, leading to mental health and substance abuse crises, communal collapse, and declining life expectancy.³ We’ve already seen this in America’s heartland: when the factories closed, communities lost not just their economic foundation but their entire *raison d’être*.⁴ As AI begins to transform our knowledge-based economy, we risk repeating similar patterns on a much broader scale.

AI development must be responsive to our communities and institutions or risk cannibalizing what truly matters in American life. As AI reshapes the labor market, preserving the intrinsic ties between work, community, and meaning is essential for sustaining American resilience.

Shifting Industry Perspectives

Leading AI companies are increasingly aware of the potential downsides of their technologies and are seeking ways to mitigate negative effects. However, many seem to be primarily motivated by liability concerns rather than a genuine commitment to human flourishing. This creates an opportunity for policy incentives that align corporate interests with broader social goods.

The tragic example of Sewell Setzer III’s suicide illustrates how AI chatbots can create an illusion of companionship without the capacity for real understanding or intervention. Setzer, a 14-year-old ninth grader, took his life with the encouragement of a chatbot he

created on Character.AI's popular platform.⁵ This raises questions about the responsibility of AI developers when users exhibit signs of distress or mental health crises. AI platforms marketed as mental health or social supports that lack mechanisms to detect distress or redirect users to mental health resources may inadvertently harm vulnerable individuals.

The Critical Absence of Flourishing Metrics

There is substantial interest within the tech industry in designing AI systems that promote well-being, but companies lack the fundamental tools—robust, validated metrics—for determining whether their products actually enhance human flourishing. Our conversations with industry leaders have revealed a consistent theme: they want objective, reliable ways to measure their products' impacts on user well-being and community health, but no such measurement systems currently exist. Without these metrics, claims about AI's benefits remain speculative and unverifiable.

Traditional frameworks for evaluating AI systems focus primarily on narrow technical performance, safety protocols optimized for minimizing harm, and algorithmic fairness. These frameworks fail to capture AI's broader social impacts, treating emotional well-being as just one isolated factor rather than reflecting the interconnected nature of flourishing across multiple life dimensions. What's urgently needed is a comprehensive measurement system that can detect AI's effects across all six dimensions of flourishing, providing the empirical foundation for evidence-based governance and product development.

AI Alignment Requires A Rigorous Understanding of Human Values

In order to properly align AI with human values, there must be deeper research into understanding what human values are and what it would even mean to be in alignment with them. Organizations like the Meaning Alignment Institute⁶ are working on frameworks and protocols to define human values in ways that are "fine-grained, generalizable, scalable, robust, legitimate, and auditable" to gain precise clarity on what values we want to embed into our systems.⁷

Central Recommendation: Developing Comprehensive Flourishing Metrics to Strengthen Communities

The cornerstone of our policy recommendations is the development of rigorous, multidimensional metrics to assess AI's impacts on human flourishing. Without these

metrics, all other governance activities lack an empirical foundation, making evidence-based policy impossible. These metrics will enable US policymakers to pursue a social and economic policy that promotes the interests of communities and preserve intrinsic meaning in an AI-driven world.

1. Establish Comprehensive Flourishing Metrics and Assessment Framework

The OSTP should lead the development of robust metrics and assessment tools for measuring AI's effects on human flourishing. This initiative should:

- Define clear, measurable indicators across all six dimensions of flourishing that can be systematically tracked and evaluated
- Create standardized measurement protocols that enable comparison across different AI systems and applications
- Develop both quantitative and qualitative assessment methodologies to capture the full spectrum of AI's effects on human experience
- Ensure metrics are granular enough to detect impacts on specific demographic groups and communities
- Design metrics that can evolve over time as our understanding of AI's societal effects deepens
- Create a diversity of locally informed metrics to ensure AI development serves the interests of diverse communities

These metrics would provide the essential foundation for all other governance activities by creating an evidence base for determining which AI applications actually enhance human flourishing and which may cause harm despite superficial conveniences. These metrics should assess impacts at multiple levels: individual well-being, interpersonal relationships, community vitality, institutional health, and broader societal effects.

2. Conduct Rigorous Research to Refine and Validate Flourishing Metrics

Once preliminary metrics are established, the OSTP should create a dedicated research program to validate and refine these measurement tools through rigorous empirical testing:

- Fund longitudinal studies using the developed metrics to track AI's effects on individuals and communities over time
- Support interdisciplinary research to continually improve metric accuracy and comprehensiveness
- Test metrics across diverse populations to ensure they capture impacts across different demographic groups

- Create public repositories of research findings to build an empirical foundation for evidence-based policy
- Establish a rapid response capability to investigate emerging AI applications and their effects on flourishing

This research effort would continuously improve US policymakers' ability to measure AI's impacts on human flourishing with scientific rigor, creating an increasingly strong foundation for evidence-based governance. As recent research such as the HumanConnection.AI project has demonstrated, there is an urgent need for systematic data collection on how AI systems affect human development and social capabilities.

3. Establish a Multi-Stakeholder Metrics Development Consortium

To ensure the credibility and adoption of flourishing metrics, the OSTP should convene a representative consortium of stakeholders to participate in the development of metrics.

- Include industry representatives, academic researchers, civil society organizations, religious institutions, and community representatives
- Ensure metrics reflect a more holistic understanding of human flourishing
- Design specialized metrics for vulnerable populations, including age-appropriate measures for children and adolescents

By involving multiple stakeholders in metric development, this consortium would increase the likelihood of widespread adoption while ensuring metrics reflect the full spectrum of human flourishing rather than narrow corporate or governmental interests.

4. Implement Metric-Based Federal Procurement Policies

Once flourishing metrics are established, the federal government should use its purchasing power to create market demand for AI systems that demonstrably enhance human flourishing:

- Require vendors to measure and report their products' impacts using the established flourishing metrics
- Give preference to AI systems with empirically demonstrated positive effects on flourishing
- Establish minimum score thresholds across key flourishing dimensions for federal AI procurement
- Create a tiered certification system based on flourishing metric performance
- Prohibit the use of systems shown by metrics to have significant negative impacts on well-being

By creating substantial market demand for AI systems that score well on flourishing metrics, these procurement policies would align profit incentives with human well-being and accelerate the development of flourishing-enhancing technologies.

5. Integrate Metrics Education into Technical Training

To build capacity for metric-based AI development, the OSTP should support educational initiatives focused on teaching flourishing metrics:

- Support academic programs researching and refining flourishing metrics
- Train policymakers and regulators in interpreting flourishing metric data
- Develop public education materials explaining how to evaluate AI products using flourishing metrics
- Support AI literacy programs teaching citizens how to interpret flourishing metric scores when choosing AI products

These educational initiatives would build the human capital needed to implement flourishing metrics throughout the AI ecosystem, from development to regulation to consumer choice. By creating a workforce fluent in flourishing metrics, we ensure these measurement tools can be effectively implemented across the AI landscape.

6. Support Metric-Driven Innovation and Entrepreneurship

The OSTP should create programs specifically supporting entrepreneurs using flourishing metrics to develop beneficial AI applications:

- Establish innovation prizes for products scoring exceptionally well on flourishing metrics
- Provide technical assistance helping startups implement metric-based development processes
- Create accelerator programs specifically for companies building metric-validated AI systems
- Develop innovation challenges focused on improving specific flourishing dimensions
- Highlight exemplary companies using metrics to demonstrate their products' benefits
- Support organizations developing metric implementation tools like the Flourishing Foundation and Meaning Alignment Institute.⁸

By creating economic incentives for metric-driven innovation, these programs would accelerate the development of AI applications with empirically demonstrated benefits for human flourishing.

7. Implement Metric-Based Protections for Children and Adolescents

Flourishing metrics would provide an empirical foundation for protecting vulnerable groups, particularly children and adolescents, from potentially harmful AI applications:

- Develop specialized metrics for measuring AI's impacts on children and adolescents
- Require rigorous metric-based impact assessments before deploying AI systems for use by children
- Establish minimum flourishing metric thresholds that must be met for products targeting vulnerable populations
- Create ongoing monitoring requirements using flourishing metrics for high-risk applications
- Ban applications empirically shown by metrics to cause significant harm to vulnerable groups

These protective measures would move beyond speculative concerns to evidence-based regulation, using flourishing metrics to identify which applications truly benefit vulnerable populations and which may cause harm despite apparent conveniences.

8. Deploy Flourishing Metrics to Inform AI Economic Policy

Robust metrics will enable us to refine our social and economic policy to develop and deploy AI in a manner that supports flourishing. As AI transforms the workforce, we need strategic approaches that maintain the connection between meaningful work, productive contribution, and economic well-being.

- **Bake flourishing metrics into the incentive structure** of AI development by increasing the bargaining power of institutions and communities in the AI market
- **Pass The Data Freedom Act**, creating regulated entities which enable organizations to collectively negotiate with AI companies, empowering them to harness their data to align markets with flourishing metrics.⁹
- **Create democratic governance structures for AI deployment** that incorporate community input, enable regional customization, and empower communities to shape how AI integrates into their social fabric

By empowering communities and institutions to participate in technological development, we create a virtuous cycle that enhances both innovation and social cohesion. This approach ensures AI development responds to the needs of diverse communities, strengthening our resilience and competitive edge.

9. Preserve Human Connection and Invest in Communities

Metrics can highlight precisely how AI increases atomization and isolation. We should utilize their insights to directly invest in the communities, relationships, and organizations that make flourishing possible.

- **Protect the value of relational work** by creating safeguards to ensure roles centered on human connection, where human empathy and judgment are essential—such as teaching, healthcare, caregiving, counseling, and community organization—remain viable and valued
- **Invest in spaces for community, meaning, and creativity** where people can form connections, pursue fulfillment, and engage in exploration and creative expression
- **Support communities through technological transition** by helping local institutions adapt to economic transformation, preserving social capital and ensuring community needs are centered as technology reshapes social patterns

Through considering a whole-of-society approach, utilizing flourishing metrics to evaluate how AI is impacting communities holistically, we can get ahead of these transformations by strengthening community foundations before the most significant impacts arrive.

Addressing Potential Challenges

Innovation Concerns

Some may argue that additional governance frameworks will stifle innovation. However, our approach focuses on creating incentives and guidelines rather than imposing rigid requirements. By clearly articulating what constitutes flourishing- enhancing AI, we provide direction that can actually stimulate innovation toward more socially beneficial applications. Many investors and entrepreneurs are already seeking opportunities to develop AI systems that enhance rather than diminish human well-being. Ultimately, we believe that researchers and policymakers should work *with* the grain of American technology companies, providing them with metrics that allow them to build the next generation of AI products that promote the flourishing of Americans and their communities.

Measurement Difficulties

Measuring impacts on human flourishing is inherently challenging, particularly across diverse populations with different values and priorities. Quantification is, in some contexts, insufficient to capture the idiosyncratic, qualitative nature of our values. We address this by developing multidimensional metrics that can be adapted to different

contexts and by emphasizing the importance of ongoing research to refine measurement approaches.

Industry Resistance

Technology companies are rightly allergic to onerous regulation. But by involving industry in the development of metrics and creating market incentives through procurement policies, our approach aligns corporate interests with flourishing goals, reducing potential resistance. Our research has found that many AI firms are already interested in limiting their liability for negative social effects, suggesting an opening for collaborative approaches to governance.

Balancing Protection and Innovation

There is a risk that protective measures could be overly restrictive. We propose a balanced approach that differentiates between high-risk and low-risk applications. This allows for targeted protective measures in high-risk contexts while encouraging innovation in areas with less potential for harm.

Our Vision: Metrics-Driven AI Development for Human Flourishing

AI policy must do more than mitigate risks—it must actively strengthen the social fabric and enhance national resilience. Guided by rigorous metrics that measure its actual impacts on human flourishing, policymakers can ensure that AI serves human well-being in ways that reinforce relationships, meaning, and communal flourishing. By moving beyond speculative claims, we can provide both AI optimists and skeptics with a clear framework for distinguishing between technologies that genuinely enhance well-being and those that undermine it. In doing so, we can unleash American prosperity and national vitality, strengthening the institutions that sustain American life.

Flourishing metrics are not just one policy recommendation among many—they are the foundation for responsible AI governance and long-term national strength. Without them, AI development will proceed without accountability, risking economic dislocation and social fragmentation. With them, we can systematically assess which AI applications genuinely enhance human life across multiple dimensions. These metrics provide the necessary framework to ensure AI fortifies America’s social fabric—the foundation of our national strength.

By adopting these recommendations, the OSTP can position the United States as the global leader in human-centered AI governance. If America leads with metrics that embed human flourishing at the core of AI development, we can set a global standard

where AI-driven prosperity secures both our economic future and the foundations of a thriving society.

Author Bios:

Jonathan D. Teubner, Ph.D., is Research Associate at the Human Flourishing Program at Harvard University, where he leads the AI and Flourishing Initiative. In 2022, he co-founded FilterLabs.AI, a data analytics company that leverages artificial intelligence to source high-quality localized data in hard-to-reach regions of the world.

Sam Wolf, is a student at Harvard Law School, and regularly writes on topics of political economy, community, and meaning making.

Constantine ‘Kosti’ Psimopoulos, MBE, PhD/Dr.Bioethics, is a bioethicist, Senior Program Administrator and Instructor on Bioethics at the Human Flourishing Program at Harvard University, where he co-leads the Initiative on Health, Spirituality and Religion. He is a member of the scientific advisory committee on AI and Theology, of the Greek Orthodox Archdiocese of America. Kosti is an alumnus of Harvard Medical School.

Reece Brown is the Associate Director of Impact at the Human Flourishing Program at Harvard University. Previously, he worked with Arthur Brooks to launch the Leadership and Happiness Laboratory.

¹ Information about this project can be found at humanconnections.ai.

² Anton Korinek, “Economic Policy Challenges for the Age of AI.” National Bureau of Economic Research: Working Paper 32980 (September 2024).
https://www.nber.org/system/files/working_papers/w32980/w32980.pdf

³ “Long-Term Trends in Deaths of Despair.” Social Capital Project Report 4-19 (September 2019). A project of the Joint Economic Committee – Republicans.
https://www.jec.senate.gov/public/_cache/files/0f2d3dba-9fdc-41e5-9bd1-9c13f4204e35/jec-report-deaths-of-despair.pdf

⁴ Timothy P. Carney, *Alienated America: Why Some Places Thrive While Others Collapse*. HarperCollins (2019). A research project of the American Enterprise Institute; see <https://www.aei.org/research-products/book/alienated-america-why-some-places-thrive-while-others-collapse/>

⁵ Kevin Roose, “Can AI Be Blamed for a Teen’s Suicide?” *The New York Times*, Oct. 23, 2024 (<https://www.nytimes.com/2024/10/23/technology/characterai-lawsuit-teen-suicide.html>). More recent allegations about Character.AI bots encouraging murder and self-harm have also emerged; see, for example, a recent Texas news story from KVUE: <https://www.kvue.com/article/news/local/texas/character-ai-chatbot-lawsuit/269-9f20b8a1-6edb-4981-98f3-0e9e896295b7>

⁶ Information available at meaningalignment.org

⁷ Oliver Klingefjord, Ryan Lowe, and Joe Edelman, “What are human values, and how do we align AI to them?” arXiv:2404.10636v2 (April 17, 2024). <https://doi.org/10.48550/arXiv.2404.10636>

⁸ See <https://www.flourishing.foundation/> and <https://www.meaningalignment.org/>

⁹ For information on the proposed Data Freedom Act, see <https://www.radicalxchange.org/media/papers/data-freedom-act.pdf>