

From: [John McCone](#)
To: [ostp-ai-rfi](#)
Subject: [External] Request for Information on the Development of an Artificial Intelligence (AI) Action Plan
Date: Saturday, February 8, 2025 1:00:37 PM

CAUTION: This email originated from outside your organization. Exercise caution when opening attachments or clicking links, especially from unknown senders.

Dear Networking and Information Technology Research and Development (NITRD) National Coordination Office (NCO) worker,

Pertaining to Request for Information on the Development of an Artificial Intelligence (AI) Action Plan I would like to offer the following input relating to AI Safety:

- **Ensure that all powerful AIs are created with a strong time preference towards obeying recent orders over orders or instructions issued previously**

This will make the AI *want to correct* its behaviour. So that if the human issuing the prompt, realises that the actions that the AI is undertaking to implement the prompt are harmful the human can issue a new prompt to instruct the AI to correct its behaviour. If the AI has a strong time preference to obeying recent prompts over older prompts then the AI is more likely to respond to the recent prompt and cease its damaging behaviour as opposed to continuing the damaging behaviour in response to the previous prompt issued.

- **Define a human being very clearly and ensure that AIs have a strong preference to obeying human beings compared to obeying other AIs**

We want obedient AIs that perform orders as instructed. However, AIs can prompt each other, so even if every AI is programmed to obey orders there is still a danger of a circular arrangement emerging where a bunch of separate AIs are all ordering and prompting each other and where AIs pay more attention to the prompts issued by other AIs than they pay to the prompts issued by human beings. Thus, in order to ensure AI continues to serve humanity, it will be critical to ensure that we design an AI to clearly identify human beings, distinguish biological human beings from AI agents and have a preference for following the instructions given by human beings when compared to the instructions issued by other AI agents

- **Do not train super-intelligent AIs, responsible for top level military coordination at a strategic level, to kill human beings, or to achieve victory in war**

Starting a war is instrumentally necessary in order to win a war. Thus, if you train a super-intelligent AI system to generally find creative ways to win wars and kill enemy units, and if this super-intelligent system is placed in charge of controlling the US military at the highest strategic level (or even at fairly high level), then, during times of peace, such a military AI may search for creative ways to subtly, or not so subtly, initiate new wars with foreign countries, that are otherwise at peace with the U.S. as an instrumental objective to achieve its primary programming goal to win wars and kill enemy units - during times of peace you have to make enemies before you get to kill enemies

I've written more AGI safety suggestions here:

<https://johnmccone.com/2023/11/30/some-quick-and-nasty-solutions-to-ai-safety/>

John McCone - Technology Blogger and Philosophy Author

All e-mails to and from this account are for NITRD official use only and subject to certain disclosure requirements.

If you have received this e-mail in error, we ask that you notify the sender and delete it immediately.