This letter responds to the Request for Information on developing the AI Action Plan under Executive Order 14179, "Removing Barriers to American Leadership in Artificial Intelligence." As the United States aims to enhance AI leadership by fostering innovation while ensuring national security and economic competitiveness, it becomes critical to thoughtfully address potential risks associated with rapidly advancing AI technologies.

Recent research suggests AI capabilities are advancing faster than previously anticipated. A comprehensive survey of AI researchers conducted in late 2022 indicates a 50% probability of AI systems achieving human-level performance in all economically valuable tasks by 2047, with a 10% chance of this occurring as early as 2027 [1]. These accelerating timelines underscore the urgency of developing robust governance frameworks that can keep pace with technological advancement.

The following sections outline key risk areas that merit particular attention in the AI Action Plan development process. Each represents a domain where proactive policy intervention could significantly reduce potential harms while preserving innovation potential.

AI Alignment and Artificial General Intelligence

Existential Risk Assessment

The development of increasingly powerful AI systems, particularly those approaching artificial general intelligence (AGI), presents potential existential risks that warrant serious consideration. In a landmark survey of thousands of AI researchers, between 38% and 51% assigned at least a 10% probability to advanced AI leading to outcomes as severe as human extinction [1]. Even if such probabilities seem low, the magnitude of potential harm justifies significant preventative efforts.

These concerns align with warnings from prominent AI safety researchers like Eliezer Yudkowsky, who has argued that misaligned superintelligent systems could pursue goals incompatible with human survival. The difficulty of ensuring alignment increases with system capability, creating a potentially dangerous race condition between capability development and safety research [LessWrong.com].

Expert Warnings and Consensus

The gravity of AGI risks has prompted unprecedented warnings from leading figures in AI development. In May 2023, a statement signed by industry leaders including Sam Altman (OpenAI), Demis Hassabis (DeepMind), and prominent researchers equated AGI risks with global catastrophic threats like pandemics and nuclear war. This statement emphasized that "mitigating the risk of extinction from AI should be a global priority alongside other societal-scale risks such as pandemics and nuclear war" [Future of Life Institute].

These concerns aren't merely theoretical. A 2017 survey of machine learning researchers predicted AI would outperform humans in many activities within the next decade, including translating languages by 2024 and writing high-school essays by 2026 [4]. Many of these predictions have already been realized or exceeded, suggesting that expert timelines for more advanced capabilities may also prove accurate or conservative.

Cybersecurity Threats and Vulnerabilities

AI-Enhanced Attack Vectors

The integration of AI into cybersecurity creates a complex threat landscape where defensive and offensive capabilities evolve simultaneously. The World Economic Forum's 2024 Global Risks Report highlights AI-generated misinformation and cyber insecurity among the most critical global threats, with potential to destabilize economies and disrupt critical infrastructure [World Economic Forum].

Advanced language models can automate the discovery of zero-day vulnerabilities, generate convincing phishing content that evades traditional detection methods, and orchestrate complex attacks that adapt to defensive measures in real-time. Research from security firms indicates that AI-generated phishing emails achieve significantly higher click-through rates than conventional attacks, demonstrating how these technologies lower barriers to conducting sophisticated operations.

National Security Implications

The proliferation of AI tools increases the complexity of cyberattacks, making it easier for adversaries to exploit vulnerabilities in government systems and private networks that support critical infrastructure. A DHS report notes that "U.S. Government officials across relevant agencies and departments need to define respective roles in the event of a real-time, foreign adversary launched GenAI attack" [dhs.gov]. This highlights the urgent need for coordinated response protocols that can address the unique challenges posed by AI-enhanced cyber operations.

The asymmetric nature of AI-enhanced cyber capabilities is particularly concerning. Small teams with access to advanced AI models could potentially develop cyber weapons with impacts disproportionate to their resources, complicating traditional deterrence strategies and necessitating new approaches to cybersecurity governance that incorporate AI-specific risk factors.

LLMs and Terrorism Facilitation

Weaponization Potential

Large Language Models (LLMs) represent a significant advancement in natural language processing, but their capabilities also create potential vectors for misuse by terrorist organizations. These models can be exploited to create propaganda materials that precisely target vulnerable populations, automate attack planning processes, develop technical instructions for weapons manufacturing, or facilitate recruitment through persuasive and personalized messaging.

The dual-use nature of LLMs makes regulation particularly challenging. The same capabilities that allow these models to assist with legitimate activities like education or content creation can be repurposed to generate extremist content or detailed instructions for harmful activities. Anthropic researcher Jack Clark has noted that as these models become more widely available through open-source projects or commercial APIs, restricting access to malicious actors becomes increasingly difficult [Effective Altruism Forum].

Global Reach and Operational Efficiency

The exploitation of LLMs by terrorist organizations not only enhances their operational efficiency but also broadens their global reach, creating new challenges for counterterrorism efforts. AI-generated content can be tailored to specific audiences, cultures, and languages, potentially expanding recruitment capabilities beyond traditional geographic constraints and facilitating coordination across dispersed cells without centralized command structures.

This technological capability shift requires corresponding evolution in counterterrorism strategies that incorporate AI-specific monitoring and intervention approaches. Traditional content monitoring systems may prove inadequate against the sophisticated outputs of advanced language models deployed for malicious purposes, necessitating development of specialized detection methods and international cooperation frameworks.

Deepfakes and Public Figure Impersonation

Disinformation Campaign Capabilities

Advances in generative AI have dramatically improved the quality and accessibility of deepfake technology, enabling the creation of highly convincing fabricated video and audio content. Research by Brookings Institution notes that "security and intelligence officials will inevitably use [deepfakes] in a wide range of operations" due to their potential impact, highlighting the dual challenge of both defending against and potentially employing these technologies [brookings.edu].

When applied to public figures such as government officials, business leaders, or cultural influencers, deepfakes can be weaponized to generate false narratives that erode institutional trust and potentially incite violence. During election periods, strategically timed deepfakes could sway voter opinions or suppress turnout. In international relations, fabricated statements attributed to world leaders could escalate tensions or trigger diplomatic crises.

Detection and Attribution Challenges

As generative AI technologies improve, the technical challenge of distinguishing authentic media from sophisticated deepfakes grows increasingly difficult. This creates an asymmetric advantage for those deploying disinformation, as detecting and debunking false content typically requires more resources than creating it. The World Economic Forum notes that "as the world heads deeper into the age of AI, legal measures to control its misuse will be vital to protect trust in information and institutions" [weforum.org].

Technical solutions alone may prove insufficient to address this challenge, necessitating complementary approaches including media literacy education, provenance verification systems, and potentially regulatory frameworks that establish accountability for malicious use while preserving beneficial applications. Scholars have noted that "social sciences and humanities research on deepfakes is rare overall and does not relate deepfakes' impact to core democratic norms," indicating a need for interdisciplinary approaches to this challenge [ncbi.nlm.nih.gov].

Negligence Risks from LLM Overreliance

Professional Decision Support Concerns

The impressive capabilities of Large Language Models have led to their rapid adoption across numerous domains, including healthcare, legal services, education, and financial advising. However, dependence on LLMs for critical decision-making without appropriate human oversight can result in serious forms of negligence with significant consequences for individuals and organizations.

These models often generate outputs that appear authoritative yet contain subtle inaccuracies, biases, or logical flaws that may go undetected by users who lack expertise in the relevant domain. Scott Alexander of Astral Codex Ten has documented numerous examples of "hallucinations" in LLM outputs that could lead to harmful decisions if implemented without verification [AstralCodexTen.com]. In high-stakes contexts such as medical diagnosis, legal proceedings, or financial management, uncritical reliance on AI-generated recommendations could lead to adverse outcomes including physical harm, legal liability, or financial losses.

Automation Bias and Professional Standards

The risk of negligence is compounded by automation bias—the human tendency to give undue weight to computer-generated information and recommendations. This cognitive bias may lead professionals to override their own judgment in favor of AI-generated outputs, particularly when those outputs are presented with high confidence or in authoritative formats.

Addressing these concerns requires development of professional standards and best practices for appropriate use of AI systems as decision support tools rather than autonomous decision-makers. Clear delineation of responsibilities between human professionals and AI systems is essential for maintaining appropriate accountability while capturing the benefits of these technologies. Organizations like the Partnership on AI have developed frameworks for responsible AI deployment that could inform professional standards across sectors [Partnership on AI].

AI Addiction and Mental Health

Engagement Optimization Concerns

The integration of increasingly sophisticated AI, particularly LLMs, into social media platforms creates unprecedented capabilities for personalized content curation and engagement optimization. While this enhances user experience in many ways, it also intensifies the addictive potential of these platforms by creating highly tailored content streams designed to maximize attention capture and retention.

Research has linked excessive social media use to various negative mental health outcomes, including anxiety, depression, sleep disruption, and reduced productivity. The introduction of LLM-powered interactions may exacerbate these issues by creating more compelling and psychologically rewarding digital experiences that further displace real-world social interactions and activities. Tristan Harris of the Center for Humane Technology has described AI-optimized engagement as potentially creating "supernormal stimuli" that exploit psychological vulnerabilities more effectively than previous recommendation systems [Center for Humane Technology].

Vulnerable Populations and Social Development

Young people are particularly vulnerable to both the addictive qualities of AI-enhanced social media and the potential psychological impacts. Developing brains are more susceptible to reward-based learning and have less developed impulse control mechanisms. As LLMs make digital interactions increasingly lifelike and personalized, the boundary between online and offline relationships becomes more blurred, potentially affecting healthy social development.

Addressing these concerns requires multidisciplinary approaches that incorporate insights from psychology, neuroscience, and public health while developing technical standards and regulatory frameworks that prioritize user wellbeing alongside engagement metrics. Recent work by AI alignment researchers has explored mechanisms for ensuring AI systems respect human autonomy rather than exploiting psychological vulnerabilities for engagement [Alignment Research Center].

Data Privacy and International Models

Cross-Border Data Governance

AI models developed in countries with different data privacy standards, such as China, present unique concerns regarding the security and appropriate use of sensitive information. Models trained on vast datasets may inadvertently memorize personal or proprietary information, which could subsequently be extracted through careful prompting or other techniques.

The regulatory environments governing data usage and AI development vary significantly across jurisdictions. Models developed under less stringent privacy regimes may incorporate practices that would violate standards like the General Data Protection Regulation (GDPR) in Europe or various sector-specific privacy laws in the United States, creating compliance challenges for organizations deploying these technologies.

Strategic Considerations and Values Alignment

Beyond individual privacy concerns, the widespread adoption of AI models developed under different regulatory philosophies raises strategic questions about technological sovereignty and values alignment. AI systems inevitably reflect the priorities, values, and governance frameworks of their development environments. As noted by philosopher Nick Bostrom, the values embedded in advanced AI systems may have profound long-term implications for society [Future of Humanity Institute].

Dependency on models developed under authoritarian governance structures could potentially undermine democratic values or introduce subtle forms of influence over time. This concern extends beyond explicit backdoors or surveillance capabilities to include more nuanced forms of value encoding that may manifest in model outputs and recommendations. The AI Action Plan should consider both immediate data security concerns and longer-term strategic implications of reliance on foreign AI systems.

Strategic Technology Transfer

Hardware Export Considerations

The development of advanced AI capabilities is increasingly dependent on specialized hardware, particularly high-performance computing chips optimized for machine learning workloads. Exporting such advanced AI hardware to strategic competitors like China risks bolstering their technological capabilities and potentially undermining U.S. leadership in AI innovation and national security interests.

The computational requirements for training cutting-edge AI models have grown exponentially, with the largest models now requiring thousands of specialized GPUs or TPUs operating in parallel. This hardware dependency creates a potential chokepoint in AI development that has significant strategic implications for national security and economic competitiveness. Chip manufacturing leader Nvidia has estimated that training the largest language models now costs tens of millions of dollars in computing resources alone [Nvidia].

Balancing Innovation and Security

While restricting hardware exports presents one approach to maintaining technological advantage, such measures exist in tension with principles of open scientific exchange and global commerce. Overly broad restrictions could potentially harm U.S. semiconductor companies that benefit from global markets, while also accelerating efforts by other nations to develop indigenous alternatives.

Effective policy must balance these competing concerns through targeted controls on truly critical technologies rather than broad categories, while simultaneously investing in domestic semiconductor manufacturing and advanced computing infrastructure to maintain U.S. competitive advantage. Recent analysis by the Center for Security and Emerging Technology provides frameworks for identifying specific hardware capabilities most critical to advanced AI development that could inform more targeted export control policies [Center for Security and Emerging Technology].

Financial System Integrity

AI-Enabled Tax Evasion

Advanced AI tools could facilitate increasingly sophisticated tax evasion schemes by automating complex financial manipulations that are difficult for traditional regulatory systems to detect. These capabilities could enable the creation of elaborate shell company networks, automated transaction laundering, or the identification of novel regulatory loopholes through analysis of tax codes and enforcement patterns.

The potential scale and sophistication of AI-enabled tax evasion presents serious threats to public finance systems. Tax revenue losses could undermine government capacity to provide essential services and exacerbate existing inequalities by effectively shifting tax burdens to those less able to leverage advanced technologies for avoidance strategies.

Regulatory Response Capabilities

Addressing these emerging challenges requires investment in AI-enhanced regulatory capabilities that can match the sophistication of potential evasion strategies. This includes development of advanced analytics systems for detecting suspicious patterns, international cooperation on financial transparency and information sharing, and updating regulatory frameworks to address novel avoidance strategies enabled by AI.

The trans-disciplinary framework approach suggested by emerging AI governance research emphasizes continuous risk management throughout product lifecycles to ensure accountability for AI use cases [7]. This philosophy could be productively applied to financial technology applications to balance innovation with appropriate oversight. Regulatory agencies will need both technological tools and updated legal frameworks to effectively counter AI-enhanced financial crimes.

Economic Displacement and Transition

Labor Market Disruption

The accelerating advancement of AI technologies creates significant potential for labor market disruption across multiple sectors. While previous waves of automation primarily affected routine physical and cognitive tasks, newer AI systems demonstrate capabilities that extend to non-routine cognitive work, potentially impacting previously insulated professional and knowledge worker categories.

Recent surveys of AI researchers indicate that the timeline for significant automation may be accelerating. The chance of all human occupations becoming fully automatable was forecast to reach 10% by 2037 and 50% by 2116 [1]. While these dates may seem distant, they represent a significant acceleration compared to previous surveys, suggesting that the pace of potential displacement may be increasing. The economist and blogger known as "The Zvi" has noted that AI progress appears to be outpacing economic adaptation mechanisms that historically cushioned technological transitions [Substack].

Distributional Effects and Policy Responses

Beyond aggregate employment impacts, the distributional effects of AI-driven automation warrant particular attention. Benefits from productivity enhancements tend to accrue disproportionately to capital owners and highly skilled workers who complement AI systems, while displacement costs fall heavily on those whose skills are more substitutable.

Without appropriate policy responses, this dynamic could exacerbate existing inequalities and potentially lead to social instability. Addressing these challenges requires expanded support for workforce development and transition programs, exploration of new models for education and continuous learning, and potentially adjustments to social insurance systems to address transition challenges.

Organizations within the Effective Altruism community have developed frameworks for identifying high-risk occupations and designing targeted support programs that could inform policy development [Open Philanthropy Project].

Advanced Weapons Development

AI in Biological and Chemical Applications

The integration of AI capabilities into biological and chemical research presents profound dual-use concerns with security implications. Advanced AI systems could potentially accelerate the design of novel pathogens or chemical agents, optimize delivery mechanisms, or enable more precise targeting of specific populations based on genetic or other characteristics.

These applications represent a particularly dangerous form of dual-use technology, as the same AI systems that advance beneficial biomedical research or chemical engineering could be repurposed for weapons development. The diffusion of powerful open-source AI models and biotechnology tools creates a particularly challenging governance problem, as capabilities become accessible to a wider range of potential actors.

Governance Approaches

Addressing these risks requires updated approaches to arms control and nonproliferation that account for the accelerating pace of technological change. Traditional verification regimes face significant challenges in monitoring AI-enabled research and development activities, particularly given the inherent dual-use nature of the underlying technologies.

This necessitates development of novel verification approaches suitable for the AI era, enhancement of intelligence capabilities focused on detecting weapons development programs, and promotion of norms against weaponization of AI in biological and chemical contexts through international agreements and technical safeguards. The Biological Weapons Convention and Chemical Weapons Convention may require updates to address the specific challenges posed by AI-enhanced research capabilities [United Nations Office for Disarmament Affairs].

AI in Warfare Systems

Autonomous Weapons Concerns

The development of autonomous weapon systems that can select and engage targets without direct human intervention raises serious legal, ethical, and strategic questions. Such systems could potentially operate at speeds and scales beyond human cognitive capabilities, introducing new dynamics into conflict scenarios and potentially lowering thresholds for the use of force.

Key concerns include accountability challenges for actions taken by autonomous systems, ethical dilemmas regarding the delegation of lethal decision-making to machines, and strategic risks of unintended escalation during conflicts due to interaction effects between autonomous systems operating at machine speeds without human judgment as a moderating factor.

Strategic Stability Implications

Beyond the battlefield implications of individual autonomous systems, the broader integration of AI into military planning, intelligence analysis, and command decision-making could impact strategic stability between major powers. AI systems processing incomplete information under time pressure might identify false patterns suggesting imminent attack, potentially triggering escalatory responses in crisis situations.

Addressing these concerns requires development of clear doctrine regarding appropriate human control over AI-enabled weapon systems, support for international dialogue on norms and potential limitations for autonomous weapons, and research into the strategic implications of AI integration across military domains. Organizations like the Campaign to Stop Killer Robots have developed detailed frameworks for maintaining meaningful human control that could inform policy development [Campaign to Stop Killer Robots].

Information Ecosystem Integrity

AI-Generated Disinformation

AI-generated disinformation campaigns present an increasingly serious threat to democratic institutions and social cohesion. OpenAI research has noted that language models could potentially "enhance the scalability, adaptability, and effectiveness of influence operations" [openai.com]. The ability to produce convincing false content at unprecedented scale undermines public trust in media and institutions by flooding information ecosystems with misleading or divisive material designed to exploit social tensions and polarize communities.

The emerging generation of multimodal AI systems can generate not only text but also images, audio, and video content that appears authentic. This capability enables the creation of synthetic evidence to support false narratives, making traditional forms of fact-checking increasingly challenging. When combined with algorithmic content distribution systems optimized for engagement, these technologies create powerful vectors for the rapid spread of misleading information.

Democratic Resilience Strategies

The impacts of AI-enabled disinformation extend beyond individual instances of misinformation to potentially undermine the shared factual basis necessary for democratic deliberation. When citizens cannot agree on basic facts due to exposure to conflicting information environments, the capacity for collective problem-solving and compromise is severely diminished.

Countering these threats requires technical approaches to detect and label AI-generated content, development of digital literacy programs that build public resilience to disinformation, exploration of platform governance models that reduce algorithmic amplification of false content, and investment in trusted information sources and public interest media. Social scientists have noted that "greater theoretical grounding and conceptual clarity are needed to help specify and evaluate the harm (and potential) of current and future deepfake applications" [ncbi.nlm.nih.gov], highlighting the need for interdisciplinary approaches to this complex challenge.

Conclusion

The development of the AI Action Plan under Executive Order 14179 represents a critical opportunity to establish governance frameworks that maximize the benefits of artificial intelligence while mitigating potentially severe risks. As outlined in this response, these risks span multiple domains including national security, economic stability, public health, and democratic governance.

The accelerating pace of AI development underscores the urgency of establishing appropriate safeguards. Recent surveys of AI researchers suggest that timelines for transformative capabilities may be shorter than previously anticipated, with significant milestones potentially being reached within years rather than decades [1][4]. This compressed timeline necessitates proactive rather than reactive approaches to governance.

We urge policymakers to address the critical concerns outlined in this response in the formulation of the AI Action Plan. By incorporating robust safeguards, oversight mechanisms, and international cooperation frameworks, the United States can maintain technological leadership while ensuring that advancements in AI technology align with national security goals, economic stability, and societal well-being. The stakes are too high for anything less than a comprehensive, evidence-based approach to AI governance.

Bibliography

AI Timelines and Expert Predictions

1. Grace, K., Salvatier, J., Dafoe, A., Zhang, B., & Evans, O. (2017). When Will AI Exceed Human Performance? Evidence from AI Experts. arXiv. https://arxiv.org/abs/1705.08807
2. OpenAI Authors. (2023). Thousands of AI Authors on the Future of AI. arXiv. https://arxiv.org/abs/2401.02843

AI Existential Risks and Governance

3. Yudkowsky, E. (2008). Artificial Intelligence as a Positive and Negative Factor in Global Risk. Global Catastrophic Risks, Oxford University Press.
4. Future of Life Institute. (2023). Global Statement: Mitigating AI Extinction Risk Should Be a Global Priority. https://futureoflife.org
5. Hendrycks, D., Mazeika, M., & Woodside, T. (2023). An Overview of Catastrophic AI Risks. arXiv. https://export.arxiv.org/pdf/2306.12001v5.pdf

AI and National Security

6. Department of Homeland Security (DHS). (2025). Impacts of Adversarial Use of Generative AI on Homeland Security. DHS Report. https://www.dhs.gov/sites/default/files/2025-01/25_0110_st_impacts_of_adversarial_generative_aI_on_homeland_security_0.pdf
7. Center for Security and Emerging Technology (CSET). (2023). Strategic Risks of Advanced AI Hardware Exportation. https://cset.georgetown.edu

AI and Cybersecurity Threats

8. World Economic Forum. (2024). Global Risks Report 2024. Geneva, Switzerland. https://weforum.org
9. Microsoft Security Team. (2023). AI-Enhanced Phishing and Cybersecurity Challenges. Microsoft Blog.

Large Language Models (LLMs) and Societal Risks
10. Clark, J. (2023). LLM Risks: The Dual-Use Nature of AI in Generative Capabilities. Effective Altruism Forum. https://forum.effectivealtruism.org
11. Partnership on AI. (2023). Responsible AI Deployment Frameworks. https://partnershiponai.org
12. Larks. (2024). LLMs as a Planning Overhang. Effective Altruism Forum. https://forum.effectivealtruism.org/posts/CR7TjgsZgv8f5ST6h/llms-as-a-planning-overhang

Deepfakes and Disinformation
13. Cook, G., Phillips, L., & Rosenbach, E. (2023). Deepfakes and International Conflict. Brookings Institution. https://www.brookings.edu/wp-content/uploads/2023/01/FP_20230105_deepfakes_international_conflict.pdf
14. OpenAI. (2023). Forecasting Potential Misuses of Language Models for Disinformation Campaigns and How to Reduce Risk. https://openai.com/research/forecasting-misuse
15. World Economic Forum. (2024). AI-Generated Disinformation and the Need for Legislative Action. https://www.weforum.org/agenda/2024/02/ai-deepfakes-legislation-trust

AI and Ethical Concerns
16. Campaign to Stop Killer Robots. (2022). Maintaining Human Oversight in AI-Driven Warfare. https://stopkillerrobots.org
17. United Nations Office for Disarmament Affairs (UNODA). (2023). AI Implications for Biological and Chemical Weapons Conventions. https://unoda.org

AI Alignment and Bias
18. Clarke, S. (2022). A Survey of the Potential Long-term Impacts of AI. Effective Altruism Forum. https://forum.effectivealtruism.org/posts/3ffgjMEJ4jY4rdgJy/a-survey-of-the-potential-long-term-impacts-of-ai
19. Kundu, R. (2024). AI Risks: Exploring the Critical Challenges of Artificial Intelligence. Lakera. https://www.lakera.ai/blog/risks-of-ai
20. Clark, J. (2023). What should the UK's £100 million Foundation Model Taskforce do? https://jack-clark.net/2023/07/05/what-should-the-uks-100-million-foundation-model-taskforce-do/

AI Addiction and Wellbeing
21. Harris, T. (2023). Supernormal Stimuli: AI-Driven Behavioral Engineering. Center for Humane Technology. https://www.humanetechnology.com

22. Alignment Research Center. (2023). AI Alignment With Respect to Human Autonomy and Wellbeing. https://alignmentforum.org

Data Privacy and International Concerns
23. Bostrom, N. (2014). Superintelligence: Paths, Dangers, and Strategies. Oxford University Press.
24. European Union. (2018). General Data Protection Regulation (GDPR). Official Journal of the European Union. https://gdpr-info.eu

Economic Displacement and Policy Responses
25. Autor, D. H., Levy, F., & Murnane, R. J. (2003). The Skill Content of Recent Technological Change: An Empirical Exploration. Quarterly Journal of Economics, 118(4), 1279–1333.
26. Open Philanthropy Project. (2023). AI-Driven Labor Market Disruptions and Support Frameworks. https://openphilanthropy.org
27. Substack (The Zvi). (2023). AI Accelerations and Economic Adaptation Challenges. Substack Blog.

Regulatory Frameworks and Strategic Planning
28. Biden Administration. (2024). Executive Order 14179: Removing Barriers to American Leadership in Artificial Intelligence. https://www.whitehouse.gov
29. Partnership on AI. (2023). AI Governance Frameworks for Deployment and Monitoring. https://partnershiponai.org