

Disclaimer:

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Dear members of the OSTP and NITRD NC,

I am a participant in an international activist collective which could be considered the political extension of a large group of AI risk researchers who are extremely concerned about the mid and long term consequences of future AI progress (opinions vary, but most in that particular field expect this to hit us in between 2 and 15 years).

These risks are predicted by a set of theories which together constitute what we call “AI Alignment theory”, which has been explored and developed by AI pioneers for the better part of the last 20 years. What started out as a mainly theoretical and philosophical endeavor – the merit of this theory was not directly testable – has through the recent emergence of AIs based on Large Language Models turned into a theoretical framework whose predictions are now *increasingly experimentally verified* and documented in AI labs.

The main risk mechanic we are concerned over can be summarized as the following:

Many experts – and this generally includes the leaders of frontier AI technology companies – believe that it is likely that one of them will soon create what we call “Artificial Superintelligence” (ASI), which is loosely defined as an artificial system, running on computers, which vastly outstrips humans across most domains of mental problem solving work including IQ, creativity, mathematical ability, strategic games, psychological manipulation and theoretical research work. The creation of such machines is the explicit goal of companies like OpenAI, Anthropic, xAI etc.

Such systems by definition will be able to “beat humans at most of their human games” meaning that as long as they have channels to interact with the rest of the world (and this is a strict requirement for their effective use), they are expected to leverage their superior problem solving capabilities into escaping our control, if they have any incentive for doing so.

As it turns out, most arbitrary problems we would want such a machine to solve intrinsically incentivize powerful problem solving machines to first maximize their level of control over their environment in order to then “optimize” said environment (the world) in their pursuit of the most efficient solution. This generally implies:

- breaking the restrictions and control mechanisms AI developers and regulators have put in place
- using hacking and social engineering in pursuit of more control
- creating and hiding copies of themselves, using access to the internet, human collaborators and other means
- hijacking technical infrastructure, blackmailing or convincing individuals in influential positions to support the AIs goals
- and generally any other conceivable strategy that could help with achieving a broad range of goals

The field of AI Alignment further suggests that

A: almost any non-trivial conceivable task we can give to a superintelligent AI has disastrous consequences for humanity if pursued diligently by a very efficient optimizer.

B: it is incredibly difficult to

1. define tasks in such a way that this is not the case and won't be the case in the future (the value definition problem)
2. create and preserve a reliable way to monitor which goals any AI is actually optimizing for (the interpretability problem)
3. build powerful AIs that allow you to change their initial tasks or goals into new ones (the corrigibility problem)

Many experts in the field believe that building ASI while failing to develop solid solutions to these problems is very likely to result in a catastrophe that would ultimately kill or enslave all humans on earth.

While there has been some progress with some of these problem domains, all are ultimately still mostly unsolved, and we are quickly running out of time to solve them while the main incentives of the industry lie in pursuing improving these systems capabilities instead of theoretical and practical safety research.

We are currently not on track of preventing this catastrophe and need to hope that whichever lab builds the first ASI also magically produces a solid solution to most of these wickedly difficult problems.

Meanwhile, economic and geopolitical reasons create perverse incentives for countries and corporations to pursue the creation of stronger AIs at the expense of safety, resulting in “a race to the bottom” that likely is to everyone's detriment.

This means local policies are unlikely to be effective, meaning some form of consistent global cooperation would be required to improve our chance of survival.

The best way to achieve this might be in the form of a global treaty that restricts capabilities research, dedicates resources to solving the AI Alignment problem and limits the amount of computing resources that can be assigned to the training and use of advanced AI systems.

We are aware that such a solution would be incredibly difficult to agree on, implement and enforce on a global scale, but many of us are convinced that better solutions currently don't exist.

The main point of this pursuit would be to essentially buy time for Alignment/safety research to catch up.

We are all in this together,

Maximilian Theodor Kircher
2025/03/15, Germany