

Nathaniel Sharpe

The rapid development of artificial intelligence has brought immense potential for scientific progress, economic growth, and societal benefit. However, it also carries unprecedented risks—none more critical than the development of artificial superintelligence (ASI), which, if misaligned or uncontrolled, could result in catastrophic outcomes. The current trajectory of AI development is reckless, competitive, and uncoordinated. If left unchecked, the race to achieve ASI will not reward the "winner"—it will simply bring forward the moment when humanity loses control. We urge the AI Action Plan to prioritize the following policy interventions:

- 1. Prohibit the Development of Artificial Superintelligence and Self-Improving AI** The most immediate and non-negotiable safeguard is a global ban on the development of ASI and AI systems capable of recursively improving themselves. AI models that can autonomously enhance their own intelligence pose an existential risk because they could escape human oversight, develop unforeseen capabilities, and alter their own objectives. Current AI developers cannot reliably predict the capabilities of the systems they train before deployment, and post-hoc evaluations are insufficient to ensure safety. We must explicitly prohibit AI that: Improves its own architecture or code without human oversight. Possesses capabilities that exceed general human intelligence across a broad set of domains. Can break out of its training or operational environment. Just as the international community has established frameworks to prevent the proliferation of nuclear and biological weapons, we must categorically prevent the creation of ASI until proven control mechanisms exist.
- 2. Establish a Licensing System for AI Development and Deployment** AI systems should not be developed in a regulatory vacuum. The AI Action Plan must implement a strict licensing regime for both AI models and the compute infrastructure that supports them:
 - Training Licenses:** No AI model above a designated capability threshold should be trained without explicit government approval and oversight.
 - Compute Licenses:** Data centers and cloud computing providers must be required to track AI workloads and report compliance with safety regulations.
 - Application Licenses:** Deploying AI systems in high-risk domains must require ongoing evaluation to ensure safety constraints remain effective over time.A failure to implement such oversight mechanisms will allow AI capabilities to scale unchecked, outpacing regulatory and security measures.
- 3. Shift from Competitive AI Development to Cooperative International Safeguards** A core assumption in recent AI policy rhetoric is that the United States must "win" the AI race against China. This framing is dangerously misguided. AI safety is not a zero-sum game—an uncontrollable ASI developed in any country is an existential threat to all of humanity. The AI Action Plan should:
 - Engage in direct diplomatic efforts** to ensure China, the EU, and other AI-leading nations commit to a coordinated slowdown in frontier AI development.
 - Establish a global AI safety agency** with verification mechanisms similar to those used in nuclear arms control, including:
 - Compute tracking via semiconductor manufacturers.
 - International inspections of AI research labs and data centers.
 - Energy monitoring to detect unauthorized AI training runs.It is crucial that policymakers recognize that competition fuels reckless acceleration. The only way to ensure AI safety is through an enforceable, internationally verifiable slowdown—not through an arms race that increases the risk of losing control.

Conclusion The most dangerous assumption we can make is that we can "safely" continue to advance AI without binding international safeguards. The risks of ASI are not merely hypothetical; leading AI researchers and industry executives themselves have publicly acknowledged the lack of control mechanisms for frontier models. The AI Action Plan must prioritize prohibition, licensing, and international cooperation over laissez-faire development.

Once the intelligence and capability of AI surpasses human oversight, we do not get a second chance to correct our mistakes. Thank you for considering these recommendations. Sincerely,
Nate Sharpe