

Policy Recommendation: Adopting the Self-Alignment Framework (SAF) for AI Ethics Governance

Introduction:

The rapid advancement of artificial intelligence presents both transformative opportunities and critical ethical challenges for the United States. To maintain U.S. leadership in AI while safeguarding public values, we recommend the adoption of the Self-Alignment Framework (SAF) as a governance model for AI ethics. SAF is an open-source, structured system for aligning intelligent behavior with human values, designed to ensure AI systems self-regulate their conduct in an ethical manner. By embedding a continuous closed-loop alignment process within AI development and deployment, SAF addresses issues like algorithmic bias, misinformation propagation, and value drift without imposing excessive external regulatory burdens. This proposal outlines how SAF works, how it mitigates key AI risks, and how it aligns with U.S. policy goals of responsible innovation and global AI leadership.

SAF Overview: A Self-Regulating, Closed-Loop Framework for Long-Term Alignment

In any complex system—whether mechanical, biological, or conceptual—maintaining alignment with a guiding principle requires a robust feedback loop. The Self-Alignment Framework (SAF) employs just such a loop, ensuring that individuals or organizations consistently act in accordance with their deepest values.

1. Core Principles of a Closed-Loop System:

In traditional control theory, a closed-loop system monitors its output (results) and uses that information to steer the system toward a target state, known as the set point. The Self-Alignment Framework does something very similar:

- A **set point (Values)** defines the target state.
- A **controller (Intellect + Will)** interprets feedback and decides how to act.
- **Sensors (Conscience + Spirit)** measure alignment between action and the set point, both short-term and long-term.
- Feedback is looped back to the controller, prompting adjustments or corrections whenever needed.

What makes SAF distinctive is its human-centered application. It translates the mechanics of a feedback loop into an ethical and introspective domain, ensuring that

moral principles are not only clearly defined but also actively maintained.

2. Values: The Guiding North Star:

At the heart of SAF lies a set of **Values**—the foundational ethical principles that define what “alignment” means. Think of Values as the system’s compass:

- Purpose: They determine what the “correct” or desired state should look like.
- Stability: They provide consistency and clarity, serving as the reference point against which every action is measured.
- Moral Anchor: They keep the framework grounded, preventing it from drifting arbitrarily in response to external pressures or fleeting impulses.

In SAF, Values can mature over time through insights from the **Spirit** component, making the framework adaptive when genuinely needed.

3. Intellect: Analysis and Discernment:

Intellect is the analytical engine that takes in information, compares it to the Values, and makes reasoned judgments on how to act:

- Observation & Synthesis: It gathers relevant facts, weighs possible outcomes, and interprets any feedback from the system.
- Comparison to Values: Intellect scrutinizes various courses of action against the established Values, looking for potential mismatches or ethical conflicts.
- Decision Making: Once the analysis is complete, Intellect formulates the most value-aligned option.

4. Will: From Decision to Action:

Where Intellect is analytical, **Will** is action-oriented:

- Execution: Will “pushes the button,” transforming decisions into tangible steps and real-world effects.
- Commitment: It also represents the tenacity to see actions through, even in the face of obstacles.
- Ethical Fortitude: By upholding the Values in practice, Will ensures that ethical principles do not stay theoretical but become lived reality.

5. Conscience: The Immediate Feedback Sensor:

Conscience serves as the short-term feedback sensor:

- Real-Time Check: Conscience evaluates whether the action just taken aligns with the Values.
- Internal Signal: It may manifest as a sense of peace when actions are consistent with Values—or discomfort, guilt, and uncertainty when they are not.
- Motivator for Correction: If there is misalignment, Conscience alerts both Intellect

and Will to reconsider or adjust future actions.

6. Spirit: Long-Term Feedback and System Evolution:

Spirit scans the horizon of time:

- **Long-Range Monitoring:** Spirit looks beyond isolated decisions, identifying patterns of behavior to ensure sustained ethical integrity.
- **Early Warning System:** Repeated small misalignments can accumulate and lead to “ethical drift.” Spirit’s long-term perspective can detect these subtle shifts before they become entrenched.
- **Adaptive Refinement:** In cases where the Values themselves might be incomplete or need updating, Spirit can trigger a re-evaluation of those very principles.

Addressing Key AI Challenges: Bias, Misinformation, and Value Drift:

SAF directly tackles several priority AI ethics issues through its design, using the closed loop system to correct errors.

- **Mitigating Bias:** SAF helps AI systems recognize and correct biased decision pathways by continuously referencing core values like fairness and nondiscrimination.
- **Countering Misinformation:** When AI reasoning is compromised by false or misleading information, SAF’s Intellect and Conscience work in tandem to identify inconsistencies with truth-oriented values.
- **Preventing Value Drift:** SAF is expressly designed to combat this through its dual feedback layers. The Conscience component offers immediate course correction when minor misalignments occur, while the Spirit component provides continuous long-term oversight.

Governance Without Excessive Regulatory Burden:

A major advantage of SAF as an AI ethics governance model is that it can be implemented as a self-regulatory system, reducing the need for heavy-handed external regulation.

- **Structured Ethical Governance:** SAF provides a universal framework for decision-making and governance that organizations can adopt internally.
- **Reduced Compliance Costs:** By adopting SAF, businesses can streamline their ethical compliance efforts.
- **Certification and Standards:** We also propose establishing a SAF certification program in partnership with standards bodies.

Benefits of SAF for Stakeholders:

- **AI Developers & Researchers:** SAF provides developers with a clear ethical design template.
- **Businesses & AI Providers:** Companies that adopt SAF stand to gain a

competitive edge in trust and compliance.

- **Policymakers & Regulators:** From a governance perspective, SAF offers a scalable oversight model.

Alignment with U.S. AI Policy Goals:

Adopting SAF as part of the national AI Action Plan will directly support the United States' twin objectives of maintaining AI leadership and promoting responsible innovation.

Recommendations and Conclusion:

1. Integrate SAF into the National AI Action Plan.
2. Encourage Industry Adoption via Incentives.
3. Develop Certification and Oversight Mechanisms.
4. Ongoing Monitoring and Research.

Conclusion:

By adopting the Self-Alignment Framework, the United States can establish a forward-looking governance model that marries innovation with responsibility. SAF represents a pragmatic approach to AI ethics – one that leverages self-governance within AI systems to uphold our values and laws. We urge policymakers to include SAF in the national AI strategy as a means to foster responsible AI innovation at scale, ensuring a future where AI systems act as reliable partners in advancing economic prosperity, social well-being, and the values we cherish as a nation.

For more details and documentation visit the <https://selfalignmentframework.com> website.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.