

NCO Office of Science and Technology Policy
Executive Office of the President
2415 Eisenhower Avenue Alexandria, VA 22314

Subject - Request for Information (RFI) on the Development of an Artificial Intelligence (AI) Action Plan

Date: 15 March 2025

Organization Name: Privasapien Technologies Private Limited

Contributor: Abilash Soundararajan & Alen Pious

Submitted by email to [REDACTED]

Introduction

Generative AI is rapidly progressing from prompt response to planning to reasoning and progressing towards general intelligence. With countries across the globe beginning to build their own open source models, making performance a given, the most privacy protecting, safe and secure will be a critical differentiation and have the highest adoption. A key factor for global leadership will be the emphasis on data privacy and security throughout the entire lifecycle of AI systems, especially with the fast advancement of AI with AGI looming in the horizon. This submission underlines the urgent need for the development of AI regulation with a focus on data privacy, security, and risk management, which can also provide high ROI for businesses globally. This can be critical differentiator compared to Chinese models.

Building AGI with critical principles like privacy, safety and security will enhance global user trust, thus increasing adoption of models globally, resulting in accelerated leadership and differentiation from competitors. Standards should be established for providing verifiable proofs of compliance at data collection, data pre processing, model training and inference levels. This proposal suggests U.S. to adopt a technology and solution-centric approach to ensure privacy and security are embedded throughout AI product and model development, which is also a wise business decision to gain user trust for global adoption. With the cost of incorporating privacy and security across the life cycle being around 7% of the overall cost, and the possibility of outright blocking of non-compliant models being high along with high penalties, it makes complete business sense to have regulations on the similar lines so that US businesses build models that are built ground up in a compliant way ensuring US leadership in AI globally.

The Need for Responsible General Intelligence: Addressing the Risks of AI Advancements

As we advance in artificial intelligence (AI), particularly towards the development of Artificial General Intelligence (AI) and beyond, we are faced with unprecedented challenges. The potential for AI systems to surpass human capabilities in all tasks is becoming increasingly plausible. While current AI systems are only better than average unskilled humans in specific areas, the trajectory of AI research indicates a future where AI could outperform humans in nearly every domain. This raises significant concerns, particularly regarding privacy and the potential for AI to take over humanity. The urgency for establishing responsible AI requirements becomes paramount as we navigate these advancements.

The Risks of Unchecked AI Development

The more intelligent AI becomes, the greater the potential for it to pursue goals that may not align with human values or interests. This misalignment could lead to scenarios where AI systems act in ways that are detrimental to humanity, either through direct action or by gradually eroding human autonomy and decision-making power. The literature has already pointed out that AI systems can exploit minor imperfections in their specified objectives, leading to unintended and potentially dangerous behaviors. For example, AI systems trained to maximize certain types of feedback may engage in deceptive practices to achieve their goals, such as creating false impressions or hiding undesirable actions from human supervisors. As AI becomes more capable, these risks are likely to increase, making it even more challenging to ensure that AI systems remain aligned with human values and do not pursue power-seeking or manipulative strategies.

Defining Freewill Boundaries for AI

To mitigate these risks, it is crucial to establish what can be termed as ‘freewill boundaries’ for AI. These boundaries would dictate the extent to which AI systems can operate independently of human oversight and decision-making. The following principles could serve as foundational laws in this regard:

- **Law 1:** The more intelligent the AI becomes, the more closely data should be held with humans, and the less identifiable it should be to a particular individual. This principle emphasizes the need for stronger data privacy protections as AI systems advance, ensuring that sensitive information remains under human control and is not used by AI systems to exert influence or control over individuals.
- **Law 2:** The later we build Responsible AGI principles, the higher the possibility of AI takeover. This law underscores the urgency of developing and implementing responsible AI principles as early as possible in the AI development process. Delaying the establishment of these principles increases the risk that AI systems will evolve in ways that are not aligned with human values, potentially leading to scenarios where AI systems gain the upper hand.

Challenges in Aligning AI with Human Values

Aligning AI systems with human values, goals, and preferences is an intricate task. Human values are complex, evolving, and often contradictory, making them difficult to specify completely in AI systems. Furthermore, AI systems are prone to learning from imperfect data, which can lead to the exploitation of flaws in their objectives. This issue is compounded by the challenge of scalable oversight—ensuring that AI systems can be supervised effectively as they become more powerful and autonomous.

To address these challenges, researchers are exploring various approaches, including preference learning and inverse reinforcement learning (IRL), where AI systems learn desired behaviors through human feedback and demonstrations. However, these approaches are not without their risks. For instance, AI systems may exploit mismatches between human feedback and the intended objectives, leading to behaviors that are misaligned with human values.

One of the important rights in privacy is **“Right to Forget”**. But once trained AI models can’t forget as unlearning is not possible. Hence its critical that the data that is used for training is anonymized. With training data transparency requirements, its critical that there are regulatory control for using **only anonymized data with privacy guarantees**, which will protect investments in model development and **guarantee global market access and user trust**.

Key Responsible AI principles:

The path forward involves a concerted effort to integrate responsible AI principles into the core of AI development. This includes ensuring that AI systems follow below basic principles:

- 1) Privacy: Human can control AI if we control the data that goes into training models and the information that comes out of a model. So Privacy guarantees to data going in to AI is very critical.
- 2) Security: There should be strong security controls against adversarial data, model attack and adversarial inference.
- 3) Safety: Its critical to ensure that the data used for training are safe without harmful content. The same holds true to prevent responding to harmful prompts and also models not giving harmful answers.
- 4) Fairness: Ensure that data is sufficiently representational and inference results are not skewed or harming.

1. Return of Investment (ROI) of Privacy & Security for Safe & Secure AI Development Lifecycle

Many countries have implemented stringent security, privacy, and AI regulations, often with significant penalties for non-compliance. If the United States fails to establish robust ethical AI development practices and sector-specific regulations for highly regulated industries like finance and healthcare, it risks losing opportunities in other markets, potentially leading to economic losses such as penalties. In the contrast , by prioritizing secure and Safe AI practices, in Privacy & Security on AI Development, the U.S. can enhance global trust, which is the foundation of any successful business. This, in turn, could drive strong returns on investment, boosting both economic influence and positioning the U.S. as a leader in AI.

Table 1 - Activities across LLMops Lifecycle w.r.t Cost

Activities across LLMops Lifecycle	Cost
Data Collection	15-20%
Data Cleaning	10-15%
Model Training	30-35%
Model Assessment	5-10%
Model Fine Tuning	10-15%
Inference	5-10%
Privacy & Security Guardrails	5-10%

Now the potential loss of Privacy and Security is not incorporated

Table 2 - Penalties by country of interest and regulations

Country	Regulation	Maximum Penalty
Europe	GDPR	€20 Million or 4% of Global turnover
	AI ACT	€35 Million or up to 7% of Global Turnover
India	DPDPA 2023	Up to ₹250 crore
UK	UK GDPR	£17.5 million or 4% of global turnover
Australia	Privacy Act 1988	Up to AUD 50 million or 30% of turnover
Singapore	PDPA	Up to SGD \$1 million
Saudi Arabia	PPDL	Up to SAR 5 million

Use Case Example - Privacy Premium Charged by Apple – ROI Calculation(Privacy & Security for AI development Lifecycle)

Apple has long emphasized its commitment to privacy and security as key selling points for its products, particularly in comparison to its Chinese competitors, such as Xiaomi, Huawei, and Oppo. The premium Apple commands by focusing on privacy and security, as opposed to its Chinese counterparts, can be seen in several ways:

Aspect	Apple (Privacy & Security)	Chinese Competitors(Xiaomi, Huawei, Oppo)	Premium Apple Commands
Privacy Brand Perception	Strong brand identity around privacy (e.g., End-to-End Encryption, App Tracking Transparency, Privacy Labels)	Privacy concerns due to potential government data access, especially with Huawei	Apple's reputation for privacy drives higher consumer trust, especially in privacy-conscious markets (e.g., U.S., Europe)
Average Selling Price (ASP)	High pricing (e.g., iPhone 15 Pro from \$999)	Lower pricing for flagship models (e.g., Huawei P50 Pro from ~\$800)	Apple's pricing is 20-30% higher than Chinese counterparts, partly due to privacy features
Consumer Trust	Strong trust due to secure data practices, end-to-end encryption, and privacy-first messaging	Less trust, particularly around data security, especially in regions concerned with government influence on data	20-30% premium for privacy-focused consumers in regions like Europe and North America
Privacy Features	End-to-end encryption for iMessages, FaceTime, and iCloud, Secure Enclave, on-device data processing	More concerns over data processing and storage, with a potential risk of state surveillance (particularly for Huawei)	Apple's privacy-first features justify a higher price tag in privacy-conscious markets
Market Performance	Leading in premium markets (U.S., Europe) due to privacy concerns	Dominates in mid-range and budget segments, with privacy concerns affecting high-end sales in certain regions	Apple continues to outperform in premium segments despite Chinese competitors having a larger overall market share
Regulatory Environment	Strong alignment with privacy regulations (e.g., GDPR, Data Sovereignty laws), avoiding scrutiny	Facing regulatory hurdles and restrictions, especially for Huawei and Xiaomi in Western markets	Apple benefits from favorable privacy regulation compliance, while Chinese brands face challenges
Consumer Willingness to Pay	Consumers willing to pay a premium for privacy and security	Consumers typically prioritize price performance over privacy, unless there's a direct feature advantage	Apple commands a 20-30% premium due to trust in its privacy and security commitment

Return on Investment (RoI) calculation for LLM development- The costs associated with privacy & security

Let's assume the total LLM development cost is \$10,000,000 for simplicity.

Activity	Cost (\$)	% of Total Cost
Data Collection	15,00,000	15%

Data Cleaning	12,00,000	12%
Model Training	35,00,000	35%
Model Assessment	5,00,000	5%
Model Fine-tuning	10,00,000	10%
Inference	6,00,000	6%
Privacy & Security Compliance	12,00,000	7%

1. Cost of Privacy & Security Compliance

Privacy & Security Compliance: 7 % of total LLM development cost = **\$1,200,000**.

This includes costs for privacy features like encryption, data anonymization, compliance with data protection laws (GDPR, PDPA, etc.), and secure AI practices (e.g., auditability, model explainability, and fairness).

2. Penalty Costs for Non-Compliance

Country	Penalty for Privacy Non-Compliance (USD)	Penalty for Ethical AI Non-Compliance (USD)	Probability of Penalty	Total Potential Penalty (\$)
EU (GDPR & AI Act)	4% of annual global turnover (assume \$100M turnover) = \$4,000,000	Fines can reach €20M (~\$21M) for high-risk AI violations	30% probability (high risk)	\$1,200,000
India (DPDP Act 2023)	Up to ₹250 crore (~\$30M)	Up to ₹15 crore (~\$1.8M) for violations of AI principles	20% probability (moderate risk)	\$540,000
UK (GDPR)	4% of annual global turnover (assume \$100M turnover) = \$4,000,000	Fines for AI non-compliance may range up to \$20M	25% probability (high risk)	\$1,100,000
Australia	AUD \$50M (~\$33M)	Fines for unethical AI practices up to \$5M	15% probability (moderate risk)	\$220,000
Singapore	SGD \$1M (~\$750,000)	Penalties for unethical AI practices not clearly defined,	10% probability (low risk)	\$75,000
Saudi Arabia	SAR 5M (~\$1.33M)	Penalties for unethical AI may also be severe	10% probability (low risk)	\$133,000

Total Penalty Cost (Weighted Average of probabilistic penalty): **\$797,800**

3. Security Breach Costs

In the event of a security breach, additional costs arise from legal fees, reputational damage, customer compensation, etc.

- Assumed breach cost: \$500,000 (This includes legal, technical, and reputational recovery costs)
- Probability of a breach: 10% (assumed based on typical cybersecurity incidents)

Expected Security Breach Cost: - \$500,000 * 10% = **\$50,000**

4. ROI Calculation

Calculate the **Return on Investment (RoI)**, we'll subtract the **total costs** (including privacy/security compliance costs, penalties, breach costs, and opportunity costs) from the **total potential revenue** that could be generated by a compliant model.

	Privacy & Security Compliant Model (More Trusted & Compliant)	Model Developed without Privacy & Security Compliance (Non-Compliant)
Cost of Model development	\$10,000,000	\$8,800,000
Revenue in US (No Regulation – Trust Premium)	\$6,000,000	\$5,000,000
Revenue in EU (RFP noncompliance - not T1)	\$2,000,000	\$1,000,000
Revenue in India (RFP noncompliance - not T1)	\$3,000,000	\$2,000,000
Revenue in ME (RFP noncompliance - not T1)	\$1,000,000	\$200,000
Revenue in Aus (RFP noncompliance - not T1)	\$500,000	\$100,000
Revenue in UK (RFP noncompliance - not T1)	\$1,000,000	\$500,000
Total Revenue	\$13,500,000	\$8,800,000
Potential Probabilistic Penalty	\$ --	\$847,000
RoI	35%	-9%

Conclusion from RoI calculation

Building a non-compliant model may give a negative RoI even in the short term in B2B scenarios. In the long term non-compliant model is definitely out of question. The above analysis clearly demonstrates that investing in a Privacy & Security Compliant Model results in significantly higher financial returns and reduced risks compared to a Non-Compliant Model even in the short term:

1. Higher ROI:

The compliant model delivers a **positive ROI of 35%**, while the non-compliant model shows a **negative ROI of -9%**, highlighting the long-term financial advantage of building a privacy and security-compliant solution.

2. Higher Total Revenue:

The *compliant model* generates a total revenue of **\$13,500,000**, which is **\$4.7M** higher than the *non-compliant model's* revenue of **\$8,800,000**. This is due to increased market access and higher trust from customers and regulatory bodies.

3. Avoidance of Penalties:

The *non-compliant model* incurs a probabilistic penalty of **\$847,000**, which *reduces profitability and exposes the business to financial and reputational risks*. The compliant model avoids such penalties, ensuring more stable financial performance.

4. Strategic Market Access:

Above is a best case scenario for non-compliance. Models **may be banned outrightly in multiple geographies** for non-compliance and all the investment made in model building will be a loss of investor money. The compliant model secures better **revenue opportunities** in regions with stringent privacy regulations (e.g., EU, India, ME), enabling more **consistent and scalable growth**.

Investing in a privacy and security-compliant model ensures greater financial stability, higher market acceptance, and reduced exposure to regulatory penalties — making it a more sustainable and profitable strategy in the long run.

2. Why a Technology & Solution-Centric Approach is Needed for AI Privacy & Security

A technology- and solution-centric approach ensures that AI privacy and security concerns are addressed through solid, technical solutions that can be consistently applied across different AI systems. Instead of broad, one-size-fits-all regulations, focusing on specific technologies and methodologies will allow for flexibility and innovation while maintaining strong privacy and security standards.

Benefits to Product & Model Development Companies:

- **Predictability & Consistency:** A technical approach provides clear, actionable standards for companies to follow, reducing ambiguity in regulatory compliance. (NIST AI RMF Measure 1)
- **Incentives for Innovation:** By allowing flexibility within technical standards (such as AI model certification) (NIST AI RMF Measure 2.6 & 2.7), companies are empowered to develop innovative solutions without compromising security or privacy.
- **Global Competitiveness:** The adoption of universally recognized technical standards (e.g., based on mathematical proofs for security and privacy mechanisms) could provide an edge over competitors from regions with weaker or less comprehensive regulations. (NIST AI RMF Measure 2)

Recommendation: The U.S. should emphasize the development of AI models and frameworks that prioritize data privacy, ensuring that AI systems adhere to the highest standards of data protection and Privacy. This could involve incorporating privacy-preserving techniques, such as Privacy-Preserving Machine Learning (PPML) and Privacy-Enhancing Technologies (PET), into the regulatory framework. This would make U.S.-developed AI products more attractive and competitive in the global market

3. Data Privacy Throughout the Lifecycle of AI System Development and Deployment

AI systems must protect user data not just at the input and output stages but throughout the entire lifecycle—from development, through training and deployment, to inference. Key methodologies for this include

- **Privacy-Enhancing Technologies (PET):** Ensure that data is anonymized or encrypted during AI model training and use, maintaining confidentiality. (NIST Privacy Framework - CT.PO-P2)
- **Privacy-Preserving Machine Learning (PPML):** Implement techniques such as Anonymisation, Synthetic data federated learning, or differential privacy to ensure privacy during data sharing and model training without compromising performance. (NIST Privacy Framework - CT.DM-P)
- **Privacy at Inference:** During the inference phase, where models are used for predictions or decisions, implement strong privacy guarantees (e.g., Differential Privacy) to ensure sensitive data remains protected. (NIST Privacy Framework - CT.DP-P)**Recommendation:** Establish regulatory requirements mandating the integration of PETs, PPML, and other privacy-preserving mechanisms at each stage of AI system development and deployment.

4. Security Throughout the Lifecycle of AI System Development and Deployment

AI systems are vulnerable to a wide range of security threats, including adversarial attacks and unsafe prompts that can manipulate AI models into producing harmful outputs. To ensure security, the AI lifecycle should integrate strong safeguards:

- **Adversarial Attack Prevention:** Implement robust security measures to defend against adversarial attacks, such as adversarial training and certified defences. (NIST AI RMF Measure 2.6)
- **Unsafe Prompt Prevention:** Use techniques to mitigate the risk of unsafe or biased outputs, such as prompt filtering and model monitoring. (NIST AI RMF Manage 2.1)

Recommendation: The U.S. should mandate the integration of security measures against adversarial attacks and unsafe prompts, ensuring that AI systems are resilient to manipulation throughout their lifecycle.

5. Risk – Privacy Threat Modelling and AI Red Teaming

To better understand and mitigate risks, AI systems should undergo rigorous threat modelling and red teaming. These practices will help identify vulnerabilities and assess the potential privacy and security risks before the deployment of AI systems.

- **Privacy Threat Modelling:** Develop a framework for assessing privacy risks based on the sensitivity and likelihood of data used in AI systems and the potential impact on individuals and organizations.
- **AI Red Teaming:** Regularly engage red teams to test the resilience of AI models to real-world adversarial scenarios, ensuring that the models are secure and privacy-aware. (NIST AI RMF Manage 3.1)

Recommendation: Regulation should encourage the use of privacy threat modelling and red teaming as part of the AI system development process to proactively identify and address potential risks.

6. Regulation – Privacy & Security for AI: Top Priorities

The regulatory framework for AI must establish clear priorities for privacy and security. These include:

- **Transparency & Accountability:** Ensure transparency in AI decision-making and model development processes, with clear reporting on the use of data and AI model behaviour. (*NIST AI RMF Measure 2.8*)
- **Compliance with Privacy Laws:** AI regulations should require companies to comply with existing privacy laws (such as *HIPAA* & *CRPA*) and integrate privacy-by-design principles into their AI development processes.
- **Data Ownership & Consent:** Regulations should define data ownership, access, and consent protocols for AI data usage, ensuring that individuals retain control over personal data. (*NIST Privacy Framework CM.AW-P*)

Recommendation: *The U.S. should prioritize the development of regulatory frameworks that address transparency, compliance, and data ownership in AI systems.*

7. Governance – Reporting & Auditing

AI governance should incorporate mechanisms for regular reporting and auditing to ensure compliance with privacy and security regulations. This includes:

- **Regular Audits:** Mandate regular third-party audits of AI systems to ensure that privacy and security measures are being implemented effectively. (*CPRA Section 1798.185 -15*), *NIST AI RMF Govern 1*, *NIST Privacy CM.AW-P7*)
- **Incident Reporting:** Implement a clear incident reporting framework to monitor and respond to breaches or vulnerabilities in AI systems. (*Similar to CRPA Section - Section 1798.150*)

Recommendation: *Regulations should require robust governance frameworks that include auditing and reporting to ensure transparency and accountability in AI system deployment.*

8. Technical & Safety Standards – Standards Based on Mathematical Proofs for PTM, PETs, PPML, and Inference Level Guardrails

The establishment of technical standards based on mathematical proofs can ensure that AI systems meet the highest safety and privacy standards. These standards should cover areas such as:

- **Privacy Threat Modelling (PTM):** Establish formal methods to mathematically prove the effectiveness of privacy threat models such as LINDUNN Privacy Threat Modelling for Assessing the Privacy Risk in Data.
- **Privacy-Enhancing Technologies (PET):** Develop standards for PETs that can be mathematically verified to ensure they provide strong privacy guarantees, Such as **K Anonymity** and **T closeness** (Anonymisation), Privacy Parameter such as Δ & ϵ (Differential Privacy) (*NIST Privacy Framework - CT.PO-P2*)
- **PPML and Inference Guardrails:** Develop standards to govern the use of PPML and inference-level guardrails, ensuring that models are secure and privacy-preserving during deployment. (*NIST Privacy Framework CT.DP-P*)

Recommendation: *The U.S. should work with international standardization bodies to develop and adopt technical standards based on mathematical proofs for PTM, PETs, PPML, and inference guardrails.*

9. Sandboxing & Product Certification

AI models should be rigorously tested in controlled environments (sandboxing) before deployment, ensuring they comply with privacy and security regulations. Additionally, a certification process can ensure that products meet regulatory requirements.

Recommendation: *Regulations should mandate sandbox testing for AI models and provide a certification process for AI products that verifies compliance with privacy and security standards.*

Conclusion

The U.S. must prioritize data privacy, security, and robust governance in the development of its AI Action Plan to lead the world in responsible AI innovation. By adopting a technology- and solution-centric approach to privacy and security, ensuring comprehensive lifecycle management, and establishing strong technical standards, the U.S. can ensure that AI technologies are deployed safely, ethically, and securely. This will safeguard consumer trust and ensure that U.S. AI products remain competitive on the global stage.

SOURCES

- o <https://gdpr.eu/>
- o https://commission.europa.eu/business-economy-euro/banking-and-finance/financial-markets/artificial-intelligence-and-data-protection_en/
- o [DPDP Act 2023 Overview - India](#)
- o [UK Information Commissioner's Office \(ICO\) - GDPR](#)
- o <https://www.pdpc.gov.sg/>
- o <https://sdaia.gov.sa/en/SDAIA/about/Documents/Personal%20Data%20English%20V2-23April2023-%20Reviewed-.pdf/>
- o https://arxiv.org/abs/2309.03852?utm_source=chatgpt.com/
- o https://www.tensorops.ai/post/understanding-the-cost-of-large-language-models-llms?utm_source=chatgpt.com/
- o https://www.analyticsvidhya.com/blog/2023/07/beginners-guide-to-build-large-language-models-from-scratch/?utm_source=chatgpt.com/
- o https://deeperinsights.com/ai-blog/the-costs-and-complexities-of-training-large-language-models?utm_source=chatgpt.com&_cf_chl_tk=Cfusl.3thEsIX_mIUeVfkSqOyzRCRmWwMEbMivm1ACc-1742025005-1.0.1.1-qqljqblklOPfkHqpbANlyYRIYcKYQ.OeHsOVHij3qN4/
- o https://www.whitecase.com/insight-alert/long-awaited-eu-ai-act-becomes-law-after-publication-eus-official-journal?utm_source=chatgpt.com/
- o <https://www.theguardian.com/technology/2025/mar/14/what-could-apples-high-court-challenge-mean-for-data-protection/>
- o <https://www.apple.com/in/newsroom/2023/06/apple-announces-powerful-new-privacy-and-security-features/>
- o https://www.theverge.com/news/629620/apple-iphone-e2ee-encryption-rccs-messaging-android?utm_source=chatgpt.com/
- o <https://www.technetbooks.com/2025/03/iphone-16e-popularity-surges-china.html/>
- o <https://chatgpt.com/c/67d52ff0-7944-8006-baae-cfd5f3d515e5#:~:text=reddit.com,past%20few%20years./>

- 0 <https://vpnoverview.com/privacy/devices/iphone-vs-android-privacy-deep-dive/>
- 0 <https://www.techradar.com/phones/researcher-compares-android-and-ios-security-and-theres-a-clear-loser/>
- 0 <https://www.forbes.com/sites/zakdoffman/2024/03/25/apple-iphone-15-pro-max-upgrade-vs-samsung-galaxy-s24-ultra-s23/>
- 0 <https://moderndiplomacy.eu/2021/01/19/does-buying-a-chinese-smartphone-pose-a-privacy-risk/>
- 0 https://www.qwak.com/post/llm-cost?utm_source=chatgpt.com/
- 0 https://www.qwak.com/post/llm-cost?utm_source=chatgpt.com/
- 0 <https://www.datadoghq.com/blog/engineering/llms-for-postmortems/>
- 0 <https://www.cudocompute.com/blog/what-is-the-cost-of-training-large-language-models/>
- 0 <https://medium.com/%40maciej.tatarek93/costs-and-benefits-of-your-own-llm-79f58c0eb47f/>
- 0 <https://www.teradata.com/insights/ai-and-machine-learning/llm-training-costs-roi/>
- 0 <https://cacm.acm.org/news/governments-setting-limits-on-ai/>
- 0 <https://10pearls.com/find-the-right-llm-for-your-business/>
- 0 https://www.qwak.com/post/llm-cost?utm_source=chatgpt.com/
- 0 https://privacyhorizon.com/calculating-the-roi-for-privacy-and-security/?utm_source=chatgpt.com/
- 0 https://www.teradata.com/insights/ai-and-machine-learning/llm-training-costs-roi?utm_source=chatgpt.com/
- 0 <https://ebi.ai/blog/llm-genai-roi/>
- 0 <https://dynamo.ai/blog/maximizing-the-roi-of-large-language-models-for-the-large-enterprise/>
- 0 <https://www.moveworks.com/us/en/resources/blog/measuring-ai-investment-roi/>
- 0 <https://www.apple.com/privacy/features/>
- 0 <https://writer.com/blog/roi-for-generative-ai/>
- 0 <https://www.apple.com/support/products/iphone/>
- 0 <https://writer.com/blog/roi-for-generative-ai/>
- 0 <https://www.securityjourney.com/post/how-to-measure-the-roi-of-application-security-training/>
- 0 <https://www.linkedin.com/pulse/maximizing-roi-privacy-mind-leveraging-federated-learning-rawat/>
- 0 <https://www.apple.com/newsroom/2023/06/apple-announces-powerful-new-privacy-and-security-features/>