

March 15, 2025

National Science Foundation
White House Office of Science and Technology Policy (OSTP)
Attn: Faisal D'Souza, NCO
2415 Eisenhower Avenue,
Alexandria, VA 22314, USA
[REDACTED]

RE: Artificial Intelligence (AI) Action Plan

Dear Members of the National Science Foundation and White House Office of Science and Technology Policy (OSTP),

Responsible Innovation Labs (RIL) appreciates the opportunity to submit comments in response to the National Science Foundation's Request for Information (RFI) on recommendations concerning the development of an Artificial Intelligence (AI) Action Plan following President Donald Trump's Executive Order 14179 ("Removing Barriers to American Leadership in Artificial Intelligence") signed on January 23, 2025. As we all know, the United States has always been the biggest driver of technological innovation in the world. Fueled by our unmatched ingenuity, our tremendous enabling environment, and the incredible drive of our entrepreneurs and startups, America is the place to launch the next generation of AI startups.

Responsible Innovation Labs is a coalition of startups and investors who believe smart and responsible innovation is how you win. It is how great founders transform startups into enduring companies enjoying global success. RIL's mission is twofold. First, RIL accelerates founders and operators building great companies responsibly. We develop startup-specific content, provide trainings, and convene startups to talk about shared challenges and opportunities. Second, RIL is a forum for startups, large enterprise, investors, policy makers, and others to get ahead of the challenges posed by emerging technologies so that tech is worthy of the trust necessary to continue benefitting from permissionless innovation. Our efforts are aimed at nurturing the next generation of American business leadership. We strive to increase know-how when it comes to innovating creatively so that we can secure America's position as the global leader in AI and create a brighter future for every American because – not in spite – of this technological step change.

I. Introduction

The benefits of today's AI revolution are well-documented. AI is driving unparalleled benefits across our economy, enhancing productivity, efficiency, safety, customer service, and

accelerating discovery and decision-making to create tremendous value for consumers and entrepreneurs alike. Already, by automating everyday tasks - from travel planning to financial management - AI has saved countless hours and opened new avenues for economic growth, in finance, energy, and agriculture. However, concerns over AI risks remain significant.

There are real risks to American's physical, property, and economic security. AI can supercharge threats to our national cyber and critical infrastructure. Seniors and other vulnerable populations are being targeted in sophisticated scams exacerbated by AI. Children are being targeted with sexualized content with little to no transparency for parents to make decisions as to which products actually prioritize safety. For creators, confusion abounds around the application of existing intellectual property (IP) laws. To truly take advantage of the opportunities AI presents for American leadership and prosperity, it is crucial that we address these near-term risks.

We believe America must be the leader in establishing the global benchmarks for smart AI innovation.

II. Important Goals of AI Policy

As part of our public mission, RIL engages with policymakers at both state and federal levels to provide technical assistance focused on the foundational components of AI: compute power, data stewardship, and model development.

Compute Power: To ensure that AI innovation is not monopolized by large tech incumbents, it is imperative to level the playing field for startups. To this end, projects like the National AI Research Resource (NAIRR) Pilot should be advanced. This sort of investment recognizes that many researchers and students lack the necessary access to the computing power needed to fully conduct research and advance AI technology. Americans will not benefit from AI if too few players have the basic tools and know-how to accelerate discovery, innovation, and commercialization. We believe a level playing field also requires rigorous enforcement of existing laws designed to ensure that large tech incumbents are not engaged in anticompetitive conduct to crush nascent startups.

Data Stewardship: The federal government has an important role to play in promoting transparency and ensuring that AI companies are responsible stewards of data. We urge the Trump Administration and Congress to work together to enact federal privacy legislation that would provide additional transparency and choices, require companies to maintain reasonable data security, foster data minimization principles, and ensure data subject rights to access their data.

AI Action Plan Development: Prior regulatory proposals such as those at the state level have run into opposition because they focused on model development, rather than model applications and use, particularly those posing high risks to public safety and security. Government efforts

must recognize the differences between AI model applications and their development. Tailoring regulations based on risk will ensure that safety and security are prioritized without stifling innovation. Moreover, Congress and federal agencies may want to consider antitrust safe harbors to encourage self-regulation among AI companies, so that they can share information about threats and risks, similar to what the DOJ and FTC have done in the area of cybersecurity.

Existing laws apply and should be enforced in the AI space, ranging from antitrust and privacy laws to laws prohibiting unfair and deceptive practices to more specialized laws such as the Fair Credit Reporting Act (which would require adverse action notices in certain circumstances where AI is used for decision making) and the Equal Credit Opportunity Act (which prohibits credit discrimination, through AI or otherwise). Continued enforcement of these laws is crucial to AI development and the ability of startups to compete with other rule-abiding companies.

The guidance below is offered as policy makers consider a potential AI Action Plan for American AI leadership.

III. Principles for a Winning AI Action Plan

A winning AI Action Plan should incorporate the following principles:

- America wins when startups scale with an intense focus on a first order priority of product integrity. AI applications – often driven by startups – will win commercially by emphasizing the benefits of AI quality control. Mechanisms to increase explainability, assurance, transparency, and getting ahead of 2nd/3rd order impacts of AI will benefit customers and all who encounter AI because of American ingenuity. These behaviors will drive demand for American AI applications and systems in the long run. Policies should emphasize how a well-calibrated approach to AI governance enables innovation while strengthening the U.S. AI ecosystem for long-term competitiveness.
- Responsibility in innovation does not mean zeroing out risk. Startups should be entitled to safe harbors, so long as they have employed and documented common sense responsible innovation methods, including testing and risk mitigation controls. An AI Action plan should include a means to incentivize companies to develop processes to forecast and address potential harms associated with AI. Rules demanding perfect outcomes are likely to stifle interest in developing AI that benefits the American people on some of our most pressing challenges – for example, creating abundance in health care, individualizing education for all kids, and accelerating scientific discovery. Similarly, the NSF and OSTP should consider safe harbors as part of any proposed AI Action Plan, so that companies are not penalized for internal testing that would potentially find problems with AI systems.

- An AI Action Plan should promote competition, particularly among startups and small businesses, and should not entrench existing large tech players. American startups should be scaling with stage- and maturity-specific requirements for AI governance. Onerous technical and legal requirements not tied to addressing tangible harms can favor large companies with established legal teams at the expense of smaller companies that may struggle to comply. Legal requirements contemplated by the AI Action Plan should therefore be scaled to account for the needs of smaller players. This can be accomplished through tiered requirements, small business exceptions, and/or regulatory sandboxes. Competition can also be supported through efforts like the NAIRR, as well as block grants to states so that they may complement the NAIRR with research, data, and compute endeavors specific to their state economies and needs.
- To promote startup competition and avoid a state-by-state patchwork, the federal government should lead on AI policy. A federal standard could ensure a level playing field of protection for American consumers, as well as smaller startups in the AI space that may not have the resources to hire outside counsel to parse through state-by-state nuances. A national standard could form the basis of multilateral engagement by the US government to harmonize international standards such that they do not harm American startups and consumers. States would still have significant latitude to legislate in their traditional domains of state power in our constitutional system.
- Clarify the responsibilities between developers and deployers. Because deployers have the most control over how an AI system is used, developers should not bear the main responsibility for unsafe uses. Developers should, however, be obligated to provide sufficient information to deployers so that deployers can fulfill their obligations, consistent with protection of trade secrets and confidential information. These responsibilities could be modeled on the relative responsibilities between controllers and processors under privacy law, with controllers ultimately responsible for compliance with most provisions, and processors responsible for assisting controllers with their obligations.
- Promote the testing and assurance marketplace. More and more startups (and established large enterprises) are offering governance, testing, and assurance services when it comes to novel AI technologies. These market mechanisms can spur safer and more reliable AI. Supporting free market mechanisms will enable better and more durable responsible behaviors by both developers and deployers in the long run.
- Balance transparency with protection of trade secrets. Companies should not have to disclose trade secrets or confidential commercial information to the public absent compelling circumstances. An AI Action Plan should strike a balance between appropriate transparency and protecting confidential information important for AI development and American companies.

- Protect American workers impacted by AI as its velocity increases. Every day startups are setting out to improve life and commerce in our country. They are fighting to survive and solve problems. It is especially important that startups have room to develop applications that fill gaps and meet the unmet needs of all Americans. At the same time, rapid advancement of AI will reshape industries and impact labor markets. It is imperative a federal AI Action Plan get ahead of impacts on workers caused by these transformative technologies. Ignoring the threats will encourage fear mongering and backlash for AI innovators notwithstanding upside benefits from these technologies. Among other examples, the U.S. Trade Adjustment Assistance Program (TAA Program) helps workers who have been adversely affected by foreign trade and displaced by globalization. A similar framework could be established to provide financial stability and reskilling opportunities for those adversely impacted as AI reshapes our economy.
- Strengthen NIST to Advance AI Self-Governance. For decades, the National Institute of Standards and Technology (NIST) has played a pivotal role in establishing standards and best practices that have reinforced U.S. leadership in emerging technologies, including artificial intelligence. Its rigorous, consensus-driven approach has not only enhanced domestic AI governance but has also positioned the U.S. as a global leader in responsible AI innovation. Foundational efforts such as the NIST AI Risk Management Framework have provided a flexible, risk-based foundation for AI policy, influencing international partners to align with U.S. principles. As AI continues to evolve, NIST's expertise will be critical in ensuring that AI governance frameworks remain innovation-friendly while providing a forum to address key risks. By fostering international collaboration and shaping global AI norms, NIST has ensured that American AI policy remains central to global discussions. Moving forward, leveraging NIST's leadership will be essential to deepening AI policy coordination with allies, strengthening transatlantic and Indo-Pacific AI partnerships, and ensuring that pro-innovation, risk-based regulatory approaches remain the global standard.

The U.S. AI Action Plan represents a critical opportunity to align AI policy with the realities of the startup ecosystem, where innovation flourishes when integrity, trust, and responsibility are built into the development process. This is the time for startups to enhance public trust not because they are told to do so, but because they have the freedoms and incentives to build companies to endure for the long run.

By fostering an approach that enhances governance and integrity while enabling AI's full economic and societal benefits, the U.S. can strengthen its leadership in AI, drive economic prosperity for the American worker, and ensure long-term technological progress.

RIL looks forward to engaging further on these issues and contributing to efforts that ensure American AI leadership in the years ahead.