

AI Action Plan

Submitted By :Shubham Gupta

This AI Action Plan outlines clear policy actions across key areas – *Model Development, Cybersecurity, Data Privacy, Technical & Safety Standards, Research & Development*, and *International Collaboration* – and expands on additional priorities such as energy efficiency, regulation & governance, national security, and education/workforce development. Each section includes actionable steps, measurable targets where possible, and references to existing policies and best practices to ensure the United States continues to “sustain and enhance America’s global AI dominance” in a safe and trustworthy manner. Implementation will involve public-private collaboration and a whole-of-government approach, with progress measured against clear benchmarks.

Model Development

Develop cutting-edge AI models that advance American innovation and uphold our values. This involves promoting robust model design, transparency, and evaluation practices that fuel innovation without ideological bias, aligned with the directive to keep AI development free from undue agendas. Key actions include:

- **Incentivize Advanced AI Models:** Support the development of next-generation AI systems (e.g. foundation models in language, vision, and science) by U.S. companies and researchers. Set a target for U.S. innovators to produce *at least 5 world-leading AI models annually* in critical domains. Progress can be tracked via benchmark achievements (e.g. improvements on language understanding or healthcare AI benchmarks by 10% each year).
- **Robust Testing & Transparency:** Require rigorous pre-release testing and risk assessment for any high-impact AI model. Encourage use of independent “red teams” to stress-test models for safety, fairness, and security issues before public deployment. For example, all major AI models above a certain capability threshold should undergo external security testing and publish transparency reports (such as *Model Cards* detailing training data, intended use, and limitations). *Measurable outcome:* 100% of federal AI projects and 90% of industry-led frontier models report testing results and known risks publicly.
- **Best Practices for Responsible AI:** Promote industry adoption of responsible AI best practices. This includes sharing information on risk management and mitigation techniques across the AI sector. Leading AI firms have voluntarily committed to share best

practices and information on managing AI risks (e.g. addressing attempts to bypass safeguards)– the government will build on these efforts by facilitating an AI Safety Coalition to exchange information on emerging risks and solutions.

- **Open Innovation & Interoperability:** Encourage open-source frameworks and standards that allow interoperability of AI models and tools. Reducing barriers for startups and researchers to build on prior models (while respecting IP and security) will spur innovation. *Goal:* Increase the number of openly available reference models or datasets by 20% per year, enabling wider experimentation and lowering entry hurdles for new AI developers.

Cybersecurity

Strengthen the cybersecurity of AI systems and leverage AI to enhance cyber defense. As AI becomes ubiquitous, it creates new attack surfaces – a significant portion of organizations (over 40%) have already reported AI-related security incidents such as adversarial attacks on machine learning models. This plan fortifies AI against threats and uses AI as a tool for security:

- **Secure AI Systems Against Attacks:** Develop and disseminate guidelines for making AI models resilient to adversarial attacks and data poisoning. By 2026, all federal AI systems should implement adversarial robustness testing (e.g. testing against manipulated inputs) as part of their development lifecycle. Industry will be encouraged to do the same, given that 41% of organizations have faced AI security incidents (with threats like model manipulation or data poisoning). *Measurable target:* Reduce known AI vulnerabilities (such as successful adversarial attacks in evaluation tests) by 50% within 3 years across critical systems.
- **AI for Cyber Defense:** Invest in AI-driven cybersecurity tools to protect critical infrastructure and networks. AI can help detect intrusions and anomalies faster than traditional methods. The plan sets a goal to *cut cybersecurity incident response times by 50%* using AI analytics and to deploy advanced AI monitoring on 100% of federal civilian networks by 2026. Support public-private partnerships (e.g. with cybersecurity firms and academia) to develop algorithms that identify and counter new threats in real-time.
- **Protect AI Assets and IP:** Mandate strong cybersecurity and insider-threat safeguards for sensitive AI assets like model weights and training data. AI models, especially large ones, are strategic assets; losing them to theft or espionage would be a national security risk. Companies are already committing to protect proprietary model weights and to enable third-party vulnerability reporting– the U.S. government will reinforce this by establishing a confidential AI incident reporting mechanism (possibly via the Cybersecurity and Infrastructure Security Agency) so that AI developers can report breaches or model

compromises anonymously and get support. Success will be measured by increased reporting and reduction in successful theft attempts.

- **Cybersecurity Workforce & Training:** Expand training for cybersecurity professionals in AI and for AI practitioners in cybersecurity. The goal is to train at least 5,000 cybersecurity experts in AI-specific security techniques (like securing ML pipelines) and vice versa, by 2028. This cross-training ensures a workforce capable of securing AI systems and using AI to secure other systems. Agencies will also incorporate AI-security exercises (e.g. “red team/blue team” drills with AI attack scenarios) into their regular cybersecurity training by 2025.

Data Privacy and Security

Ensure AI development and deployment uphold rigorous data privacy standards and protect sensitive information. Building public trust is crucial for AI adoption. Key measures:

- **Robust Privacy Standards for AI:** Require that AI systems, especially those handling personal or sensitive data, adhere to privacy-by-design principles. Developers should integrate techniques like data anonymization, encryption, and differential privacy during model training to safeguard individual data. All federal AI projects involving personal data will undergo Privacy Impact Assessments and implement recommended protections. *Target:* Achieve 100% compliance with privacy assessments for federal AI systems by 2026, and encourage the private sector to match this standard for high-risk AI applications.
- **Promote Privacy-Enhancing Technologies (PETs):** Accelerate the research and use of PETs in AI, such as federated learning (which allows AI models to train on data without centralizing it) and secure multi-party computation. These technologies enable AI innovation while minimizing raw data sharing. The plan will launch grand challenges or pilot programs (in domains like healthcare and finance) to apply PETs at scale, aiming for a 50% increase in the use of PETs in government AI systems within 3 years.
- **Data Security and Governance:** Establish clear data governance frameworks for AI. This includes setting standards for data quality, integrity, and lineage (provenance) so that AI models are trained on secure and reliable data. Agencies will inventory and securely open up government data sets for AI R&D, with appropriate access controls and privacy safeguards (aligning with the principle of enhancing access to high-quality data while maintaining security). We will measure progress by the number of new high-value government data sets released for AI use (with a goal of at least 20 major new datasets per year) and zero significant data breaches from those resources.

- **Strengthen Public Trust and Transparency:** To address public concerns, mandate transparency measures for AI systems that use personal data. For example, if an AI system impacts an individual (in loan decisions, hiring, etc.), there should be an explanation or documentation of how it uses data and protects privacy. Agencies will develop easy-to-understand user consent and notification guidelines for AI services. Industry will be encouraged to adopt voluntary transparency labels for AI products, indicating if and how personal data is used. Success will be reflected in improved public sentiment (as measured by surveys) and trust in AI – e.g., increasing the percentage of Americans who trust AI with their data from current low levels to over 50% in the next five years.

Technical and Safety Standards

Establish and champion technical standards and safety best practices to ensure AI systems are reliable, safe, and compliant with American values. Standardizing how we evaluate and manage AI risks will remove uncertainty for innovators while protecting the public. The U.S. will build on frameworks like the NIST AI Risk Management Framework, which provides guidelines to incorporate trustworthiness into AI design. Key initiatives:

- **Adopt AI Risk Management Frameworks:** Encourage widespread adoption of NIST’s AI Risk Management Framework (RMF) across federal agencies and industry. This voluntary framework outlines characteristics of trustworthy AI (such as validity, security, accountability, and transparency) and methods to minimize potential negative impacts on civil liberties while maximizing benefits. *Policy action:* All federal agencies will use the NIST AI RMF (or a comparable standard) as guidance when procuring or deploying AI systems, starting within one year. The private sector will be engaged through NIST-led workshops to tailor the framework for different industries. *Measurable outcome:* By 2025, 75% of Fortune 500 companies developing AI will report alignment with NIST’s AI guidelines or similar industry standards, up from an initial baseline.
- **Develop and Update Safety Standards:** In collaboration with industry and international bodies, accelerate the development of technical standards for AI safety, performance, and interoperability. This includes standards for testing AI algorithms (for example, standard benchmark suites for verifying an AI system’s accuracy and failure modes) and for certifying AI in high-stakes uses (like healthcare, transportation). The plan calls for publishing at least 3 new technical standards per year via organizations like IEEE, ISO/IEC JTC1, or NIST’s collaborative programs. These standards will address things like AI system robustness against attacks, explainability of AI decisions, and fail-safe mechanisms to handle AI errors. Agencies like NIST will also work internationally to ensure U.S. perspectives shape global standards (avoiding a patchwork of conflicting rules).

- **Certification and Audit Mechanisms:** Establish mechanisms to test and certify AI systems before wide deployment, especially in safety-critical sectors. For example, create an AI Safety Certification Program (potentially via a public-private partnership) that can evaluate AI models for known risks (bias, security holes, reliability) similar to how Underwriters Laboratories certifies electrical appliances. Aim to have all high-risk AI systems (as defined by usage in areas like medicine, automotive, etc.) undergo independent assessment by 2026. Additionally, promote third-party algorithm audits and red-team exercises as a norm. *Outcome:* Reduced incidents of AI failures causing harm – track and aim for a year-over-year decrease in reported AI malfunctions or accidents in regulated sectors.
- **Continuous Improvement & Research:** Support R&D in tools that improve AI safety (e.g., AI systems that can explain their reasoning, techniques for better model validation). Create feedback loops: incidents or near-misses involving AI will be analyzed by a national AI Safety Board (or existing bodies like the National AI Advisory Committee) to recommend new best practices or standards. Success in this area will be measured by the breadth of safety measures implemented (for instance, an increase in the number of AI systems that provide *explanations for their outputs*, or the proliferation of “ethical AI” reviews in organizations). By continuously updating technical guidelines, the U.S. ensures AI systems are not only innovative but also trustworthy by design, reinforcing public confidence and international credibility.

Regulation and Governance

Develop a balanced regulatory and governance framework for AI that protects public interests and national values without stifling innovation. The U.S. approach will be risk-based and flexible, focusing on high-risk AI applications with appropriate oversight, while allowing space for experimentation and growth in low-risk areas. This aligns with the Administration’s directive to reduce regulatory barriers that unnecessarily hinder AI deployment, even as we enact sensible guardrails for safety and ethics. Key steps:

- **Light-Touch, Risk-Based Regulation:** Rather than broad one-size-fits-all rules, pursue a regulatory strategy that distinguishes between uses of AI. High-risk AI systems (for example, AI in medical diagnosis, autonomous vehicles, or loan decisions) should meet stricter requirements for safety, fairness, and transparency. Lower-risk applications (like AI in spam filtering or routine business analytics) would face minimal regulation, guided by voluntary standards. This approach contrasts with more prescriptive regimes abroad – e.g., the EU’s draft AI Act imposes stringent rules on “high-risk” AI systems – and is intended to maintain U.S. competitiveness. *Action:* Within 12 months, issue an AI Governance Framework document that outlines categories of AI risk and recommended

regulatory or oversight actions for each category. This will guide sector-specific regulators (like FDA, DOT, FTC) in updating rules for AI in their domains.

- **Update and Harmonize Laws:** Conduct a comprehensive review of existing laws and regulations to identify those that impede AI development or are insufficient to address AI-driven outcomes. For example, assess whether current product liability laws adequately cover harms caused by AI decisions, or if intellectual property rules need clarification for AI-generated content. Remove or modernize regulations that unintentionally block AI innovation (a direct mandate of the President’s AI leadership strategy), such as outdated restrictions on data use, while filling gaps in consumer protection and anti-discrimination enforcement as they relate to AI. *Measurable goal:* Propose or enact updates to 100% of relevant federal regulations identified as outdated or hindering AI within 2 years, and develop new guidelines for at least 5 emerging AI issue areas (e.g. autonomous drones, AI in HR, etc.).
- **Regulatory Sandboxes and Pilots:** Encourage innovation-friendly governance by establishing “regulatory sandboxes” for AI – controlled environments where companies can pilot new AI technologies under the supervision of regulators without immediately being subject to all regulations. These sandboxes (potentially hosted by agencies like the FTC or sectoral regulators) will allow experimentation with, say, AI-driven healthcare diagnostics or fintech algorithms, while ensuring oversight to manage risks. The plan aims to launch sandbox programs in key sectors (health, finance, transportation) by next year, with *at least 20 companies/teams participating* and generating data to inform smarter future regulations.
- **Ethical Governance and Advisory Bodies:** Strengthen governance by incorporating ethical oversight in AI development. Agencies and companies will be encouraged to form AI Ethics Boards or advisory committees that include diverse stakeholders (technologists, ethicists, community representatives). These bodies can review AI initiatives for compliance with ethical guidelines and societal values. The federal government will support the ongoing National AI Advisory Committee and similar bodies to continually advise on AI policy, ensuring that governance keeps pace with technological advances. We will also promote the OECD AI Principles – which urge that AI be innovative, trustworthy, and respect human rights and democratic values – as a foundation for our governance approach. Success here is measured by the establishment of clear, trusted processes for AI oversight (e.g., a majority of federal agencies with significant AI projects have an internal AI governance policy and advisory process by 2025) and by maintaining public trust in AI as usage expands.

Research and Development (R&D)

Expand support for AI research and development to drive technological breakthroughs and maintain America's scientific and economic edge. The Federal Government will prioritize AI R&D in budgets and programs, in line with the national AI strategy to invest in the "Industries of the Future." This echoes long-standing policy objectives to promote sustained AI R&D collaboration across government, industry, and academia. Key components of this R&D agenda:

- **Increase Federal R&D Investment:** While avoiding specific dollar amounts, the plan calls for significantly ramping up AI R&D funding across agencies (NSF, DARPA, DOE, NIH, etc.). The focus is on *long-term, high-reward research* that industry might not undertake alone. Priority areas include fundamental AI algorithms, next-gen computing (quantum AI, neuromorphic chips), and applied research in fields like healthcare, agriculture, climate, and education. We set a quantifiable goal to double the number of federally funded AI research projects (grants, programs) over the next 5 years. Agencies will report on the number of AI projects and publications yearly to track progress. This fulfills the objective of promoting sustained investment to generate AI breakthroughs that enhance economic and national security.
- **National AI Research Infrastructure:** Build out the infrastructure needed to support cutting-edge AI R&D. This includes high-performance computing resources, cloud platforms, and testbed environments accessible to researchers nationwide. For example, the plan will establish at least three AI research testbeds (for instance, a large-scale autonomous vehicle testing environment, a secure federated learning network for health data, and a simulation platform for military AI testing). We will also make high-quality federal datasets and tools more accessible for AI training – expanding initiatives like *Data.gov* and domain-specific data trusts. *Measured outcome:* By 2026, the number of researchers (academic or small-business) who access national AI research infrastructure (computing or data) should increase by 50%, democratizing R&D beyond well-funded tech labs.
- **Public-Private Research Partnerships:** Encourage joint research endeavors between government, academia, and industry. Building on the success of programs like the National AI Research Institutes, the plan will create new AI innovation hubs that bring together university researchers, startups, and industry partners to tackle grand challenges (e.g., developing AI for clean energy or pandemic response). Our goal is to launch *8 new multi-sector AI institutes or centers* within the next 2 years, a 50% increase in such collaborative entities, covering all major regions of the country to spur local innovation ecosystems. These partnerships will ensure that research results transition into practical applications quickly, fueling economic growth.

- **Metrics and Evaluation of Progress:** Establish clear metrics for evaluating U.S. AI research leadership. For instance, track the number of top-tier AI publications by U.S. researchers, patents filed in AI, and international AI competitions won by U.S.-led teams. The plan aspires to achieve a *10% year-over-year increase* in AI patent applications and keep the U.S. at the forefront of key international AI benchmarks (e.g., ranking #1 in competitive tasks like vision recognition, or maintaining the majority of the world’s most powerful AI computing installations in the U.S.). These quantifiable indicators will help monitor the effectiveness of R&D investments and signal areas needing more attention.

Energy Efficiency and Sustainability

Mitigate the energy consumption and environmental impact of AI systems as their use grows. Advanced AI, especially large-scale model training and deployment, can be extremely energy-intensive – for example, training a single state-of-the-art model like GPT-3 (175 billion parameters) consumed about *1,287 MWh of electricity and produced 500+ tons of CO2 emissions*, comparable to the annual emissions of 100+ cars. This action plan integrates energy efficiency as a core priority to ensure AI advancements align with sustainability goals:

- **Green AI Research Initiatives:** Invest in R&D for energy-efficient AI algorithms and hardware. This means developing new techniques that require less computation (and thus less power) for training and inference, as well as supporting the design of specialized AI chips that are more power-efficient. Goals include achieving a 5-10x improvement in energy efficiency for training large AI models by 2030 (for example, through better algorithms or hardware accelerators). The government will launch programs (possibly via DOE and NSF) focused on “Green AI” – measuring researchers not just on accuracy of models, but also on energy usage. *Measurable target:* Publish energy usage metrics with all major federal AI research projects and incentivize projects that reduce compute requirements for the same performance.
- **Sustainable AI Infrastructure:** Promote the use of renewable energy and advanced cooling in data centers that power AI computations. Federal data centers and facilities used for AI will transition to clean energy sources in line with broader sustainability commitments (e.g., aiming for 100% carbon-free power for these centers by 2030). For the private sector, encourage and potentially offer fast-track permitting or recognition for AI data centers that are energy-efficient (e.g., using state-of-the-art cooling, waste-heat recovery) and predominantly powered by renewables. We will track the carbon footprint of AI operations: the aim is to reduce CO2 emissions per AI training task by 50% over the next five years, through efficiency gains and green energy adoption. Progress will be reported via an annual *AI Energy Impact Report*.

- **Energy Transparency and Standards:** Develop standards and best practices for reporting the energy usage and carbon impact of AI models. Just as we have model cards for transparency about data and performance, we will promote “Energy Cards” or similar documentation that companies and researchers can provide, detailing the compute resources and electricity used for training major models. This transparency will raise awareness and drive competition to optimize energy use. The plan will work with industry and IEEE/ISO to standardize how AI energy consumption is measured. By 2025, our goal is that all major AI research papers or product releases voluntarily include an energy impact statement (with metrics like kilowatt-hours consumed). This echoes the growing international call for sustainable AI practices
- **AI for Climate Solutions:** Leverage AI as a tool to improve energy efficiency across the economy. Support projects where AI optimizes power grids, improves energy storage, or reduces waste in manufacturing and transportation. For example, set a target to deploy AI-driven energy management systems in *at least 50% of federal buildings* by 2026 to cut energy use. By showing how AI can contribute to climate goals (e.g., smarter grids reducing blackouts, AI in agriculture lowering water and fertilizer needs), we reinforce a positive feedback loop: AI helping to solve the very environmental challenges it can exacerbate. Success will be measured in tangible savings – for instance, petroleum use reduced, megawatts saved, or emissions avoided due to AI-augmented optimizations, as reported by agencies and partners.

National Security

Implement AI initiatives that bolster national security, defense, and protection of critical technologies. The United States faces strategic competitors investing heavily in AI for military and intelligence purposes. Maintaining leadership in AI is not just an economic issue but a national security imperative. This plan focuses on both adopting AI for defense and shielding our AI advantage from adversaries:

- **AI Integration in Defense:** Rapidly expand the integration of AI systems into appropriate defense and intelligence operations to enhance capabilities. This includes applications such as AI-driven analysis of intelligence data, autonomous systems for surveillance or logistical support, and decision-support tools for commanders. The Department of Defense (DoD) should aim to be “AI-ready” across core functions by 2025, as recommended by the National Security Commission on AI. *Action:* Scale up successful pilot programs (e.g., predictive maintenance of equipment using AI, or cybersecurity threat detection as mentioned earlier) into fully deployed solutions across the armed services. Set a quantifiable objective that, within 3 years, at least 20 major defense programs incorporate AI and demonstrate improved performance or cost savings (for

example, reduced equipment downtime, faster intelligence processing). Regular progress updates will be provided to ensure these integrations are on track.

- **Protect Critical AI Technology:** Strengthen measures to prevent the transfer or theft of critical U.S. AI technology and expertise to adversarial nations. As stated in prior policy, maintaining our lead requires protecting the American AI technology base from acquisition by strategic competitors and adversaries. The U.S. will use export controls, investment screening (via CFIUS), and cyber protections to safeguard key assets like advanced AI chips, sensitive datasets, and intellectual property. *Policy action:* Update export control lists annually to cover the most advanced AI-related hardware and algorithms, balancing security with the need for a robust domestic industry. Also, enhance counter-intelligence efforts against industrial espionage targeting AI (e.g., espionage directed at tech companies or research labs). A measurable outcome is reducing instances of illicit technology transfer – tracked by enforcement actions or prevented attempts – and ensuring no significant loss of cutting-edge AI source code or models to hostile actors.
- **AI Workforce for National Security:** Dramatically expand AI training and recruitment within national security agencies (DoD, intelligence community, homeland security). The talent deficit in government is one of the greatest impediments to U.S. AI readiness. We will create specialized Digital Corps or AI career tracks to attract AI engineers and scientists into public service for national security. For example, establish an AI ROTC-style program: fund students through AI-related degrees in exchange for a term of service in defense or public sector roles. Aim to train at least 1,000 military and civilian personnel with advanced AI skills (through programs like the DoD’s JAIC initiatives or scholarships) by 2025, and continue scaling that number. Additionally, promote exchange programs where private sector AI experts can do short-term stints in government and vice versa, to cross-pollinate expertise.

International Collaboration

Engage with allies, partners, and multilateral forums to promote a global environment that supports responsible AI innovation and upholds democratic values. AI is a frontier not bound by national borders; U.S. leadership in AI must be exercised on the world stage to shape norms, standards, and partnerships. In accordance with policy, the U.S. will “promote an international environment that supports American AI research and innovation and opens markets for American AI industries”, while also protecting our interests. Major strategies:

- **Lead in Global AI Governance:** Take an active role in international forums setting the rules and norms for AI. This includes the Global Partnership on AI (GPAI) – a multi-country initiative to guide the responsible development and use of AI in line with human rights

and democratic values– as well as engagement with the G7 (e.g. the G7’s AI code of conduct), G20, and United Nations discussions on AI. The U.S. will use its influence to champion frameworks that emphasize trustworthy and human-centric AI globally (echoing principles similar to the OECD AI Principles). *Action:* Within the next year, convene a high-level Global Summit on AI with allies to align strategies on AI safety, trade, and security (building on existing efforts). Measure progress by concrete outcomes such as the adoption of international AI standards that reflect U.S. input, or new agreements on data sharing and research collaboration among like-minded nations.

- **Harmonize Standards and Regulations:** Work with allied nations to align AI standards and regulatory approaches to the extent possible, reducing fragmentation. Through bodies like the OECD and bilateral engagements (e.g. the U.S.-EU Trade and Technology Council’s AI working group), promote interoperability – so an AI system developed in the U.S. can be trusted and used in Europe, and vice versa, without compromising safety. For instance, collaborate on defining what constitutes “high-risk AI” and appropriate safeguards, so that democratic nations present a unified front versus more permissive or authoritarian standards. A practical step is to develop joint testing methodologies or certification processes recognized across countries (perhaps a mutual recognition agreement for AI certifications). Success would be reflected in fewer conflicts between regulatory regimes and in U.S. companies being able to export AI systems abroad seamlessly, bolstering our industry.
- **Joint Research and Talent Exchange:** Expand international collaboration in AI R&D. Increase funding for joint research projects between U.S. labs and those in allied countries (for example, co-funding AI research centers tackling global challenges like pandemic prediction or climate modeling). Create talent exchange programs that allow top AI researchers and students from partner countries to work/study in the U.S. and vice versa, building a network of experts aligned with democratic values. For measurable impact, aim to sponsor *at least 100 international AI research fellows* in U.S. institutions each year and double the number of collaborative publications between U.S. and allied researchers within five years. Such collaboration will not only advance science but also strengthen relationships and standardize ethical practices.
- **Global Market Access and Competition:** Work diplomatically to open markets for U.S. AI products and services. Push back against protectionism or unfair practices that could exclude U.S. companies (for instance, advocating against unjustified data localization requirements that hamper cloud-based AI services). Ensure that trade agreements include provisions on digital trade and AI that favor free flow of data with trust (consistent with privacy/security). Meanwhile, coordinate with allies on monitoring and responding to

strategic competitors' AI advancements – sharing assessments of adversaries' capabilities and coordinating export controls or sanctions when necessary to maintain a strategic edge. The outcome should be a robust international market where U.S. AI firms can compete and a coalition that collectively addresses any malicious uses of AI (such as disinformation campaigns or AI-enabled cyber attacks from adversaries).

- **Capacity Building for Developing Nations:** As part of international leadership, help build AI capacity in developing countries in a way that aligns with responsible use. This can involve providing technical assistance or tools for using AI in areas like agriculture, health, and education. By offering partnership (instead of allowing adversarial powers to fill the void with potentially exploitative tech), the U.S. can foster goodwill and ensure international AI progress is inclusive. For example, extend initiatives like AI for Good and work with organizations (World Bank, UNDP) to fund projects where U.S. tech companies and universities mentor or support AI deployments for social good abroad. Over time, this will enlarge the global community of AI practitioners who share our standards and bolster U.S. soft power.

Education and Workforce Development

Cultivate a skilled AI workforce and educate citizens to thrive in an AI-driven economy. Talent is the bedrock of AI leadership – the U.S. faces a global competition for AI expertise and must invest in both developing home-grown talent and attracting the best minds worldwide. This section outlines how we will expand education, training, and workforce initiatives to ensure a robust pipeline of AI-capable workers, while also creating opportunities for all Americans to benefit from the AI boom:

- **K-12 and STEM Education Expansion:** Integrate computer science and AI concepts into K-12 curricula to spark interest and build foundational skills early. By 2026, the plan is for every student to have access to some form of AI or computer science education before graduating high school (whether through coursework, after-school programs, or online modules). Federal grants and challenges (in line with the 2018 STEM Education Strategic Plan) will support states in teacher training and curriculum development for AI literacy. We also encourage project-based learning (e.g., students building simple AI projects) to demystify AI. *Metric:* Double the number of high schools offering AI-related courses or clubs in the next three years, and increase enrollment in high school computer science classes (as a proxy) significantly year-over-year.
- **Higher Education and Vocational Training:** Greatly increase the output of AI specialists from universities and training programs. As noted by the NSCAI, the number of U.S.-born students pursuing advanced AI degrees has stagnated since 1990– we must reverse this

trend. The plan will fund scholarships, fellowships, and targeted expansion of university programs in AI, machine learning, and data science. For example, establish a “National AI Fellowship” program to support an additional 500 domestic graduate students in AI-related fields annually (with a focus on diversity and underrepresented communities). Partner with community colleges and technical schools to create AI certification and apprenticeship programs for those not pursuing 4-year degrees, enabling pathways into AI technician roles (e.g., data annotation specialists, AI model testers). Aim to retrain workers from declining industries by enrolling at least 100,000 American workers in AI-skills training programs (online or in person) over five years. *Outcome:* A 50% increase in the number of AI and data science degrees/certificates awarded per year by 2028, helping meet industry demand.

- **Attract and Retain Global Talent:** Make the U.S. the most attractive destination for AI researchers and engineers worldwide. Recognizing that cultivating domestic talent and recruiting international talent are both essential to sustain U.S. AI leadership, we will streamline immigration pathways for individuals with AI expertise. This includes expanding the eligibility for *O-1 visas* (for individuals with extraordinary ability) in AI fields, clearing green card backlogs for STEM Ph.D. graduates, and potentially creating a new “AI Talent Visa.” The goal is to retain top foreign AI students graduating from U.S. universities (who might otherwise be recruited by competitors) and draw experienced professionals from abroad. We aim for a net gain of AI talent each year (as measured by jobs filled and companies started by immigrants in AI). Additionally, encourage U.S. AI experts working overseas to return or collaborate through “brain circulation” initiatives (like jointly funded labs).
- **Lifelong Learning and AI Literacy:** Implement programs to ensure the broader U.S. workforce can adapt alongside AI. This involves upskilling workers whose jobs may be transformed by AI automation – offering short courses or micro-credentials in using AI tools in one’s occupation (for example, AI in marketing, AI in manufacturing). The Department of Labor, in partnership with industry, will launch an AI Skills Toolkit and outreach campaign to help small businesses and workers understand how to leverage AI, not fear it. We set a target to provide AI awareness or basic training to 500,000 workers (across various sectors) within 5 years, which will be tracked by enrollment numbers in such programs. Moreover, promote digital literacy for all citizens: libraries, community centers, and online platforms will disseminate easy-to-understand materials about what AI is and how it affects daily life, so the public can engage with AI knowledgeably. This will help mitigate fear and misinformation, ensuring society at large is prepared for AI’s impacts on work and life.

By executing these strategic actions in Model Development, Cybersecurity, Privacy, Standards, R&D, Energy, Governance, National Security, International Collaboration, and Education/Workforce, the United States will foster an environment where AI innovation flourishes and American leadership is strengthened. This Action Plan emphasizes measurable outcomes – from doubling research initiatives to increasing skilled graduates and reducing security incidents – to track progress. It aligns with the Administration’s vision of removing unnecessary obstacles for AI developers while instituting responsible guardrails for safety and ethics. Through coordinated implementation (led by the White House’s science and security advisors in partnership with agencies and stakeholders), we will ensure that the transformative power of AI is harnessed in a manner that secures economic prosperity, upholds our democratic values, and protects national security for years to come.