

Attention To: Faisal D'Souza
Networking and Information Technology Research and Development
(NITRD) National Coordination Office (NCO), National Science Foundation

Snowflake Inc. ("Snowflake") Official Response To:

Request for Information on the Development of an Artificial Intelligence (AI) Action Plan

As the United States embarks on the development of a new National AI Action Plan, Snowflake welcomes the opportunity to contribute and collaborate. AI will drive economic growth, benefit the workforce, and improve the delivery of services across the commercial and public sectors. The potential impact of artificial intelligence and data is promising for economic growth and delivery across all industries.

The barriers to realizing this growth and ensuring lasting and safe adoption should not be underestimated. Balancing efficiency, performance, regulation, investments, privacy, and security are just a few of the areas to consider. Snowflake stands by the goals of growth and efficiency- these two areas have differentiated us in our customers' minds and decisions.

Our position as a global platform leader provides us with a detailed understanding of how thousands of enterprises are leveraging AI. This submission offers select insights for the development of the Action Plan. In a decade we have become the leading force on mobilizing the world's data- this is our charge. We support the goal of ensuring American leadership in AI.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

We appreciate the opportunity to respond to this RFI. We welcome the opportunity to speak further in upcoming sessions. Should you have any questions or wish to learn more, please reach out to Matthew Rose at [REDACTED] ¹

¹Snowflake submits this paper for informational purposes only and with the understanding that it will not create any legally binding obligations on either party, and that binding obligations will arise only upon execution of definitive agreements that are mutually acceptable to both parties.

US ARTIFICIAL INTELLIGENCE ACTION PLAN

Snowflake recognizes the importance of a coherent approach to AI (Artificial Intelligence) in order to improve economic growth, service interaction and the future of work. The United States' AI Action Plan should carefully choose initiatives that will create leverage, accelerate progress, and position us as the global leader in this race.

We continue to observe the importance of laying the foundations that enable AI. To that end, we highly recommend the United States include two critical areas in its Action Plan:

- **First, the United States should take lead in the development of secure, efficient, and performant AI Data Clouds.**
- **Second, the United States should enable AI data collaboration.**

The United States' AI Data Clouds

Governments are using AI to improve their services, products, workforce, and operations. Organizations of all sizes who have a scalable data foundation are able to adopt and scale AI: those who do not are falling behind. As a nation, the United States is positioned for expansion given multiple advantages, such as the availability of cloud infrastructure, computing, enterprise technologies, and policy.

As AI becomes the foundation of modern business strategy, one truth has become abundantly clear- enterprise AI cannot succeed without a robust data strategy. A data strategy isolated in silos is equally ineffective; siloed and incorrect data is costing companies up to 30% of their annual revenue.² Additionally, organizations are failing to use up to 97% of their data, leaving vast opportunities for innovation and growth on the table.

Organizations are increasingly recognizing the need for unified data platforms, but not all organizations can or want to centralize their data and governance. This is where interoperability becomes crucial; ensuring responsible data access and human oversight, powering everything from day-to-day decisions to advanced AI applications. Creating useful AI applications will not be achieved without solving the data silo issue.

The United States Government ("US Government") set supportive national policy with items like the Foundations for the Evidence-Based Policymaking Act, OPEN Government Data Act, Paperwork Reduction Act, and Removing Barriers to American Leadership in Artificial Intelligence (AI) Executive Order. The infrastructure has generally been normalized with advancements like hybrid public clouds. Investments from hyperscalers like Google, Amazon Web Services and Microsoft have made cloud services available across the United States and key international markets. In parallel to this, the compute required to train and provide

²IDC Market Research and Gartner.

inference for AI is also improving. These improvements are being realized and go beyond chip advancements. Foundational models are becoming more and more available- enterprises are now training industry and domain specific models with their own proprietary data.

The Data Layer remains a key lagging component; In stark contrast to other necessary foundational components, the development of a normalized data layer is absent. Moreover, the US Government is lagging behind other countries. Just within the Five Eyes (FVEY) nations, the UK, Canada, and Australia have set national goals and initiatives to establish national, industry, research, or department-equivalent data clouds.³ For example, New Zealand is outpacing others in government and is sharing data seamlessly across ministries. Retail, financial services, health and life sciences, and marketing are emerging industries where corporate and enterprise organizations are evolving beyond data collaboration for AI development. Research and academic institutions are often seeking to expand efficient partnerships for data access; access to government data is especially useful and important.

These findings are documented across the whole-of-government, Department, and programmatic levels. US Government organizations like the Government Accountability Office, Chief Data Officer Council, and Congress are in; the Federal CDO council reiterated findings that data discovery, inventory, collaboration and use remains a major challenge.⁴

The AI Action Plan should include initiatives to address this gap and formulate AI Data Clouds. The following initiatives could achieve progress towards this goal:

- **Informed Federated Learning Infrastructure:** The US Government must rapidly develop an industry-informed infrastructure to support federated learning. Some will indubitably involve the government operationally while others need balanced government regulation. Within six months, it should convene a panel of industry leaders to define use cases, address barriers, and guide the implementation of multi-cloud solutions and policy for secure AI model training across sectors.
- **Rapid AI Development:** To accelerate AI development, it is important to adopt a dynamic funding approach. This involves initially supporting a broad range of models, followed by rigorous evaluation to identify high-potential technologies. The selection process for the best-in-breed will ensure that the remaining resources are strategically invested in scaling impactful solutions.

³<https://www.canada.ca/en/government/system/digital-government/digital-government-innovations/responsible-use-ai/gc-ai-strategy-overview.html> <https://ised-isde.canada.ca/site/advisory-council-artificial-intelligence/en> <https://webarchive.nla.gov.au/awa/20220816053410/https://www.industry.gov.au/data-and-publications/australias-artificial-intelligence-action-plan> <https://www.gov.uk/government/publications/ai-opportunities-action-plan/ai-opportunities-action-plan>

⁴https://resources.data.gov/assets/documents/2021_DSWG_Recommendations_and_Findings_508.pdf

- **Secure Data Cloud for Model Research:** Several US adversaries and threat actors now publish their own foundational models, making AI applications available to US citizens with nefarious intent. A secure data cloud that allows the US Government to review these models and applications will be important. Similar sandboxes exist and have been proven invaluable in domains like cybersecurity.
- **National AI Data Clouds:** The US Government should develop federated National AI Data Clouds (Environments) that will facilitate collaboration, resource sharing, and data access. These National AI Data Clouds should be connected through innovative data sharing capabilities and concepts such as data clean rooms. This will enable secure cross sectoral data sharing and analysis while protecting sensitive information. Data clouds for research, transportation, education, testing and evaluation, and national security are well documented domains of need. A unified approach of federated data clouds will set the US apart from on-going international sovereign equivalents.

Expanding AI Data Collaboration

Data is the fuel that is powering the age of intelligence- without data, there is no AI. Useful AI applications require access to high-quality data: it follows that since AI represents unprecedented opportunity, data has become a precious resource. Real historical data, the likes of which the US Government has locked away on its many disparate networks, is more valuable than gold. Unifying this data for US Government employee use by marrying it with publicly available data will allow all key personnel to make better, faster decisions that serve US interests.

We have seen that data is not only an asset, but it is a reusable one. This central asset is driving growth and innovation: it improves services across health, commercial systems, and government. Data assets are also being referenced in companies that seek government loans. The US Government is not alone in this recognition- the United Kingdom, United Arab Emirates, China, Canada, Australia, and several others describe how data will drive their economic growth.⁵

We maintain one of the largest data marketplaces in the world. When data is categorized by industry segment, Snowflake's customers have shown that public sector data has an outsized impact on business and the development of AI. An anecdote showing this is the weather data market, which was projected to reach \$635M in 2022 and is expected to grow by 9% resulting in a \$1.5B market by 2029. The United Kingdom is one of several countries who have realized this and are moving towards advantage. The UK MET weather office has expanded in their data sharing, monetization and development from data to data products. Categories of valuable US Government data sets include non-personally identifiable sets and include categories like weather, transportation, government contracts, and policies. Many data

⁵World Governments Summit 2025- Report & Country AI Strategies from listed Countries

providers begin to share their data for cost recoverability- ensuring they can house the valuable and growing asset.

The US Government provided publicly available data repositories, such as Data.gov & Census Opportunity Project, that rely on data sharing methodologies like APIs. These methods are antiquated. APIs do provide automated access to the underlying data. They offer the ability for data replication and usage based off of the provided format and schema. The data providers are unable to gather telemetry on who is using the data; the lack of data telemetry is in part to data aggregators who are offering curation services in a hub and spoke model. Every time a data consumer copies the data, they must pay for the storage and subsequent curation for their use. Moreover, the data is then no longer current and canonical. For example: the FDIC's Bankfind API is especially useful to the financial services industry, but this data set must be curated with other government data sets in a common join table for AI and other applications. The US Government is unable to determine the true value of this to the market and specific consumers because of the old methodology.

Newer methodologies of data sharing provide access and do not replicate the data, thus eliminating the creation of silos. Users can now access data and allow for seamless joins regardless of the infrastructure such as On-Premise, AWS, Azure, GCP. This method will reduce the cost and improve access needed for AI development. It also addresses bi-partisan legislation cutting down on egress costs by eliminating the costs altogether. Moreover, it addresses the national approach to cybersecurity as data is not duplicated- it avoids another attack surface for bad actors.

AI Development and value requires diversified data sources. Probabilistic models, automated decisions are influencing citizens' lives increasingly and the quality and accuracy of data is paramount. AI enabled decisions need as much input as possible. Getting a second opinion is common practice for humans- increasing data sources is the equivalent for AI.

Data diversity is one means of mitigating the risk of AI models capturing internal bias and "hallucinating" or, plainly speaking, making mistakes. The following approaches apply:

- **Break down internal silos to access cross-functional sources:** Historically, data was isolated within applications or systems distributed across an organization. Modern data architectures must support discovery and access across organizational boundaries.
- **Transform unstructured data to expand available internal data:** To ensure that all data is made available, organizations must adopt tools to transform unstructured data into usable formats. Documents, emails, images, videos and voice recordings provide valuable input for training.
- **Collaborate with partners to access different data sources:** Data collaboration enables organizations to expand access to data across their business ecosystem.

- **Acquire and access third-party external data sources:** There is a growing realization that AI models capture historical biases. Expanding data sources can provide a bigger picture, reducing blind spots.
- **Create synthetic data:** Another approach is to create synthetic data to balance representation. If bias is expected or observed, new data can be created to increase under-represented characteristics.

The AI Action Plan should include initiatives to expand data collaboration, and thus AI development. The following initiatives could achieve progress towards this goal:

- **High Impact Datasets:** The US Government should identify at least 50 high-impact data sets that will make an outsized impact on AI Development and make them available to researchers and innovators. The data should be canonical and made available through a method that supports a myriad of infrastructures in order to reduce the cost burden.
- **US Sovereign Wealth Fund & Data Assets:** The development of the Sovereign Wealth Fund offers an opportunity for the US to create revenue from government data assets. The US should invest in its own data assets and remunerate the use and access. This will, at a minimum, recover the cost of storage and maintaining them; additional funding can be used to improve the data assets to make them more attractive for business.
- **Data Discoverability Initiative:** AI developers currently spend an inordinate amount of time and money finding, inspecting, and subsequently curating data sets from APIs/other sources to train models and develop AI applications. Improvements on data sets routinely do not make it back to the owner based on this model. The AI Action Plan should include an initiative for AI developers- both private, public, and academic- to locate data sets and inspect and inform data owners.
- **International Data Sharing Initiative:** The US Government should identify and create an international data sharing arrangement to hasten the development of AI. The international data sharing initiative would go beyond previous policy agreements and would allow access for developers to research.

AI ACTION PLAN INTEREST AREAS

Hardware and Chips

The availability of flexible and efficient computing environment(s) for organizations is crucial to any AI Action Plan. In the near term, we believe the availability of chips and hardware for some will be limited. In part, this is why Snowflake makes compute available to customers of all types through cloud service providers. A multicloud approach allows for access to compute resources where they are needed.

The ideal architecture is suited to minimize latency and maximize performance, ensuring that organizations can quickly and easily access the computing resources required to access and control its data. The AI Action Plan should consider approaches that abstract the hardware/chip layer. When abstraction is created, users can route their work to the optimal location. Methods that tie workloads to hardware should be avoided because they are inefficient and lack flexibility.

Data Centers

The hardware landscape supporting AI is rapidly changing: the accelerating growth of data centers is driven by the increasing importance of data in economic growth. This growth varies in deployment models, including sovereign clouds to public clouds. Snowflake's expansion is partially realized through US bilateral and multilateral efforts. The goal should be to unlock the data regardless of its location- this will be required for useful AI applications. To achieve this, an AI-powered data cloud must be scalable to support various workloads across different deployment methods, including Sovereign clouds, Public clouds, Hybrid environments, and Multicloud deployments.

An AI data cloud that has pervasive availability across all clouds is critical because it allows workloads to be shifted where the hardware is best and/or available. The tradespace between this concept and current state is time and money for things like migrations or network connectivity. Organizations should have the choice to run their workloads in the cloud provider, physical location, and isolation they need and prefer. Therefore, solutions that allow for flexibility in which data centers are used, how they are interconnected, and the ease of change to meet emerging needs will be more advantageous.

Energy Consumption and Efficiency

Energy consumption and efficiency are critically important to the success of an AI Action Plan. Driving efficiency is our company's north star, and we believe that this north star should be prioritized and incentivized within the AI Action Plan.

We consider efficiency across all aspects of AI; inefficiencies will drive increased energy consumption, maintenance, and cost. The infrastructure and services that will enable AI



development is varied, but we have found solutions that offer a fully managed, multi-tenant solution are uniquely positioned for efficiency. This model allows us to share resources (including GPUs) between multiple users at a fine-grained level, maximizing the efficiency of resource utilization and consumption. Charging organizations and users based on their consumption encourages them to use the most efficient, and therefore lowest cost, algorithms and methods. It also correlates with the value they receive from the system.

Finally, a managed service model is unique when coupled with automatic upgrades to the most efficient hardware and data centers. As an example, we have migrated customers from x86 chips to ARM chips seamlessly, receiving a significant boost to efficiency without customers having to make any decisions or change any of their code.

Model Development

Model Development is the epicenter of AI utilization: without a consistent method of refining the data, enhancements and insights stall. The AI Action Plan should support deployments that allow for multilevel modeling (MLM) in an efficient manner. One example of this would be users having the capability to engineer and preprocess ideal functions, allowing for accelerated training of patterns, which accelerates AI development.

Open Source Development

Open source development offers several benefits, such as transparency and community collaboration. In parallel, open source development comes with certain downsides that must be weighed against the benefits. These include areas like security vulnerabilities, reliability on community support, and compatibility. In contrast, proprietary models and development tools offer other different benefits and challenges. Broad development and growth will require emerging organizations, and some will not have the ability to negotiate for proprietary access. Based on these factors, we believe that a structured approach to open source and proprietary development should be used for the AI Action Plan.

The US Government should carefully consider the areas that are inappropriate for open source development, such as unique national security use cases. Snowflake is aware that there is no broad consensus on whether open sourcing AI development will be helpful or harmful vis-a-vis the United States national interests. A US Government assessment in this area in collaboration with global technology firms could be beneficial.

We have developed a method by which open source models can be hosted while still protecting data. Methods such as copying the open source models and hosting them in a data cloud that provides data protection are viable options in this space.

Application and Use (either in the private sector or by government)

Data remaining stagnant without applying it to business applications and government programs is inefficient. Value is optimized when AI applications use the data that is most appropriate and current to their use case. Data access must be secure and controlled by the data owner.

Platforms can offer unique capabilities that enable AI applications by providing a unified, scalable, and secure environment. This environment is used for data processing, storage, and AI model development, allowing users to build and deploy AI applications directly on their data without moving it.

Explainability and Assurance of AI Model Outputs

Users should be able to understand both the outputs and a model's behavior within the AI infrastructure that they are implementing. Based on this, our approach to AI explainability is designed to help users understand and explain their model's behavior by computing and analyzing data values. The goal is to improve model development by identifying issues and refinement capabilities. Providing detailed information about AI systems, including data sources and algorithms, helps provide transparency in these systems and how they are utilized to create the necessary outputs.

Cybersecurity

Applying artificial intelligence in a cybersecurity context allows organizations to strengthen their security posture with sophisticated solutions for identifying vulnerabilities, detecting and mitigating potential attacks, and enhancing incident response. Machine learning techniques and AI algorithms can help organizations successfully counter increasingly sophisticated and frequent cyber threats. Adaptability along with the ability to process and analyze data concurrently is what makes AI-enabled cybersecurity solutions so powerful.

This adaptability comes in many forms, including the below:

- Trained on the latest security intelligence, these tools can continuously adapt and evolve to counter new types of threats.
- Self-learning systems can learn normal patterns of user behavior and detect anomalies and variations that may indicate potential insider threats or account compromises.
- AI algorithms become more effective over time as they operate within a specific business environment.

We have aggressively invested in building a world class and differentiated data driven Security Operations Program with a focus on Threat Detection, Threat Intelligence, Incident Response, and Offensive Security. We proactively monitor for active threats to Snowflake systems, data, users, IP, and valued partners.



Regulation and Governance

Our stance on regulation and governance is clear: organizations and users retain ownership of their data. This stance should be supported and reinforced by the AI Action Plan.

To support this, we provide tooling that enables easy access to data, while also implementing object-based security controls to ensure compliance with relevant regulations. For enhanced security, governance frameworks with object-based controls are incorporated.

The solutions used by Snowflake are designed to address key governance elements, including:

- **Compliance:** meeting regulatory requirements
- **Security:** protecting data from unauthorized access
- **Privacy:** safeguarding sensitive information
- **Interoperability:** enabling seamless data exchange
- **Accessibility:** providing secure and controlled access to data

These elements were built in close collaboration with organizations across numerous industries and can be considered for the AI Action Plan & initiatives.

Technical and Safety Standards

International standards from bodies like International Organization for Standardization (ISO) and National Institute of Standards & Technology (NIST) provide solid frameworks for trusted and secure AI. However, most standards bodies are lagging in their adoption of advancing technological advancements and applicability on new frameworks. A one-size-fits all approach is potentially creating less secure and inefficient responses and audits, while also restricting innovation. The field of AI is advancing rapidly and standards bodies should consider a risk-based approach and collaboration process for AI developers.

Education and Workforce

To achieve lasting economic growth, it is crucial to work with groups in the industry to build AI literacy and provide relevant skills for US citizens. A multisectoral partnership approach with educational institutions, industries, and the government to promote hands-on skills and learning should be considered. Solutions that assist users, regardless of their technical background or coding experience, should be identified and promoted. Snowflake is investing in and advocating for more education on AI and data.⁶ Our Million Minds global program aims to train and certify 1 million students and professionals in high-growth and emerging markets by 2029 on the Snowflake AI Data Cloud, thus arming them with the knowledge and resources to pursue and expand careers in the technology industry. Snowflake welcomes the White House's partnership to make this available in the areas of greatest need or greatest impact.

⁶<https://www.snowflake.com/en/news/press-releases/snowflake-launches-one-million-minds-one-platform-program-investing-20-million-toward-ai-upskilling/>

Innovation and Competition

The US Government should consider adopting initiatives and solutions that are operational in like-scale industries. Fast moving industries and first-mover enterprises are driven by competition, and too often the government procurement processes reinforce past contract performance over technological performance. “Bake-off” arrangements could be expanded: often innovative projects at one agency are not shared across other agencies with similar workloads. Additionally, the US Government should consider technologies and solutions that offer continuous improvements- such as fully managed offerings- instead of service and SKU/feature based methodologies.

It is important to work with technologies that provide a seamless, always up-to-date experience for users while also delivering ever-increasing value through rapid development and continual innovation. Organizations should deploy new releases regularly with no downtime or disruption of service, ensuring the government is running on the most-recent release with access to the latest features.

Snowflake embraces being a leader in the industry and understands that competition is a component to driving innovation within the market. We continue to strive for excellence in a way that directly impacts our customers with a consistent and efficient solution.

Intellectual Property

The US Government should carefully consider intellectual property arrangements that do not limit growth and investment. Restrictive intellectual property clauses in acquisitions create increased costs for the government, while limiting the pool of solutions. Additionally, specific domains like national security or classified workloads may require unique arrangements.

We acknowledge that both Intellectual Property and Privacy laws and policies are not settled in the context of content for AI/ML and AI data collaboration. Our data marketplace- enabled by cross-sectoral organizations data sharing for their benefit on their terms- highlights the growing desire for balanced approaches that include new methods like data clean rooms.

Snowflake clearly outlines its policies surrounding Intellectual Property within its Terms of Service and its AI Terms of Service.⁷ Ultimately our customers retain their rights to their data as well as any modifications made while utilizing Snowflake. This is particularly important for AI usage, and is highly recommended for any AI tools the government wishes to implement, as it ensures that the primary usage of the AI tool does not affect the government’s control or ownership of its data. Notably, Snowflake commits that it will not use customer data or

⁷<https://www.snowflake.com/en/legal/terms-of-service/> and <https://www.snowflake.com/en/legal/optional-offerings/offering-specific-terms/ai-features/ai-terms/>

customer usage data (in each case, including any AI inputs or outputs) to train any model made available for use by other parties.

Procurement

The US Government remains a sizable buyer and can use its position to further the development of AI. The procurement of AI, applications and related services within the government is challenging. The challenges associated with procurement are well documented, with policies that create a “customer of one” within the government being a direct example.

Procurement that allows for scalable performance, lowered or shared operations and maintenance, and support multiple deployment architectures are ideal. Snowflake has observed that several Software as a Service providers are moving towards a consumption model. Consumption based pricing is more applicable for AI solutions than legacy methods of government procurement. The US Government should incentive performance and cost cutting improvements on procured solutions. To date, it is noticeably uncommon that technology improvements are prevented or disincentivized by government buyers.

International Collaboration

International collaboration is required to achieve the promise of growth, and will be needed beyond research and development. Snowflake’s platform was primarily designed for data collaboration, and our unique design allows Snowflake the ability to share data across Regions, allowing for international collaboration via controlled and secure methods. International cooperation and collaboration can be enabled by new technologies where data residency and sovereignty can be achieved. The AI Action Plan should ensure initiatives can build, share, and connect applications AI models both internally and externally.

The US Government’s AI Action Plan should not follow restrictive regulatory models. Organizations will unlock growth when they are able to use their data in new and unique ways. International initiatives like the EU Data Governance Act (DGA) are regulatory models where “novel data intermediaries” are introduced. The DGA is a cross-sectoral instrument that aims to regulate the reuse of publicly/held, protected data, by both boosting data sharing through the regulation of novel data intermediaries and encouraging the sharing of data for altruistic purposes. Both personal and non-personal data are in scope of the DGA, and wherever personal data is concerned, the General Data Protection Regulation (GDPR) applies. In addition to the GDPR, inbuilt safeguards will increase trust in data sharing and reuse, a prerequisite to making more data available on the market.⁸

With all of this in mind, the introduction of “novel data intermediaries” is misguided and unnecessary. A modern data architecture would not require an additional layer of technology or

⁸<https://digital-strategy.ec.europa.eu/en/policies/data-governance-act-explained>



an additional party to the data sharing: a data architecture should enable real-time direct access to data sources, secured through robust access and usage controls.

Export Controls

AI and infrastructure providers should comply with U.S. and international trade laws, which control where the company can send or receive its products and services, along with verifying to whom it can sell them. Technologies that allow for data residency to remain within regulated environments is critical to ensuring the needed measures are met.

The AI Action Plan should consider a balanced approach that weighs issues like national security and international economic growth. We recognize that our platform architecture and approach offers unique flexibility to adopt new controls.



ABOUT SNOWFLAKE

Snowflake's founders started from scratch and built a solution that would harness the immense power available in the public cloud. They created the Data Cloud- a global network where thousands of organizations mobilize data with near-unlimited scale, concurrency, and performance. Inside the Data Cloud, organizations unite their siloed data, easily discover and securely share governed data, and execute diverse analytic workloads. Wherever data or users live, Snowflake delivers a single and seamless experience across multiple public clouds. With the addition of Cortex AI to our solution suite, Snowflake is now able to empower customers to make better decisions, faster. Snowflake's platform is the engine that powers and provides access to this AI Data Cloud, creating a solution for applications, collaboration, cybersecurity, data engineering, data lake, data science, data warehousing, and unistore. Snowflake's vision is a world with unlimited access to governed data, so every organization can tackle the challenges and opportunities of today and reveal the possibilities of tomorrow.

Globally, Snowflake delivers the most performant data platform to over 11,000 organizations, powering 6.2 Billion average daily queries, with an innovative approach to delivering to governments and customers.

Snowflake continues to demonstrate strong financial performance and market presence including its most recent fiscal year (FY 2025). Snowflake reported a total revenue of \$986.8 million for the fourth quarter of FY 2025, reflecting a 27% year-over-year growth and full year revenue of \$3.5B, reflecting a growth of 29.2% compared to the previous fiscal year. With a growing number of Forbes Global 2000 customers (745 as of fiscal 2025), Snowflake has solidified its reputation as a trusted partner for large enterprises. ⁹

Snowflake is dedicated to serving regulated industries and environments. We are continuously expanding our portfolio of security and compliance certifications in cooperation with our customers. To help US, federal, state, and local agencies meet security and compliance standards, Snowflake's platform has achieved StateRAMP High, FedRAMP Moderate, FedRAMP High and Department of Defense Impact Level 4 (IL4) and is currently in the process of obtaining the Impact Level 5 authorization. We support regulated workloads subject to International Traffic in Arms Regulations, Criminal Justice Information Services, Internal Revenue Service 1075, Family Education and Privacy Act, and Federal Acquisition Regulation (FAR) and Defense FAR Supplement safeguarding requirements across certain of its U.S. government-designated and commercial regions.

⁹See Snowflake Investor Relations, available at investors.snowflake.com.

THE SNOWFLAKE PLATFORM

The Snowflake AI Data Cloud Platform is a robust, fully featured solution to process data that includes data warehousing, data lake, data engineering, data science and data sharing capabilities. The robustness of the platform and our ability to consistently innovate is a result of our architecture. Snowflake was founded in August 2012 as a cloud-native, SQL-compliant, relational analytical database platform, delivered as a fully-managed software as a service (SaaS). Snowflake founders recognized the challenges of legacy architectures and realized that creating a better solution required rethinking how data platforms and database architectures were built.

That's why the Snowflake team decided to write every line of code anew in the core data platform engine. By building an entirely new data platform and relational database management system from the ground up, they could deliver a dynamic infrastructure with instant, disruption-free scalability and performance at any cloud scale, all at a fraction of the cost of traditional systems. Our approach is unique and is tailored for cloud and the integration of AI.

Snowflake was built in the cloud, for the cloud, and offers virtually unlimited storage, and compute. Snowflake is a massively parallel processing (MPP) database that is fully relational, ACID-compliant, and processes standard SQL natively without translation or simulation.

The Snowflake AI and Data cloud platform can:

- **Store any type of business data:** Natively handles diverse types of data without requiring complex transformations before loading that data into the data platform.
- **Instantly scale for flexible performance and concurrency:** Able to scale up and instantly scale down at any time without disruption. It can scale out to as many use cases as needed without disruption.
- **Manage services:** The platform automatically handles management and infrastructure so that users can focus on getting value from their data instead of tuning infrastructure.
- **Fit seamlessly with existing skills and tools:** Snowflake leverages industry-standard data types, open interfaces, and common languages to enable innovation and insights, without incurring technical debt or skills gaps.
- **Offer unparalleled intelligence, natively:** Snowflake offers AI-powered data driven insights, runs analytical workflows on unstructured data, develops agentic apps, and trains models using both structured and unstructured data — all with minimal operational overhead and end-to-end governance.