

From: [Stampy's AI Safety Info](#)
To: [ostp-ai-rfi](#)
Subject: [External] AI Action Plan
Date: Monday, March 17, 2025 9:19:30 AM

CAUTION: This email originated from outside your organization. Exercise caution when opening attachments or clicking links, especially from unknown senders.

Name of the organization filing this comment: Stampy, fiscally sponsored by [Ashgro Inc.](#)

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

We would like to submit several proposals for your consideration. First, a proposal for increasing AI security: "Foom Liability". That is, a liability for rapidly improving AI systems which pose a risk of human extinction. Second, a Robots-Took-Most-Jobs Insurance, to deal with advanced AI automating many jobs. Third, a proposal to aid AI innovation by making AI models open-weights once they are several years behind the frontier.

Section 1: Foom liability.

The body of this proposal was copied from Robin Hanson's essay on the topic. [1]

"I instead prefer robust policies, ones we can expect to promote total welfare for a wide range of parameter values. We want policies that will give big benefits if foom risk is high, but impose low costs if foom risk is low. In that spirit, let me suggest as a compromise a particular apparently-robust policy for dealing with AI foom risk.

If you recall, the foom scenario requires an AI system that **a**) is tasked with improving itself. It finds a quite unusually lumpy innovation for that task that is **a**) secret **b**) huge **c**) improves well across a very wide scope of tasks, and **d**) continues to create rapid gains over many orders of magnitude of ability. By assumption, this AI then improves fast. It somehow **e**) become an agent with a **f**) wide scope of values and actions, **g**) its values (what best explains its choices) in effect change radically over this growth period, yet **h**) its owners/builders do not notice anything of concern, or act on such concerns, until this AI becomes able to either hide its plans and actions well or to wrest control of itself from its owners and resist opposition. After which it just keeps growing, and then acts on its radically-changed values to kill us all.

Given how specific is this description, it seems plausible that for every extreme scenario

like this there are many more “near miss” scenarios which are similar, but which don’t reach such extreme ends. For example, where the AI tries but fails to hide its plans or actions, where it tries but fails to wrest control or prevent opposition, or where it does these things yet its abilities are not broad enough for it to cause existential damage. So if we gave sufficient liability incentives to AI owners to avoid near-miss scenarios, with the liability higher for a closer miss, those incentives would also induce substantial efforts to avoid the worst-case scenarios.

In liability law today, we hold familiar actions like car accidents to a negligence standard; you are liable if damage happens and you were not sufficiently careful. But for unfamiliar actions, where it seems harder to judge proper care levels, such as for using dynamite or having pet tigers, we hold people to a strict liability standard. As it is hard to judge proper care levels, it makes sense to use strict liability there.

Also, if there is a big chance that a harm might happen yet not result in a liability penalty, it makes to add extra “punitive” damages, for extra discouragement. Finally, when harms might be larger than the wealth of responsible parties, it makes sense to require liability insurance. That is, to make people at substantial risk of liability prove that they could pay damages if they were found liable.

Thus I suggest that we consider imposing extra liability for certain AI-mediated harms, make that liability strict, and add punitive damages according to the formulas $D = (M+H) \cdot F^N$. Here D is the damages owed, H is the harm suffered by victims, $M > 0, F > 1$ are free parameters of this policy, and N is how many of the following eight conditions contributed to causing harm in this case: self-improving, agentic, wide scope of tasks, intentional deception, negligent owner monitoring, values changing greatly, fighting its owners for self-control, and stealing non-owner property.

If we could agree that some sort of cautious policy like this seems prudent, then we could just argue over the particular values of M, F . [And we'd] want not much more than D to be paid to those harmed, the rest to be paid to a much more diffuse recipient."

Further reading:

1) Weil, Gabriel, Tort Law as a Tool for Mitigating Catastrophic Risk from Artificial Intelligence (January 13, 2024). Available at SSRN: <https://ssrn.com/abstract=4694006> or <http://dx.doi.org/10.2139/ssrn.4694006>

Section 2: Robots-Took-Most-Jobs-Insurance

The body of this proposal was copied from Robin Hanson's essay on the topic. [2]

"While far from obvious, it is plausible that someday machines may displace most all

humans from their jobs. It is also not crazy to think that this might happen relatively suddenly, and without much warning. Which seems a pretty big problem, given that most people don't own much more besides their ability to earn wages. Without some sort of charity or insurance, they might starve. Thus each of us seems well-advised to try to set up some insurance re this risk, instead of just hoping to rely on charity.

One needs to set up insurance well before problems are realized or revealed. Yet it can be hard to motivate people to insure against risks that seem too unlikely or remote in time.

Which might make the current moment an ideal time to consider this. Many are now worrying loudly about AI, saying that AI might soon take all the jobs, or worse. And yet I'm pretty sure that most investors see this risk as actually still pretty remote and unlikely.

Maybe making now a great time to set this up.

I say "insurance" but, as this risk is easily measured and widely shared, we wouldn't need to use the usual insurance industry. Here is a simple plan:

- A) Carefully define the key event E of "automation suddenly takes most jobs". Maybe "labor force participation rate falls from >35% to <10% in <10 years by date D".
- B) Collect some diversified financial assets A likely to retain substantial value both after such an event, and also if that event never happens. Such as global stock, bond, and real estate index funds.
- C) Split these assets A into "A if E" and "A if not E". This split can be done with no risk. We should always see prices satisfy $p(A) = p(A \text{ if } E) + p(A \text{ if not } E)$.
- D) Sell "A if E" assets to workers. The more they buy, the better insured they are. They could buy these slowly over time, instead of all at once.
- E) Sell "A if not E" assets to any willing investors. Buyers of this are in effect selling robots-take-most-jobs insurance.
- F) When the event E happens, "A if E" assets turn into A assets, which can then be sold off to pay for ex-worker living expenses. If E seems likely to happen soon, "A if E" assets can also be sold then to pay for living expenses of early job losers.
-

G) As date D approaches, workers switch to buying assets with later dates D'.

- H) If date D comes without E ever having happened, "A if not E" assets turn into A assets, giving their investors a higher return than if they'd just bought A.
- I) At all times the price ratios $p(A \text{ if } E)/p(A)$ for various dates D warn us all via a probability distribution over time of when robots might take most jobs.

And that's it. If people vary in their tolerance for the risk of E, then there are gains from trade in having some people hold "A if E", while others hold "A if not E". And thus there is value to be released by splitting assets A into "A if E" and "A if not E". Yes, regulations may now stupidly prevent selling such split assets to workers; most of the work here may be to overturn such regulations.

Sure, there would be some work to do to advise workers of how well they could expect to live when holding how much of each kind of A asset. Workers should prefer global portfolio assets A, to insure against regional risks, and should consider the risk-return tradeoff re different kinds of assets A.

But those are minor issues; the main priority is to get workers to hold such assets. Maybe tech firms could signal their concern about AI by buying such assets for their employees. And then maybe cities or regions could buy some on behalf of their citizens. (Note: due to regional risks, planning to tax some local citizens to support others can fail badly.)

Note that whether this plan is a good idea doesn't depend much on the chances or timing of machines taking most jobs. The lower are these chances, the cheaper is this insurance, making it still a wise precaution.

Concretely, The cost of insuring against a 1% per year risk of workers losing a \$30K US median wage job, is ~1% of that wage, or \$300/yr."

All e-mails to and from this account are for NITRD official use only and subject to certain disclosure requirements.

If you have received this e-mail in error, we ask that you notify the sender and delete it immediately.