

AI Action Plan

Attn: Faisal D'Souza, NCO (National Coordination Office)
2415 Eisenhower Avenue Alexandria, VA 22314

Dear Mr. D'Souza and Colleagues,

We appreciate the opportunity to provide input on the development of the AI Action Plan. Our comments below highlight critical areas we believe are essential to include in the US AI Action Plan, to avoid falling behind, to enhance America's AI leadership, and to balance innovation, risk reduction, and strategic interests.

Trustworthiness and Internal Alignment As AI agents take on increasingly significant roles—from financial transactions to critical infrastructure—trustworthiness and internal alignment (aligned with human intent) become paramount. We must ensure that AI systems can reliably execute user-intended goals. For instance, can we trust an AI agent to move money in and out of a bank account without error? Clearly communicating the reliability of a model will be critical to broad adoption of AI technology by consumers and business-to-business, which is necessary to establish and maintain leadership. Research in internal alignment (for example, [Utility Engineering](#)) demonstrates both progress and remaining challenges in this domain. The Wall Street Journal recently reported on the challenges of getting a Chatbot to say "I don't know" instead of hallucinating (making up) an incorrect answer. We recommend that the AI Action Plan prioritize the development of tests and metrics standards to evaluate the internal alignment and trustworthiness of AI systems.

Metrics and Benchmark Standards Reliable benchmarks like [Harmbench](#), which showed OpenAI's o1-preview model resisting 77% of attacks versus DeepSeek R1 failing completely, underscore the need for standardized metrics. These metrics must emphasize accuracy and trustworthiness to reliably compare AI capabilities, incent leadership level responsible scaling, and reduce AI exploitation risks.

Threats, Risks and Capabilities from Open-Source Fine-Tuning Rapid proliferation of open-source models necessitates monitoring and managing potential threats and risks, particularly in cybersecurity, chemical, biological, and other sensitive domains. Advancing capabilities in open-source models allows for custom fine-tuning of such models to increase opportunities for misuse. Implementing metrics to track these risks, and leveraging training methods such as [Anthropic's Constitutional AI](#) and [OpenAI's deliberative alignment](#), can significantly reduce the chance of harmful misuse, which could undermine US leadership due to lost trust.

Standards, Liability Clarification, and International Cooperation Clear liability frameworks are critical for responsible AI adoption without excessive regulation. Strict liability with liability waivers (sometimes called safe harbors) for trust enhancing, risk reducing actions, such as above can accelerate leading development without burdensome regulation. Frameworks similar to those used for autonomous vehicle liability can improve transparency and risk communication. Aligning expectations will also allow for 3rd party companies to thrive by providing review and evaluation services in a common way to an actively evolving industry. Additionally, we strongly advocate US-led global standardization efforts, collaborating as needed with international partners like the EU to ensure broad, consistent adoption of robust AI practices.

Infrastructure, Hardware Resilience, and Strategic Supply Chains Addressing compute shortages, such as those encountered during ChatGPT-4.5 deployment, highlights a need for robust, scalable infrastructure. We urge inclusion of policies that support robust, scalable, and energy-efficient AI infrastructure. Additionally, ensuring hardware resilience is vital, given potential impacts of transient hardware faults on model reliability. Securing access to advanced semiconductor manufacturing, domestically (Intel's upcoming 1.8nm process, AMD's energy-efficient initiatives) and internationally (e.g., TSMC's \$100B US investment and Taiwan-based 2nm process), is critical to maintaining AI leadership. Continued support for Taiwan, including freedom of navigation and ensuring Taiwan does not fall under adversarial control, remains strategically essential to US AI leadership.

Responsible Scaling and Coordination Promoting the current practice of AI development companies publishing "[responsible scaling policies](#)" to encourage coordination among developers, users, and evaluators is critical. Standardized communication of risks and capabilities reduces national security risks and strategic risks and supports US leadership via wider adoption of trustworthy AI.

Data Integrity and Protection Against Data Poisoning High-quality, secure training data is foundational to AI reliability. Protecting training data integrity from adversarial poisoning—particularly from foreign threats—is critical. Partnerships with specialized data integrity evaluation companies like Scale AI could enhance data verification standards. Additionally, anonymization techniques should be emphasized to safeguard privacy (e.g., in medical datasets).

Strategic Alignment with US Interests We recommend clearly aligning AI development goals with US strategic interests to ensure dedicated and urgent attention from policymakers. Guidelines should highlight the importance of ensuring American AI systems are secure and aligned with democratic values. Policies must actively counter adversarial AI use and foreign efforts to dominate the field, to maintain US technological sovereignty and leadership.

In conclusion, we believe that the above areas are essential for an AI Action Plan that maximizes US AI leadership—one that advances innovation and human flourishing while ensuring robust trustworthiness, security, and accountability. By addressing internal alignment, establishing common benchmarks, clarifying standards, and reinforcing infrastructure resilience, the United States can maintain and grow its leadership in AI and promote broad, responsible adoption across the economy.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Sincerely,
Steve Kommrusch
Director of Research and Senior AI Scientist
Leela AI

Henry Minsky
Co-founder and CTO
Leela AI

John Diamant

Retired Distinguished Technologist, Software Assurance and Secure Development Chief
Technologist

Formerly with Hewlett-Packard

Dan Yee

Retired Senior Member of Technical Staff in silicon design

Formerly with Advanced Micro Devices