



Response to the NITRD NCO RFI on Development of an Artificial Intelligence (AI) Action Plan

March 14, 2025

The Transparency Coalition



This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Overview

About the Transparency Coalition

The Transparency Coalition's mission is to champion policies that ensure AI technologies are developed and used in ways which prioritize safety, transparency, and the public good.

We believe artificial intelligence has the potential to be a powerful tool for human progress if properly trained and deployed within guardrails that protect both US citizens as well as businesses adopting generative AI in their products.

We believe policies must be created to ensure model developers respect the rights of creators, reduce the potential for harmful disinformation, protect personal privacy, as well as fulfill their legal duty of care to both their business customers and end-users. This is accomplished by providing transparency and accepting accountability for the inputs and outputs of generative AI.

Key Insights

A lack of trust is stalling corporate adoption of generative AI.

In a Q4/2024 report from Deloitte on the State of Generative AIⁱ, mistrust in generative AI (Gen AI) was identified as a key barrier to business uptake:

- 35% of responding organizations said their top potential barrier to adopting Gen AI is mistakes / errors with real-world consequences.
- 29% were reluctant due to a general loss of trust due to the potential for bias, hallucinations and inaccuracies.
- 25% were worried about the liability of intellectual property contained in the training data used by the model being incorporated into their own product offerings.

These data points show that potential customers are deeply concerned about their ability to trust what went in to and comes out of Gen AI models.

These “duty of care” product liability concerns can only be addressed by the Gen AI developers *themselves*.

Potential customers wanting to deploy products that depend on large language models (LLMs) understand they are ultimately bound by the longstanding “duty of care” doctrine that underpins all modern product liability lawⁱⁱ and they are right to be concerned about putting their trust in Gen AI. Presently, the risk inherent in creating products and solutions powered by Gen AI developers passes *straight through* to the deployers and their derivative products since these LLM developers are today largely unaccountable.

Would-be corporate deployers of Gen AI have legitimate concerns on both sides of the model. Data being used to *train* the model can cause problems that range from copyright infringement to disclosure of personally identifiable information (PII). Furthermore, even a harmless model could be trained with biased data that skews its output.

Problems with content *generated* by a model can range from just giving incorrect responses that are somewhat off base to completely made-up “hallucinations” where the answers have no factual basis whatsoever. Then there is the problem of Gen AI developers not automatically labeling and providing user-viewable provenance that the content created was made with AI. Doing so would not only mitigate the impact of the unintentional disinformation; it would also help contain the spread of malicious deepfakes, social media posts and phishing emails that disinform, misinform and defraud our citizens.

Some of the very worst instances of these problems have intersected to contribute to despondent users' self-harm after being subjected to a dangerous chatbot conversation, or after learning of malicious pornographic deepfakes of themselves.ⁱⁱⁱ

Innovating with Gen AI today is too risky for most American companies.

Today, this unmet duty of care means innovating on top of a Gen AI model is too risky for all but the largest corporations that can afford to build their own trust, risk and security scaffolding around LLM outputs to mitigate risk.^{iv} In contrast, a single issue with a hallucination, copyright or PII violation could put a startup out of business, and it can take months for the underlying Gen AI developer to remedy with their current reactive approach.

America cannot continue to lead in Gen AI adoption and innovation when the rest of the AI ecosystem is concerned they will proliferate and amplify the risks that the Gen AI model developers presently pass onto them in their models without consequence or assurances.

Cybersecurity is threatened when the data supply chain is unsecured.

It is nearly impossible to guarantee data integrity and security in Gen AI implementations without strict adherence to data governance and cybersecurity best practices.

Securing the data supply chain is a foundational component of such practices and requires full disclosures on training data provenance and outputs. Only by following best practices is it possible to ensure that contaminated generative digital content is not used unwittingly to train and adversely impact future generations of models.

Emerging state regulations on Gen AI developers mandate a duty of care that helps businesses deploying AI as much as private citizens.

Several states are working on or have already enacted legislation to hold Gen AI developers accountable to their duty of care responsibilities by mandating copyright protection, personal privacy preservation, user safety protocols, and the labeling of Gen AI model outputs, which in turn helps businesses deploying Gen AI meet *their* duty of care.

Rectifying training data problems in LLMs *reactively* is very expensive for Gen AI developers and too slow a remedy for users and deployers.

While the state-of-the-art in AI changes daily, at present there is technically no way to completely remove the impact of a single data point from a Gen AI large language model without removing the data and then *retraining the entire model completely from the start*. Retraining an LLM like Open AI's ChatGPT or Anthropic's Claude can take months to complete which means the affected user whose PII was exposed or copyright infringed will

wait *months* for a remedy. This also means that an AI deployer's product leveraging a Gen AI LLM could be in violation of several state and international copyright, privacy and AI safety regulations from the moment they deploy and stay that way until the Gen AI developer retrains to resolve the underlying data issues. Retraining also cost hundreds of millions of dollars *every time* it is performed with the newest models being expected to cost more than one billion dollars to train.^v

The ecosystem definitely needs to move away from these reactive remedies, but that does *not* mean we should allow the Gen AI developers to be unaccountable for these copyright and privacy violations.

The next wave of legislation and industry action will drive *proactive* transparency from Gen AI developers.

Moving the industry to proactive and programmatic model and data transparency protects US citizens and AI deploying organizations alike by making it easy for Gen AI developers to cost-effectively address their duty of care.

A proactive approach allows content owners and users to control whether their data is used to train a model in the first place. In this manner they can prevent many, if not most, data infringements before they can happen. This in turn will lead to models that are safer to deploy with fewer risks or adverse outcomes. Which will then result in more deployers who build on foundational models to move forward with deployments. Furthermore, these legislative approaches will also facilitate the creation of Small Language Models (SLMs) that contain curated and properly licensed domain-specific data that are intended to address specific use cases/domains. A combination of these industry actions, in response to thoughtful legislation, will allow the U.S. economy to achieve the full benefits of an AI roll-out without falling victim to scaling challenges. When content owners control their data and its use as AI training material, GenAI developers and deployers will have the assurance they need to innovate without the legal uncertainty of copyright infringement or data misuse.

Our recommendations for enabling this proactive transparency make up the remainder of this document.

Recommended Actions

To enable proactive and programmatic model and data transparency, Transparency Coalition recommends the following policy actions by legislators which can in turn, be implemented by a Gen AI developer in a straightforward manner:

- **Make it clear when content is AI-generated**
 - Inform consumers when an image, video, sound, or text has been created or modified by AI. Embed that information in all AI-generated material.
- **Publish AI training data ingredient lists**
 - Developers of AI systems should be required to provide documentation for the training data used to develop an AI model.
- **Require opt-in consent to use personal data**
 - Flip the paradigm and put people in charge of their personal data. Require consumers to intentionally "opt-in" to allow tech companies to collect and use their personal information.
- **Minimize the personal data collected and kept**
 - Limit data collection to information necessary to perform a transaction or optimize a user's website experience—and nothing more.

Make it clear when content is AI-generated.

The ability to know whether an image, video, sound, or text was created by AI is critical to the healthy functioning of society. Deepfake videos and other misused AI can cause profound harm to critical components of our society, including the criminal justice system, medical industry, and electoral process.

AI Developers should always Inform consumers when an image, video, sound, or text has been created or modified by AI and attach that information in all AI-generated material.

Require the embedding of provenance metadata in AI created content

Lawmakers should require AI developers and deployers to embed provenance data within all digital objects created or modified by a Gen AI system. Provenance data is coded within metadata for the purpose of verifying the digital content's authenticity, origin, or history of modification.

Require Gen AI companies to provide a tool to detect provenance metadata in their generated content; Require large online content platforms to use it.

Gen AI companies must offer a tool to detect the embedded provenance metadata and large online platforms where content can be shared must be required to use it and make the latent provenance data available for the end user to view.

Provenance disclosure as a legal baseline and foundational requirement.

California will require this embed-and-disclose capability beginning on Jan. 1, 2026, due to the AI Transparency Act (SB 942). Other states should enact similar legislation to give their citizens the ability to sort authentic evidence from machine-made make-believe. State laws requiring AI disclosure tools set a fair and appropriate legal baseline for the tech industry, protecting ethical AI companies from bad actors offering deceptive and malicious products.

AI disclosure laws like SB 942 incentivize the creation of ethical AI tools and products while driving unethical and malicious actors out of the marketplace. They establish appropriate guardrails that encourage high-quality innovation and discourage the spread of misinformation and deceptive imagery.

Content Credentials provenance technology is available today.

Some leading companies—including Adobe, Nikon, and Microsoft—are already working with a tool known as Content Credentials which is based on an open technical specification developed and maintained by the Coalition for Content Provenance and Authenticity (C2PA). C2PA is a cross-industry standards development organization whose steering committee includes Google, Amazon, Adobe, Meta, Microsoft, and OpenAI as well as several traditional content providers like Sony and the BBC who intend to include provenance data with their news and entertainment content to establish its authenticity in the fight against disinformation and knockoffs.

Content Credentials embeds provenance data in a digital object that accompanies the content. Clicking a small pin reveals that information, as seen in figure 1




Figure 1-Adobe's implementation of Content Credentials

To be clear, this technology is available *today* and can be adopted by media companies, the Gen AI developers and content sharing platforms like Facebook and YouTube *immediately*. There is *no* barrier to these AI and media leaders adopting C2PA—they helped make the technology!

Other forms of provenance are also available to developers.

Other readily available techniques used to embed provenance in an AI-created image or video include:

- **Watermarking:** These are visible or invisible marks added to digital media files to indicate their origin or ownership. Watermarks can be textual or graphical symbols that are embedded in images, videos, or audio files. (e.g. Most people are familiar with the well-known Getty Images watermark, which allows users to view the image while preserving Getty’s copyright protections.)
- **Digital Signatures:** Digital signatures use public-key encryption to generate a unique code that is attached to a digital media file. The code can be verified by anyone who has access to the public key of the signer.
- **Blockchain:** Blockchain can be used to create and manage digital assets on the web, such as cryptocurrencies, tokens, smart contracts, etc. Blockchain can also be used to record and preserve provenance data for AI-generated content, such as source, creation process, ownership, and distribution.

Publish AI training data ingredients lists.

Developers of AI systems should be required to provide high-level documentation for the training data used to develop an AI model.

The Transparency Coalition has formulated a Data Declaration (Fig. 2) that would contain basic information about the data used to train the AI model. TCAI's Data Declaration is fully compliant with California AB 2013's requirements. California's training data disclosures will be required starting on Jan. 1, 2026. Legislators in several other states are considering introducing versions of AB 2013 for their jurisdictions in the sessions that open in Jan. 2025.

This type of auditable information set provides transparency and assurance to deployers, consumers, and regulators. It is similar to industry-standard SOC 2 reports, which assess and address the cybersecurity risks associated with software or technology services.

FIELD NAME	POSSIBLE VALUES
Data Set Name	Text
Data Set Owner	Text
Data Set Description	Text
Data Set Size	Numerical
Data Set Category	Web text, images, music, video, books
Data Set License Type	Commercial License, Proprietary, Public Domain, Fair Use claim
Data Set License Name	e.g. GPL, Apache, Creative Commons
Data Set Collection Period	Start date, End date (or Present)
Data Set Usage Period	Start date, End date (or Present)
Data Set contains personal or personally identifiable information	Yes or No
Personal Information Opt-in obtained	Yes or No
Personal Information License mechanism	EULA, Terms of Service, Privacy Policy, Click Through
Personal information anonymized prior to training	Yes or No
Data Set contains Copyrighted Information	Yes or No
License governing Copyrighted Information	Fair use, Commercial license
Synthetic Training Data use	Yes or No

Figure 2- Transparency Coalition's Data Declaration Template

To be clear, a Data Declaration is not necessarily tied to government oversight. Rather, it should become a standard component of every AI model—expected and demanded by AI system deployers as a transparent mark of quality and legal assurance.

The technology to provide Data Declarations exists already and is being standardized by D&TA.

Data and Trust Alliance's (D&TA) Data Provenance Standards define the declaration of name-value pairs for several industry use cases and the export of that metadata into several standard markup formats including JSON, XML and YAML. Furthermore C2PA, the same technology that powers Content Credentials, could be used in conjunction with

D&TA's data declarations as provenance information about the *model itself* and not just the content it creates.

Training Data is not a trade secret.

As the need for AI training data transparency gains traction in policy circles, some of the most powerful AI companies are pushing back. They argue that this basic descriptive information about the data used to train AI models constitutes proprietary information. In other words, they claim it's a trade secret. It's not.

The ingenuity in AI does not lie in the datasets. In fact, many of today's AI systems were trained on datasets like the Common Crawl, WebText2, Books1, and Wikipedia that have been shown to be problematic in the ways noted earlier.

The innovation in AI lies in the construction of the model itself. Developing an AI model requires months or years of work envisioning the system, lining up compute power, creating the algorithms, training the model, weighting the data, and building the end-user interface.

Data Declarations destroy Gen AI developer's plausible deniability.

The real fear of Gen AI developers in disclosing this most basic information is explicitly asserting whether the training data contains copyrighted materials or PII. Many of the initial "frontier" models were trained on raw, unlicensed data that may not have been legally obtained^{vi} or properly anonymized.

Require opt-IN consent to use personal data to train AI.

Flip the paradigm and put people in charge of their personal data. Require consumers to intentionally "opt-in" to allow tech companies to collect and use their personal information.

Most privacy laws in the United States operate on an opt-out basis, meaning that users are assumed to have consented to the use of their personal data unless they actively decline, i.e., opt-out. The EU's General Data Protection Regulation (GDPR), by contrast, is an opt-in mechanism. That means companies may not use personal data unless an individual actively offers consent to do so, i.e., opt-in.

State and federal lawmakers should require opt-in consent from consumers to use an individual's data to train an AI system. In other words: Companies should assume they have no right to use an individual's data unless that individual gave them affirmative consent.

Individuals are the rightful owners of their personal data. Companies should assume they cannot collect, store, or use personal data to train AI models unless given affirmative consent. Opt-in should not only be the industry standard; it should be the law of the land. California's Privacy laws were amended to cover personal data used to train AI and even go so far as to classify the insights into a user's personal data made by AI as *also* belonging to the user and other states are drafting similar legislation.

It's much easier for everyone to leave personal information out of the training data by default than to try and remove it later.

It is expensive in terms of time, money and energy consumption to retrain a frontier model just to remove the impact of one piece of improperly obtained personal data, and it's no harder than the GDPR cookie consent pop-up to implement explicit opt-out. It's the right thing to do for consumers and it will save the Gen AI developers billions in retraining costs by not including the personal data in the first place.

Minimize the personal data collected and kept.

'Data minimization' is a privacy principle that limits personal information collected from a consumer to data strictly relevant and necessary to accomplish a specific purpose.

While not strictly an "AI action" data minimization legislation prevents the collection and abuse of personal information before it starts, which *does* include Gen AI developers.

Data minimization includes limits on both data collection and data retention:

- **Collection:** Only collect personal information that is directly relevant and necessary to the purpose.
- **Retention:** Only keep the data for as long as it is necessary to fulfill the purpose. No perpetual storage.

Data minimization best practices

Data minimization best practices map directly to similar legislative requirements

- Time limitations: Establish retention policies for various data types. For example, retain email data for 90 days, retain employment data every year-end for 5 years.
- Establish periodic reviews of retained data: Identify when and how personal information is being deleted or purged.
- Establish solutions for deleting personal information upon an individual's request.
- If personal information is used as a unique identifier (such as a social security number), consider whether it is possible to use or create an alternate ID.

Respect “Do Not Train” data set and website designations

Labeling a certain set of data as “Do Not Train” should prevent all AI developers from using the data to train AI models. The effect of a Do Not Train (DNT) data designation would be similar to a copyright page in book publishing or a robots.txt file in web crawling; it announces that the following material is off-limits for use as AI training data.

Implementing “Do Not Train”

Currently the AI community is coalescing around the use of “ai.txt” as an AI version of the robots.txt directive. An ai.txt file embedded in a website’s root directory allows or denies AI developers the use of a domain’s text or media files to train AI models.

An early version of DNT data designation has been created by Spawning.ai, an independent third party AI governance start-up. Spawning maintains a Do Not Train registry and provide machine-readable opt-out tools for domain hosts. Two of the world’s biggest AI developers, Stability and HuggingFace, have partnered with Spawning and agreed to honor their DNT registry. Stability is the creator of the Gen AI image-creation system Stable Diffusion. HuggingFace is the world’s largest repository of models and datasets.

Starting in January 2025, many state legislatures will consider bills to enhance the safety and transparency of AI systems and those proposals should require Gen AI developers to honor the ‘DNT data’ designation, just as search engines today honor the directives of robots.txt as they crawl the digital world.

Transparency Coalition’s training data request prompts

As part of our mission to create AI safeguards for the greater good, Transparency Coalition is introducing new command concepts designed to infuse Gen AI systems with a new level of transparency.

Training Data Requests (TDRs) offer content creators, copyright owners, and individuals a basic level of agency in the use of their data property. We have developed two of these TDRs for consideration in AI-related bills during the upcoming 2025 state legislative sessions:

Training Data Verification Request (TDVR)

This is a mechanism by which a primary content owner submits a verified request to a Gen AI developer to inquire if their content is included in the AI model’s training dataset.

Training Data Deletion Request (TDDR)

This is a mechanism by which a primary content owner submits a verified request to a developer to delete content that was or will be included in a generative artificial intelligence training dataset.

Designating a dataset as ‘Do Not Train’ data may be done only by the primary owner of that specific data. Similarly, TDDRs may be sent only by primary content owners.

A “Primary content owner” means a person, partnership, or company that owns, in full or part, digital data, content, or objects that are subject to copyright protection. The definition is also meant to include an individual with personally identifiable information (PII) whose PII has been included in the Gen AI model’s training dataset.

Hold Gen AI Developers accountable for privacy and copyright violations by giving them a deadline to move to a proactive and scalable data removal process.

Transparency is foundational to the creation of a vibrant and safe American AI ecosystem. The current reactive approach that has arisen to remedy data set problems in Gen AI doesn’t serve anyone in the ecosystem well and needs to be addressed to prevent America’s AI innovation from stalling and to allow consumers to get a timely remedy.

Give Gen AI developers a deadline by which they must move to proactive and programmatic data and model transparency using the recommendations described in this document. This protects US citizens’ rights, mitigates the risk of AI deploying organizations and makes it easy for Gen AI developers to cost-effectively address their duty of care by not training on the problematic data in the first place.

Resources

The Transparency Coalition maintains a resource repository for legislators, policy makers, journalists, thought leaders, and researchers. The modules, articles, and guides presented here are intended to explain fundamental concepts in artificial intelligence and AI governance in accurate and non-technical language. New articles are added as the technology and language of AI evolve—and they're evolving quickly.

To learn more about the Transparency Coalition's top remedies for current risks in AI safety and transparency, see our Solutions page.

<https://www.transparencycoalition.ai/>

Endnotes

ⁱ <https://www2.deloitte.com/content/dam/Deloitte/us/Documents/consulting/us-state-of-gen-ai-q4.pdf>

ⁱⁱ <https://www.transparencycoalition.ai/learn/what-is-a-duty-of-care-and-how-does-it-apply-to-ai>

ⁱⁱⁱ <https://theconversation.com/deaths-linked-to-chatbots-show-we-must-urgently-revisit-what-counts-as-high-risk-ai-242289>

^{iv} <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>

^v <https://hai.stanford.edu/ai-index/2024-ai-index-report>

^{vi} <https://www.technologyreview.com/2024/07/02/1094508/ai-companies-are-finally-being-forced-to-cough-up-for-training-data/>