

March 15, 2025

**Comments from Wolters Kluwer on Development
of President Trump's Artificial Intelligence Action Plan**

Below are comments from Wolters Kluwer to the Networking and Information Technology Research and Development National Coordination Office (NITRD NCO) of the National Science Foundation and the White House Office of Science and Technology Policy (OSTP) in response to the recently released Request for Information (FR Doc. 2025-02305) on the development of an Artificial Intelligence (AI) Action Plan. We appreciate the opportunity to share our views.

As way of background, Wolters Kluwer is a leading global provider of clinical technology and evidence-based solutions that assist effective decision-making and outcomes across the healthcare continuum. Key solutions include UpToDate®, UpToDate® Lexidrug™, UpToDate® Patient Engagement, Medi-Span®, Senti7®, Lippincott® Solutions, Ovid®, and Health Language®. Wolters Kluwer had annual revenue in 2024 of \$6.4 billion.

This document is approved for public dissemination. The document contains no business-proprietary or confidential information. Document contents may be reused by the government in developing the AI Action Plan and associated documents without attribution.

Our comments propose concrete policy actions the Trump Administration should include in its AI Action Plan that will speed the development and adoption of generative AI-enabled (GenAI) health software used in clinical care. The specific policy actions we recommend are summed up immediately below, followed by a detailed discussion of each.

AI Action Plan Policy Recommendations Related to GenAI-enabled Health Software

1. Federal regulation of AI should be sector-specific, not one-size-fits-all.
2. FDA regulation of AI should focus on device software functions that are “agentic,” that is, capable of acting independently of the medical professional.
3. HHS should seek public feedback in 2026 on the health care industry’s experience implementing the Decision Support Interventions criterion that took effect January 1, 2025.
4. FDA’s September 2022 Final Guidance on Clinical Decision Support Software should be withdrawn and new Guidance issued that better aligns with Congress’ intent to exempt low-risk CDS software from FDA oversight.
5. To facilitate the development of AI health solutions, the AI Action Plan should include protections for identifiable health data not already protected by existing federal law.

6. The AI Action Plan should maintain the current legal liability framework that governs the use of AI-enabled software in clinical care.
7. The AI Action Plan should clarify that existing federal copyright protections for creators whose content is used to train AI models remain in place and will be rigorously enforced.

Discussion

1. Federal regulation of AI should be industry-specific, not one-size-fits-all

New federal regulations governing the use of AI should be industry-specific, accounting for the relative benefits and risks posed by AI use in each sector of the U.S. economy. The European Union's new AI law is misguided in this respect, as it takes a one-size-fits-all approach, disregarding the unique market dynamics that might differ from sector to sector. We believe this will stifle development and adoption of AI solutions in Europe, specifically in industries where a lighter regulatory approach may be advisable.

Ensuring a sector-by-sector approach to U.S. federal regulation is particularly important for the use of AI in health care, where there is significant potential for the technology to improve patient outcomes and lower costs, but also concerns for patient safety if the software does not function as designed. Federal regulations should therefore be carefully tailored to ensure the benefits of improved diagnosis, more effective treatment and enhanced clinician productivity, while properly managing the risks posed by AI. Safeguarding patient safety must be the primary goal of any regulatory framework, but engendering trust among clinician users, and preserving incentives for the private sector to invest in AI tools must also be key objectives.

2. FDA regulation of AI should focus on device software functions that are "agentic," capable of acting independently of the medical professional

Crafting a regulatory approach to AI-enabled health software depends not only on striking the correct balance between protecting patients, rewarding innovation and encouraging clinician adoption, it also must reflect the staff and resources available at federal agencies to conduct proper oversight.

In several [public statements](#) and an article published in late 2024,¹ FDA leadership expressed concerns about the Agency's ability to field sufficient staff and resources to properly review the expected wave of new AI-enabled device submissions, while also noting that hospitals and health systems may need to take more responsibility for ensuring the accuracy and safety of AI-enabled tools used in patient care. To optimize the use of FDA's staff and

¹ Haider J. Warraich, MD; Troy Tazbaz, BS; Robert M. Califf, MD *FDA Perspective on the Regulation of Artificial Intelligence in Health Care and Biomedicine* Journal of the American Medical Association doi.10.1001/jama.2024.21451 Published online October 15, 2024

resources, we recommend the Agency focus its active regulatory oversight on GenAI-enabled device software functions that pose the highest risk to patient safety. In our view, these are software functions that are “**agentic**,” meaning they can act autonomously, making care decisions and/or acting without the need for human intervention.

Though the deployment of agentic functions can add significant value and efficiency to the clinical workflow, freeing up trained human professionals to address more complex patients, they also pose significant threats to patient safety if the software does not perform as designed. Focusing FDA staff and resources on GenAI-enabled functions that are agentic also has the virtue of aligning with the Agency’s historical regulation of medical devices that operate independent of human intervention.

Remaining GenAI-enabled device software functions that are non-agentic, and thus incapable of acting independently of human intervention, would still technically remain under FDA’s regulatory authority, but be subject to enforcement discretion. FDA could provide guidance that defines the terms of such discretion. For example, such device software functions would only be granted enforcement discretion if their intended purpose was to merely assist human decision-making. The software’s design would require a trained human intermediary to make a clinical decision based on their experience and judgment, with the software itself having no agentic capability.

Additional transparency requirements could also be tied to enforcement discretion. For example, the developer would need to disclose their software’s GenAI-related source attribute information similar to the disclosures for predictive decision support interventions used with ONC-certified technology (see our next recommendation for more detail). Developer disclosures related to their risk identification and mitigation practices could also be required, with the National Institute for Standards and Technology’s AI Risk Management Framework (GenAI profile) serving as a model for such practices.

We appreciate that the FDA must act judiciously in granting enforcement discretion to any subset of medical devices subject to the Agency’s regulatory oversight. We believe the requirements outlined above related to non-agentic GenAI-enabled device software functions create sufficient safeguards to allow for safe, effective use of such solutions without the need for active FDA oversight.

We also expect market forces, specifically the hospitals, health systems and physicians who use these solutions at the point of care, will play an important complementary role in ensuring patient safety. A [survey](#) conducted in late 2024 by Wolters Kluwer found that 40% of the U.S. physicians polled were now ready to use GenAI when interacting with patients at the point-of-care. The findings reflect a rapid acceptance of the new technology more broadly, with 68% saying they have recently changed their views and are now more likely to think that GenAI would be beneficial to healthcare.

However, the survey also revealed that physicians will not blindly follow GenAI recommendations and that they would be selective when evaluating and implementing GenAI tools. Before relying on GenAI tools for assistance in clinical decision-making, 91% of respondents indicated they would need to know that GenAI-sourced materials were created by doctors and medical experts. Similarly, 89% reported that they would require vendors to be transparent about where information came from, who created it, and how it was sourced. If these survey results are indicative of the wider provider community, and we believe they are, clinicians will not use GenAI-enabled tools at the point of care unless they understand how such solutions were developed and trust their output.

3. *HHS should seek public feedback early in 2026 on the health care industry's experience implementing the Decision Support Interventions criterion that took effect January 1, 2025*

HHS's Office of the Assistant Secretary for Technology Policy (ASTP) finalized its Health Data, Technology and Interoperability (HTI-1) rule in December 2023. In that rule, ASTP finalized several changes to the ONC Health IT Certification program, including replacing the then-current *Clinical Decision Support* certification criterion with a new *Decision Support Interventions* (DSI) criterion. The new criterion was designed to address the use of AI in certified health systems and modules. To promote trust among medical professionals and spur adoption of AI tools used at the point of care, the new criterion includes transparency (or "source attribute") disclosures for DSI used in or with certified technology. The final rule also made the distinction between "evidence-based" and "predictive" DSI, with product examples illustrating the difference. In ASTP's view, GenAI solutions would likely meet the definition of a predictive DSI if used with certified technology, thereby triggering additional disclosures related to risk mitigation practices. The new DSI criterion took effect January 1, 2025.

We appreciate that the Trump Administration is actively searching for existing regulations to rescind, thereby lightening the administrative and compliance burden on private entities, a goal we generally support in this context. Regarding the new DSI criterion, however, the Administration may want to take a wait-and-see approach.

Compliance with the DSI criterion is technically voluntary for developers, not a legal requirement. Moreover, we consider its burden manageable, merely requiring disclosures of information related to the design, development, testing, validation and ongoing monitoring of the software's performance and, in the case of Predictive DSIs, risk mitigation practices. Even in the absence of the DSI criterion, such disclosures would almost certainly be requested (and likely required to be made) by hospitals, health systems and clinics before they commit to purchasing and deploying DSI solutions. As the results of our recent physician survey reveal (see our previous section), clinicians will not use GenAI-enabled tools in patient care unless they understand how such solutions were developed and trust their output.

Should the Trump Administration decide to leave the DSI criterion in place for now, HHS might consider issuing a Request for Information (RFI) at the 12- or 18-month mark of implementation (i.e., January or July 2026). The RFI could seek public input on whether the 31 required source attributes are excessive and assess if the disclosed information effectively fosters trust and adoption among clinician end users.

4. *FDA's September 2022 Final Guidance on Clinical Decision Support Software should be withdrawn and new Guidance issued that better aligns with Congress' intent to exempt low-risk CDS software from FDA oversight*

In December 2016, Congress passed the 21st Century Cures Act, which, among other things, carved out low-risk clinical decision support software from the definition of 'medical device,' thereby exempting such software from FDA oversight. Specifically, the statute added section 520(o) to the Food, Drug and Cosmetic (FD&C) Act, which excludes certain software functions from the device definition if four criteria are met:

- 1) The software function is not intended to acquire, process or analyze a medical image or a signal from an invitro diagnostic device or a pattern from a signal acquisition system;
- 2) The software function is intended for the purpose of displaying, analyzing or printing medical information about a patient or other medical information;
- 3) The software function is intended for the purpose of supporting or providing recommendations to a health care professional about the prevention, diagnosis or treatment of a disease or condition;
- 4) The software function is intended for the purpose of enabling such health professional to independently review the basis for such recommendations that such software presents so that it is not the intent that such health care professional rely primarily on any of such recommendations to make a clinical diagnosis or treatment decision regarding an individual patient.

In September 2022, FDA issued final policy guidance on Clinical Decision Support (CDS) Software, which provided the Agency's interpretation of the 21st Century Cures Act statutory language noted above. While not ostensibly about artificial intelligence, FDA's guidance includes product examples that use algorithms to generate clinical recommendations and therefore impacts the Agency's approach to regulating AI.

In our view, FDA's interpretation did not adhere to the plain language of the 21st Century Cures Act, choosing instead to discuss concepts such as the "time-critical nature" of the medical decision and "automation bias," which FDA defined as the "propensity of humans to over-rely on a suggestion from an automated system"

In the guidance, FDA noted that automation bias was more likely to occur if the software provided a user with a "single, specific, selected output or solution rather than a list of options or complete information for the user to consider." FDA also said the "time-critical

nature” of the medical decision might also increase the risk of automation bias “because there is not sufficient time for the user to adequately consider other information.” Thus, in FDA’s view, software that provides a specific output or directive and/or addresses a time-critical decision would be interpreted as “directing the (health care professional) to take action and substituting for their judgment,” and therefore likely subject to FDA regulation as a medical device.

First, it is important to note that neither “automation bias,” nor the “time-critical nature” of the medical decision appear in the 21st Century Cures Act, either as statutory language or even as concepts to be considered by FDA. To the contrary, as the language above shows, the statute was quite clear in emphasizing that the “intended” purpose or design of the underlying software should be the primary factor to consider in each of the 4 criteria related to product exemption.

Second, while we agree it is important for GenAI-enabled software to present clinicians with recommendations incorporating all appropriate options, and to do so in a way that allows clinicians to adequately consider all relevant information, we do not agree that GenAI-enabled software that presents only a single recommendation and/or supports time-critical medical decisions necessarily increases the risk of automation bias.

There are many instances in clinical care when compelling scientific and medical evidence indicates that there is only a single selected output or solution that is correct for the patient. For example, if a particular infection can be treated by multiple antibiotics, but a patient is allergic to all but one, then it would be inappropriate to suggest anything other than the one antibiotic that is safe for that patient. Similarly, if one treatment has overwhelmingly been shown to be the best approach for treating a particular condition, then suggesting alternatives could lead to a patient receiving suboptimal care. GenAI-enabled software should be designed to provide a range of options for the clinician when appropriate, but if there is only one best option dictated by the evidence and the patient’s situation, then only one should be presented.

In addition, many CDS solutions assign confidence levels to their output, thereby giving the clinician additional reason to either accept or question a single recommendation. Confidence in a medical recommendation is often communicated in medical practice guidelines and elsewhere using well-defined methodology, such as GRADE. Note also that using confidence levels with GenAI-enabled software will require assurance that the output grade reflects the accumulated medical evidence and the conviction of the medical community about an appropriate course of action, and not the software’s belief that it found and incorporated the appropriate content when making the recommendation.

FDA has recently spoken to the issue of regulating CDS software that issues a single recommendation in its response to a Citizen’s Petition challenging its CDS Guidance.² There, the

² FDA’s January 17, 2025, Response to the CDS Coalition, denying their Citizen’s Petition (Docket Number FDA-2023-P-0422)

Agency argued that its decision to continue regulating such software is consistent with the language in Section 520(o) of the FD&C Act because Congress used the plural, “recommendations,” in describing Criteria 3 and 4 of the exempting language. Thus, any software that may deliver a single recommendation would not be exempt from the Agency’s oversight.

We strongly disagree with this interpretation. A more accurate read is that Congress understood that CDS software is designed to provide many recommendations over the course of its lifecycle to many clinicians covering multiple patients, addressing thousands of care events that require a corresponding number of clinical decisions. Most CDS output will provide clinicians with multiple options, but in some clinical situations as we outlined above, there is only one right recommendation. Even CDS software used as little as 3 times per day in a small clinical practice could be classified as having issued “recommendations,” even if on those 3 occasions the output was a single recommendation that represented the best evidence-based option for the patient.

Regarding the time-sensitivity of medical decision-making, in practice, identifying such decisions is not as clear as it might seem due to numerous factors, including the patient, their presentation, and the setting in which treatment is being delivered. For example, the management of a patient with gastrointestinal bleeding may be emergent if the patient is severely bleeding and is in shock, whereas it may be less urgent if the patient had a single episode of bleeding an hour earlier and is hemodynamically stable.

The ramifications of FDA’s overly narrow interpretation of Section 520(o) have led to fewer CDS solutions being exempt from FDA oversight. New product innovations and features, possibly including AI-enablement, could be (and likely already have been) abandoned, thus depriving providers and patients of enhanced functionality that could improve health outcomes and lower costs.

Dr. Scott Gottlieb, who served as Commissioner of FDA during President Trump’s first term, recently raised similar concerns with FDA’s interpretation of Section 520(o) and how it might stifle the use of AI in CDS software. In an article that appeared in the online version of the *Journal of the American Medical Association* in early February,³ Dr. Gottlieb said:

“If these tools are classified as medical devices merely because they...possess analytical capabilities that are so comprehensive and intelligent that clinicians are likely to accept their analyses in full, then nearly any AI tool embedded in an EMR (electronic medical record) could fall under regulation. The risk is that EMR developers may attempt to circumvent regulatory uncertainty by omitting these features from their software. This could deny health care clinicians access to AI tools that have the potential to transform the productivity and safety of medical care.”

³ *JAMA Health Forum*. 2025;6(2): e250289. doi:10.1001/jamahealthforum.2025.0289

Dr. Gottlieb went on to conclude that “(i)f these AI tools are designed to augment the information available to clinicians and do not provide autonomous diagnoses or treatment decisions, they should not be subject to premarket review,” a point that comports with Congress’ intent to loosen oversight of low-risk CDS software (and also lends support to our second policy recommendation, above, that FDA focus its regulatory resources on GenAI-enabled software that is “agentic,” capable of acting autonomously without human intervention).

We therefore urge the Trump Administration to include in its AI Action Plan a directive to the FDA’s Center for Device and Radiological Health to review its FDA Final Guidance on Clinical Decision Support Software with a view to ultimately rescinding it and issuing new Guidance that more closely reflects Congress’ intent to exempt low-risk CDS software from FDA regulation.

5. To facilitate the development of AI health solutions, the AI Action Plan should include protections for identifiable health data not already protected by existing federal law

Identifiable patient health data might be the most important resource for successful development and advancement of AI-related health tools. Such data is critical for training, testing and validating the performance of new algorithms deployed in both back-office administrative tasks or used at the patient’s bedside. But development of new AI-enabled software also poses a risk to privacy if AI developers do not put sufficient safeguards in place to protect such data.

Health care has an existing federal privacy statute, the Health Insurance Portability and Accountability Act (HIPAA), which sets standards for safeguarding protected health information (PHI) and places limits on the use and disclosure of such information without the patient’s consent. HIPAA’s narrow scope leaves potential gaps in the protection of patient personal data, particularly regarding the third-party use of AI to process patient personal data. For example, unless an AI developer is producing an AI solution for a specific HIPAA “covered entity” (e.g., hospital) or “business associate” (e.g., medical lab), they would likely not be subject to HIPAA requirements to protect PHI.

The Federal Trade Commission (FTC) recently clarified the Health Breach Notification Rule’s (HBNR) applicability to health apps and other technologies and entities not traditionally covered by HIPAA and who engage in the processing of personal health records (PHR). The HBNR requires notification of customers and the FTC following a breach of unsecured personally identifiable health information but does not specifically mandate the implementation of safeguards to protect patient health data. Like HIPAA, HBNR has a limited scope and applicability, which fails to mitigate the risk of identifiable health data not being safeguarded in the context of AI.

It is unlikely the Trump Administration can act unilaterally to address these gaps without the involvement of Congress. As such, we recommend the Administration’s AI Action Plan

articulate the following principles: a) Individuals should maintain control of their health data, including whether it is used in developing, training, testing and validating AI algorithms to be used in health care delivery. This would include the ability to opt-out; b) Developers of AI solutions intended for broad commercial sale should be required to implement privacy protections for any identifiable, individual identifiable health data in their possession or control; and c) Clarify that the FTC will continue to actively enforce the Health Breach Notification Rule.

6. The Trump Administration should maintain the current legal liability framework that governs the use of AI-enabled software in clinical care

The development, deployment and use of AI tools in patient care should remain subject to existing federal and state laws related to medical device liability, medical malpractice and other relevant civil liability.

There have been suggestions made that clinicians should be granted a liability “safe harbor” when using AI at the point of care. We think this would be a mistake, particularly because for the foreseeable future, many AI tools developed for bedside care will be designed to merely assist the medical professional make a clinical decision. This design assumes the professional will use their experience and medical judgment to make the final decision regarding patient care, relying on an AI-generated care recommendation as only one of many data points for consideration. This is entirely appropriate as trained medical professionals are licensed to practice medicine by the states, but AI tools used by these professionals are not. Giving protection from liability would also make clinicians more vulnerable to automation bias, which is of great concern to FDA as discussed above in Recommendation 4.

While we fully encourage widespread adoption and use of AI-enabled tools at the point of care, we also believe it important that clinicians maintain a healthy skepticism of such tools, continue to scrutinize the provenance of their source data, remain curious as to their underlying rationale and provide robust feedback on software performance. Such an approach will optimize the use of such tools to improve decision-making, while improving the tools themselves. On the other hand, a liability safe harbor for clinicians would discourage such practices and ultimately slow the adoption of safe and effective AI in health care.

7. The AI Action Plan should clarify that existing federal copyright protections for creators whose content is used to train AI models remain in place and will be rigorously enforced.

To date, concerns expressed about unauthorized use of copyrighted content in training large language models primarily pertain to the works of authors, artists, musicians and actors. But medical content creators such as researchers, professional societies, and clinical decision support developers would also be negatively impacted should AI developers ignore or violate federal copyright laws by using protected materials in the development of their products.

We strongly urge the AI Action Plan to clarify to AI developers that their use of copyrighted materials can only be done with the holder’s permission, and that federal

copyright protections will be vigorously enforced.

We appreciate the opportunity to comment. If NITRD NCO or OSTP have questions or would like to discuss our comments in more detail, please contact Bob Hussey at

[REDACTED]