

Wu-chang Feng

The non-stop and rapid pace of change in technology and AI has made education and workforce development challenging. Courses developed years ago are now obsolete, leaving students ill-prepared for a job market that has seen rapid change over the last several years. In our experience attempting to deliver a course applying advances in generative AI towards solving problems in cybersecurity (Labs are available at <https://codelabs.cs.pdx.edu>), we have seen that a different approach is necessary when attempting to prepare students for an AI-enabled future. We believe that such an approach will allow the US to remain competitive in the technological future. Our approach embodies the main aspects below: 1) Creativity. AI has reduced the importance of the mechanics of technology. Impactful solutions are now driven as much by the ability for someone to imagine an application as it does actually being able to build one. Towards this end, an approach where students learn a breadth of topics and then practice creatively building AI-enabled applications that apply them will be paramount. 2) Skepticism. AI and reports of its capabilities can lead many to believe that models can solve problems well beyond their ability to do so. It is important that students are continually exposed to problems with known solutions that AI can not solve at all or can not reliably solve. 3) Change. New AI models are released and deprecated over a period of months. Tasks that models failed at before, may be reliably solved with newer models. It is important that students develop an agile mindset, understanding that a narrative they believed about AI previously, may no longer apply currently. 4) Security. Most developers of AI applications have no formal training in cybersecurity. As a result, simple vulnerabilities such as prompt injection, excessive agency, insecure tools, insecure output handling, and over-reliance are proliferating the marketplace of AI applications. It is essential that students get practice exploiting vulnerable AI applications in order for us to ensure that our AI future isn't controlled by adversaries able to exploit insecure applications. 5) Cost. One of the problems with AI applications is that models are being used to solve problems that they are not trained to solve or that can be solved more efficiently with alternate approaches. Rather than throw everything at the model to perform, it is essential that students understand and be able to apply the range of technological tools, both AI and non-AI ones, to build robust, cost-effective solutions. For more information about our approach, a paper can be found here: <https://arxiv.org/abs/2501.10900> and a talk on the approach given at the 2024 NSF Workshop on LLMs for Network Security can be found here: <https://www.youtube.com/watch?v=CTRGYAA-MtE> Work supported by NSF Award #2335633